

Drifting Games

Combining expert advice

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:
expert 1: 1,0,1,0,1,1,1

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- **N** experts make predictions:
 - expert 1: 1,0,1,0,1,1,1
 - expert 2: 0,0,0,1,1,0,1

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:
 - expert 1: 1,0,1,0,1,1,1
 - expert 2: 0,0,0,1,1,0,1
 - ...

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:
 - expert 1: 1,0,1,0,1,1,1
 - expert 2: 0,0,0,1,1,0,1
 - ...
 - expert N : 1,0,0,1,1,1,0

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:
 - expert 1: 1,0,1,0,1,1,1
 - expert 2: 0,0,0,1,1,0,1
 - ...
 - expert N : 1,0,0,1,1,1,0
- Algorithm makes prediction: 1,0,1,1,0,1,1

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:
 - expert 1: 1,0,1,0,1,1,1
 - expert 2: 0,0,0,1,1,0,1
 - ...
 - expert N : 1,0,0,1,1,1,0
- Algorithm makes prediction: 1,0,1,1,0,1,1
- Assumption: there exists an expert which makes at most k mistakes

Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:
 - expert 1: 1,0,1,0,1,1,1
 - expert 2: 0,0,0,1,1,0,1
 - ...
 - expert N : 1,0,0,1,1,1,0
- Algorithm makes prediction: 1,0,1,1,0,1,1
- Assumption: there exists an expert which makes at most k mistakes
- Goal: make least mistakes under the assumption (no statistics!)

Analysis using meta-experts

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert = base-expert + (sequence of locations of mistakes).

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert = base-expert + (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert = base-expert + (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert = base-expert + (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10
 - No mistake: 0,0,0,0,0,0,0,0,0,0

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert = base-expert + (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10
 - No mistake: 0,0,0,0,0,0,0,0,0,0
 - 1 mistake: 1,2,3,...,10

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert = base-expert + (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10
 - No mistake: 0,0,0,0,0,0,0,0,0,0
 - 1 mistake: 1,2,3,...,10
 - 2 mistakes: (1,2),(1,3),(1,4),..., (9,10)

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert = base-expert + (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10
 - No mistake: 0,0,0,0,0,0,0,0,0,0
 - 1 mistake: 1,2,3,...,10
 - 2 mistakes: (1,2),(1,3),(1,4),..., (9,10)
- If the number of mistakes is at most 2, then there is a meta-expert that makes no mistakes.

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert = base-expert + (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10
 - No mistake: 0,0,0,0,0,0,0,0,0,0
 - 1 mistake: 1,2,3,...,10
 - 2 mistakes: (1,2),(1,3),(1,4),..., (9,10)
- If the number of mistakes is at most 2, then there is a meta-expert that makes no mistakes.
- We can use the halving algorithm on the set of meta experts.

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert= base-expert+ (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10
 - No mistake: 0,0,0,0,0,0,0,0,0,0
 - 1 mistake: 1,2,3,...,10
 - 2 mistakes: (1,2),(1,3),(1,4),..., (9,10)
- If the number of mistakes is at most 2, then there is a meta-expert that makes no mistakes.
- We can use the halving algorithm on the set of meta experts.

Number of mistakes is smallest m such that $\frac{1}{2^{m+1}} \binom{m+1}{\leq k-i} = \frac{1}{2^{m+1}} \sum_{j=0}^{k-i} \binom{m+1}{j} < 1$

Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert= base-expert+ (sequence of locations of mistakes).
- Suppose a sequence of length $m=10$ and one of the base experts makes at most $k=2$ mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10
 - No mistake: 0,0,0,0,0,0,0,0,0,0
 - 1 mistake: 1,2,3,...,10
 - 2 mistakes: (1,2),(1,3),(1,4),..., (9,10)
- If the number of mistakes is at most 2, then there is a meta-expert that makes no mistakes.
- We can use the halving algorithm on the set of meta experts.

Number of mistakes is smallest m such that $\frac{1}{2^{m+1}} \binom{m+1}{\leq k-i} = \frac{1}{2^{m+1}} \sum_{j=0}^{k-i} \binom{m+1}{j} < 1$

- Naive implementation requires explicitly maintaining each meta-expert.

Reduction to chip game

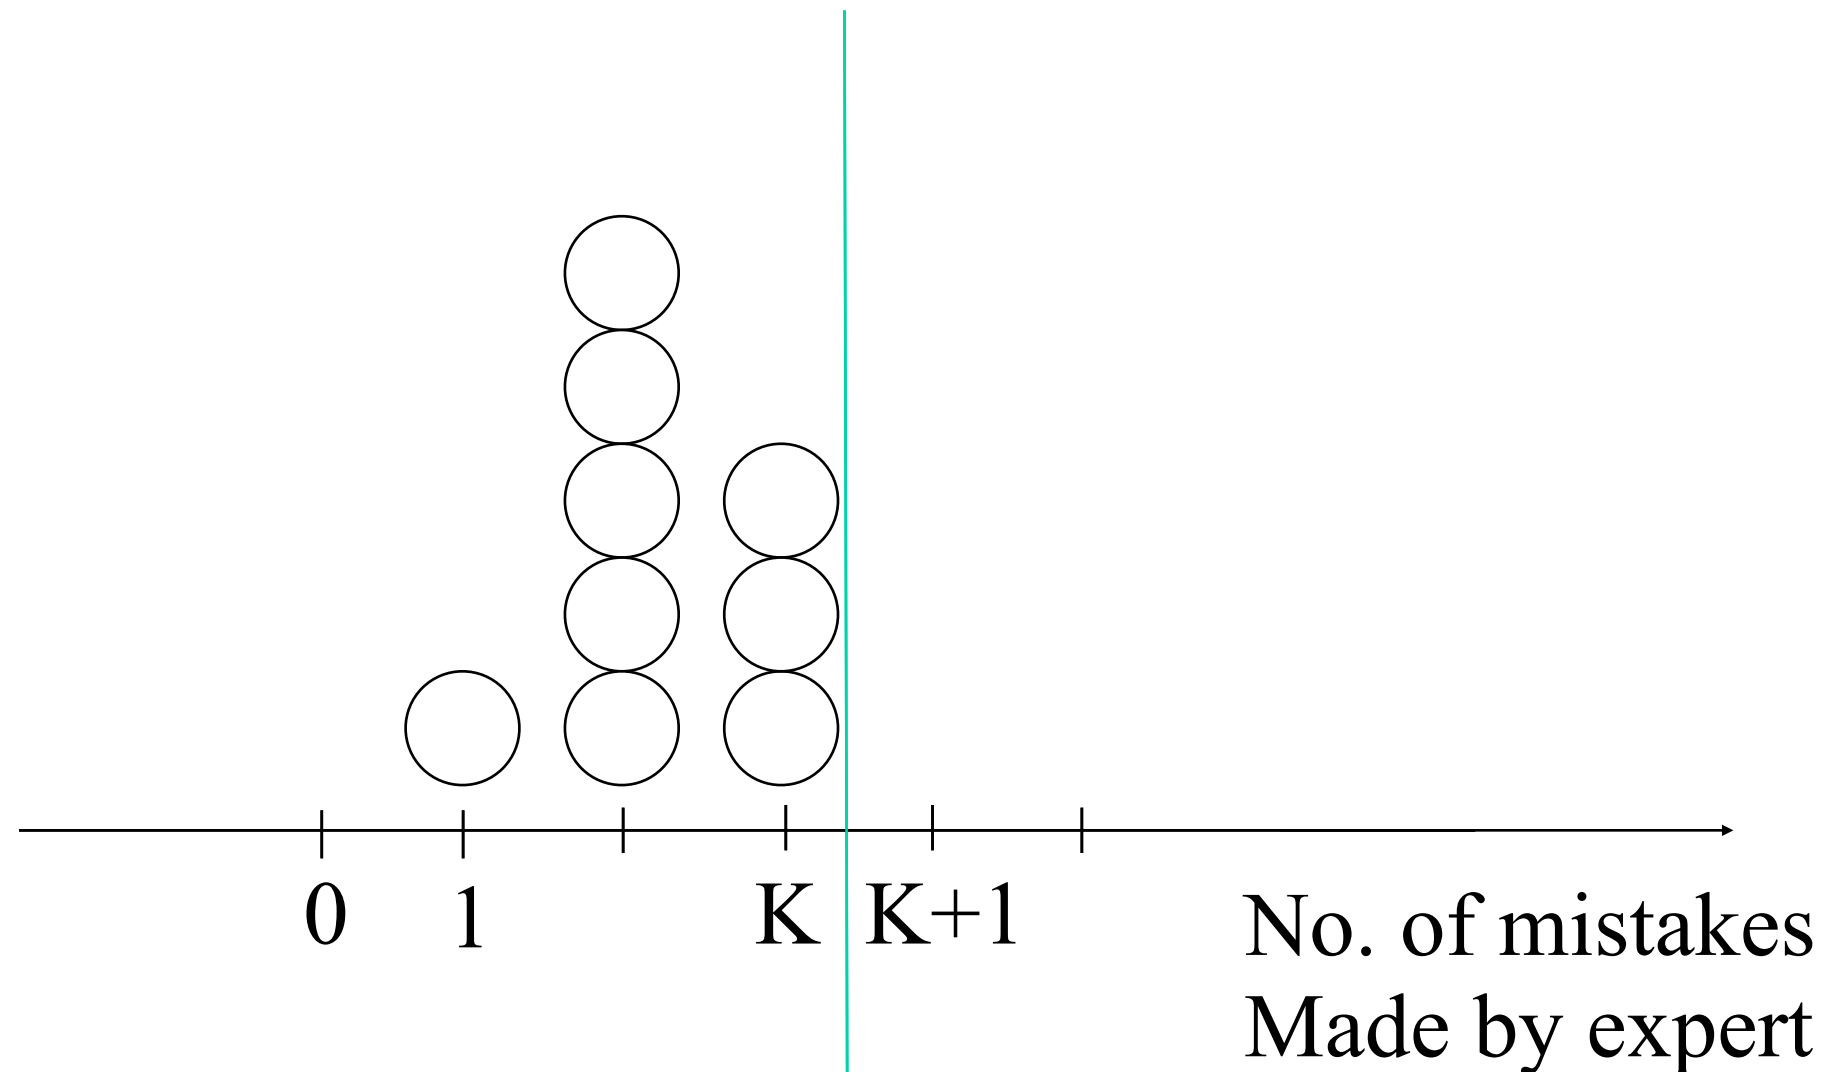
Reduction to chip game

- **Chip** = (base)-expert, **bin** = number of mistakes made so far


Reduction to chip game

- **Chip** = (base)-expert, **bin** = number of mistakes made so far
- Game step = algorithm's prediction is incorrect.

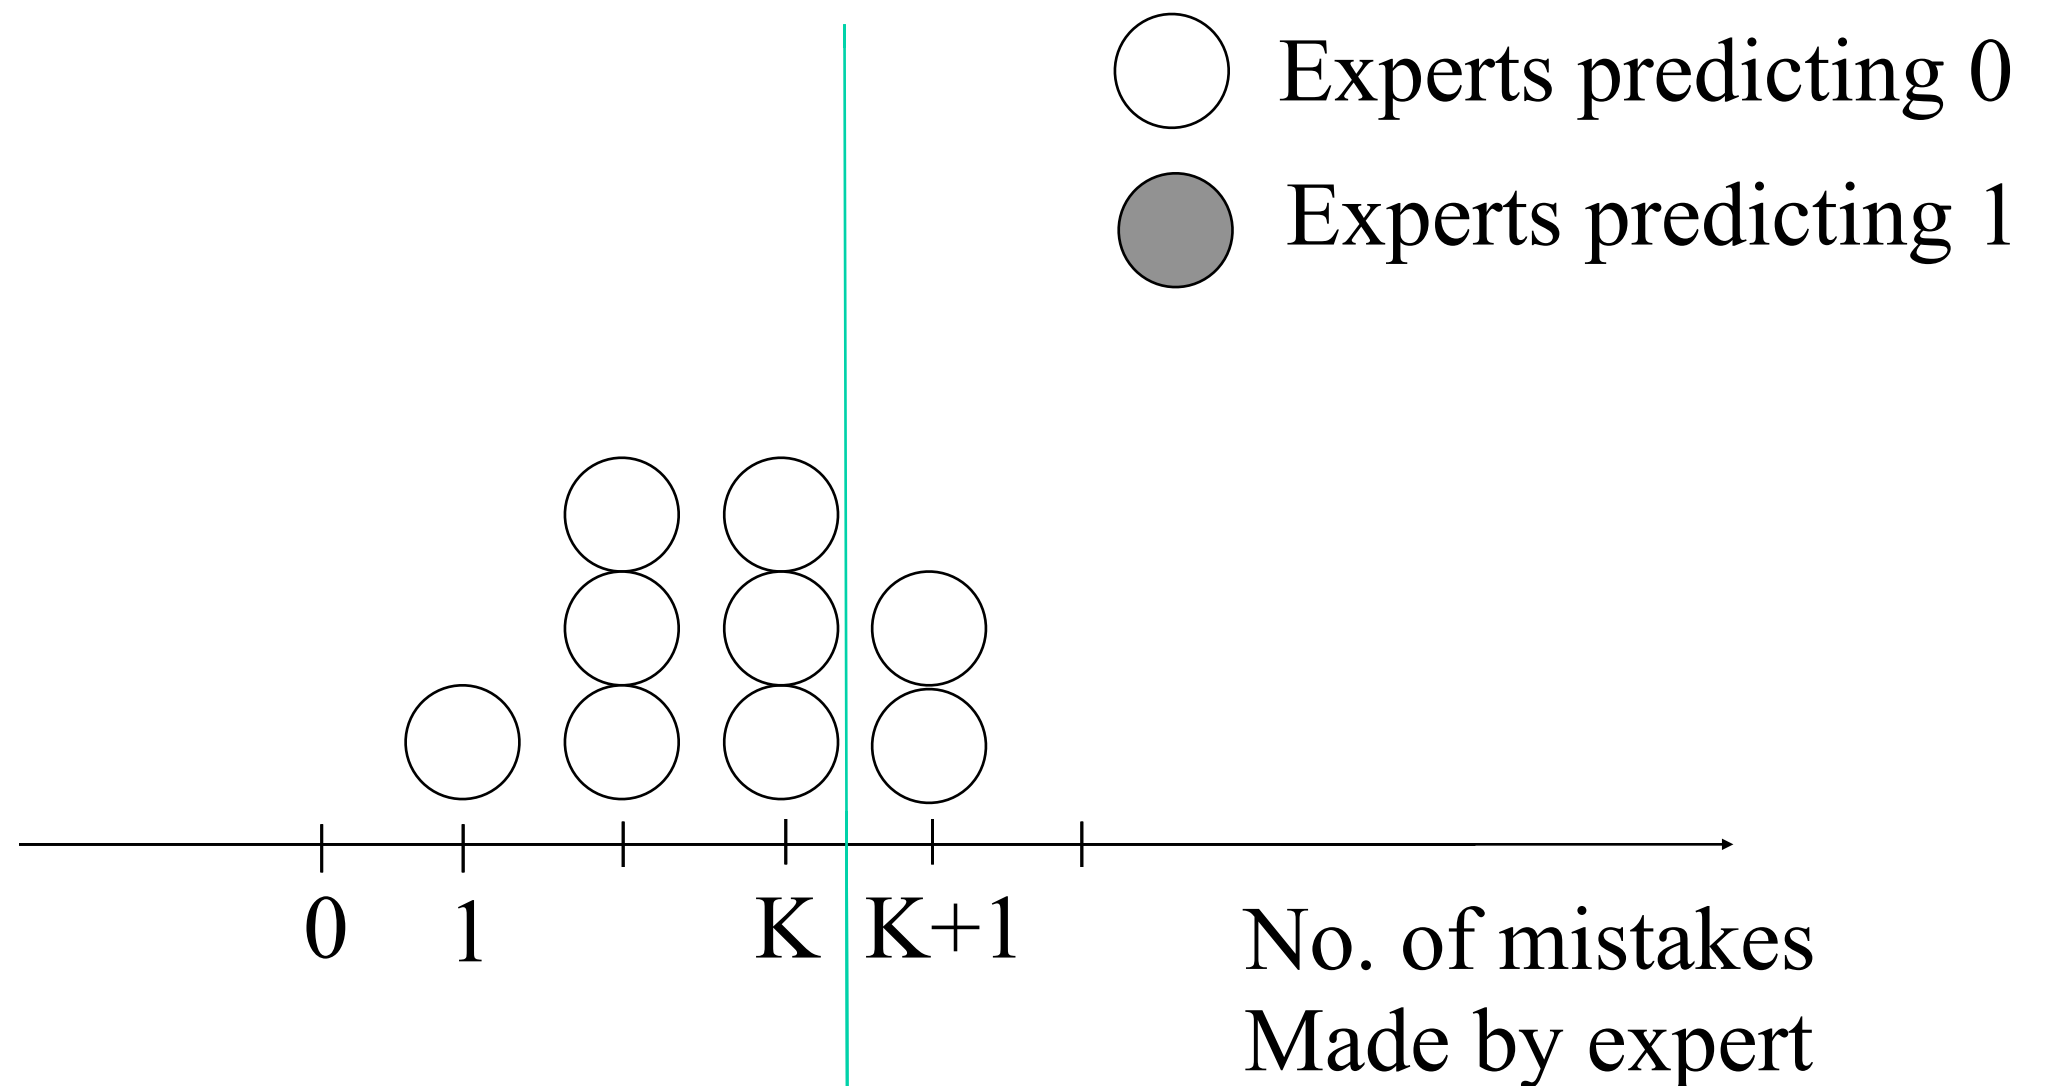
Chip game for expert advice




Chip game for expert advice



Chip game for expert advice



Chip game for expert advice



20 questions with k lies
[Ulam's game with lies]

20 questions with k lies

[Ulam's game with lies]

- Player 1 chooses secret x from $1, \dots, N$

20 questions with k lies

[Ulam's game with lies]

- Player 1 chooses secret x from $1, \dots, N$
- Player 2 asks “is x in the set A ?”

20 questions with k lies


[Ulam's game with lies]

- Player 1 chooses secret x from $1, \dots, N$
- Player 2 asks “is x in the set A ?”
- Player 1 answers yes/no , can lie at most k times.

20 questions with k lies

[Ulam's game with lies]


- Player 1 chooses secret x from $1, \dots, N$
- Player 2 asks “is x in the set A ?”
- Player 1 answers yes/no , can lie at most k times.



20 questions with k lies

[Ulam's game with lies]

- Player 1 chooses secret x from $1, \dots, N$
- Player 2 asks “is x in the set A ?”
- Player 1 answers yes/no , can lie at most k times.
- Game ends when player 2 can determine x .



20 questions with k lies

[Ulam's game with lies]

- Player 1 chooses secret x from $1, \dots, N$
- Player 2 asks “is x in the set A ?”
- Player 1 answers yes/no , can lie at most k times.
- Game ends when player 2 can determine x .
- A smart player 1 aims to have more than one $\text{consistent } x$ for as long as possible.

20 questions with k lies

[Ulam's game with lies]

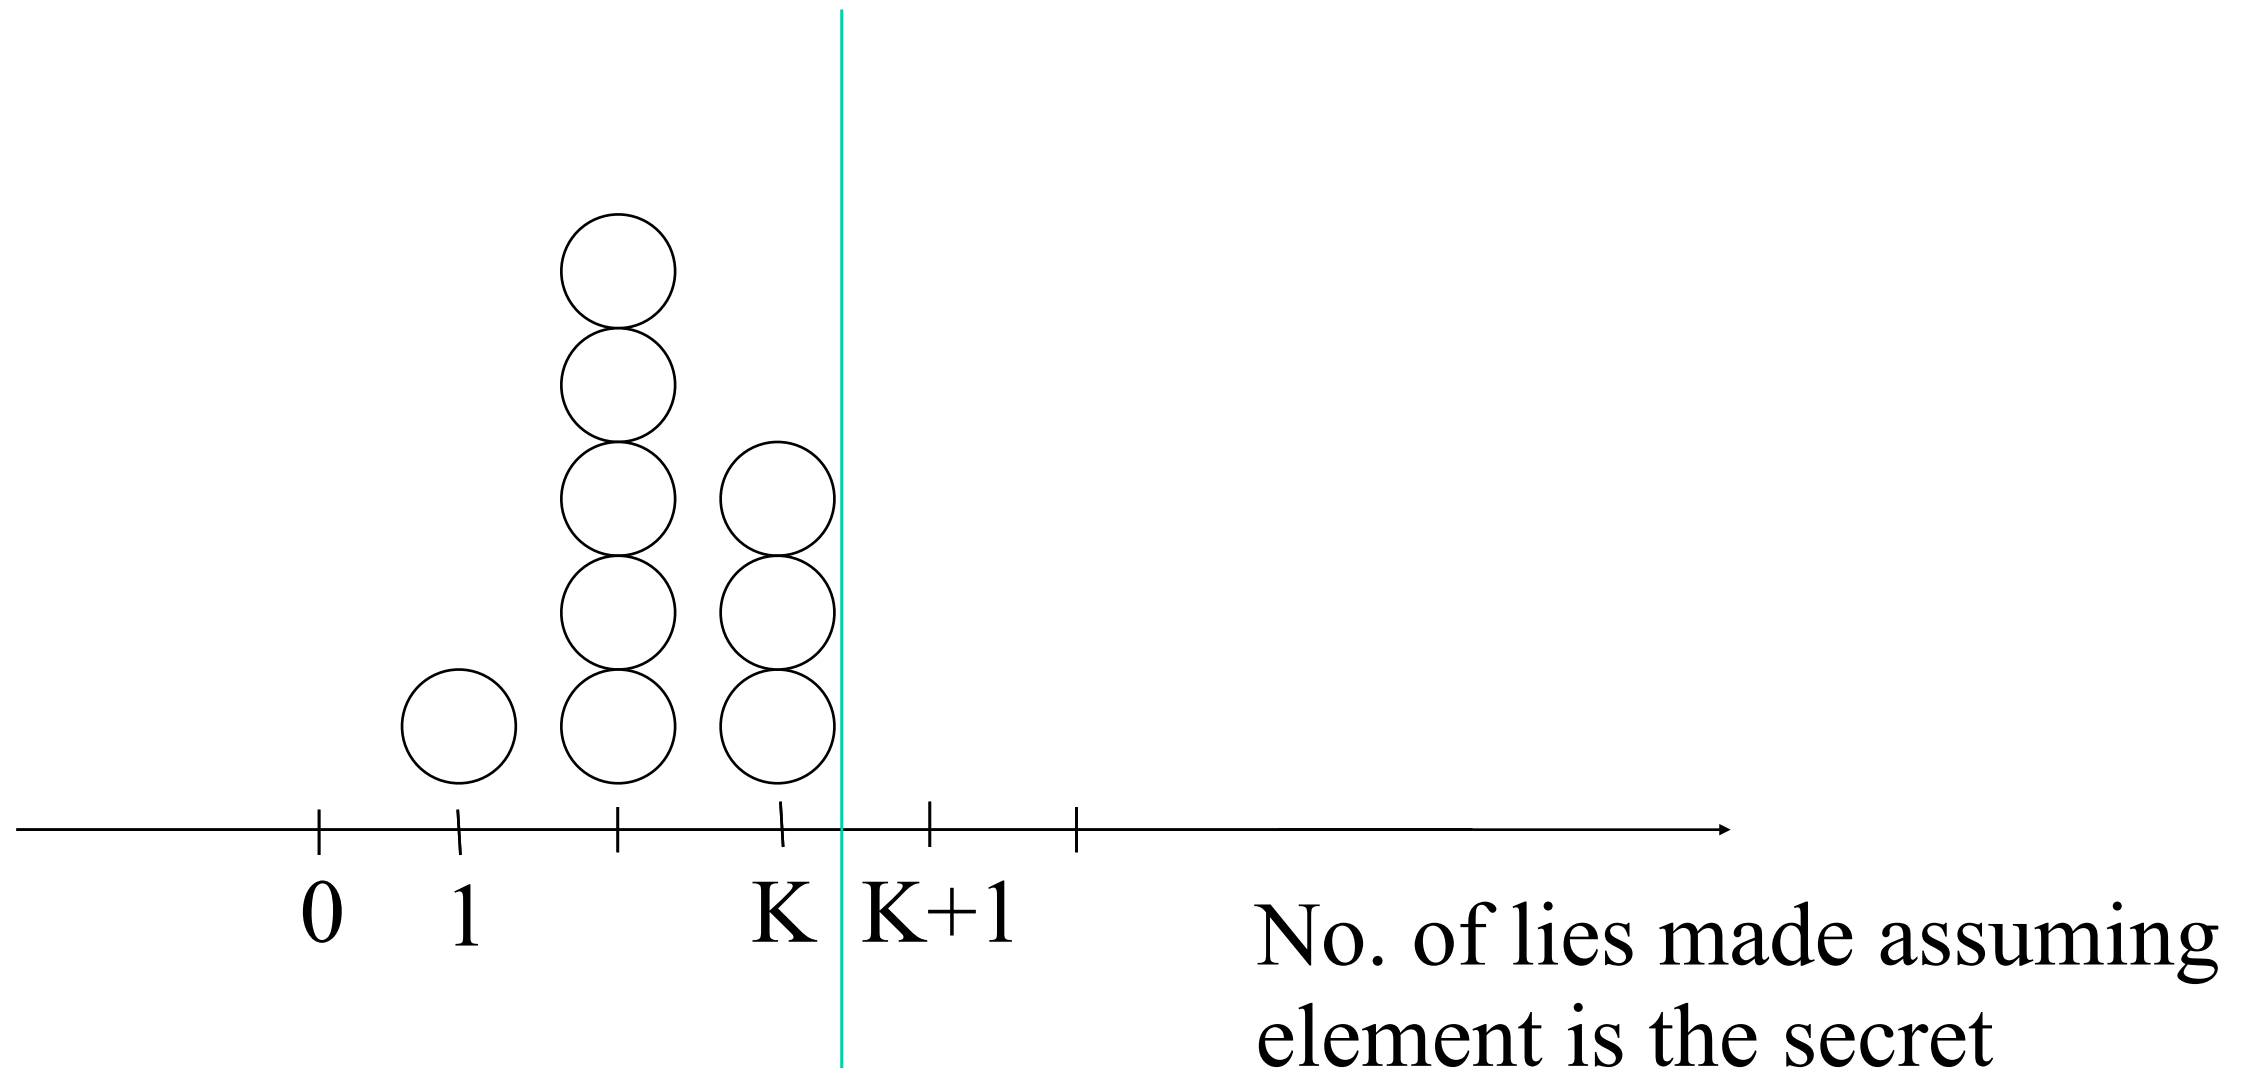
- Player 1 chooses secret x from $1, \dots, N$
- Player 2 asks “is x in the set A ?”
- Player 1 answers yes/no , can lie at most k times.
- Game ends when player 2 can determine x .
- A smart player 1 aims to have more than one $\text{consistent } x$ for as long as possible.
- Chip = element i

20 questions with k lies


[Ulam's game with lies]

- Player 1 chooses secret x from $1, \dots, N$
- Player 2 asks “is x in the set A ?”
- Player 1 answers yes/no , can lie at most k times.
- Game ends when player 2 can determine x .
- A smart player 1 aims to have more than one $\text{consistent } x$ for as long as possible.
- Chip = element i
- Bin = number of lies made regarding element


Chip game for 20 questions



Chip game for 20 questions




Chip game for 20 questions



Chip game for 20 questions

○ element in A
● element not in A



Game ends when there
is only one chip on this side

Simple case – all chips in bin 1

Simple case – all chips in bin 1

- 21 questions without lies

Simple case – all chips in bin 1

- 21 questions without lies
- Combining experts where one is perfect

Simple case – all chips in bin 1

- 21 questions without lies
- Combining experts where one is perfect
- Splitting a cookie

Simple case – all chips in bin 1

- 21 questions without lies
- Combining experts where one is perfect
- Splitting a cookie
- Optimal strategies:
 - Player 1: split chips into two equal parts
 - Player 2: choose larger part

Simple case – all chips in bin 1

- 21 questions without lies
- Combining experts where one is perfect
- Splitting a cookie
- Optimal strategies:
 - Player 1: split chips into two equal parts
 - Player 2: choose larger part
- Note problem when number of chips is odd

Binomial weights strategies

Binomial weights strategies

- What should **chooser** do if parts are not equal?

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.
- Fraction of bin **i** that will remain in bins **1..k** when **m** iterations remain

$$\Phi(i, k, m) \doteq \frac{1}{2^m} \binom{m}{\leq k-i} = \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.
- Fraction of bin **i** that will remain in bins **1..k** when **m** iterations remain

$$\Phi(i, k, m) \doteq \frac{1}{2^m} \binom{m}{\leq k-i} = \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

- Configuration: $\vec{n} = \langle n_0, n_1, \dots, n_k \rangle$

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.
- Fraction of bin **i** that will remain in bins **1..k** when **m** iterations remain

$$\Phi(i, k, m) \doteq \frac{1}{2^m} \binom{m}{\leq k-i} = \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

- Configuration: $\vec{n} = \langle n_0, n_1, \dots, n_k \rangle$
- Potential: The number of chips that will remain in bins **0..k** when **m** iterations remain and both sides play optimally:

$$\Psi(\vec{n}, k, m) = \sum_{i=0}^k n_i \Phi(i, k, m)$$

Combining experts, the binary prediction case

Combining experts, the binary prediction case

- Goal is to predict a binary sequence, making as few mistakes as possible.

Combining experts, the binary prediction case

- Goal is to predict a binary sequence, making as few mistakes as possible.
- There are N experts.

Combining experts, the binary prediction case

- Goal is to predict a binary sequence, making as few mistakes as possible.
- There are N experts.
- All predictions are binary and deterministic.

Combining experts, the binary prediction case

- Goal is to predict a binary sequence, making as few mistakes as possible.
- There are N experts.
- All predictions are binary and deterministic.
- A-priori knowledge: there is an expert that never makes more than k mistakes.

Combining experts, the binary prediction case


- Goal is to predict a binary sequence, making as few mistakes as possible.
- There are N experts.
- All predictions are binary and deterministic.
- A-priori knowledge: there is an expert that never makes more than k mistakes.
- $k=0$ corresponds to the halving algorithm.

Combining experts as a drifting game


[Cesa-Bianchi, Freund, Helmbold, Warmuth 96]

Binary instantaneous loss $l_i^t, l_A^t \in \{0,1\}$


Bin s contains all experts for which $L_i^t = s$




The game lattice



Initial configuration $t=1$



Experts predictions $t=1$

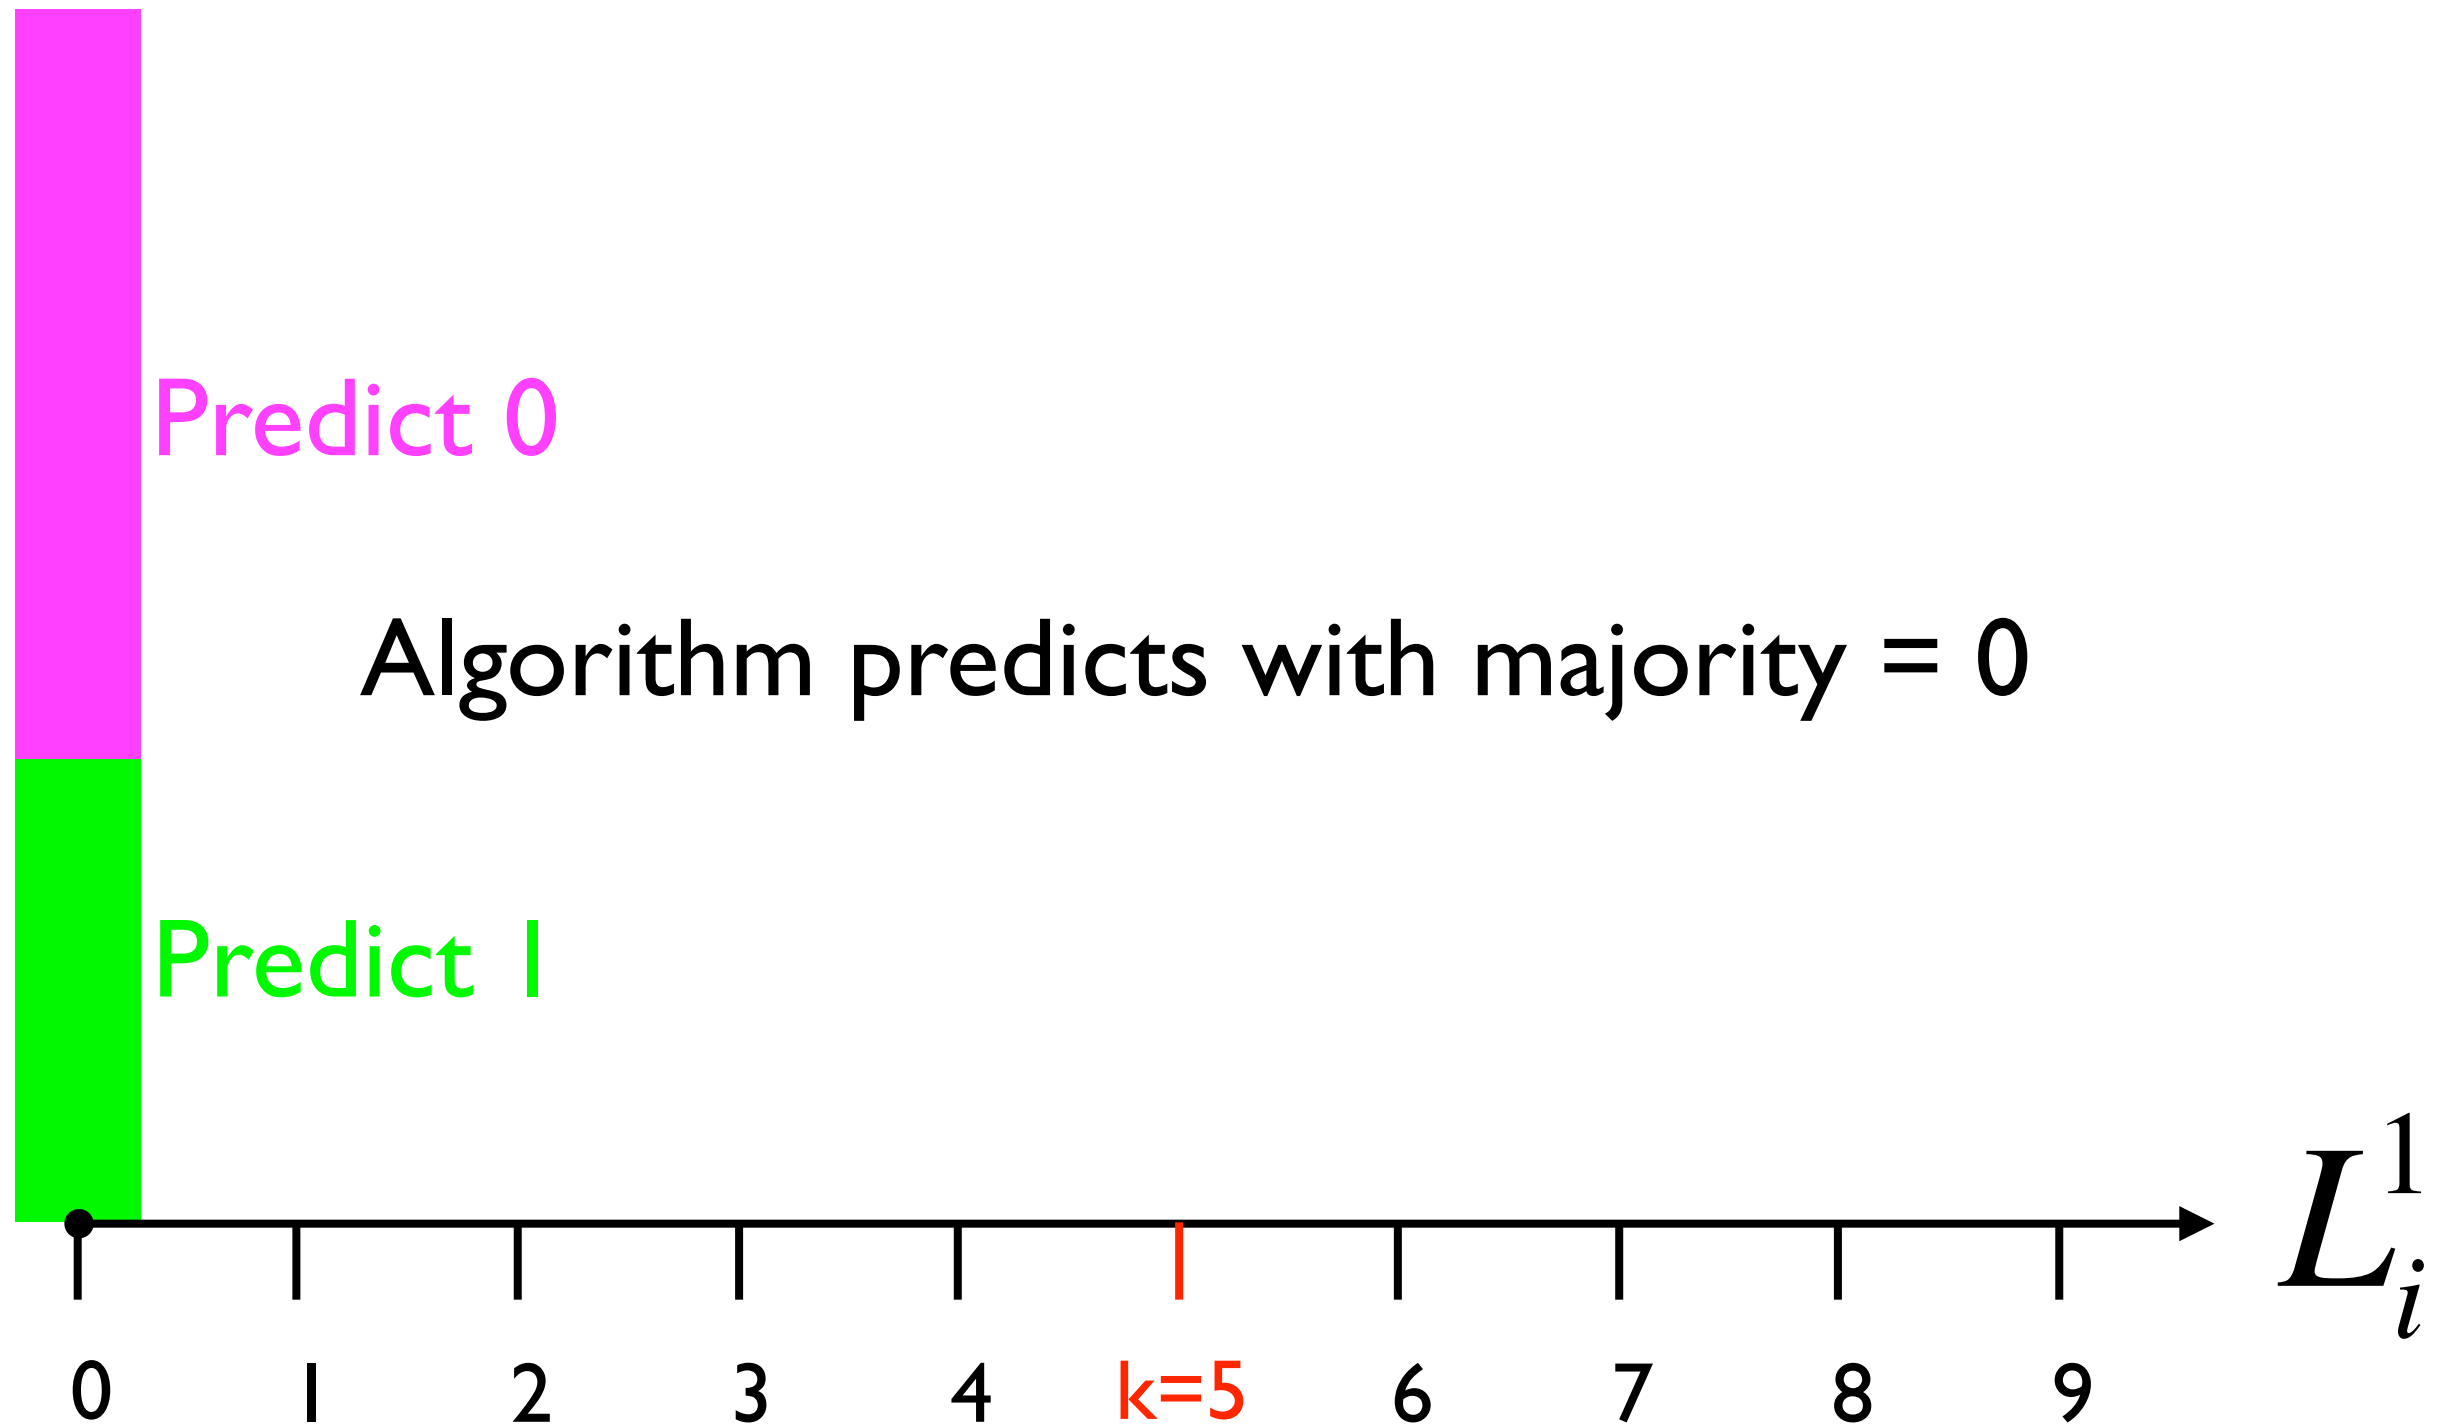


Experts predictions $t=1$

Predict 0

Algorithm predicts with majority = 0

Predict 1



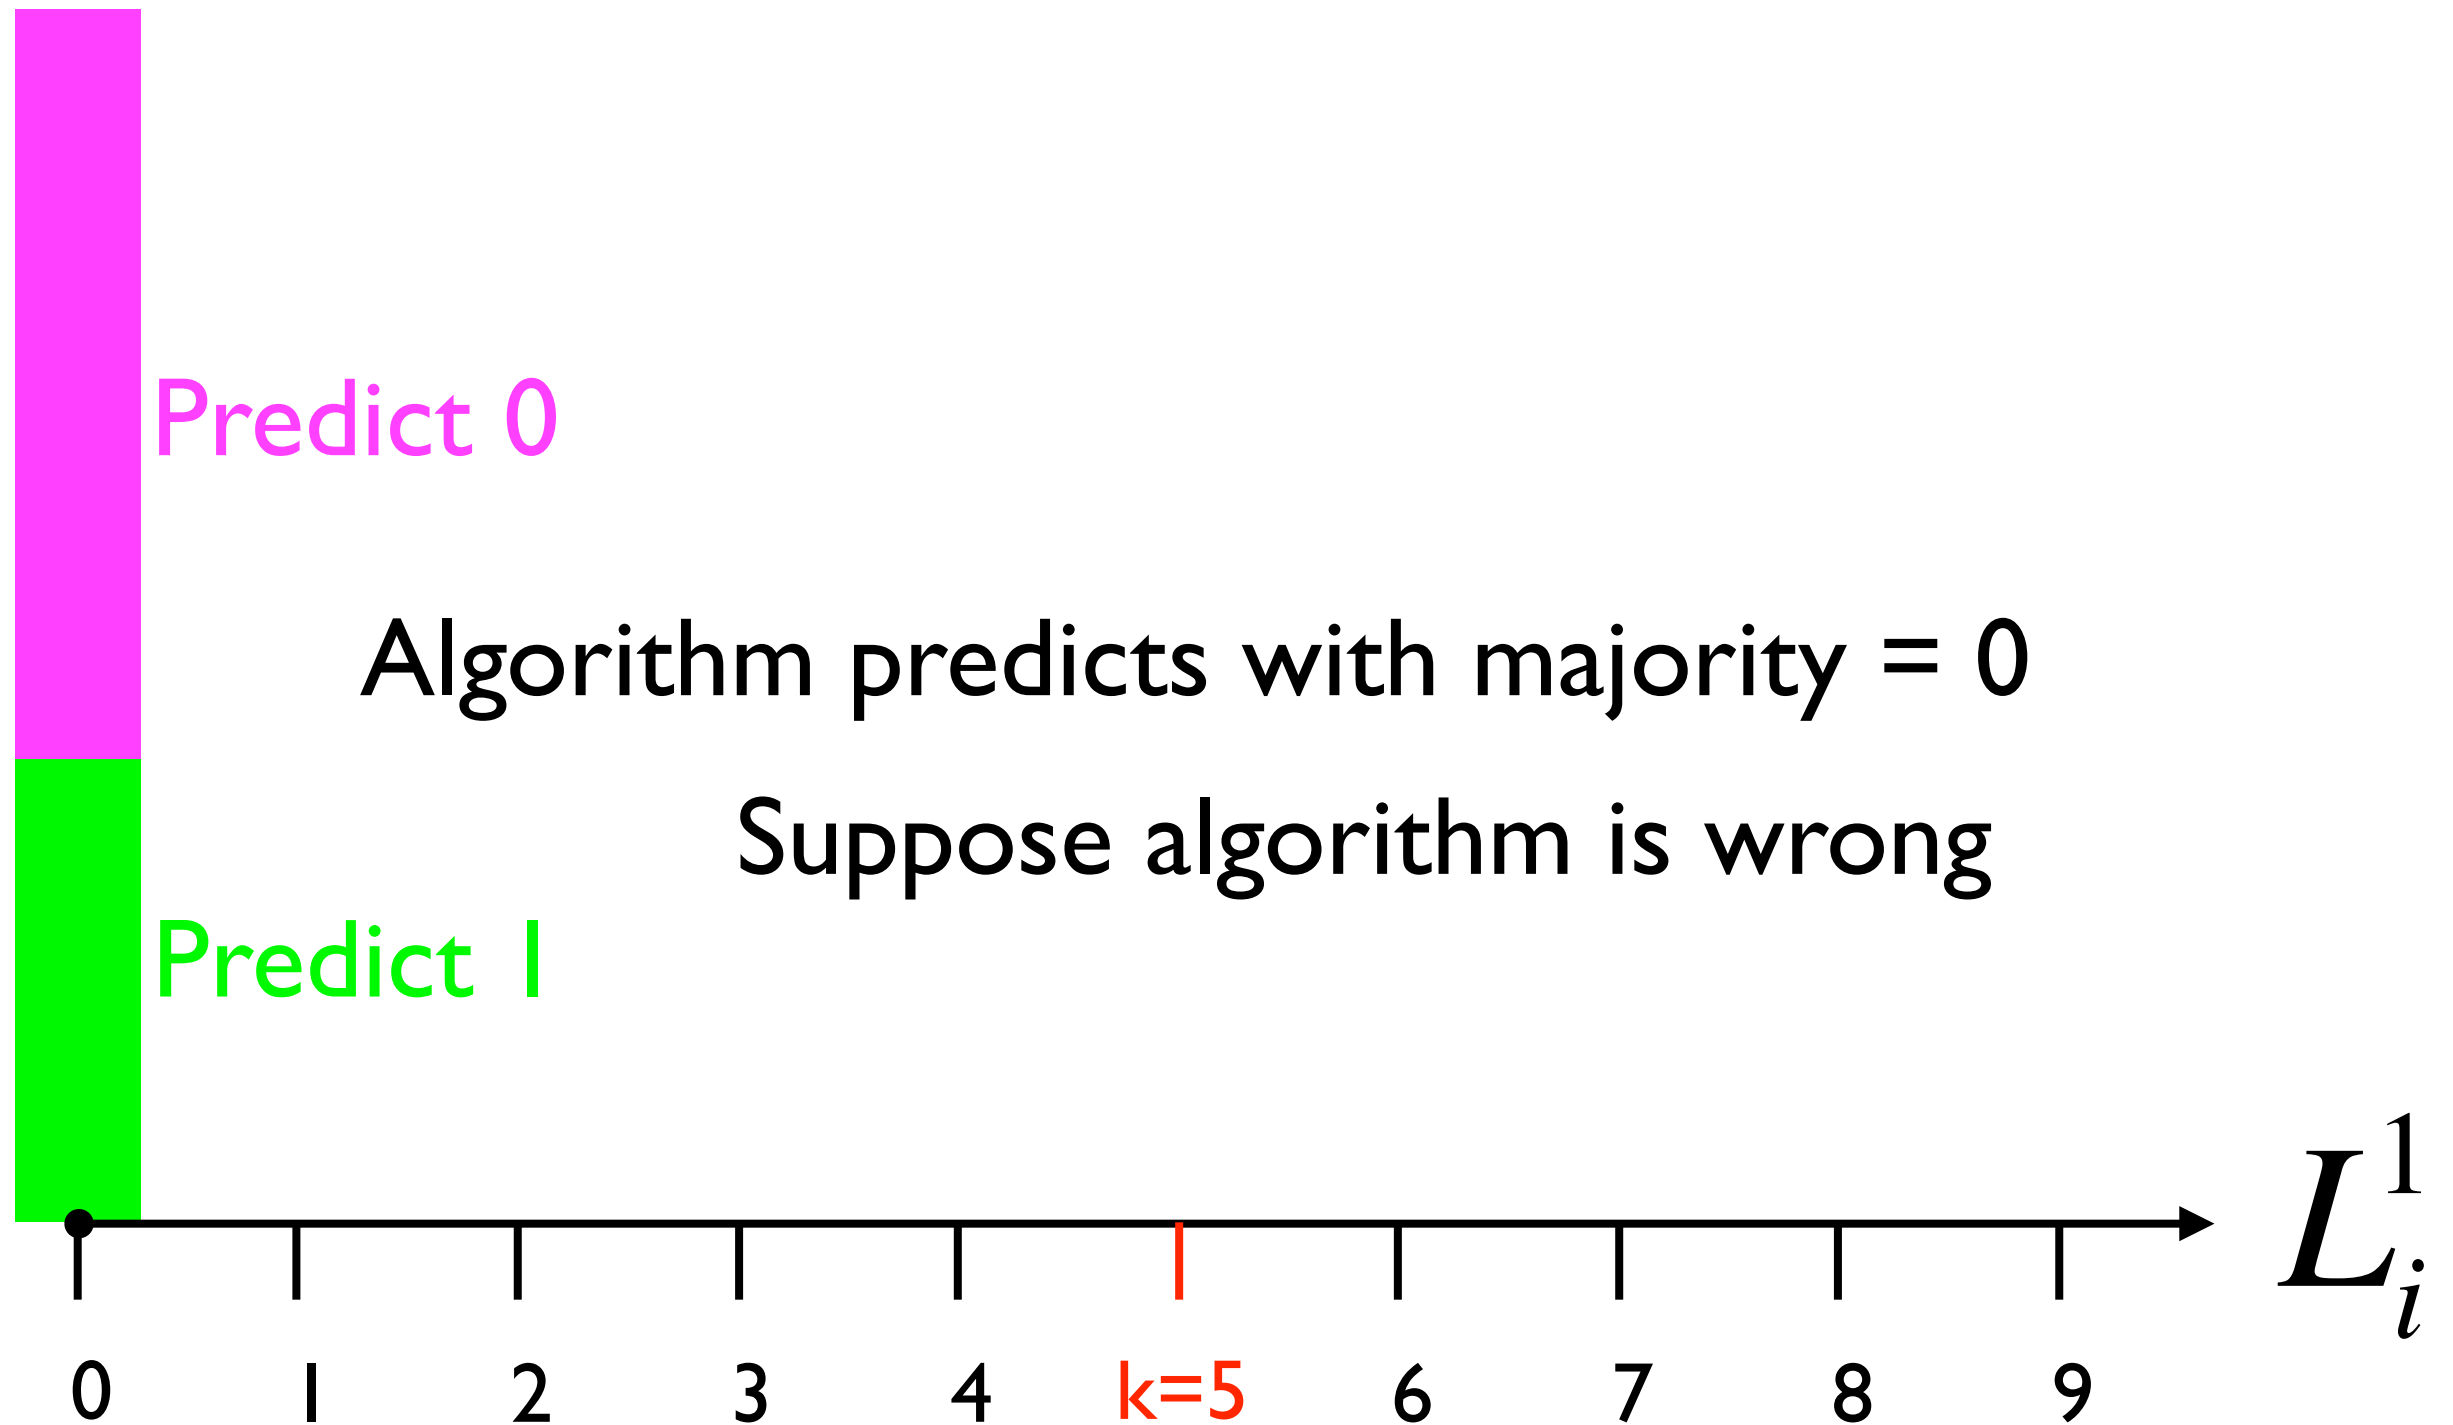
Experts predictions $t=1$


Predict 0

Algorithm predicts with majority = 0


Suppose algorithm is wrong

Predict 1






configuration at $t=2$




Experts predictions $t=2$

How should the algorithm
predict? 0 or 1?



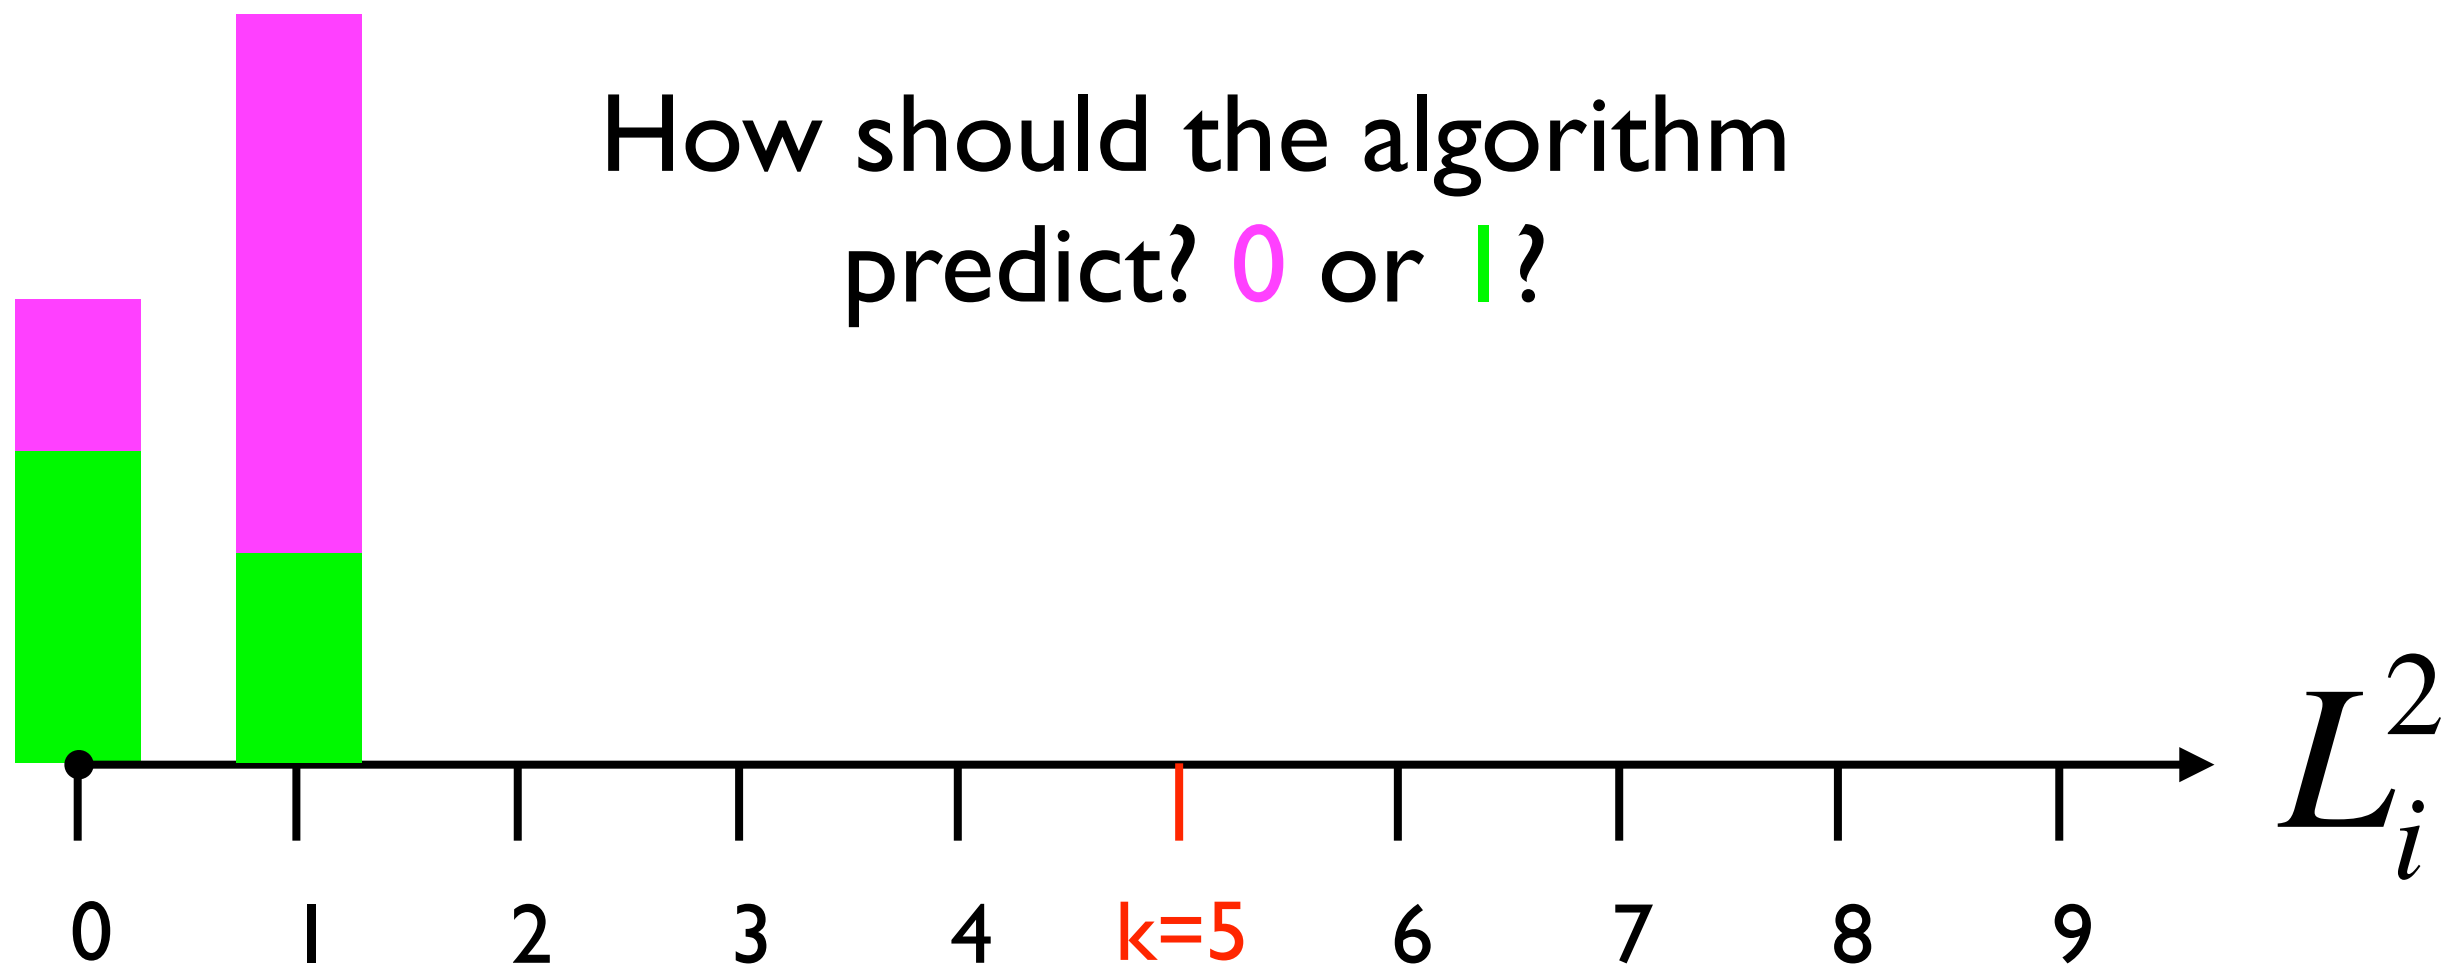
configuration $t=3$

Prediction is 0 and outcome is 1




Experts predictions $t=2$

How should the algorithm
predict? 0 or 1?




configuration $t=3$

Prediction is 1 and outcome is 0




If an error will lead to this configuration
then an error is not possible
 \Rightarrow this is a safe prediction



If an error will lead to this configuration
then an error is not possible
 \Rightarrow this is a safe prediction

Algorithm's goal is to get to an illegal configuration
with the smallest number of mistakes.



Helping the adversary.

Helping the adversary.

- Assume that the set of experts is continuous, arbitrarily divisible.

Helping the adversary.


- Assume that the set of experts is continuous, arbitrarily divisible.
- a-priori knowledge: $1/N$ fraction of the expert “mass” have cumulative loss at most k

Helping the adversary.

- Assume that the set of experts is continuous, arbitrarily divisible.
- a-priori knowledge: $1/N$ fraction of the expert “mass” have cumulative loss at most k
- Find algorithm with the tightest uniform upper bound on the cumulative loss.

If an error will lead to this configuration
then an error is not possible
 \Rightarrow this is a safe prediction

Algorithm's goal is to get to an illegal configuration
with the smallest number of mistakes.



An optimal adversarial strategy

An optimal adversarial strategy

- Split each bin to two equal parts. Algorithm's prediction is always incorrect.

An optimal adversarial strategy

- Split each bin to two equal parts. Algorithm's prediction is always incorrect.
- Equivalently: predictions of each expert are IID 0,1 with probabilities $1/2, 1/2$

An optimal adversarial strategy

- Split each bin to two equal parts. Algorithm's prediction is always incorrect.
- Equivalently: predictions of each expert are IID 0,1 with probabilities $1/2, 1/2$
- Same adversarial strategy was used to prove general lower bound on BLG

Optimal prediction strategy

Optimal prediction strategy

- Assume that adversary will play optimally from the next iteration until the end of the game.

Optimal prediction strategy

- Assume that adversary will play optimally from the next iteration until the end of the game.
- Choose as the next configuration the one that would end the game faster.

Optimal prediction strategy

- Assume that adversary will play optimally from the next iteration until the end of the game.
- Choose as the next configuration the one that would end the game faster.
- Relevant only when adversary plays sub-optimally, when adversary plays optimally the two next configurations are identical.

Binomial weights strategies

Binomial weights strategies

- What should **chooser** do if parts are not equal?

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.
- Fraction of bin **i** that will remain in bins **1..k** when **m** iterations remain

$$\phi(i, k, m) \doteq \frac{1}{2^m} \binom{m}{\leq k-i} = \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.
- Fraction of bin **i** that will remain in bins **1..k** when **m** iterations remain

$$\phi(i, k, m) \doteq \frac{1}{2^m} \binom{m}{\leq k-i} = \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

- Configuration: $\vec{n} = \langle n_0, n_1, \dots, n_k \rangle$

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.
- Fraction of bin **i** that will remain in bins **1..k** when **m** iterations remain

$$\phi(i, k, m) \doteq \frac{1}{2^m} \binom{m}{\leq k-i} = \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

- Configuration: $\vec{n} = \langle n_0, n_1, \dots, n_k \rangle$
- Potential: The number of chips that will remain in bins **0..k** when **m** iterations remain and both sides play optimally:

$$\Psi(\vec{n}, k, m) = \sum_{i=0}^k n_i \phi(i, k, m)$$

Response to suboptimal play

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
 - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
 - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).
- When split is not even, the chooser might be able to improve bound:

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
 - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).
- When split is not even, the chooser might be able to improve bound:
 - Online learning: $m := \min \{q \in N : \Psi(\vec{n}, k, q) < 1\} - 1$

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
 - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).
- When split is not even, the chooser might be able to improve bound:
 - Online learning: $m := \min \{q \in N : \Psi(\vec{n}, k, q) < 1\} - 1$
 - Ulam's game with k lies: $m := \max \{q \in N : \Psi(\vec{n}, k, q) \geq 1\} - 1$

Optimality of strategy

Optimality of strategy

- When chips are indivisible there can be situations where optimal play by the splitter is impossible.

Optimality of strategy

- When chips are indivisible there can be situations where optimal play by the splitter is impossible.
- [Spencer95] If number of chips is sufficiently large, equal-potential splits are always possible.

Optimality of strategy

- When chips are indivisible there can be situations where optimal play by the splitter is impossible.
- [Spencer95] If number of chips is sufficiently large, equal-potential splits are always possible.
- A sufficient number of chips is $\Omega(2^{2^k})$


Number of chips to infinity

Number of chips to infinity

Replace individual chips by chip **mass**


Number of chips to infinity

Replace individual chips by chip **mass**



Number of chips to infinity

Replace individual chips by chip **mass**



Optimal splitter strategy:

Split **each** bin into **two equal** parts

Equivalence to a random walk

Equivalence to a random walk

- Both sides playing optimally is equivalent to each chip performing an independent random walk.

Equivalence to a random walk

- Both sides playing optimally is equivalent to each chip performing an independent random walk.
- **Potential** = **probability** of a chip in bin **i** ending in bins **0..K** after **m** iteration

Equivalence to a random walk

- Both sides playing optimally is equivalent to each chip performing an independent random walk.
- **Potential** = **probability** of a chip in bin **i** ending in bins **0..K** after **m** iteration
- **Weight** = difference between the potentials of a chip in its two possible locations on the following iteration.

Equivalence to a random walk

- Both sides playing optimally is equivalent to each chip performing an independent random walk.
- **Potential** = **probability** of a chip in bin **i** ending in bins **0..K** after **m** iteration
- **Weight** = difference between the potentials of a chip in its two possible locations on the following iteration.
- Chooser's optimal strategy: choose set with smaller (larger) **weight**

Equivalence to a random walk

- Both sides playing optimally is equivalent to each chip performing an independent random walk.
- **Potential** = **probability** of a chip in bin **i** ending in bins **0..K** after **m** iteration
- **Weight** = difference between the potentials of a chip in its two possible locations on the following iteration.
- Chooser's optimal strategy: choose set with smaller (larger) **weight**
- **Worst case behavior of experts is to perform a random walk!**

Response to suboptimal play

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
 - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
 - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).
- When split is not even, the chooser might be able to improve bound:

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
 - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).
- When split is not even, the chooser might be able to improve bound:
 - online learning: $m := \min \{q \in N : \Psi(\vec{n}, k, q) < 1\} - 1$

Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
 - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
 - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).
- When split is not even, the chooser might be able to improve bound:
 - online learning: $m := \min \{q \in N : \Psi(\vec{n}, k, q) < 1\} - 1$
 - Ulam's game with k lies: $m := \max \{q \in N : \Psi(\vec{n}, k, q) \geq 1\} - 1$

Optimality of strategy

Optimality of strategy

- When chips are indivisible there can be situations where optimal play by the splitter is impossible.

Optimality of strategy

- When chips are indivisible there can be situations where optimal play by the splitter is impossible.
- [Spencer95] If number of chips is sufficiently large, equal-potential splits are always possible.

Optimality of strategy

- When chips are indivisible there can be situations where optimal play by the splitter is impossible.
- [Spencer95] If number of chips is sufficiently large, equal-potential splits are always possible.
- A sufficient number of chips is $\Omega(2^{2^k})$


Number of chips to infinity

Number of chips to infinity

Replace individual chips by chip **mass**


Number of chips to infinity

Replace individual chips by chip **mass**



Number of chips to infinity

Replace individual chips by chip **mass**




Optimal splitter strategy:

Split **each** bin into **two equal** parts

Example potential and weight

$m=5; k=10$



Boosting

Boosting

- A method for improving classifier accuracy

Boosting

- A method for improving classifier accuracy
- **Weak Learner**: a learning algorithm generating rules *slightly* better than random guessing.

Boosting

- A method for improving classifier accuracy
- **Weak Learner**: a learning algorithm generating rules **slightly** better than random guessing.
- **Basic idea**: **re-weight** training examples to force weak learner into different parts of the space.

Boosting

- A method for improving classifier accuracy
- **Weak Learner**: a learning algorithm generating rules **slightly** better than random guessing.
- **Basic idea**: **re-weight** training examples to force weak learner into different parts of the space.
- **Combine** weak rules by a **majority vote**.

Boost by Majority (BBM)

[Freund 95]

Boost by Majority (BBM)

[Freund 95]

- game between a **booster** and a weak **learner**.

Boost by Majority (BBM)

[Freund 95]

- game between a **booster** and a weak **learner**.
- Boosting generates a simple (unweighted) majority rule over weak learners.

Boost by Majority (BBM)

[Freund 95]

- game between a **booster** and a weak **learner**.
- Boosting generates a simple (unweighted) majority rule over weak learners.
- **T** Number of iterations is **set in advance**

Boost by Majority (BBM)

[Freund 95]

- game between a **booster** and a weak **learner**.
- Boosting generates a simple (unweighted) majority rule over weak learners.
- **T** Number of iterations is **set in advance**
- On iteration **$t=1..T$**

Boost by Majority (BBM)

[Freund 95]

- game between a **booster** and a weak **learner**.
- Boosting generates a simple (unweighted) majority rule over weak learners.
- **T** Number of iterations is set in advance
- On iteration **$t=1..T$**
 - **booster** assigns weights to the training examples.

Boost by Majority (BBM)

[Freund 95]

- game between a **booster** and a weak **learner**.
- Boosting generates a simple (unweighted) majority rule over weak learners.
- **T** Number of iterations is set in advance
- On iteration **$t=1..T$**
 - **booster** assigns weights to the training examples.
 - **learner** chooses a rule whose error wrt the chosen weights is smaller than **$1/2 - \gamma$**

Boost by Majority (BBM)

[Freund 95]


- game between a **booster** and a weak **learner**.
- Boosting generates a simple (unweighted) majority rule over weak learners.
- **T** Number of iterations is set in advance
- On iteration **$t=1..T$**
 - **booster** assigns weights to the training examples.
 - **learner** chooses a rule whose error wrt the chosen weights is smaller than **$1/2 - \gamma$**
 - Rule is added to majority rule

Boost by Majority (BBM)

[Freund 95]


- game between a **booster** and a weak **learner**.
- Boosting generates a simple (unweighted) majority rule over weak learners.
- **T** Number of iterations is set in advance
- On iteration **$t=1..T$**
 - **booster** assigns weights to the training examples.
 - **learner** chooses a rule whose error wrt the chosen weights is smaller than **$1/2 - \gamma$**
 - Rule is added to majority rule
- Goal of booster is to minimize number of errors of final majority rule.

BBM as a drifting game




BBM as a drifting game

- Chips = examples




BBM as a drifting game

- Chips = examples
- bin **s** contains the examples for which the difference between the number of correct and incorrect base rules is **s**




The continuous chip limit




The continuous chip limit

- Number of training examples increases to infinity.




The continuous chip limit


- Number of training examples increases to infinity.
- Alternatively - think of examples as a probability mass with probability measure μ defined on it.



The boosting game lattice




The boosting game lattice




$\gamma = 0.1$

Initial configuration



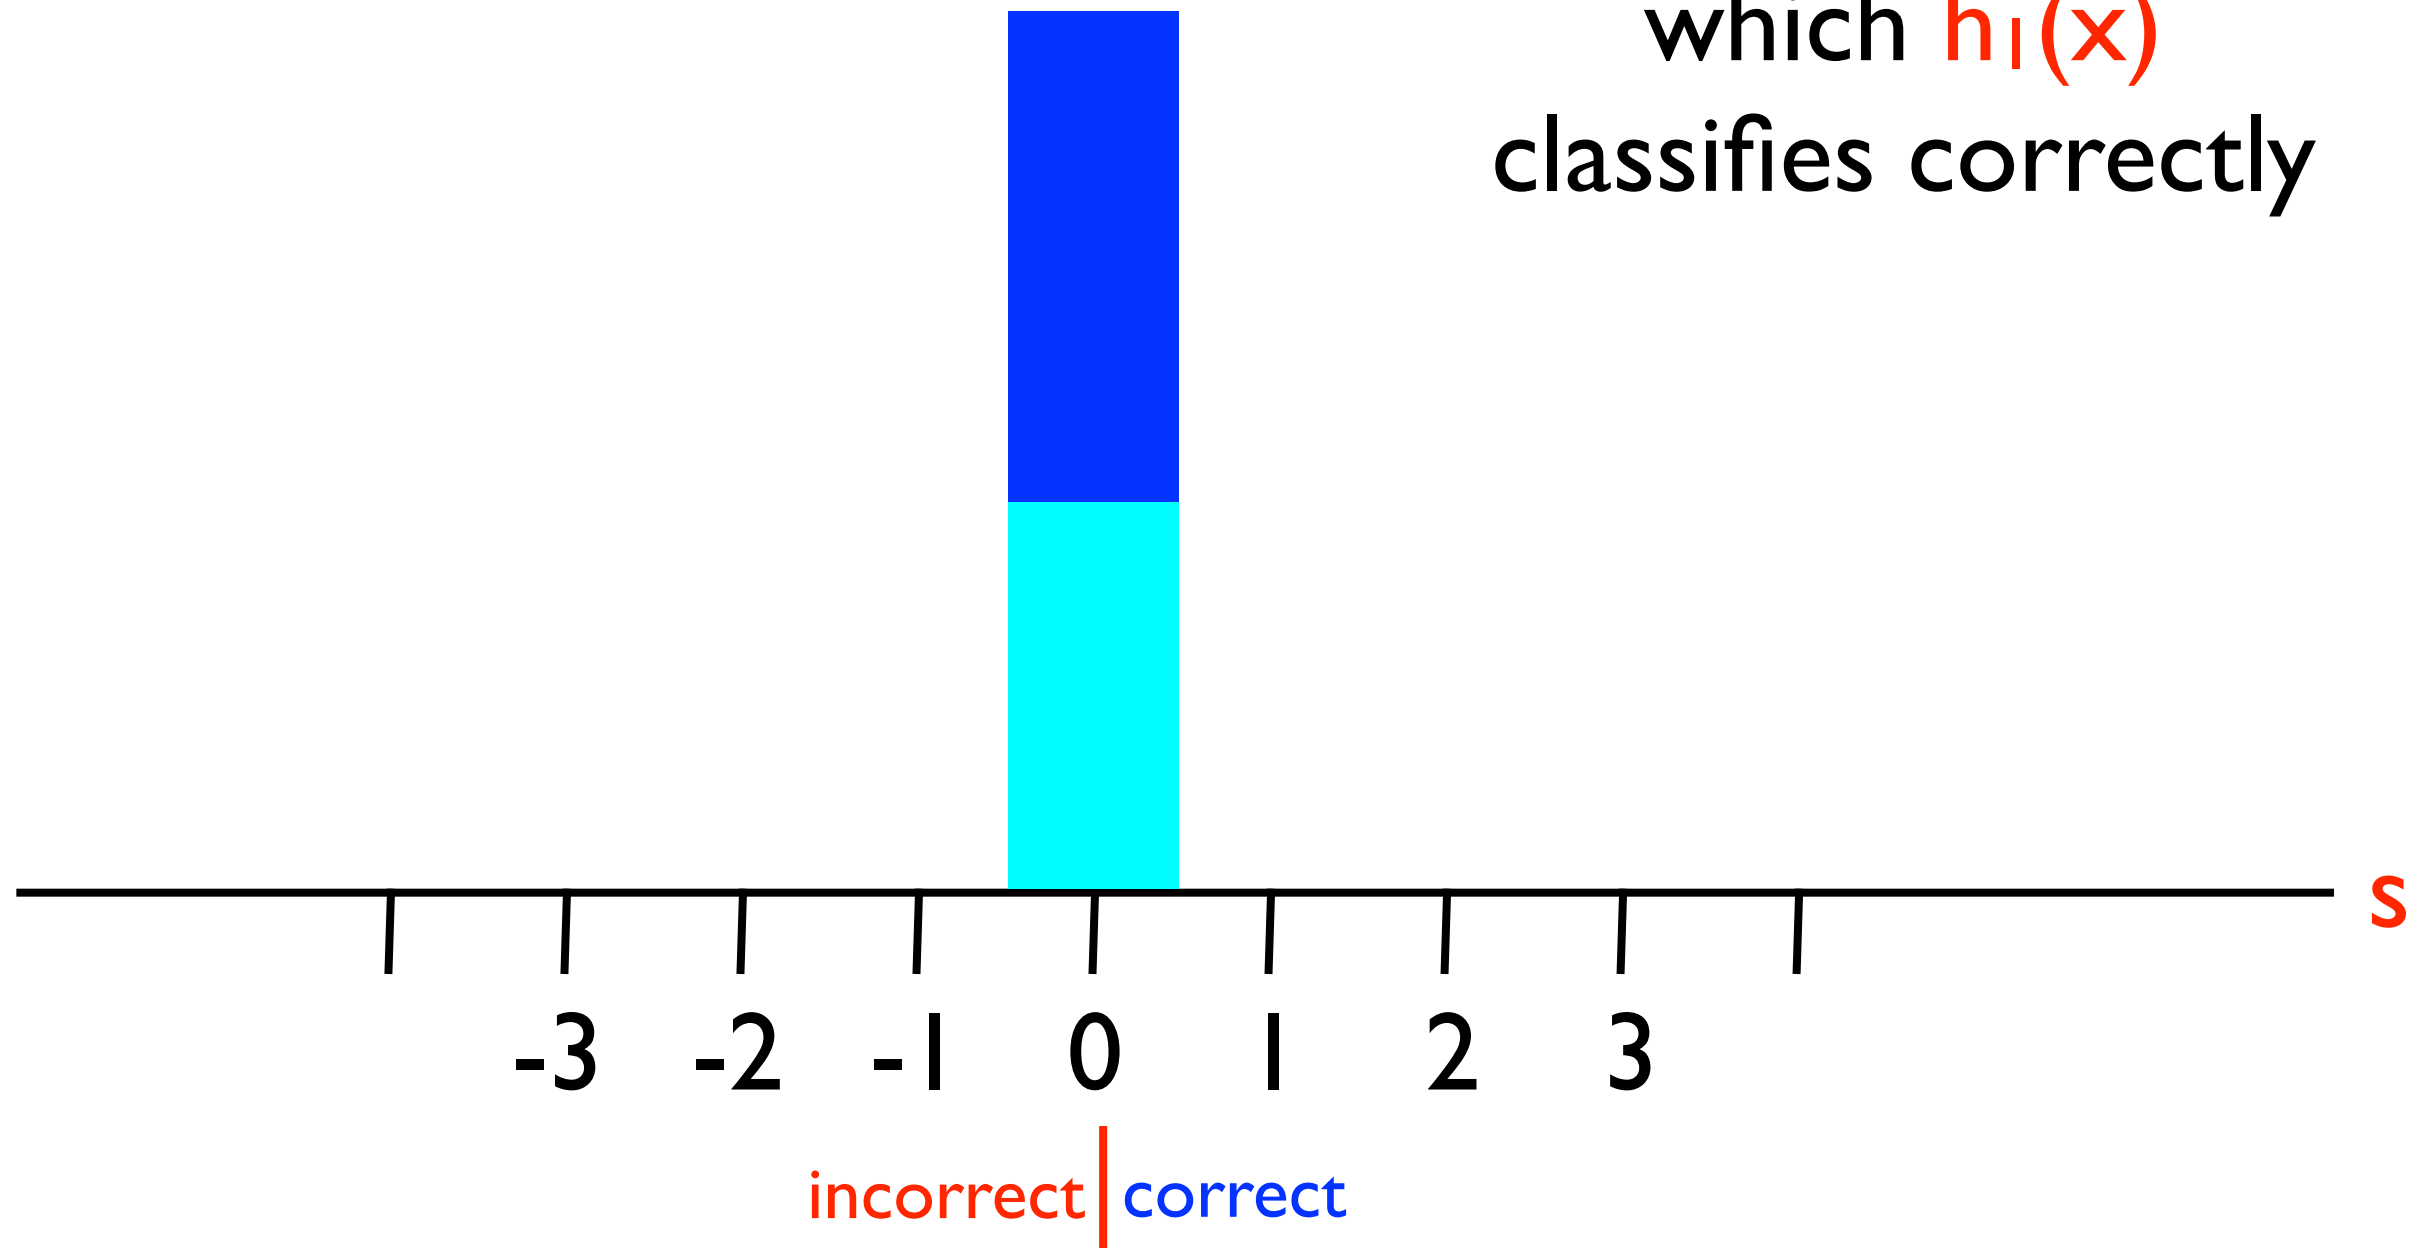
$\gamma = 0.1$

Booster
assigns weights
to
examples




$\gamma = 0.1$

Weak learner
chooses subset
with weight $1/2 + \gamma$
which $h_1(x)$
classifies correctly




$\gamma = 0.1$

Weak learner
chooses subset
with weight $1/2 + \gamma$
which $h_1(x)$
classifies correctly




$\gamma = 0.1$

Booster
assigns weights
to
examples




$\gamma = 0.1$

Weak learner
chooses subset
with weight $1/2 + \gamma$
which $h_2(x)$
classifies correctly




$\gamma = 0.1$

Weak learner
chooses subset
with weight $1/2 + \gamma$
which $h_2(x)$
classifies correctly




$\gamma = 0.1$

Booster
assigns weights
to
examples




$\gamma = 0.1$

Weak learner
chooses subset
with weight $1/2 + \gamma$
which $h_3(x)$
classifies correctly




$\gamma = 0.1$

Weak learner
chooses subset
with weight $1/2 + \gamma$
which $h_3(x)$
classifies correctly



$\gamma = 0.1$




$\gamma = 0.1$

Majority[$h_3(x) + h_3(x) + h_3(x)$]

incorrect

Majority[$h_3(x) + h_3(x) + h_3(x)$]

correct



Weak Learner's min/max strategy

Weak Learner's min/max strategy

- **AdOpt** - Choose $1/2 + \gamma$ from each bin to be correct.

Weak Learner's min/max strategy

- **AdOpt** - Choose $1/2 + \gamma$ from each bin to be correct.
- Equivalently: prediction of each base rule on each example is chosen independently at random

$$P(h_t(x)=y) = 1/2 + \gamma$$

Potential

Potential

Total potential: $\Psi(t, \text{configuration})$ - μ -prob of the examples on which the final majority vote is incorrect given the configuration after iteration t is configuration and on the remaining steps the learner plays AdOpt .

Potential

Total potential: $\Psi(t, \text{configuration})$ - μ -prob of the examples on which the final majority vote is incorrect given the configuration after iteration t is configuration and on the remaining steps the learner plays AdOpt .

$\Psi(0, \text{all at origin})$ = Initial potential.

Potential

Total potential: $\Psi(t, \text{configuration})$ - μ -prob of the examples on which the final majority vote is incorrect given the configuration after iteration t is configuration and on the remaining steps the learner plays AdOpt .

$\Psi(0, \text{all at origin})$ = Initial potential.

$\Psi(T, \text{configuration})$ = Training error of final majority rule.

Potential

Total potential: $\Psi(t, \text{configuration})$ - μ -prob of the examples on which the final majority vote is incorrect given the configuration after iteration t is configuration and on the remaining steps the learner plays AdOpt .

$\Psi(0, \text{all at origin})$ = Initial potential.

$\Psi(T, \text{configuration})$ = Training error of final majority rule.

Boosting algorithm chooses weights so that the total potential does not increase.

Potential

Total potential: $\Psi(t, \text{configuration})$ - μ -prob of the examples on which the final majority vote is incorrect given the configuration after iteration t is configuration and on the remaining steps the learner plays AdOpt .

$\Psi(0, \text{all at origin})$ = Initial potential.

$\Psi(T, \text{configuration})$ = Training error of final majority rule.

Boosting algorithm chooses weights so that the total potential does not increase.

Initial potential \geq final training error.

Bin Potential: $\phi(t,s)$ - fraction of examples
in bin s after iteration t on which the final majority
rule will be incorrect assuming AdOpt play

Bin Potential: $\phi(t,s)$ - fraction of examples
in bin s after iteration t on which the final majority
rule will be incorrect assuming AdOpt play

$f(t,s)$ The configuration

The μ -prob of examples in bin s after iteration t

Bin Potential: $\phi(t,s)$ - fraction of examples in bin s after iteration t on which the final majority rule will be incorrect assuming AdOpt play

$f(t,s)$ The configuration

The μ -prob of examples in bin s after iteration t

The potential of a configuration: $\Psi(f,t) = \sum_i f(t,s)\phi(t,s)$

Bin Potential: $\phi(t,s)$ - fraction of examples in bin s after iteration t on which the final majority rule will be incorrect assuming AdOpt play

$f(t,s)$ The configuration

The μ -prob of examples in bin s after iteration t

The potential of a configuration: $\Psi(f,t) = \sum_i f(t,s)\phi(t,s)$

$\phi(t,s)$ does not depend on the configuration

Bin Potential: $\phi(t,s)$ - fraction of examples in bin s after iteration t on which the final majority rule will be incorrect assuming AdOpt play

$f(t,s)$ The configuration

The μ -prob of examples in bin s after iteration t

The potential of a configuration: $\Psi(f,t) = \sum_i f(t,s) \phi(t,s)$

$\phi(t,s)$ does not depend on the configuration

$$\phi(t,s) = \text{Binom}\left(T-t, \frac{T-t-s}{2}, \frac{1}{2} + \gamma\right); \quad \text{Binom}(n,k,p) \doteq \sum_{j=0}^{\lfloor k \rfloor} \binom{n}{j} p^j (1-p)^{n-j}$$

Bin Potential: $\phi(t,s)$ - fraction of examples in bin s after iteration t on which the final majority rule will be incorrect assuming AdOpt play

$f(t,s)$ The configuration

The μ -prob of examples in bin s after iteration t

The potential of a configuration: $\Psi(f,t) = \sum_i f(t,s) \phi(t,s)$

$\phi(t,s)$ does not depend on the configuration

$$\phi(t,s) = \text{Binom}\left(T-t, \frac{T-t-s}{2}, \frac{1}{2} + \gamma\right); \quad \text{Binom}(n,k,p) \doteq \sum_{j=0}^{\lfloor k \rfloor} \binom{n}{j} p^j (1-p)^{n-j}$$

$$\phi(t-1,s) = \left(\frac{1}{2} - \gamma\right) \phi(t,s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t,s+1)$$

Bin Potential: $\phi(t,s)$ - fraction of examples in bin s after iteration t on which the final majority rule will be incorrect assuming AdOpt play

$f(t,s)$ The configuration

The μ -prob of examples in bin s after iteration t

The potential of a configuration: $\Psi(f,t) = \sum_i f(t,s) \phi(t,s)$

$\phi(t,s)$ does not depend on the configuration

$$\phi(t,s) = \text{Binom}\left(T-t, \frac{T-t-s}{2}, \frac{1}{2} + \gamma\right); \quad \text{Binom}(n,k,p) \doteq \sum_{j=0}^{\lfloor k \rfloor} \binom{n}{j} p^j (1-p)^{n-j}$$

$$\phi(t-1,s) = \left(\frac{1}{2} - \gamma\right) \phi(t,s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t,s+1)$$

$$\phi(T,s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases}$$

Bin Potential: $\phi(t,s)$ - fraction of examples in bin s after iteration t on which the final majority rule will be incorrect assuming AdOpt play

$f(t,s)$ The configuration

The μ -prob of examples in bin s after iteration t

The potential of a configuration: $\Psi(f,t) = \sum_i f(t,s) \phi(t,s)$

$\phi(t,s)$ does not depend on the configuration

$$\phi(t,s) = \text{Binom}\left(T-t, \frac{T-t-s}{2}, \frac{1}{2} + \gamma\right); \quad \text{Binom}(n,k,p) \doteq \sum_{j=0}^{\lfloor k \rfloor} \binom{n}{j} p^j (1-p)^{n-j}$$

$$\phi(t-1,s) = \left(\frac{1}{2} - \gamma\right) \phi(t,s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t,s+1)$$

$$\phi(T,s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \phi(0,0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)$$

Definitions

Definitions

edge: $d(t, s) \doteq \mu(h_t(x) = y | (x, y) \text{ in bin } s \text{ after iteration } t) - \left(\frac{1}{2} + \gamma\right)$

Definitions

edge: $d(t, s) \doteq \mu(h_t(x) = y | (x, y) \text{ in bin } s \text{ after iteration } t) - \left(\frac{1}{2} + \gamma\right)$

Example Weights: $w(t, s) \doteq \phi(t + 1, s - 1) - \phi(t + 1, s + 1)$

Definitions

edge: $d(t, s) \doteq \mu(h_t(x) = y | (x, y) \text{ in bin } s \text{ after iteration } t) - \left(\frac{1}{2} + \gamma\right)$

Example Weights: $w(t, s) \doteq \phi(t + 1, s - 1) - \phi(t + 1, s + 1)$

Theorem

If, for step t , $\sum_s d(t, s)w(t, s) \geq 0$ then $\Psi(t + 1) \leq \Psi(t)$

Definitions

edge: $d(t, s) \doteq \mu(h_t(x) = y | (x, y) \text{ in bin } s \text{ after iteration } t) - \left(\frac{1}{2} + \gamma\right)$

Example Weights: $w(t, s) \doteq \phi(t + 1, s - 1) - \phi(t + 1, s + 1)$

Theorem

If, for step t , $\sum_s d(t, s)w(t, s) \geq 0$ then $\Psi(t + 1) \leq \Psi(t)$

Corollary

If $\forall t$ [weighted error of $h_t(x)$] $\leq 1/2 - \gamma$

Then

Initial potential \geq final training error.

Proof of Theorem

Proof of Theorem

$$f(t+1, s) = f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right)$$

Proof of Theorem

$$f(t+1, s) = f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right) \right] \phi(t+1, s)$$

Proof of Theorem

$$f(t+1, s) = f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right) \right] \phi(t+1, s)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s) \left(\left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1) \right) - d(t, s) (\phi(t+1, s-1) - \phi(t+1, s+1)) \right]$$

Proof of Theorem

$$f(t+1, s) = f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right) \right] \phi(t+1, s)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s) \left(\left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1) \right) - d(t, s) (\phi(t+1, s-1) - \phi(t+1, s+1)) \right]$$

$$\phi(t, s) = \left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1)$$

Proof of Theorem

$$f(t+1, s) = f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right) \right] \phi(t+1, s)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s) \left(\left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1) \right) - d(t, s) (\phi(t+1, s-1) - \phi(t+1, s+1)) \right]$$

$$\phi(t, s) = \left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1)$$

$$w(t, s) \doteq \phi(t+1, s-1) - \phi(t+1, s+1)$$

Proof of Theorem

$$f(t+1, s) = f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right) \right] \phi(t+1, s)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s) \left(\left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1) \right) - d(t, s) (\phi(t+1, s-1) - \phi(t+1, s+1)) \right]$$

$$\phi(t, s) = \left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1)$$

$$w(t, s) \doteq \phi(t+1, s-1) - \phi(t+1, s+1)$$

$$\Psi(f, t+1) = \sum_s [f(t, s) \phi(t, s) + d(t, s) w(t, s)] = \Psi(f, t) - \sum_s d(t, s) w(t, s)$$

Proof of Theorem

$$f(t+1, s) = f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s-1) \left(\frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left(\frac{1}{2} - \gamma - d(t, s+1) \right) \right] \phi(t+1, s)$$

$$\Psi(f, t+1) = \sum_s \left[f(t, s) \left(\left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1) \right) - d(t, s) (\phi(t+1, s-1) - \phi(t+1, s+1)) \right]$$

$$\phi(t, s) = \left(\frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left(\frac{1}{2} - \gamma \right) \phi(t+1, s-1)$$

$$w(t, s) \doteq \phi(t+1, s-1) - \phi(t+1, s+1)$$

$$\Psi(f, t+1) = \sum_s [f(t, s) \phi(t, s) + d(t, s) w(t, s)] = \Psi(f, t) - \sum_s d(t, s) w(t, s)$$

$$\text{If } \sum_s d(t, s) w(t, s) \geq 0 \quad \text{then } \Psi(f, t+1) \leq \Psi(f, t)$$

Theorem about BBM

Theorem about BBM

setting the boosting weights at iteration **t** to be

$$w(t,s) = \binom{T-t}{\left\lfloor \frac{T-t-s+1}{2} \right\rfloor} \left(\frac{1}{2} + \gamma \right)^{\left\lfloor \frac{T-t-s+1}{2} \right\rfloor} \left(\frac{1}{2} - \gamma \right)^{\left\lceil \frac{T-t-s+1}{2} \right\rceil}$$

Theorem about BBM

setting the boosting weights at iteration **t** to be


$$w(t,s) = \binom{T-t}{\left\lfloor \frac{T-t-s+1}{2} \right\rfloor} \left(\frac{1}{2} + \gamma\right)^{\left\lfloor \frac{T-t-s+1}{2} \right\rfloor} \left(\frac{1}{2} - \gamma\right)^{\left\lceil \frac{T-t-s+1}{2} \right\rceil}$$

guarantees


Initial potential $= \Psi(0) \geq \Psi(1) \geq \dots \geq \Psi(T) =$ training error of $\text{sign}\left(\sum_{t=1}^T h_t(x)\right)$

$$\varepsilon = \Psi(0) = \psi(0,0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)$$


$t=1$




$t=11$




t=21




t=31




$t=41$




t=51




t=61




t=71




t=81




t=91



t=101



BBM/Logitboost/Adaboost



High level summary

High level summary

- The worst case adversary splits each bin into:
 $1/2 - \gamma$ incorrect / $1/2 + \gamma$ correct


High level summary

- The worst case adversary splits each bin into:
 $1/2 - \gamma$ incorrect / $1/2 + \gamma$ correct
- Alternative interpretation: Random walk with IID steps.







High level summary

- The worst case adversary splits each bin into:
 $1/2 - \gamma$ incorrect / $1/2 + \gamma$ correct
- Alternative interpretation: Random walk with IID steps.
- Algorithm is derived as optimal response to this simple worst-case adversary.

Drifting games (in 2d)








Drifting games (in 2d)



$$\sum_i \|\vec{w}_i\| = 1$$



Drifting games (in 2d)



$$\sum_i \|\vec{w}_i\| = 1$$

$$\sum_i \vec{w}_i \vec{x}_i \geq \delta > 0$$

Drifting games (in 2d)



Target
region




The allowable steps

B = the set of all allowable step

Normal **B** = minimal set that spans the space. (~basis)

Regular **B** = a symmetric regular set. (~orthonormal basis)

Regular step set in \mathbb{R}^2



The min/max solution

[Schapire99]

- A potential defined by a min/max recursion

$$\phi_T(\mathbf{s}) = L(\mathbf{s})$$

$$\phi_{t-1}(\mathbf{s}) = \min_{\mathbf{w} \in \mathbb{R}^d} \sup_{\mathbf{z} \in B} (\phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w} \cdot \mathbf{z} - \delta \|\mathbf{w}\|)$$

$$\phi_0(0) = \text{the value of the game}$$

- Shepherd's strategy

$$\mathbf{w}_i^t = \arg \min_{\mathbf{w}} \sup_{\mathbf{z} \in B} (\phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w} \cdot \mathbf{z} - \delta \|\mathbf{w}\|)$$

The solution simplifies when $\delta \rightarrow 0$


If B is normal, and δ is sufficiently small then $\exists \mathbf{w}^*$ such that

$$\phi_{t-1}(\mathbf{s}) = \phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w}^* \cdot \mathbf{z} - \delta \|\mathbf{w}^*\|$$

for all $\mathbf{z} \in B$ (and all $t = 1, 2, \dots, \mathbf{s} \in \mathbb{R}^d$)

Implies that: \mathbf{w}^* is the “local slope” at $\phi_t(\mathbf{s})$, i.e.

$$\phi_t(\mathbf{s} + \mathbf{z}_i) = C + \mathbf{w}^* \mathbf{z}_i ; \quad C \doteq \frac{\sum_{j=0}^d \phi_t(\mathbf{s} + \mathbf{z}_j)}{d+1}$$



and that

$$\phi_{t-1}(\mathbf{s}) = C - \delta \|\mathbf{w}^*\|$$

Increasing the number of steps

- Consider T steps in a unit time
- Drift δ should scale like $1/T$
- Step size $O(1/T)$ gives game to shepherd
- Step size $o(1/\sqrt{T})$ keeps game balanced

The solution when $T \rightarrow \infty$

The local slope becomes the gradient

$$\mathbf{w}^* = \nabla \phi_\tau(\mathbf{s})$$

The recursion becomes a PDE

$$\begin{aligned} \frac{\partial \phi_\tau(\mathbf{s})}{\partial \tau} &= -\frac{1}{2} \sum_{k=1}^d \frac{\partial^2 \phi_\tau(\mathbf{s})}{\partial^2 s_k} + \delta \|\mathbf{w}^*\| \\ &= -\frac{1}{2} \Delta \phi_\tau(\mathbf{s}) + \delta \|\nabla \phi_\tau(\mathbf{s})\| \end{aligned}$$

Same PDE describes time development of
Brownian motion with drift

Plan of talk

- Label noise and convex loss functions.
- Boost by Majority and drifting games.
- **Boosting in continuous time.**
- RobustBoost
- Experimental results.

Why is BBM not practical?

Why is BBM not practical?

- BBM needs to know ε, γ before starting.

Why is BBM not practical?

- BBM needs to know ε, γ before starting.

$$T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$$

Why is BBM not practical?

- BBM needs to know ϵ, γ before starting.

$$T = \frac{1}{\gamma^2} \ln \frac{1}{\epsilon}$$

- Adaboost = adaptive boosting, Adapts to the sequence $\gamma_1, \gamma_2, \gamma_3, \dots$

Why is BBM not practical?

- BBM needs to know ϵ, γ before starting.

$$T = \frac{1}{\gamma^2} \ln \frac{1}{\epsilon}$$

- Adaboost = adaptive boosting, Adapts to the sequence $\gamma_1, \gamma_2, \gamma_3, \dots$
- No need to set parameters in advance.

Why is BBM not practical?

- BBM needs to know ϵ, γ before starting.

$$T = \frac{1}{\gamma^2} \ln \frac{1}{\epsilon}$$

- Adaboost = adaptive boosting, Adapts to the sequence $\gamma_1, \gamma_2, \gamma_3, \dots$
 - No need to set parameters in advance.
 - generates a weighted majority rule.

Why is BBM not practical?

- BBM needs to know ϵ, γ before starting.

$$T = \frac{1}{\gamma^2} \ln \frac{1}{\epsilon}$$

- Adaboost = adaptive boosting, Adapts to the sequence $\gamma_1, \gamma_2, \gamma_3, \dots$
 - No need to set parameters in advance.
 - generates a weighted majority rule.
 - Decide when to stop using cross-validation.

Why is BBM not practical?

- BBM needs to know ε, γ before starting.

$$T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$$

- Adaboost = adaptive boosting, Adapts to the sequence $\gamma_1, \gamma_2, \gamma_3, \dots$
 - No need to set parameters in advance.
 - generates a weighted majority rule.
 - Decide when to stop using cross-validation.
- How can we make BBM adaptive?

Letting time step decrease to zero.

Letting time step decrease to zero.

- Number of iterations required by BBM: $T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$

Letting time step decrease to zero.

- Number of iterations required by BBM: $T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$
- Keep ε fixed and let $\gamma \rightarrow 0, T \rightarrow \infty$

Letting time step decrease to zero.

- Number of iterations required by BBM: $T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$
- Keep ε fixed and let $\gamma \rightarrow 0, T \rightarrow \infty$
- The same weak rule is added many times, until its advantage falls below γ .

Letting time step decrease to zero.

- Number of iterations required by BBM: $T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$
- Keep ε fixed and let $\gamma \rightarrow 0, T \rightarrow \infty$
- The same weak rule is added many times, until it's advantage falls below γ .
- yields **adaptive** boosting and a **weighted** majority rule.

Letting time step decrease to zero.

- Number of iterations required by BBM: $T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$
- Keep ε fixed and let $\gamma \rightarrow 0, T \rightarrow \infty$
- The same weak rule is added many times, until it's advantage falls below γ .
 - yields **adaptive** boosting and a **weighted** majority rule.
- In the limit, adversary uses random walk in continuous time = Brownian Motion.


Letting time step decrease to zero.

- Number of iterations required by BBM: $T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$
- Keep ε fixed and let $\gamma \rightarrow 0, T \rightarrow \infty$
- The same weak rule is added many times, until it's advantage falls below γ .
 - yields **adaptive** boosting and a **weighted** majority rule.
- In the limit, adversary uses random walk in continuous time = Brownian Motion.
- Instead of $t=1,2,\dots,T$ use $t=1/T,2/T,\dots,1$

The game lattice

$T=3, t=1,2,3$

s




Using step $\Delta s = \pm \frac{1}{T}$

$$T = 1, \Delta t = 1, \Delta s = \pm 1$$

Using step $\Delta s = \pm \frac{1}{T}$

$$T = 1, \Delta t = 1, \Delta s = \pm 1$$




Using step $\Delta s = \pm \frac{1}{T}$

$$T = 3, \Delta t = \frac{1}{3}, \Delta s = \pm \frac{1}{3}$$

Using step $\Delta s = \pm \frac{1}{T}$

S

$$T = 3, \Delta t = \frac{1}{3}, \Delta s = \pm \frac{1}{3}$$




Using step $\Delta s = \pm \frac{1}{T}$

$$T = 9, \Delta t = \frac{1}{9}, \Delta s = \pm \frac{1}{9}$$

Using step $\Delta s = \pm \frac{1}{T}$

S


$$T = 9, \Delta t = \frac{1}{9}, \Delta s = \pm \frac{1}{9}$$



Using step $\Delta s = \pm \frac{1}{T}$

S

$$T = 9, \Delta t = \frac{1}{9}, \Delta s = \pm \frac{1}{9}$$




Looks fine but $\text{var}(s) = T \frac{1}{T^2} = \frac{1}{T} \xrightarrow{T \rightarrow \infty} 0$

Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 1, \Delta t = 1, \Delta s = \pm 1$$


Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 1, \Delta t = 1, \Delta s = \pm 1$$



Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 1, \Delta t = 1, \Delta s = \pm 1$$




Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 3, \Delta t = \frac{1}{3}, \Delta s = \pm \frac{1}{\sqrt{3}}$$

Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 3, \Delta t = \frac{1}{3}, \Delta s = \pm \frac{1}{\sqrt{3}}$$


S



Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 3, \Delta t = \frac{1}{3}, \Delta s = \pm \frac{1}{\sqrt{3}}$$


s




Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 9, \Delta t = \frac{1}{9}, \Delta s = \pm \frac{1}{3}$$


Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$



Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$



Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$



Potentials in continuous time

Potentials in continuous time

- **Discrete time:** Equations relating time t to time $t+1$ based on random walks.

Potentials in continuous time

- **Discrete time:** Equations relating time t to time $t+1$ based on random walks.
- **Continuous time:**
Differential Equations describing the density evolution for Brownian motion with drift (known as the Kolmogorov forward and backward equations).

Example: From BBM to Brownboost

Example: From BBM to Brownboost

Potential function for BBM:

Example: From BBM to Brownboost

Potential function for BBM:

$$\phi(t-1, s) = \left(\frac{1}{2} - \gamma\right) \phi(t, s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t, s+1)$$

Example: From BBM to Brownboost

Potential function for BBM:

$$\phi(t-1, s) = \left(\frac{1}{2} - \gamma\right)\phi(t, s-1) + \left(\frac{1}{2} + \gamma\right)\phi(t, s+1)$$
$$\phi(T, s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases}$$

Example: From BBM to Brownboost

Potential function for BBM:

$$\phi(t-1, s) = \left(\frac{1}{2} - \gamma\right) \phi(t, s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t, s+1)$$
$$\phi(T, s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)$$

Example: From BBM to Brownboost

Potential function for BBM:

$$\begin{aligned}\phi(t-1, s) &= \left(\frac{1}{2} - \gamma\right) \phi(t, s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t, s+1) \\ \phi(T, s) &= \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)\end{aligned}$$

Potential function for Brownboost:

Example: From BBM to Brownboost

Potential function for BBM:

$$\begin{aligned}\phi(t-1, s) &= \left(\frac{1}{2} - \gamma\right) \phi(t, s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t, s+1) \\ \phi(T, s) &= \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)\end{aligned}$$

Potential function for Brownboost:

$$\frac{\partial}{\partial t} \phi(t, s) = -\frac{1}{2} \frac{\partial^2}{\partial s^2} \phi(t, s) - 2\sqrt{\beta} \frac{\partial}{\partial s} \phi(t, s)$$

Example: From BBM to Brownboost

Potential function for BBM:

$$\begin{aligned}\phi(t-1, s) &= \left(\frac{1}{2} - \gamma\right) \phi(t, s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t, s+1) \\ \phi(T, s) &= \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)\end{aligned}$$

Potential function for Brownboost:

$$\frac{\partial}{\partial t} \phi(t, s) = -\frac{1}{2} \frac{\partial^2}{\partial s^2} \phi(t, s) - 2\sqrt{\beta} \frac{\partial}{\partial s} \phi(t, s)$$

Boundary conditions:

Example: From BBM to Brownboost

Potential function for BBM:

$$\begin{aligned}\phi(t-1, s) &= \left(\frac{1}{2} - \gamma\right) \phi(t, s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t, s+1) \\ \phi(T, s) &= \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)\end{aligned}$$

Potential function for Brownboost:

$$\frac{\partial}{\partial t} \phi(t, s) = -\frac{1}{2} \frac{\partial^2}{\partial s^2} \phi(t, s) - 2\sqrt{\beta} \frac{\partial}{\partial s} \phi(t, s)$$

Boundary conditions:

$$\phi(1, s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases}$$

Example: From BBM to Brownboost

Potential function for BBM:

$$\phi(t-1, s) = \left(\frac{1}{2} - \gamma\right) \phi(t, s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t, s+1)$$
$$\phi(T, s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)$$

Potential function for Brownboost:

$$\frac{\partial}{\partial t} \phi(t, s) = -\frac{1}{2} \frac{\partial^2}{\partial s^2} \phi(t, s) - 2\sqrt{\beta} \frac{\partial}{\partial s} \phi(t, s)$$

Boundary conditions:

$$\phi(1, s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0)$$

properties of the potential

properties of the potential

$$\psi(i,r) = \frac{\psi(i,r-1) + \psi(i+1,r-1)}{2}$$

properties of the potential

$$\psi(i, r) = \frac{\psi(i, r-1) + \psi(i+1, r-1)}{2}$$

End of game

$$\psi(i, 0) = \begin{cases} 1 & i \leq k \\ 0 & i > k \end{cases}$$

Illegal configuration:

$$\Psi(\text{configuration}) = \sum_{j=1}^k f(i) \psi(i, 0) < \frac{1}{N}$$

properties of the potential

$$\psi(i, r) = \frac{\psi(i, r-1) + \psi(i+1, r-1)}{2}$$

End of game

$$\psi(i, 0) = \begin{cases} 1 & i \leq k \\ 0 & i > k \end{cases}$$

Illegal configuration:

$$\Psi(\text{configuration}) = \sum_{j=1}^k f(i) \psi(i, 0) < \frac{1}{N}$$

Beginning of game

Number of errors if adversary always plays optimally $r-1$, where r is the smallest integer for which

$$\psi(0, r) < \frac{1}{N}$$

BW Prediction algorithm

BW Prediction algorithm

- **Initialization:** set r to be the number of errors against optimal adversary.

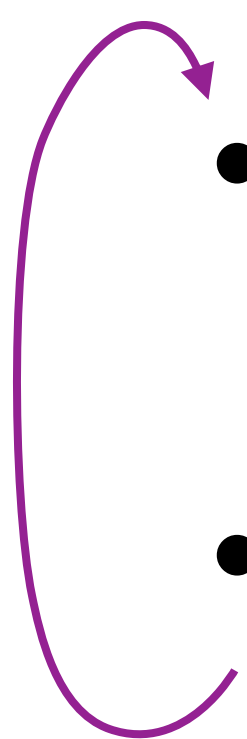
BW Prediction algorithm

- **Initialization:** set r to be the number of errors against optimal adversary.
- Given expert predictions: choose prediction that will result (assuming error) in a lower-potential configuration.

BW Prediction algorithm

- **Initialization:** set r to be the number of errors against optimal adversary.
- Given expert predictions: choose prediction that will result (assuming error) in a lower-potential configuration.
- Decrease r if possible.

BW Prediction algorithm

- **Initialization:** set r to be the number of errors against optimal adversary.
 - Given expert predictions: choose prediction that will result (assuming error) in a lower-potential configuration.
 - Decrease r if possible.
- 

Main properties

Main properties

- If algorithm is followed, the potential of the configuration never increases - is always $\leq 1/N$

Main properties

- If algorithm is followed, the potential of the configuration never increases - is always $\leq 1/N$
- Algorithm is min/max optimal.

Main properties

- If algorithm is followed, the potential of the configuration never increases - is always $\leq 1/N$
- Algorithm is min/max optimal.
 - Removing assumption that expert set is divisible
min/max optimality holds if $N > 2^{2^k}$
 - Based on relation to Ulam's game with k lies [Spencer 92]

Alternative Representation

Alternative Representation

The difference between the two configurations can be represented as a weighted sum

Alternative Representation

The difference between the two configurations can be represented as a weighted sum

$$\Psi(\text{configuration 1}) - \Psi(\text{configuration 0}) = \sum_{j=0}^k f(i)w(i,r)$$

$$w(i,r) = \psi(i+1,r-1) - \psi(i,r-1) = \frac{1}{2^{r-1}} \binom{r-1}{k-i}$$

Alternative Representation

The difference between the two configurations can be represented as a weighted sum

$$\Psi(\text{configuration 1}) - \Psi(\text{configuration 0}) = \sum_{j=0}^k f(i)w(i,r)$$

$$w(i,r) = \psi(i+1,r-1) - \psi(i,r-1) = \frac{1}{2^{r-1}} \binom{r-1}{k-i}$$

The optimal prediction is according to the the sign of this weighted sum.

The BW algorithm

The BW algorithm

- Better error bound than exponential weights.

The BW algorithm

- Better error bound than exponential weights.
- A-priori assumption that one of the experts has loss at most k , we want a bound on the regret without any a priori assumptions.

The BW algorithm

- Better error bound than exponential weights.
- A-priori assumption that one of the experts has loss at most k , we want a bound on the regret without any a priori assumptions.
- Instantaneous loss is restricted to $\{0, 1\}$, we want it to be any number in $[-1, +1]$.

Design of NormalHedge

Design of NormalHedge

- BW: potential function depends on **loss** and **number of remaining mistakes**

Design of NormalHedge

- BW: potential function depends on **loss** and **number of remaining mistakes**
- Normal-Hedge: Potential function based on **regret** and **variance of the positive regrets**

The NormalHedge potential

The NormalHedge potential

$$\text{Potential: } \psi(r, c) = \begin{cases} \exp\left(\frac{r^2}{2c}\right) & \text{if } r \geq 0 \\ 1 & \text{if } r \leq 0 \end{cases}$$

The NormalHedge potential

$$\text{Potential: } \psi(r, c) = \begin{cases} \exp\left(\frac{r^2}{2c}\right) & \text{if } r \geq 0 \\ 1 & \text{if } r \leq 0 \end{cases}$$

$$\text{Weight: } w(r, c) = \frac{\partial}{\partial r} \psi(r, c) = \begin{cases} \frac{r}{c} \exp\left(\frac{r^2}{2c}\right) & \text{if } r \geq 0 \\ 0 & \text{if } r \leq 0 \end{cases}$$

NormalHedge algorithm

NormalHedge algorithm

for $t=0,1,2,\dots$

NormalHedge algorithm

for $t=0,1,2,\dots$

if $\forall i, R_i^t \leq 0$ then $w_i^t = 1 / N$

NormalHedge algorithm

for $t=0,1,2,\dots$

if $\forall i, R_i^t \leq 0$ then $w_i^t = 1 / N$

else

NormalHedge algorithm

for $t=0,1,2,\dots$

if $\forall i, R_i^t \leq 0$ then $w_i^t = 1 / N$

else
set $c(t)$ so that $\frac{1}{N} \sum_{i=1}^N \psi(R_i^t, c(t)) = e$

NormalHedge algorithm

for $t=0,1,2,\dots$

if $\forall i, R_i^t \leq 0$ then $w_i^t = 1 / N$

else

set $c(t)$ so that $\frac{1}{N} \sum_{i=1}^N \psi(R_i^t, c(t)) = e$

$w_i^t = w(R_i^t, c(t))$

NormalHedge algorithm

for $t=0,1,2,\dots$

if $\forall i, R_i^t \leq 0$ then $w_i^t = 1 / N$

else

set $c(t)$ so that $\frac{1}{N} \sum_{i=1}^N \psi(R_i^t, c(t)) = e$

$$w_i^t = w(R_i^t, c(t))$$

Incur instantaneous losses: $\langle l_1^t, l_2^t, \dots, l_N^t \rangle$

NormalHedge algorithm

for $t=0,1,2,\dots$

if $\forall i, R_i^t \leq 0$ then $w_i^t = 1 / N$

else

set $c(t)$ so that $\frac{1}{N} \sum_{i=1}^N \psi(R_i^t, c(t)) = e$

$$w_i^t = w(R_i^t, c(t))$$

Incur instantaneous losses: $\langle l_1^t, l_2^t, \dots, l_N^t \rangle$

$$\text{Algorithm loss: } l_A^t = \frac{\sum_{i=1}^N w_i^t l_i^t}{\sum_{i=1}^N w_i^t}$$

NormalHedge algorithm

for $t=0,1,2,\dots$

if $\forall i, R_i^t \leq 0$ then $w_i^t = 1 / N$

else

set $c(t)$ so that $\frac{1}{N} \sum_{i=1}^N \psi(R_i^t, c(t)) = e$

$$w_i^t = w(R_i^t, c(t))$$

Incur instantaneous losses: $\langle l_1^t, l_2^t, \dots, l_N^t \rangle$

$$\text{Algorithm loss: } l_A^t = \frac{\sum_{i=1}^N w_i^t l_i^t}{\sum_{i=1}^N w_i^t}$$


$$\text{Update regrets: } R_i^{t+1} = R_i^t + l_A^t - l_i^t$$

Illustrative Example




Illustrative Example

● Expert



Illustrative Example

- Expert
- Algorithm




Illustrative Example

● Expert

● Algorithm

$$\text{---} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




Illustrative Example

● Expert

● Algorithm

— $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




Illustrative Example

● Expert

● Algorithm

— $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




Illustrative Example

● Expert

● Algorithm

— $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




Illustrative Example

● Expert

● Algorithm

— $\exp(\eta G)$

$$\text{---} \left\{ \begin{array}{ll} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{array} \right.$$




Normal-Hedge Performance bound

[Chaudhuri, Freund & Hsu 2009]

The regret of NormalHedge is upper bounded by

$$O\left(\sqrt{T \ln N} + \ln^3 N\right)$$

Performance on flip-flop



Summary

- Drifting games is a method for **deriving** new potential functions for online learning and boosting.
- Performed by working backwards in time, starting from final loss function and working backwards.
- Adversarial strategy: random walk / normal distribution.
- Yields Brown-boost
- Yields Normal-Hedge
- Both Brownboost and normal-hedge have difficult open problems.