

# Drifting Games

# Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- $N$  experts make predictions:
  - expert 1: 1,0,1,0,1,1,1
  - expert 2: 0,0,0,1,1,0,1
  - ...
  - expert  $N$ : 1,0,0,1,1,1,0
- Algorithm makes prediction: 1,0,1,1,0,1,1
- Assumption: there exists an expert which makes at most  $k$  mistakes
- Goal: make least mistakes under the assumption (no statistics!)

# Analysis using meta-experts

- Consider only the iterations on which the algorithm makes a mistake.
- Expand each original (base) expert into meta-expert= base-expert+ (sequence of locations of mistakes).
- Suppose a sequence of length  $m=10$  and one of the base experts makes at most  $k=2$  mistakes.
- Mistake no.: 1,2,3,4,5,6,7,8,9,10
  - No mistake: 0,0,0,0,0,0,0,0,0,0
  - 1 mistake: 1,2,3,...,10
  - 2 mistakes: (1,2),(1,3),(1,4),..., (9,10)
- If the number of mistakes is at most 2, then there is a meta-expert that makes no mistakes.
- We can use the halving algorithm on the set of meta experts.


Number of mistakes is smallest  $m$  such that  $\frac{1}{2^{m+1}} \binom{m+1}{\leq k-i} = \frac{1}{2^{m+1}} \sum_{j=0}^{k-i} \binom{m+1}{j} < 1$

- Naive implementation requires explicitly maintaining each meta-expert.

# Reduction to chip game

- **Chip** = (base)-expert, **bin** = number of mistakes made so far
- Game step = algorithm's prediction is incorrect.

# Chip game for expert advice




# 20 questions with $k$ lies

## [Ulam's game with lies]

- Player 1 chooses secret  $x$  from  $1, \dots, N$
- Player 2 asks “is  $x$  in the set  $A$ ?”
- Player 1 answers  $\text{yes/no}$ , can lie at most  $k$  times.
- Game ends when player 2 can determine  $x$ .
- A smart player 1 aims to have more than one  $\text{consistent } x$  for as long as possible.
- $\text{Chip}$  = element  $i$
- $\text{Bin}$  = number of lies made regarding element

# Chip game for 20 questions

○ element in A  
● element not in A



Game ends when there  
is only one chip on this side

# Simple case – all chips in bin 1

- 21 questions without lies
- Combining experts where one is perfect
- Splitting a cookie
- Optimal strategies:
  - Player 1: split chips into two equal parts
  - Player 2: choose larger part
- Note problem when number of chips is odd



# Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.
- Fraction of bin **i** that will remain in bins **1..k** when **m** iterations remain

$$\Phi(i, k, m) \doteq \frac{1}{2^m} \binom{m}{\leq k-i} = \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

- Configuration:  $\vec{n} = \langle n_0, n_1, \dots, n_k \rangle$
- Potential: The number of chips that will remain in bins **0..k** when **m** iterations remain and both sides play optimally:

$$\Psi(\vec{n}, k, m) = \sum_{i=0}^k n_i \Phi(i, k, m)$$

# Combining experts, the binary prediction case


- Goal is to predict a binary sequence, making as few mistakes as possible.
- There are  $N$  experts.
- All predictions are binary and deterministic.
- A-priori knowledge: there is an expert that never makes more than  $k$  mistakes.
- $k=0$  corresponds to the halving algorithm.

# Combining experts as a drifting game


[Cesa-Bianchi, Freund, Helmbold, Warmuth 96]

Binary instantaneous loss  $l_i^t, l_A^t \in \{0,1\}$


Bin  $s$  contains all experts for which  $L_i^t = s$



# The game lattice



Initial configuration  $t=1$




Experts predictions  $t=1$


Predict 0

Algorithm predicts with majority = 0


Suppose algorithm is wrong

Predict 1





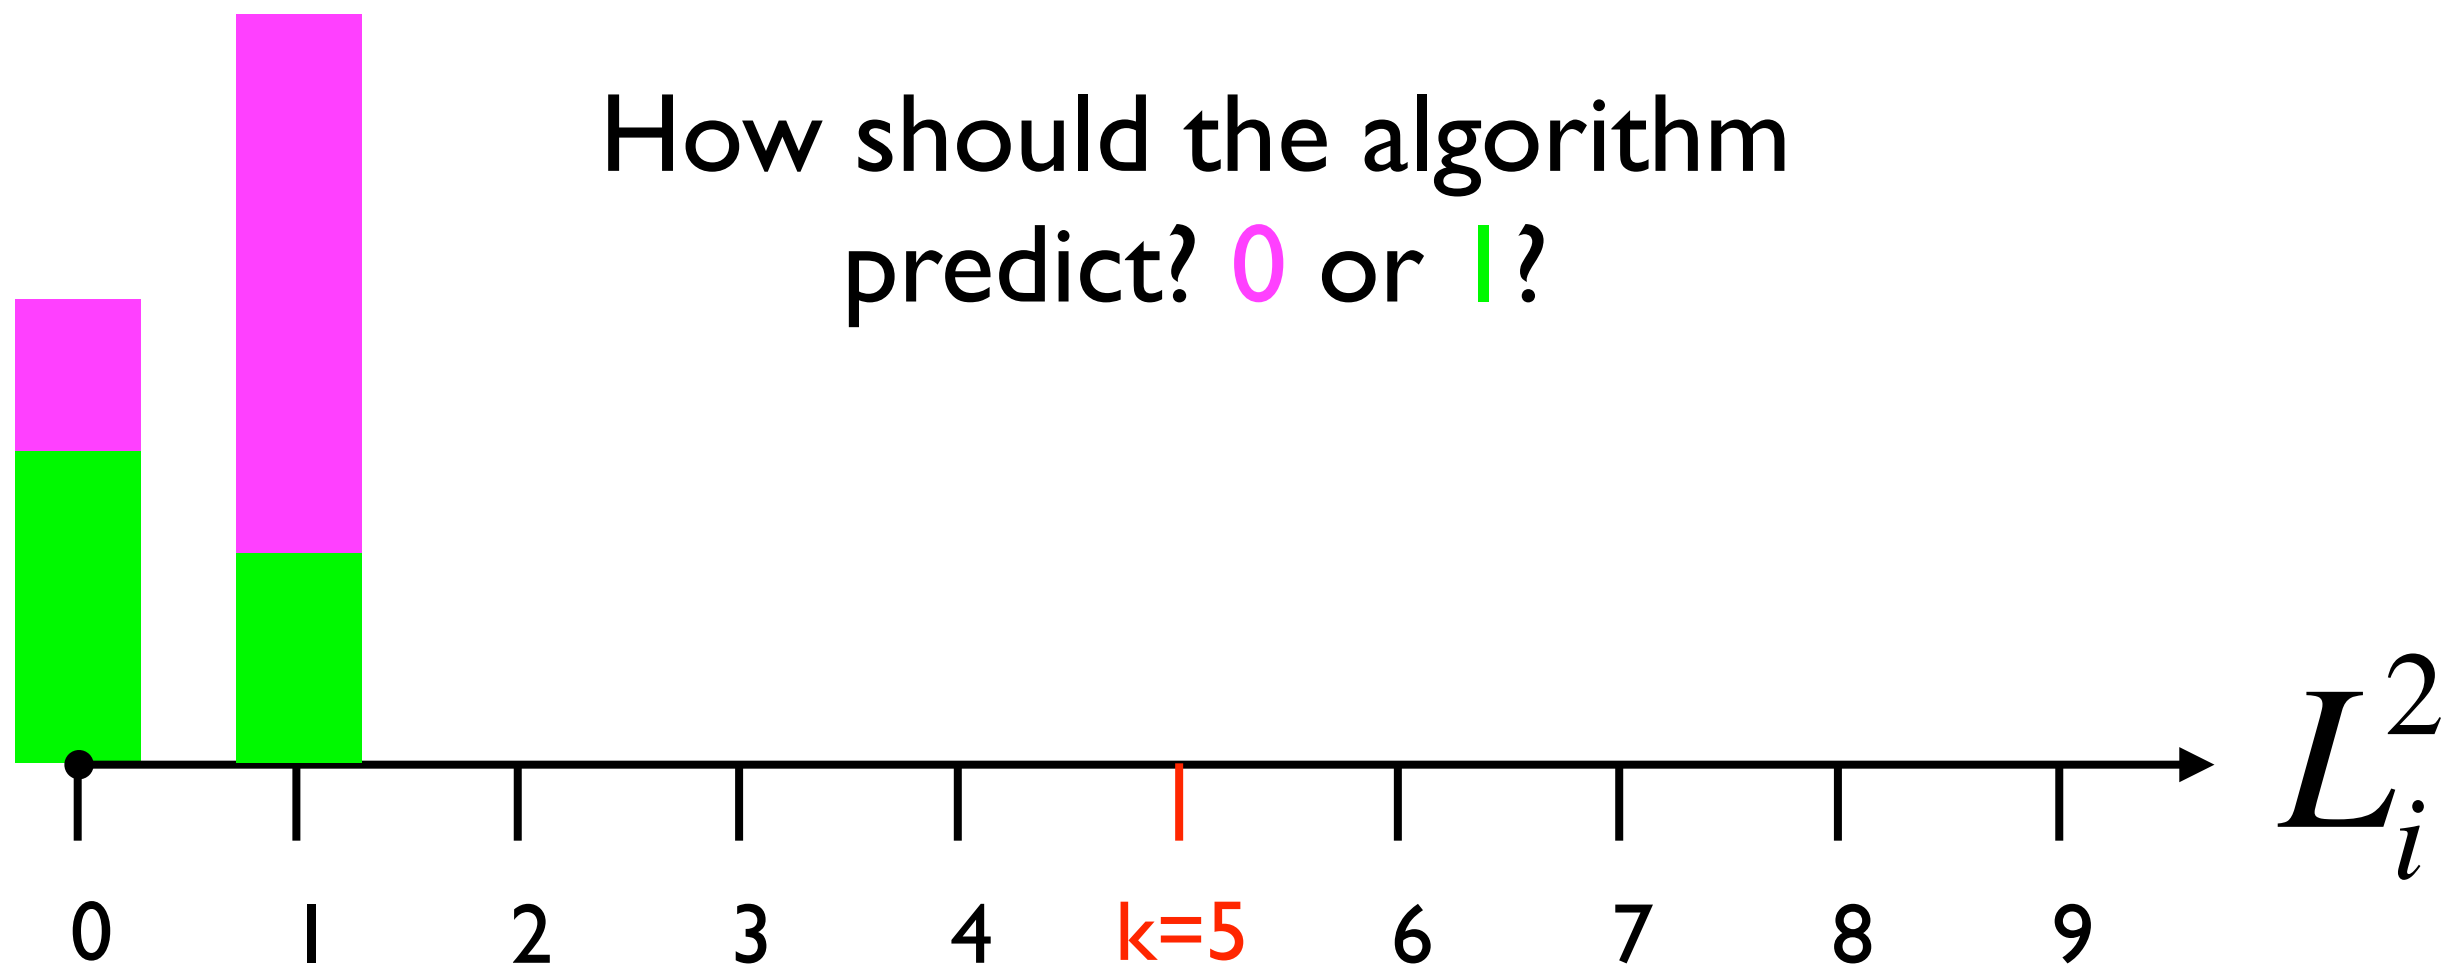
configuration at  $t=2$






Experts predictions  $t=2$

How should the algorithm  
predict? 0 or 1?




configuration  $t=3$

Prediction is 0 and outcome is 1




Experts predictions  $t=2$

How should the algorithm  
predict? 0 or 1?




configuration  $t=3$

Prediction is 1 and outcome is 0



If an error will lead to this configuration  
then an error is not possible  
 $\Rightarrow$  this is a safe prediction

Algorithm's goal is to get to an illegal configuration  
with the smallest number of mistakes.




# Helping the adversary.

- Assume that the set of experts is continuous, arbitrarily divisible.
- a-priori knowledge:  $1/N$  fraction of the expert “mass” have cumulative loss at most  $k$
- Find algorithm with the tightest uniform upper bound on the cumulative loss.

If an error will lead to this configuration  
then an error is not possible  
 $\Rightarrow$  this is a safe prediction

Algorithm's goal is to get to an illegal configuration  
with the smallest number of mistakes.



# An optimal adversarial strategy

- Split each bin to two equal parts. Algorithm's prediction is always incorrect.
- Equivalently: predictions of each expert are IID 0,1 with probabilities  $1/2, 1/2$
- Same adversarial strategy was used to prove general lower bound on BLG



# Optimal prediction strategy

- Assume that adversary will play optimally from the next iteration until the end of the game.
- Choose as the next configuration the one that would end the game faster.
- Relevant only when adversary plays sub-optimally, when adversary plays optimally the two next configurations are identical.

# Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into two **equal** parts (might not be possible when number of chips is finite).
- Future configurations independent of chooser's choice.
- Fraction of bin **i** that will remain in bins **1..k** when **m** iterations remain

$$\phi(i, k, m) \doteq \frac{1}{2^m} \binom{m}{\leq k-i} = \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

- Configuration:  $\vec{n} = \langle n_0, n_1, \dots, n_k \rangle$
- Potential: The number of chips that will remain in bins **0..k** when **m** iterations remain and both sides play optimally:

$$\Psi(\vec{n}, k, m) = \sum_{i=0}^k n_i \phi(i, k, m)$$

# Response to suboptimal play


- The chooser part so that the potential of the next configuration will be:
  - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
  - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).
- When split is not even, the chooser might be able to improve bound:
  - Online learning:  $m := \min \{q \in N : \Psi(\vec{n}, k, q) < 1\} - 1$
  - Ulam's game with k lies:  $m := \max \{q \in N : \Psi(\vec{n}, k, q) \geq 1\} - 1$

# Optimality of strategy

- When chips are indivisible there can be situations where optimal play by the splitter is impossible.
- [Spencer95] If number of chips is sufficiently large, equal-potential splits are always possible.
- A sufficient number of chips is  $\Omega(2^{2^k})$

# Number of chips to infinity

Replace individual chips by chip **mass**



Optimal splitter strategy:

Split **each** bin into **two equal** parts

# Equivalence to a random walk

- Both sides playing optimally is equivalent to each chip performing an independent random walk.
- **Potential** = **probability** of a chip in bin **i** ending in bins **0..K** after **m** iteration
- **Weight** = difference between the potentials of a chip in its two possible locations on the following iteration.
- Chooser's optimal strategy: choose set with smaller (larger) **weight**
- **Worst case behavior of experts is to perform a random walk!**

# Response to suboptimal play

- The chooser part so that the potential of the next configuration will be:
  - Predictor in online learning: Minimal weighting. Make game short (fewer mistakes).
  - Secret holder in Ulam's game: Maximal weighting. Make game long (More guesses).
- When split is not even, the chooser might be able to improve bound:
  - online learning:  $m := \min \{q \in N : \Psi(\vec{n}, k, q) < 1\} - 1$
  - Ulam's game with k lies:  $m := \max \{q \in N : \Psi(\vec{n}, k, q) \geq 1\} - 1$


# Optimality of strategy

- When chips are indivisible there can be situations where optimal play by the splitter is impossible.
- [Spencer95] If number of chips is sufficiently large, equal-potential splits are always possible.
- A sufficient number of chips is  $\Omega(2^{2^k})$



# Number of chips to infinity

Replace individual chips by chip **mass**

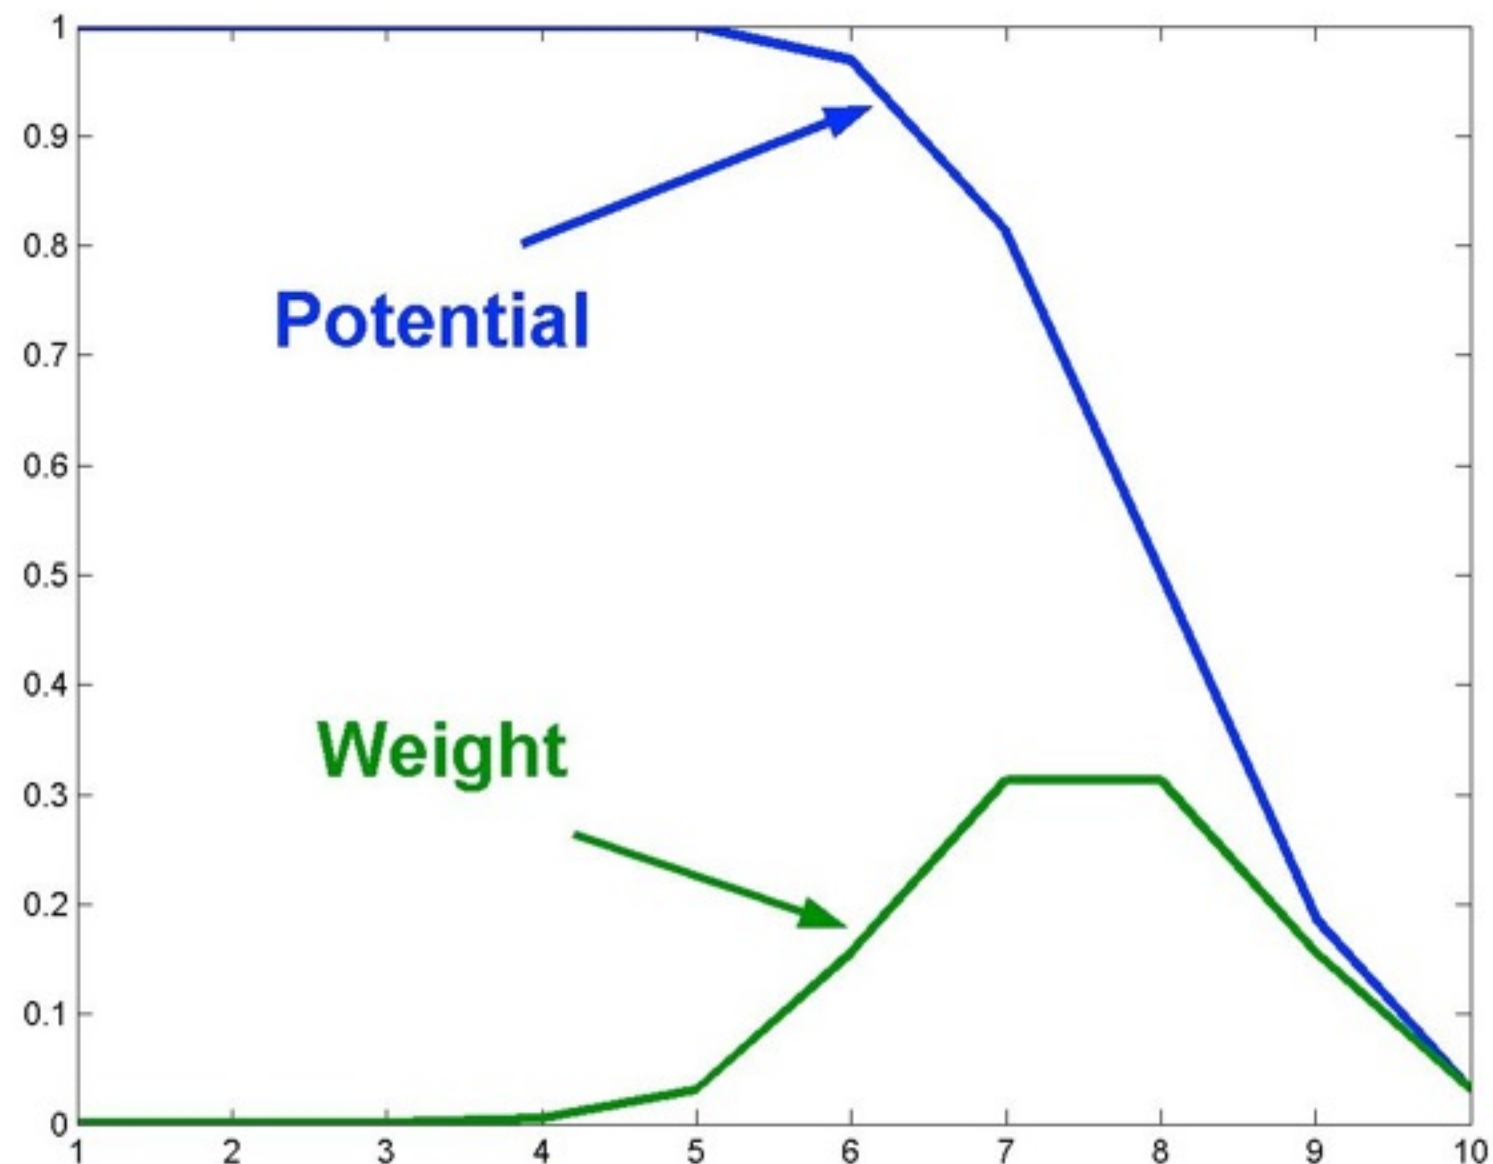


Optimal splitter strategy:

Split **each** bin into **two equal** parts

# Example potential and weight

$m=5; k=10$



# Boosting

- A method for improving classifier accuracy
- **Weak Learner**: a learning algorithm generating rules **slightly** better than random guessing.
- **Basic idea**: **re-weight** training examples to force weak learner into different parts of the space.
- **Combine** weak rules by a **majority vote**.


# Boost by Majority (BBM)

[Freund 95]

- game between a **booster** and a weak **learner**.
- Boosting generates a simple (unweighted) majority rule over weak learners.
- **T** Number of iterations is set in advance
- On iteration  **$t=1..T$** 
  - **booster** assigns weights to the training examples.
  - **learner** chooses a rule whose error wrt the chosen weights is smaller than  **$1/2 - \gamma$**
  - Rule is added to majority rule
- Goal of booster is to minimize number of errors of final majority rule.


# BBM as a drifting game

- Chips = examples
- bin **s** contains the examples for which the difference between the number of correct and incorrect base rules is **s**




# The continuous chip limit

- Number of training examples increases to infinity.
- Alternatively - think of examples as a probability mass with probability measure  $\mu$  defined on it.




# The boosting game lattice



$\gamma = 0.1$


Initial configuration






$\gamma = 0.1$

Booster  
assigns weights  
to  
examples




$\gamma = 0.1$

Weak learner  
chooses subset  
with weight  $1/2 + \gamma$   
which  $h_1(x)$   
classifies correctly




$\gamma = 0.1$

Weak learner  
chooses subset  
with weight  $1/2 + \gamma$   
which  $h_1(x)$   
classifies correctly




$\gamma = 0.1$

Booster  
assigns weights  
to  
examples



$\gamma = 0.1$

Weak learner  
chooses subset  
with weight  $1/2 + \gamma$   
which  $h_2(x)$   
classifies correctly




$\gamma = 0.1$

Weak learner  
chooses subset  
with weight  $1/2 + \gamma$   
which  $h_2(x)$   
classifies correctly




$\gamma = 0.1$

Booster  
assigns weights  
to  
examples



$\gamma = 0.1$


Weak learner  
chooses subset  
with weight  $1/2 + \gamma$   
which  $h_3(x)$   
classifies correctly





$\gamma = 0.1$

Weak learner  
chooses subset  
with weight  $1/2 + \gamma$   
which  $h_3(x)$   
classifies correctly




$\gamma = 0.1$

Majority[ $h_3(x) + h_3(x) + h_3(x)$ ]

incorrect

Majority[ $h_3(x) + h_3(x) + h_3(x)$ ]

correct



# Weak Learner's min/max strategy

- **AdOpt** - Choose  $1/2 + \gamma$  from each bin to be correct.
- Equivalently: prediction of each base rule on each example is chosen independently at random

$$P(h_t(x)=y) = 1/2 + \gamma$$

# Potential

Total potential:  $\Psi(t, \text{configuration})$  -  $\mu$ -prob of the examples on which the final majority vote is incorrect given the configuration after iteration  $t$  is  $\text{configuration}$  and on the remaining steps the learner plays  $\text{AdOpt}$ .

$\Psi(0, \text{all at origin})$  = Initial potential.

$\Psi(T, \text{configuration})$  = Training error of final majority rule.

Boosting algorithm chooses weights so that the total potential does not increase.

Initial potential  $\geq$  final training error.

Bin Potential:  $\phi(t,s)$  - fraction of examples in bin  $s$  after iteration  $t$  on which the final majority rule will be incorrect assuming AdOpt play

$f(t,s)$  The configuration

The  $\mu$ -prob of examples in bin  $s$  after iteration  $t$

The potential of a configuration:  $\Psi(f,t) = \sum_i f(t,s) \phi(t,s)$

$\phi(t,s)$  does not depend on the configuration

$$\phi(t,s) = \text{Binom}\left(T-t, \frac{T-t-s}{2}, \frac{1}{2} + \gamma\right); \quad \text{Binom}(n,k,p) \doteq \sum_{j=0}^{\lfloor k \rfloor} \binom{n}{j} p^j (1-p)^{n-j}$$

$$\phi(t-1,s) = \left(\frac{1}{2} - \gamma\right) \phi(t,s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t,s+1)$$

$$\phi(T,s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \phi(0,0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)$$

## Definitions

edge:  $d(t, s) \doteq \mu(h_t(x) = y | (x, y) \text{ in bin } s \text{ after iteration } t) - \left(\frac{1}{2} + \gamma\right)$

Example Weights:  $w(t, s) \doteq \phi(t + 1, s - 1) - \phi(t + 1, s + 1)$

## Theorem

If, for step  $t$ ,  $\sum_s d(t, s)w(t, s) \geq 0$  then  $\Psi(t + 1) \leq \Psi(t)$

## Corollary

If  $\forall t$  [weighted error of  $h_t(x)$ ]  $\leq 1/2 - \gamma$

Then

Initial potential  $\geq$  final training error.

# Proof of Theorem

$$f(t+1, s) = f(t, s-1) \left( \frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left( \frac{1}{2} - \gamma - d(t, s+1) \right)$$

$$\Psi(f, t+1) = \sum_s \left[ f(t, s-1) \left( \frac{1}{2} + \gamma + d(t, s-1) \right) + f(t, s+1) \left( \frac{1}{2} - \gamma - d(t, s+1) \right) \right] \phi(t+1, s)$$

$$\Psi(f, t+1) = \sum_s \left[ f(t, s) \left( \left( \frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left( \frac{1}{2} - \gamma \right) \phi(t+1, s-1) \right) - d(t, s) (\phi(t+1, s-1) - \phi(t+1, s+1)) \right]$$

$$\phi(t, s) = \left( \frac{1}{2} + \gamma \right) \phi(t+1, s+1) + \left( \frac{1}{2} - \gamma \right) \phi(t+1, s-1)$$

$$w(t, s) \doteq \phi(t+1, s-1) - \phi(t+1, s+1)$$

$$\Psi(f, t+1) = \sum_s [f(t, s) \phi(t, s) + d(t, s) w(t, s)] = \Psi(f, t) - \sum_s d(t, s) w(t, s)$$

$$\text{If } \sum_s d(t, s) w(t, s) \geq 0 \quad \text{then } \Psi(f, t+1) \leq \Psi(f, t)$$

# Theorem about BBM

setting the boosting weights at iteration **t** to be

$$w(t,s) = \binom{T-t}{\left\lfloor \frac{T-t-s+1}{2} \right\rfloor} \left( \frac{1}{2} + \gamma \right)^{\left\lfloor \frac{T-t-s+1}{2} \right\rfloor} \left( \frac{1}{2} - \gamma \right)^{\left\lceil \frac{T-t-s+1}{2} \right\rceil}$$


guarantees

Initial potential  $= \Psi(0) \geq \Psi(1) \geq \dots \geq \Psi(T) =$  training error of  $\text{sign}\left(\sum_{t=1}^T h_t(x)\right)$


$$\varepsilon = \Psi(0) = \psi(0,0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)$$




$t=1$




$t=11$




t=21




t=31




$t=41$




t=51



t=61




$t=71$







t=81




t=91



t=101









# BBM/Logitboost/Adaboost



# High level summary

- The worst case adversary splits each bin into:  
 $1/2 - \gamma$  incorrect /  $1/2 + \gamma$  correct
- Alternative interpretation: Random walk with IID steps.
- Algorithm is derived as optimal response to this simple worst-case adversary.


# Drifting games (in 2d)



$$\sum_i \|\vec{w}_i\| = 1$$

$$\sum_i \vec{w}_i \vec{x}_i \geq \delta > 0$$

# Drifting games (in 2d)



Target  
region




# The allowable steps

**B** = the set of all allowable step

Normal **B** = minimal set that spans the space. (~basis)

Regular **B** = a symmetric regular set. (~orthonormal basis)

Regular step set in  $\mathbb{R}^2$





# The min/max solution

[Schapire99]

- A potential defined by a min/max recursion

$$\phi_T(\mathbf{s}) = L(\mathbf{s})$$

$$\phi_{t-1}(\mathbf{s}) = \min_{\mathbf{w} \in \mathbb{R}^d} \sup_{\mathbf{z} \in B} (\phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w} \cdot \mathbf{z} - \delta \|\mathbf{w}\|)$$

$$\phi_0(0) = \text{the value of the game}$$

- Shepherd's strategy

$$\mathbf{w}_i^t = \arg \min_{\mathbf{w}} \sup_{\mathbf{z} \in B} (\phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w} \cdot \mathbf{z} - \delta \|\mathbf{w}\|)$$

The solution simplifies when  $\delta \rightarrow 0$


If  $B$  is normal, and  $\delta$  is sufficiently small then  $\exists \mathbf{w}^*$  such that

$$\phi_{t-1}(\mathbf{s}) = \phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w}^* \cdot \mathbf{z} - \delta \|\mathbf{w}^*\|$$

for all  $\mathbf{z} \in B$  (and all  $t = 1, 2, \dots, \mathbf{s} \in \mathbb{R}^d$ )

Implies that:  $\mathbf{w}^*$  is the “local slope” at  $\phi_t(\mathbf{s})$ , i.e.

$$\phi_t(\mathbf{s} + \mathbf{z}_i) = C + \mathbf{w}^* \mathbf{z}_i ; \quad C \doteq \frac{\sum_{j=0}^d \phi_t(\mathbf{s} + \mathbf{z}_j)}{d+1}$$



and that

$$\phi_{t-1}(\mathbf{s}) = C - \delta \|\mathbf{w}^*\|$$

# Increasing the number of steps

- Consider  $T$  steps in a unit time
- Drift  $\delta$  should scale like  $1/T$
- Step size  $O(1/T)$  gives game to shepherd
- Step size  $o(1/\sqrt{T})$  keeps game balanced

# The solution when $T \rightarrow \infty$

The local slope becomes the gradient

$$\mathbf{w}^* = \nabla \phi_\tau(\mathbf{s})$$

The recursion becomes a PDE

$$\begin{aligned} \frac{\partial \phi_\tau(\mathbf{s})}{\partial \tau} &= -\frac{1}{2} \sum_{k=1}^d \frac{\partial^2 \phi_\tau(\mathbf{s})}{\partial^2 s_k} + \delta \|\mathbf{w}^*\| \\ &= -\frac{1}{2} \Delta \phi_\tau(\mathbf{s}) + \delta \|\nabla \phi_\tau(\mathbf{s})\| \end{aligned}$$

Same PDE describes time development of  
Brownian motion with drift

# Plan of talk

- Label noise and convex loss functions.
- Boost by Majority and drifting games.
- **Boosting in continuous time.**
- RobustBoost
- Experimental results.

# Why is BBM not practical?

- BBM needs to know  $\varepsilon, \gamma$  before starting.

$$T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$$

- Adaboost = adaptive boosting, Adapts to the sequence  $\gamma_1, \gamma_2, \gamma_3, \dots$ 
  - No need to set parameters in advance.
  - generates a weighted majority rule.
  - Decide when to stop using cross-validation.
- How can we make BBM adaptive?


# Letting time step decrease to zero.

- Number of iterations required by BBM:  $T = \frac{1}{\gamma^2} \ln \frac{1}{\varepsilon}$
- Keep  $\varepsilon$  fixed and let  $\gamma \rightarrow 0, T \rightarrow \infty$
- The same weak rule is added many times, until it's advantage falls below  $\gamma$ .
  - yields **adaptive** boosting and a **weighted** majority rule.
- In the limit, adversary uses random walk in continuous time = Brownian Motion.
- Instead of  $t=1,2,\dots,T$  use  $t=1/T,2/T,\dots,1$

# The game lattice

$T=3, t=1,2,3$


$s$





# Using step $\Delta s = \pm \frac{1}{T}$


$$T = 1, \Delta t = 1, \Delta s = \pm 1$$



Using step  $\Delta s = \pm \frac{1}{T}$

**S**


$$T = 3, \Delta t = \frac{1}{3}, \Delta s = \pm \frac{1}{3}$$



# Using step $\Delta s = \pm \frac{1}{T}$

**S**

$$T = 9, \Delta t = \frac{1}{9}, \Delta s = \pm \frac{1}{9}$$




Looks fine but  $\text{var}(s) = T \frac{1}{T^2} = \frac{1}{T} \rightarrow 0$   
 $T \rightarrow \infty$

Using step  $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 1, \Delta t = 1, \Delta s = \pm 1$$


**s**




# Using step $\Delta s = \pm \frac{1}{\sqrt{T}}$

$$T = 3, \Delta t = \frac{1}{3}, \Delta s = \pm \frac{1}{\sqrt{3}}$$

**s**



Using step  $\Delta s = \pm \frac{1}{\sqrt{T}}$



# Potentials in continuous time

- **Discrete time:** Equations relating time  $t$  to time  $t+1$  based on random walks.
- **Continuous time:**  
**Differential Equations** describing the density evolution for Brownian motion with drift (known as the Kolmogorov forward and backward equations).

# Example: From BBM to Brownboost

Potential function for BBM:

$$\begin{aligned}\phi(t-1, s) &= \left(\frac{1}{2} - \gamma\right) \phi(t, s-1) + \left(\frac{1}{2} + \gamma\right) \phi(t, s+1) \\ \phi(T, s) &= \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0) = \text{Binom}\left(T, \frac{T}{2}, \frac{1}{2} + \gamma\right)\end{aligned}$$

Potential function for Brownboost:

$$\frac{\partial}{\partial t} \phi(t, s) = -\frac{1}{2} \frac{\partial^2}{\partial s^2} \phi(t, s) - 2\sqrt{\beta} \frac{\partial}{\partial s} \phi(t, s)$$

Boundary conditions:

$$\phi(1, s) = \begin{cases} 0 & s > 0 \\ 1 & s \leq 0 \end{cases} \quad \varepsilon = \phi(0, 0)$$



# properties of the potential

$$\psi(i, r) = \frac{\psi(i, r-1) + \psi(i+1, r-1)}{2}$$

## End of game

$$\psi(i, 0) = \begin{cases} 1 & i \leq k \\ 0 & i > k \end{cases}$$

Illegal configuration:

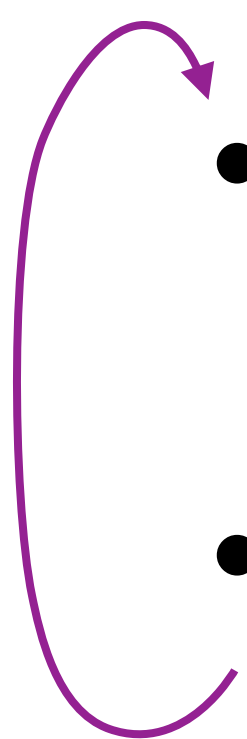
$$\Psi(\text{configuration}) = \sum_{j=1}^k f(i) \psi(i, 0) < \frac{1}{N}$$

## Beginning of game

Number of errors if adversary always plays optimally  $r-1$ , where  $r$  is the smallest integer for which

$$\psi(0, r) < \frac{1}{N}$$

# BW Prediction algorithm

- **Initialization:** set  $r$  to be the number of errors against optimal adversary.
  - Given expert predictions: choose prediction that will result (assuming error) in a lower-potential configuration.
  - Decrease  $r$  if possible.
- 

# Main properties

- If algorithm is followed, the potential of the configuration never increases - is always  $\leq 1/N$
- Algorithm is min/max optimal.
  - Removing assumption that expert set is divisible  
min/max optimality holds if  $N > 2^{2^k}$
  - Based on relation to Ulam's game with  $k$  lies [Spencer 92]

# Alternative Representation

The difference between the two configurations can be represented as a weighted sum

$$\Psi(\text{configuration 1}) - \Psi(\text{configuration 0}) = \sum_{j=0}^k f(i)w(i,r)$$

$$w(i,r) = \psi(i+1,r-1) - \psi(i,r-1) = \frac{1}{2^{r-1}} \binom{r-1}{k-i}$$

The optimal prediction is according to the the sign of this weighted sum.

# The BW algorithm

- Better error bound than exponential weights.
- A-priori assumption that one of the experts has loss at most  $k$ , we want a bound on the regret without any a priori assumptions.
- Instantaneous loss is restricted to  $\{0, 1\}$ , we want it to be any number in  $[-1, +1]$ .

# Design of NormalHedge

- BW: potential function depends on **loss** and **number of remaining mistakes**
- Normal-Hedge: Potential function based on **regret** and **variance of the positive regrets**

# The NormalHedge potential

$$\text{Potential: } \psi(r, c) = \begin{cases} \exp\left(\frac{r^2}{2c}\right) & \text{if } r \geq 0 \\ 1 & \text{if } r \leq 0 \end{cases}$$

$$\text{Weight: } w(r, c) = \frac{\partial}{\partial r} \psi(r, c) = \begin{cases} \frac{r}{c} \exp\left(\frac{r^2}{2c}\right) & \text{if } r \geq 0 \\ 0 & \text{if } r \leq 0 \end{cases}$$

# NormalHedge algorithm

for  $t=0,1,2,\dots$

if  $\forall i, R_i^t \leq 0$  then  $w_i^t = 1 / N$

else

set  $c(t)$  so that  $\frac{1}{N} \sum_{i=1}^N \psi(R_i^t, c(t)) = e$

$$w_i^t = w(R_i^t, c(t))$$

Incur instantaneous losses:  $\langle l_1^t, l_2^t, \dots, l_N^t \rangle$

$$\text{Algorithm loss: } l_A^t = \frac{\sum_{i=1}^N w_i^t l_i^t}{\sum_{i=1}^N w_i^t}$$

$$\text{Update regrets: } R_i^{t+1} = R_i^t + l_A^t - l_i^t$$




# Illustrative Example

● Expert

● Algorithm

—  $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




# Illustrative Example

● Expert

● Algorithm

—  $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




# Illustrative Example

● Expert

● Algorithm

—  $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




# Illustrative Example

● Expert

● Algorithm

—  $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




# Normal-Hedge Performance bound

[Chaudhuri, Freund & Hsu 2009]

The regret of NormalHedge is upper bounded by

$$O\left(\sqrt{T \ln N} + \ln^3 N\right)$$

# Performance on flip-flop



# Summary

- Drifting games is a method for **deriving** new potential functions for online learning and boosting.
- Performed by working backwards in time, starting from final loss function and working backwards.
- Adversarial strategy: random walk / normal distribution.
- Yields Brown-boost
- Yields Normal-Hedge
- Both Brownboost and normal-hedge have difficult open problems.