

# Mixable losses and Tracking the best Expert

Yoav Freund

January 22, 2014

## Outline

- Review

- The general prediction game

- Some useful loss functions

- Vovk's algorithm

- mixable loss functions

- The convexity condition

- Log loss

- Square loss

  - Square loss using simple averaging

- Summary table

- Switching Experts

- An inefficient algorithm

- The fixed-share algorithm

- The variable-share algorithm

## The log-loss game

- ▶ Prediction algorithm  $A$  has access to  $N$  experts.
- ▶ The following is repeated for  $t = 1, \dots, T$ 
  - ▶ Experts generate predictive distributions:  $\mathbf{p}_1^t, \dots, \mathbf{p}_N^t$
  - ▶ Algorithm generates its own prediction  $\mathbf{p}_A^t$
  - ▶  $\mathbf{c}^t$  is revealed.
- ▶ **Goal:** minimize regret:

$$-\sum_{t=1}^T \log p_A^t(\mathbf{c}^t) + \min_{i=1, \dots, N} \left( -\sum_{t=1}^T \log p_i^t(\mathbf{c}^t) \right)$$

## The online Bayes Algorithm

- ▶ Total loss of expert  $i$

$$L_i^t = - \sum_{s=1}^t \log p_i^s(c^s); \quad L_i^0 = 0$$

- ▶ Weight of expert  $i$

$$w_i^t = w_i^1 e^{-L_i^{t-1}} = w_i^1 \prod_{s=1}^{t-1} p_i^s(c^s)$$

- ▶ Freedom to choose initial weights.

$$w_i^1 \geq 0, \sum_{i=1}^N w_i^1 = 1$$

- ▶ Prediction of algorithm  $A$

$$\mathbf{p}_A^t = \frac{\sum_{i=1}^N w_i^t \mathbf{p}_i^t}{\sum_{i=1}^N w_i^t}$$

## Cumulative loss vs. Final total weight

Total weight:  $W^t \doteq \sum_{i=1}^N w_i^t$

$$\frac{W^{t+1}}{W^t} = \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t} = \frac{\sum_{i=1}^N w_i^t p_i^t(c^t)}{\sum_{i=1}^N w_i^t} = p_A^t(c^t)$$

$$-\log \frac{W^{t+1}}{W^t} = -\log p_A^t(c^t)$$

$$-\log W^{T+1} = -\log \frac{W^{T+1}}{W^1} = -\sum_{t=1}^T \log p_A^t(c^t) = L_A^T$$

**EQUALITY** not bound!

## Vovk's general prediction game

$\Gamma$  - prediction space.  $\Omega$  - outcome space.

On each trial  $t = 1, 2, \dots$

1. Each expert  $i \in \{1 \dots N\}$  makes a prediction  $\gamma_i^t \in \Gamma$
2. The learner, after observing  $\langle \gamma_1^t \dots \gamma_N^t \rangle$ ,  
makes its own prediction  $\gamma^t$
3. Nature chooses an outcome  $\omega^t \in \Omega$
4. Each expert incurs loss  $\ell_i^t = \lambda(\omega^t, \gamma_i^t)$   
The learner incurs loss  $\ell_A^t = \lambda(\omega^t, \gamma^t)$

## Achievable loss bounds

- ▶  $L_A \doteq \sum_{t=1}^T \ell_A^t$  - total loss of algorithm
- ▶  $L_i \doteq \sum_{t=1}^T \ell_i^t$  - total loss of expert  $i$
- ▶ **Goal:** find an algorithm which guarantees that

$$(a, c) \in [0, \infty), \quad L_A \leq aL_{\min} + c \ln N$$


For any sequence of events.

- ▶ We say that the pair  $(a, c)$  is **achievable**.

## The set of achievable bounds

- ▶ Fix loss function  $\lambda : \Omega \times \Gamma \rightarrow [0, \infty)$
- ▶ The pair  $(a, c)$  is *achievable* if there exists *some* prediction algorithm such that for *any*  $N > 0$ , *any* set of  $N$  prediction sequences and *any* sequence of outcomes

$$L_A \leq aL_{\min} + c \ln N$$





## Some useful loss functions

- Outcomes:  $\omega^1, \omega^2, \dots, \omega^t \in [0, 1]$
- Predictions:  $\gamma^1, \gamma^2, \dots, \gamma^t \in [0, 1]$

## Log loss (Entropy loss)



$$\lambda_{\text{ent}}(\omega, \gamma) = \omega \ln \frac{\omega}{\gamma} + (1 - \omega) \ln \frac{1 - \omega}{1 - \gamma}$$

- ▶ When  $q_t \in \{0, 1\}$  Cumulative log loss = coding length  $\pm 1$
- ▶ If  $P[\omega_t = 1] = q$ , optimal prediction  $\gamma^t = q$
- ▶ Unbounded loss.
- ▶ Not symmetric  $\exists p, q \lambda(p, q) \neq \lambda(q, p)$ .
- ▶ No triangle inequality  
 $\exists p_1, p_2, p_3 \lambda(p_1, p_3) > \lambda(p_1, p_2) + \lambda(p_2, p_3)$

## Square loss (Breier Loss)



$$\lambda_{\text{sq}}(\omega, \gamma) = (\omega - \gamma)^2$$

- ▶  $P[\omega^t = 1] = q, P[\omega^t = 0] = 1 - q,$   
optimal prediction  $\gamma^t = q$
- ▶ Bounded loss.
- ▶ Defines a metric (symmetric and triangle ineq.)
- ▶ Corresponds to regression.

## Hellinger Loss



$$\lambda_{\text{hel}}(\omega, \gamma) = \frac{1}{2} \left( (\sqrt{\omega} + \sqrt{\gamma})^2 + (\sqrt{1-\omega} + \sqrt{1-\gamma})^2 \right)$$

- ▶ If  $P[\omega^t = 1] = q$ ,  $P[\omega^t = 0] = 1 - q$ ,  
optimal prediction  $\gamma^t = q$
- ▶ Loss is bounded.
- ▶ Defines a metric.
- ▶  $\lambda_{\text{hel}}(p, q) \approx \lambda_{\text{ent}}(p, q)$  when  $p \approx q$  and  $p, q \in (0, 1)$

## Absolute loss



$$\lambda(\omega, \gamma) = |\omega - \gamma|$$

- ▶ Probability of making a mistake if predicting 0 or 1 using a biased coin
- ▶ If  $P[\omega^t = 1] = q$ ,  $P[\omega^t = 0] = 1 - q$ , then the optimal prediction is

$$\gamma^t = \begin{cases} 1 & \text{if } q > 1/2, \\ 0 & \text{otherwise} \end{cases}$$

## Structureless bounded loss

- ▶ Prediction is a distribution  $\gamma = \langle p_1, \dots, p_N \rangle$ ,  $p_i \geq 0$ ,  
 $\sum_{i=1}^N p_i = 1$
- ▶ Outcome is a loss vector  $\omega = \langle \omega_1, \dots, \omega_N \rangle$ ,  $0 \leq \omega_i \leq 1$
- ▶ Loss is the dot product:  $\lambda_{\text{dot}}(\omega, \gamma) = \gamma \cdot \omega$
- ▶ Corresponds to the hedging game.
- ▶ For hedge loss the regret is  $\Omega(\sqrt{T \log N})$ .
- ▶ For the log loss the regret is  $O(\log N)$
- ▶ Which losses behave like **entropy loss** and which behave like **hedge loss**?

## Some technical requirements

- ▶ There should be a **topology** on the prediction set  $\Gamma$  such that
- ▶  $\Gamma$  is compact.
- ▶  $\forall \omega \in \Omega$ , the function  $\gamma \rightarrow \lambda(\omega, \gamma)$  is **continuous**
- ▶ There is a **universally reasonable prediction**  
 $\exists \gamma \in \Gamma, \forall \omega \in \Omega, \lambda(\omega, \gamma) < \infty$
- ▶ There is **no universally optimal prediction**  
 $\neg \exists \gamma \in \Gamma, \forall \omega \in \Omega, \lambda(\omega, \gamma) = 0$

## Vovk's meta-algorithm

- Fix an **achievable** pair  $(a, c)$  and set  $\eta = a/c$
- 1.

$$w_i^t = \frac{1}{N} e^{-\eta L_i^t}$$

Choose  $\gamma_t$  so that, for all  $\omega^t \in \Omega$ :

$$\lambda(\omega^t, \gamma^t) - c \ln \sum_i w_i^t \leq -c \ln \left( \sum_i w_i^t e^{-\eta \lambda(\omega^t, \gamma_i^t)} \right)$$

- 2. If choice of  $\gamma^t$  always exists, then the total loss satisfies:

$$\sum_t \lambda(\omega^t, \gamma^t) \leq -c \ln \sum_i w_i^{T+1} \leq a L_{\min} + c \ln N$$


- Vovk's result: **yes!** a good choice for  $\gamma_t$  always exists!



# Vovk's algorithm is the the highest achiever [Vovk95]

The pair  $(a, c)$  is achieved by some algorithm if and only if it is achieved by Vovk's algorithm.

The separation curve is  $\left\{ \left( a(\eta), \frac{a(\eta)}{\eta} \right) \mid \eta \in [0, \infty] \right\}$



## Mixable Loss Functions

- ▶ A Loss function is **mixable** if a pair of the form  $(1, c)$ ,  $c < \infty$  is achievable.

$$L_A \leq L_{\min} + c \ln N$$

- ▶ Vovk's algorithm with  $\eta = 1/c$  achieves this bound.
- ▶  $\lambda_{\text{ent}}, \lambda_{\text{sq}}, \lambda_{\text{hel}}$  are **mixable**
- ▶  $\lambda_{\text{abs}}, \lambda_{\text{dot}}$  are **not mixable**

## The convexity condition

- ▶ requirement for loss to be  $(1, 1/\eta)$  mixable
- ▶  $\forall \langle (\gamma_1, W_1), \dots, (\gamma_N, W_N) \rangle$   
 $\exists \gamma \in \Gamma$   
 $\forall \omega \in \Omega$ :

$$\lambda(\omega, \gamma) - \frac{1}{\eta} \ln \sum_i W_i \leq -\frac{1}{\eta} \ln \left( \sum_i W_i e^{-\eta \lambda(\omega, \gamma_i)} \right)$$

- ▶ Can be re-written as:


$$e^{-\eta \lambda(\omega, \gamma)} \geq \sum_i \left( \frac{W_i}{\sum_j W_j} \right) e^{-\eta \lambda(\omega, \gamma_i)}$$

- ▶ Equivalently - the image of the set  $\Gamma$  under the mapping  $F(\gamma) = \langle e^{-\eta \lambda(\omega, \gamma)} \rangle_{\omega \in \Omega}$  is concave.

## convexity condition: Pictorially

- **Example:** Suppose  $\Omega = \{0, 1\}$ ,  $\Gamma = [0, 1]$ . then

$$F(\gamma) = \left\langle e^{-\eta\lambda(0,\gamma)}, e^{-\eta\lambda(1,\gamma)} \right\rangle$$



## Vovk Algorithm for log loss

- ▶ The log loss is mixable with  $\eta = 1$
- ▶ The image of  $[0, 1]$  through  $F(\gamma) = \langle e^{-\eta\lambda(0,\gamma)}, e^{-\eta\lambda(1,\gamma)} \rangle$  is a straight line segment.
- ▶ The **only** satisfactory prediction is

$$\gamma = \frac{\sum_i w_i \gamma_i}{\sum_i w_i}$$

- ▶ We are back to the online Bayes algorithm.

## Vovk algorithm for square loss

- ▶ The square loss is mixable with  $\eta = 2$ .
- ▶ Prediction must satisfy

$$1 - \sqrt{-\frac{1}{2} \ln \sum_i V_i^t e^{-2(1-p_i^t)^2}} \leq p^t \leq \sqrt{-\frac{1}{2} \ln \sum_i V_i^t e^{-2(p_i^t)^2}}$$

where  $V_i^t = \frac{W_i^t}{\sum_s W_i^s}$ .



$$L_A \leq L_{\min} + \frac{1}{2} \ln N$$

## Simple prediction for square loss

- ▶ We can use the prediction

$$\gamma = \frac{\sum_i W_i \gamma_i}{\sum_i W_i}$$

- ▶ But in that case we must use  $\eta = 1/2$  when updating the weights.
- ▶ Which yields the bound

$$L_A \leq L_{\min} + 2 \ln N$$

## Summary of bounds for mixable losses

### TRACKING THE BEST EXPERT


Loss Functions:	$c$ values: ( $\eta = 1/c$ )	
	$\text{pred}_{\text{wmean}}(v, x)$	$\text{pred}_{\text{Vovk}}(v, x)$
$L_{\text{sq}}(p, q)$	2	$1/2$
$L_{\text{ent}}(p, q)$	1	1
$L_{\text{hel}}(p, q)$	1	$1/\sqrt{2}$

Figure 2.  $(c, 1/c)$ -realizability:  $c$  values for loss and prediction function pairings



## Switching experts setup

- ▶ **Usually:** compare algorithm's total loss to total loss of the best expert.
- ▶ **Switching experts:** compare algorithm's total loss to total loss of **best expert sequence** with  **$k$  switches**.
- ▶



## An inefficient algorithm

- ▶ Fix:
  - ▶  $l$  - sequence length
  - ▶  $k$  - number of switches
  - ▶  $n$  - number of experts
- ▶ Consider one **partition-expert** per sequence of switching experts.
- ▶ No. of **partition-experts** :  $\binom{l}{k-1} n(n-1)^k = O\left(n^{k+1} \left(\frac{el}{k}\right)^k\right)$
- ▶ The log-loss regret is at most  $(k+1) \log n + k \log \frac{l}{k} + k$
- ▶ Requires maintaining  $O\left(n^{k+1} \left(\frac{el}{k}\right)^k\right)$  weights.

## generalization to mixable losses

- ▶ In this lecture we assume loss function is **mixable**.
- ▶ There is an exponential weights algorithm with learning rate  $\eta$  that achieves (in the non-switching case) a bound

$$L_A \leq \min_i L_i + \frac{1}{\eta} \log n$$

- ▶ Then using the **partition-expert** algorithm for the switching-experts case we get a bound on the regret  $\frac{1}{\eta} ((k+1) \log n + k \log \frac{l}{k} + k)$

## Weight sharing algorithms

- Update weights in two stages: loss update then share update.
- Prediction uses the normalized  $s$  weights  $w_{t,i}^s / \sum_j w_{t,j}^s$
- **Loss update** is the same as always, but defines intermediate  $m$  weights:


$$w_{t,i}^m = w_{t,i}^s e^{-\eta L(y_t, x_{t,i})}$$

- **Share update**: redistribute the weights
- **Fixed-share**:

$$pool = \alpha \sum_{i=1}^n w_{t,i}^m$$

$$w_{t+1,i}^s = (1 - \alpha) w_{t,i}^m + \frac{1}{n - 1} (pool - \alpha w_{t,i}^m)$$

## The fixed-share algorithm



## Proving a bound on the fixed-share


- ▶ The relation between algorithm loss and total weight does not change because share update does not change the total weight.
- ▶ Thus we still have

$$L_A \leq \frac{1}{\eta} \sum_{i=1}^n w_{l+1,i}^s$$

- ▶ The harder question is how to lower bound  $\sum_{i=1}^n w_{l+1,i}^s$

## Lower bounding the final total weight

- Fix some switching experts sequence:



- “follow” the weight of the chosen expert  $i_t$ .
- The loss update reduces the weight by a factor of  $e^{-\eta \ell_t, i_t}$ .
- The share update reduces the weight by a factor larger than:
  - $1 - \alpha$  on iterations with no switch.
  - $\frac{\alpha}{n-1}$  on iterations where a switch occurs.

## Bound for arbitrary $\alpha$

- Combining we lower bound the final weight of the last expert in the sequence

$$w_{l+1, e_k}^s \geq \frac{1}{n} e^{-\eta L_*} (1 - \alpha)^{l-k-1} \left( \frac{\alpha}{n-1} \right)^k$$

Where  $L_*$  is the cumulative loss of the switching sequence of experts.

- Combining the upper and lower bounds we get that for any sequence

$$L_A \leq L_* + \frac{1}{\eta} \left( \ln n + (l - k - 1) \ln \frac{1}{1 - \alpha} + k \left( \ln \frac{1}{\alpha} + \ln(n - 1) \right) \right)$$



## Tuning $\alpha$

- ▶ let  $k^*$  be the best number of switches (in hind sight) and  $\alpha^* = k^*/I$
- ▶ Suppose we use  $\alpha \approx \alpha^*$  then the bound that we get is

$$L_A \leq L_* + \frac{1}{\eta} ((k+1) \ln n + (I-1)(H(\alpha^*) + D_{\text{KL}}(\alpha^* || \alpha)))$$

Where

$$H(\alpha^*) = -\alpha^* \ln \alpha^* - (1 - \alpha^*) \ln(1 - \alpha^*)$$

$$D_{\text{KL}}(\alpha^* || \alpha) = \alpha^* \ln \frac{\alpha^*}{\alpha} + (1 - \alpha^*) \ln \frac{1 - \alpha^*}{1 - \alpha}$$

- ▶ This is very close to the loss of the computationally inefficient algorithm.
- ▶ For the log loss case this is essentially optimal.
- ▶ Not so for square loss!

## What can we hope to improve?

- ▶ In the fixed-share algorithm, the weight of a suboptimal expert never decreases below  $\alpha/n$ .
- ▶ The algorithm does not concentrate only on the best expert, even if the last switch is in the distant past.
- ▶ The regret depends on the length of the sequence.

## The idea of variable-share

- ▶ Let the fraction of the total weight given to the best expert get arbitrarily close to **1**.
- ▶ we can get a regret bound that depends only on the number of switches, not on the length of the sequence.
- ▶ Requires that the loss be bounded.
- ▶ Works for **square** loss, but not for **log** loss!

## Variable-share

$$pool = \sum_{i=1}^n \left(1 - (1 - \alpha)^{\ell_{t,i}}\right) w_{t,i}^m$$

$$w_{t+1,i}^s = (1 - \alpha)^{\ell_{t,i}} w_{t,i}^m + \frac{1}{n-1} \left( pool - (1 - (1 - \alpha)^{\ell_{t,i}}) w_{t,i}^m \right)$$

If  $\ell_{t,i} = 0$ , then expert  $i$  does not contribute to the pool.  
Expert can get fraction of the total weight arbitrarily close to 1.  
Shares the weight quickly if  $\ell_{t,i} > 0$

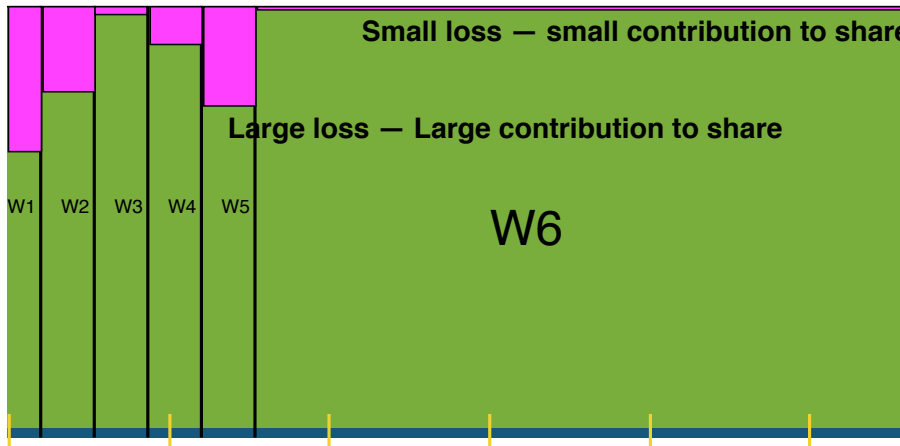
## Bound for variable share




$$\frac{1}{\eta} \ln n + \left(1 + \frac{1}{(1-\alpha)\eta}\right) L_* + k \left(1 + \frac{1}{\eta} \left(\ln n - 1 + \ln \frac{1}{\alpha} + \ln \frac{1}{1-\alpha}\right)\right)$$

- $\alpha$  should be tuned so that it is (close to)  $\frac{k}{2k+L_*}$

## Variable share figure



## An experiment using variable share



## Next Class

- ▶ Suppose the best switching sequence is repeatedly switching among a small subset of the experts  $n' \ll n$
- ▶ In the context of speech recognition - the speaker repeatedly uses a small number of phonemes.
- ▶ If we know the subset, we can pay  $\ln n'$  per switch rather than  $\ln n$
- ▶ Can track switches much more closely.
- ▶ Easy to describe an inefficient algorithm (consider all  $\binom{n}{n'}$  subsets.)
- ▶ Next class - how to do as well with just one weight per expert.