

Mixable losses and Tracking the best Expert

Yoav Freund

January 22, 2014

Outline

Review

Outline

Review

The general prediction game

Outline

- Review

- The general prediction game

- Some useful loss functions

Outline

- Review

- The general prediction game

- Some useful loss functions

- Vovk's algorithm

Outline

- Review

- The general prediction game

- Some useful loss functions

- Vovk's algorithm

- mixable loss functions

Outline

- Review

- The general prediction game

- Some useful loss functions

- Vovk's algorithm

- mixable loss functions

- The convexity condition

Outline

- Review

- The general prediction game

- Some useful loss functions

- Vovk's algorithm

- mixable loss functions

- The convexity condition

- Log loss

Outline

- Review

- The general prediction game

- Some useful loss functions

- Vovk's algorithm

- mixable loss functions

- The convexity condition

- Log loss

- Square loss

 - Square loss using simple averaging

Outline

Review

The general prediction game

Some useful loss functions

Vovk's algorithm

mixable loss functions

The convexity condition

Log loss

Square loss

Square loss using simple averaging

Summary table

Outline

Review

The general prediction game

Some useful loss functions

Vovk's algorithm

mixable loss functions

The convexity condition

Log loss

Square loss

Square loss using simple averaging

Summary table

Switching Experts

Outline

Review

The general prediction game

Some useful loss functions

Vovk's algorithm

mixable loss functions

The convexity condition

Log loss

Square loss

Square loss using simple averaging

Summary table

Switching Experts

An inefficient algorithm

Outline

Review

The general prediction game

Some useful loss functions

Vovk's algorithm

mixable loss functions

The convexity condition

Log loss

Square loss

Square loss using simple averaging

Summary table

Switching Experts

An inefficient algorithm

The fixed-share algorithm

Outline

Review

The general prediction game

Some useful loss functions

Vovk's algorithm

mixable loss functions

The convexity condition

Log loss

Square loss

Square loss using simple averaging

Summary table

Switching Experts

An inefficient algorithm

The fixed-share algorithm

The variable-share algorithm

The log-loss game

- ▶ Prediction algorithm A has access to N experts.

The log-loss game

- ▶ Prediction algorithm A has access to N experts.
- ▶ The following is repeated for $t = 1, \dots, T$

The log-loss game

- ▶ Prediction algorithm A has access to N experts.
- ▶ The following is repeated for $t = 1, \dots, T$
 - ▶ Experts generate predictive distributions: $\mathbf{p}_1^t, \dots, \mathbf{p}_N^t$

The log-loss game

- ▶ Prediction algorithm A has access to N experts.
- ▶ The following is repeated for $t = 1, \dots, T$
 - ▶ Experts generate predictive distributions: $\mathbf{p}_1^t, \dots, \mathbf{p}_N^t$
 - ▶ Algorithm generates its own prediction \mathbf{p}_A^t

The log-loss game

- ▶ Prediction algorithm A has access to N experts.
- ▶ The following is repeated for $t = 1, \dots, T$
 - ▶ Experts generate predictive distributions: $\mathbf{p}_1^t, \dots, \mathbf{p}_N^t$
 - ▶ Algorithm generates its own prediction \mathbf{p}_A^t
 - ▶ c^t is revealed.

The log-loss game

- ▶ Prediction algorithm A has access to N experts.
- ▶ The following is repeated for $t = 1, \dots, T$
 - ▶ Experts generate predictive distributions: $\mathbf{p}_1^t, \dots, \mathbf{p}_N^t$
 - ▶ Algorithm generates its own prediction \mathbf{p}_A^t
 - ▶ c^t is revealed.
- ▶ **Goal:** minimize regret:

$$-\sum_{t=1}^T \log p_A^t(c^t) + \min_{i=1, \dots, N} \left(-\sum_{t=1}^T \log p_i^t(c^t) \right)$$

The online Bayes Algorithm

- Total loss of expert i

$$L_i^t = - \sum_{s=1}^t \log p_i^s(c^s); \quad L_i^0 = 0$$

The online Bayes Algorithm

- Total loss of expert i

$$L_i^t = - \sum_{s=1}^t \log p_i^s(c^s); \quad L_i^0 = 0$$

- Weight of expert i

$$w_i^t = w_i^1 e^{-L_i^{t-1}} = w_i^1 \prod_{s=1}^{t-1} p_i^s(c^s)$$

The online Bayes Algorithm

- ▶ Total loss of expert i

$$L_i^t = - \sum_{s=1}^t \log p_i^s(c^s); \quad L_i^0 = 0$$

- ▶ Weight of expert i

$$w_i^t = w_i^1 e^{-L_i^{t-1}} = w_i^1 \prod_{s=1}^{t-1} p_i^s(c^s)$$

- ▶ Freedom to choose initial weights.

$$w_i^1 \geq 0, \sum_{i=1}^n w_i^1 = 1$$

The online Bayes Algorithm

- Total loss of expert i

$$L_i^t = - \sum_{s=1}^t \log p_i^s(c^s); \quad L_i^0 = 0$$

- Weight of expert i

$$w_i^t = w_i^1 e^{-L_i^{t-1}} = w_i^1 \prod_{s=1}^{t-1} p_i^s(c^s)$$

- Freedom to choose initial weights.

$$w_i^1 \geq 0, \sum_{i=1}^N w_i^1 = 1$$

- Prediction of algorithm A

$$\mathbf{p}_A^t = \frac{\sum_{i=1}^N w_i^t \mathbf{p}_i^t}{\sum_{i=1}^N w_i^t}$$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

$$\frac{W^{t+1}}{W^t} = \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t}$$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

$$\frac{W^{t+1}}{W^t} = \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t} = \frac{\sum_{i=1}^N w_i^t p_i^t(c^t)}{\sum_{i=1}^N w_i^t}$$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

$$\frac{W^{t+1}}{W^t} = \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t} = \frac{\sum_{i=1}^N w_i^t p_i^t(c^t)}{\sum_{i=1}^N w_i^t} = p_A^t(c^t)$$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

$$\begin{aligned}\frac{W^{t+1}}{W^t} &= \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t} = \frac{\sum_{i=1}^N w_i^t p_i^t(c^t)}{\sum_{i=1}^N w_i^t} = p_A^t(c^t) \\ -\log \frac{W^{t+1}}{W^t} &= -\log p_A^t(c^t)\end{aligned}$$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

$$\frac{W^{t+1}}{W^t} = \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t} = \frac{\sum_{i=1}^N w_i^t p_i^t(c^t)}{\sum_{i=1}^N w_i^t} = p_A^t(c^t)$$

$$-\log \frac{W^{t+1}}{W^t} = -\log p_A^t(c^t)$$

$$-\log \frac{W^{T+1}}{W^1} = -\sum_{t=1}^T \log p_A^t(c^t)$$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

$$\frac{W^{t+1}}{W^t} = \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t} = \frac{\sum_{i=1}^N w_i^t p_i^t(c^t)}{\sum_{i=1}^N w_i^t} = p_A^t(c^t)$$

$$-\log \frac{W^{t+1}}{W^t} = -\log p_A^t(c^t)$$

$$-\log \frac{W^{T+1}}{W^1} = -\sum_{t=1}^T \log p_A^t(c^t) = L_A^T$$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

$$\frac{W^{t+1}}{W^t} = \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t} = \frac{\sum_{i=1}^N w_i^t p_i^t(c^t)}{\sum_{i=1}^N w_i^t} = p_A^t(c^t)$$

$$-\log \frac{W^{t+1}}{W^t} = -\log p_A^t(c^t)$$

$$-\log W^{T+1} = -\log \frac{W^{T+1}}{W^1} = -\sum_{t=1}^T \log p_A^t(c^t) = L_A^T$$

Cumulative loss vs. Final total weight

Total weight: $W^t \doteq \sum_{i=1}^N w_i^t$

$$\frac{W^{t+1}}{W^t} = \frac{\sum_{i=1}^N w_i^t e^{\log p_i^t(c^t)}}{\sum_{i=1}^N w_i^t} = \frac{\sum_{i=1}^N w_i^t p_i^t(c^t)}{\sum_{i=1}^N w_i^t} = p_A^t(c^t)$$

$$-\log \frac{W^{t+1}}{W^t} = -\log p_A^t(c^t)$$

$$-\log W^{T+1} = -\log \frac{W^{T+1}}{W^1} = -\sum_{t=1}^T \log p_A^t(c^t) = L_A^T$$

EQUALITY not bound!

Vovk's general prediction game

Γ - prediction space. Ω - outcome space.

Vovk's general prediction game

Γ - prediction space. Ω - outcome space.

On each trial $t = 1, 2, \dots$

Vovk's general prediction game

Γ - prediction space. Ω - outcome space.

On each trial $t = 1, 2, \dots$

1. Each expert $i \in \{1 \dots N\}$ makes a prediction $\gamma_i^t \in \Gamma$

Vovk's general prediction game

Γ - prediction space. Ω - outcome space.

On each trial $t = 1, 2, \dots$

1. Each expert $i \in \{1 \dots N\}$ makes a prediction $\gamma_i^t \in \Gamma$
2. The learner, after observing $\langle \gamma_1^t \dots \gamma_N^t \rangle$, makes its own prediction γ^t

Vovk's general prediction game

Γ - prediction space. Ω - outcome space.

On each trial $t = 1, 2, \dots$

1. Each expert $i \in \{1 \dots N\}$ makes a prediction $\gamma_i^t \in \Gamma$
2. The learner, after observing $\langle \gamma_1^t \dots \gamma_N^t \rangle$,
makes its own prediction γ^t
3. Nature chooses an outcome $\omega^t \in \Omega$

Vovk's general prediction game

Γ - prediction space. Ω - outcome space.

On each trial $t = 1, 2, \dots$

1. Each expert $i \in \{1 \dots N\}$ makes a prediction $\gamma_i^t \in \Gamma$
2. The learner, after observing $\langle \gamma_1^t \dots \gamma_N^t \rangle$,
makes its own prediction γ^t
3. Nature chooses an outcome $\omega^t \in \Omega$
4. Each expert incurs loss $\ell_i^t = \lambda(\omega^t, \gamma_i^t)$
The learner incurs loss $\ell_A^t = \lambda(\omega^t, \gamma^t)$

Achievable loss bounds

► $L_A \doteq \sum_{t=1}^T \ell_A^t$ - total loss of algorithm

Achievable loss bounds

- ▶ $L_A \doteq \sum_{t=1}^T \ell_A^t$ - total loss of algorithm
- ▶ $L_i \doteq \sum_{t=1}^T \ell_i^t$ - total loss of expert i

Achievable loss bounds

- ▶ $L_A \doteq \sum_{t=1}^T \ell_A^t$ - total loss of algorithm
- ▶ $L_i \doteq \sum_{t=1}^T \ell_i^t$ - total loss of expert i
- ▶ **Goal:** find an algorithm which guarantees that

$$(a, c) \in [0, \infty), \quad L_A \leq aL_{\min} + c \ln N$$

For any sequence of events.

Achievable loss bounds

- ▶ $L_A \doteq \sum_{t=1}^T \ell_A^t$ - total loss of algorithm
- ▶ $L_i \doteq \sum_{t=1}^T \ell_i^t$ - total loss of expert i
- ▶ **Goal:** find an algorithm which guarantees that

$$(a, c) \in [0, \infty), \quad L_A \leq aL_{\min} + c \ln N$$

For any sequence of events.

- ▶ We say that the pair (a, c) is **achievable**.

The set of achievable bounds

- Fix loss function $\lambda : \Omega \times \Gamma \rightarrow [0, \infty)$

The set of achievable bounds


- ▶ Fix loss function $\lambda : \Omega \times \Gamma \rightarrow [0, \infty)$
- ▶ The pair (a, c) is *achievable* if there exists *some* prediction algorithm such that for *any* $N > 0$, *any* set of N prediction sequences and *any* sequence of outcomes

$$L_A \leq aL_{\min} + c \ln N$$

The set of achievable bounds

- ▶ Fix loss function $\lambda : \Omega \times \Gamma \rightarrow [0, \infty)$
- ▶ The pair (a, c) is *achievable* if there exists *some* prediction algorithm such that for *any* $N > 0$, *any* set of N prediction sequences and *any* sequence of outcomes

$$L_A \leq aL_{\min} + c \ln N$$



Some useful loss functions

- Outcomes: $\omega^1, \omega_2, \dots \omega^t \in [0, 1]$

Some useful loss functions

- ▶ Outcomes: $\omega^1, \omega_2, \dots \omega^t \in [0, 1]$
- ▶ Predictions: $\gamma^1, \gamma^2, \dots \gamma^t \in [0, 1]$

Log loss (Entropy loss)



$$\lambda_{\text{ent}}(\omega, \gamma) = \omega \ln \frac{\omega}{\gamma} + (1 - \omega) \ln \frac{1 - \omega}{1 - \gamma}$$

Log loss (Entropy loss)



$$\lambda_{\text{ent}}(\omega, \gamma) = \omega \ln \frac{\omega}{\gamma} + (1 - \omega) \ln \frac{1 - \omega}{1 - \gamma}$$

- When $q_t \in \{0, 1\}$ Cumulative log loss = coding length ± 1

Log loss (Entropy loss)



$$\lambda_{\text{ent}}(\omega, \gamma) = \omega \ln \frac{\omega}{\gamma} + (1 - \omega) \ln \frac{1 - \omega}{1 - \gamma}$$

- ▶ When $q_t \in \{0, 1\}$ Cumulative log loss = coding length ± 1
- ▶ If $P[\omega_t = 1] = q$, optimal prediction $\gamma^t = q$

Log loss (Entropy loss)



$$\lambda_{\text{ent}}(\omega, \gamma) = \omega \ln \frac{\omega}{\gamma} + (1 - \omega) \ln \frac{1 - \omega}{1 - \gamma}$$

- ▶ When $q_t \in \{0, 1\}$ Cumulative log loss = coding length ± 1
- ▶ If $P[\omega_t = 1] = q$, optimal prediction $\gamma^t = q$
- ▶ Unbounded loss.

Log loss (Entropy loss)



$$\lambda_{\text{ent}}(\omega, \gamma) = \omega \ln \frac{\omega}{\gamma} + (1 - \omega) \ln \frac{1 - \omega}{1 - \gamma}$$

- ▶ When $q_t \in \{0, 1\}$ Cumulative log loss = coding length ± 1
- ▶ If $P[\omega_t = 1] = q$, optimal prediction $\gamma^t = q$
- ▶ Unbounded loss.
- ▶ Not symmetric $\exists p, q \lambda(p, q) \neq \lambda(q, p)$.

Log loss (Entropy loss)



$$\lambda_{\text{ent}}(\omega, \gamma) = \omega \ln \frac{\omega}{\gamma} + (1 - \omega) \ln \frac{1 - \omega}{1 - \gamma}$$

- ▶ When $q_t \in \{0, 1\}$ Cumulative log loss = coding length ± 1
- ▶ If $P[\omega_t = 1] = q$, optimal prediction $\gamma^t = q$
- ▶ Unbounded loss.
- ▶ Not symmetric $\exists p, q \lambda(p, q) \neq \lambda(q, p)$.
- ▶ No triangle inequality
 $\exists p_1, p_2, p_3 \lambda(p_1, p_3) > \lambda(p_1, p_2) + \lambda(p_2, p_3)$

Square loss (Breier Loss)



$$\lambda_{\text{sq}}(\omega, \gamma) = (\omega - \gamma)^2$$

Square loss (Breier Loss)



$$\lambda_{\text{sq}}(\omega, \gamma) = (\omega - \gamma)^2$$

- ▶ $P[\omega^t = 1] = q, P[\omega^t = 0] = 1 - q,$
optimal prediction $\gamma^t = q$

Square loss (Breier Loss)



$$\lambda_{\text{sq}}(\omega, \gamma) = (\omega - \gamma)^2$$

- ▶ $P[\omega^t = 1] = q, P[\omega^t = 0] = 1 - q,$
optimal prediction $\gamma^t = q$
- ▶ Bounded loss.

Square loss (Breier Loss)



$$\lambda_{\text{sq}}(\omega, \gamma) = (\omega - \gamma)^2$$

- ▶ $P[\omega^t = 1] = q, P[\omega^t = 0] = 1 - q,$
optimal prediction $\gamma^t = q$
- ▶ Bounded loss.
- ▶ Defines a metric (symmetric and triangle ineq.)

Square loss (Breier Loss)



$$\lambda_{\text{sq}}(\omega, \gamma) = (\omega - \gamma)^2$$

- ▶ $P[\omega^t = 1] = q, P[\omega^t = 0] = 1 - q,$
optimal prediction $\gamma^t = q$
- ▶ Bounded loss.
- ▶ Defines a metric (symmetric and triangle ineq.)
- ▶ Corresponds to regression.

Hellinger Loss



$$\lambda_{\text{hel}}(\omega, \gamma) = \frac{1}{2} \left((\sqrt{\omega} + \sqrt{\gamma})^2 + (\sqrt{1-\omega} + \sqrt{1-\gamma})^2 \right)$$

Hellinger Loss



$$\lambda_{\text{hel}}(\omega, \gamma) = \frac{1}{2} \left((\sqrt{\omega} + \sqrt{\gamma})^2 + (\sqrt{1-\omega} + \sqrt{1-\gamma})^2 \right)$$

- ▶ If $P[\omega^t = 1] = q$, $P[\omega^t = 0] = 1 - q$,
optimal prediction $\gamma^t = q$

Hellinger Loss



$$\lambda_{\text{hel}}(\omega, \gamma) = \frac{1}{2} \left((\sqrt{\omega} + \sqrt{\gamma})^2 + (\sqrt{1-\omega} + \sqrt{1-\gamma})^2 \right)$$

- ▶ If $P[\omega^t = 1] = q$, $P[\omega^t = 0] = 1 - q$,
optimal prediction $\gamma^t = q$
- ▶ Loss is bounded.

Hellinger Loss



$$\lambda_{\text{hel}}(\omega, \gamma) = \frac{1}{2} \left((\sqrt{\omega} + \sqrt{\gamma})^2 + (\sqrt{1-\omega} + \sqrt{1-\gamma})^2 \right)$$

- ▶ If $P[\omega^t = 1] = q$, $P[\omega^t = 0] = 1 - q$,
optimal prediction $\gamma^t = q$
- ▶ Loss is bounded.
- ▶ Defines a metric.

Hellinger Loss



$$\lambda_{\text{hel}}(\omega, \gamma) = \frac{1}{2} \left((\sqrt{\omega} + \sqrt{\gamma})^2 + (\sqrt{1-\omega} + \sqrt{1-\gamma})^2 \right)$$

- ▶ If $P[\omega^t = 1] = q$, $P[\omega^t = 0] = 1 - q$,
optimal prediction $\gamma^t = q$
- ▶ Loss is bounded.
- ▶ Defines a metric.
- ▶ $\lambda_{\text{hel}}(p, q) \approx \lambda_{\text{ent}}(p, q)$ when $p \approx q$ and $p, q \in (0, 1)$

Absolute loss



$$\lambda(\omega, \gamma) = |\omega - \gamma|$$

Absolute loss



$$\lambda(\omega, \gamma) = |\omega - \gamma|$$

- ▶ Probability of making a mistake if predicting 0 or 1 using a biased coin

Absolute loss



$$\lambda(\omega, \gamma) = |\omega - \gamma|$$

- ▶ Probability of making a mistake if predicting 0 or 1 using a biased coin
- ▶ If $P[\omega^t = 1] = q$, $P[\omega^t = 0] = 1 - q$, then the optimal prediction is

$$\gamma^t = \begin{cases} 1 & \text{if } q > 1/2, \\ 0 & \text{otherwise} \end{cases}$$

Structureless bounded loss

- Prediction is a distribution $\gamma = \langle p_1, \dots, p_N \rangle$, $p_i \geq 0$,
 $\sum_{i=1}^N p_i = 1$

Structureless bounded loss

- ▶ Prediction is a distribution $\gamma = \langle p_1, \dots, p_N \rangle$, $p_i \geq 0$,
 $\sum_{i=1}^N p_i = 1$
- ▶ Outcome is a loss vector $\omega = \langle \omega_1, \dots, \omega_N \rangle$, $0 \leq \omega_i \leq 1$

Structureless bounded loss

- ▶ Prediction is a distribution $\gamma = \langle p_1, \dots, p_N \rangle$, $p_i \geq 0$,
 $\sum_{i=1}^N p_i = 1$
- ▶ Outcome is a loss vector $\omega = \langle \omega_1, \dots, \omega_N \rangle$, $0 \leq \omega_i \leq 1$
- ▶ Loss is the dot product: $\lambda_{\text{dot}}(\omega, \gamma) = \gamma \cdot \omega$

Structureless bounded loss

- ▶ Prediction is a distribution $\gamma = \langle p_1, \dots, p_N \rangle$, $p_i \geq 0$,
 $\sum_{i=1}^N p_i = 1$
- ▶ Outcome is a loss vector $\omega = \langle \omega_1, \dots, \omega_N \rangle$, $0 \leq \omega_i \leq 1$
- ▶ Loss is the dot product: $\lambda_{\text{dot}}(\omega, \gamma) = \gamma \cdot \omega$
- ▶ Corresponds to the hedging game.

Structureless bounded loss

- ▶ Prediction is a distribution $\gamma = \langle p_1, \dots, p_N \rangle$, $p_i \geq 0$,
 $\sum_{i=1}^N p_i = 1$
- ▶ Outcome is a loss vector $\omega = \langle \omega_1, \dots, \omega_N \rangle$, $0 \leq \omega_i \leq 1$
- ▶ Loss is the dot product: $\lambda_{\text{dot}}(\omega, \gamma) = \gamma \cdot \omega$
- ▶ Corresponds to the hedging game.
- ▶ For hedge loss the regret is $\Omega(\sqrt{T \log N})$.

Structureless bounded loss

- ▶ Prediction is a distribution $\gamma = \langle p_1, \dots, p_N \rangle$, $p_i \geq 0$,
 $\sum_{i=1}^N p_i = 1$
- ▶ Outcome is a loss vector $\omega = \langle \omega_1, \dots, \omega_N \rangle$, $0 \leq \omega_i \leq 1$
- ▶ Loss is the dot product: $\lambda_{\text{dot}}(\omega, \gamma) = \gamma \cdot \omega$
- ▶ Corresponds to the hedging game.
- ▶ For hedge loss the regret is $\Omega(\sqrt{T \log N})$.
- ▶ For the log loss the regret is $O(\log N)$

Structureless bounded loss

- ▶ Prediction is a distribution $\gamma = \langle p_1, \dots, p_N \rangle$, $p_i \geq 0$,
 $\sum_{i=1}^N p_i = 1$
- ▶ Outcome is a loss vector $\omega = \langle \omega_1, \dots, \omega_N \rangle$, $0 \leq \omega_i \leq 1$
- ▶ Loss is the dot product: $\lambda_{\text{dot}}(\omega, \gamma) = \gamma \cdot \omega$
- ▶ Corresponds to the hedging game.
- ▶ For hedge loss the regret is $\Omega(\sqrt{T \log N})$.
- ▶ For the log loss the regret is $O(\log N)$
- ▶ Which losses behave like **entropy loss** and which behave like **hedge loss**?

Some technical requirements

- ▶ There should be a **topology** on the prediction set \mathcal{I} such that

Some technical requirements

- ▶ There should be a **topology** on the prediction set Γ such that
- ▶ Γ is compact.

Some technical requirements

- ▶ There should be a **topology** on the prediction set Γ such that
- ▶ Γ is compact.
- ▶ $\forall \omega \in \Omega$, the function $\gamma \rightarrow \lambda(\omega, \gamma)$ is **continuous**

Some technical requirements

- ▶ There should be a **topology** on the prediction set Γ such that
- ▶ Γ is compact.
- ▶ $\forall \omega \in \Omega$, the function $\gamma \rightarrow \lambda(\omega, \gamma)$ is **continuous**
- ▶ There is a **universally reasonable prediction**
 $\exists \gamma \in \Gamma, \forall \omega \in \Omega, \lambda(\omega, \gamma) < \infty$

Some technical requirements

- ▶ There should be a **topology** on the prediction set Γ such that
- ▶ Γ is compact.
- ▶ $\forall \omega \in \Omega$, the function $\gamma \rightarrow \lambda(\omega, \gamma)$ is **continuous**
- ▶ There is a **universally reasonable prediction**
 $\exists \gamma \in \Gamma, \forall \omega \in \Omega, \lambda(\omega, \gamma) < \infty$
- ▶ There is **no universally optimal prediction**
 $\neg \exists \gamma \in \Gamma, \forall \omega \in \Omega, \lambda(\omega, \gamma) = 0$

Vovk's meta-algorithm

- Fix an **achievable** pair (a, c) and set $\eta = a/c$

Vovk's meta-algorithm

- Fix an **achievable** pair (a, c) and set $\eta = a/c$

Vovk's meta-algorithm

- ▶ Fix an **achievable** pair (a, c) and set $\eta = a/c$
- ▶ 1.

$$w_i^t = \frac{1}{N} e^{-\eta L_i^t}$$

Vovk's meta-algorithm

- ▶ Fix an **achievable** pair (a, c) and set $\eta = a/c$
- ▶ 1.

$$w_i^t = \frac{1}{N} e^{-\eta L_i^t}$$

Choose γ_t so that, for all $\omega^t \in \Omega$:

$$\lambda(\omega^t, \gamma^t) - c \ln \sum_i w_i^t \leq -c \ln \left(\sum_i w_i^t e^{-\eta \lambda(\omega^t, \gamma_i^t)} \right)$$

Vovk's meta-algorithm

- ▶ Fix an **achievable** pair (a, c) and set $\eta = a/c$
- ▶ 1.

$$W_i^t = \frac{1}{N} e^{-\eta L_i^t}$$

Choose γ_t so that, for all $\omega^t \in \Omega$:

$$\lambda(\omega^t, \gamma^t) - c \ln \sum_i W_i^t \leq -c \ln \left(\sum_i W_i^t e^{-\eta \lambda(\omega^t, \gamma_i^t)} \right)$$

- 2. If choice of γ^t always exists, then the total loss satisfies:

$$\sum_t \lambda(\omega^t, \gamma^t) \leq -c \ln \sum_i W_i^{T+1} \leq a L_{\min} + c \ln N$$

Vovk's meta-algorithm

- Fix an **achievable** pair (a, c) and set $\eta = a/c$
- 1.

$$W_i^t = \frac{1}{N} e^{-\eta L_i^t}$$

Choose γ_t so that, for all $\omega^t \in \Omega$:

$$\lambda(\omega^t, \gamma^t) - c \ln \sum_i W_i^t \leq -c \ln \left(\sum_i W_i^t e^{-\eta \lambda(\omega^t, \gamma_i^t)} \right)$$

- 2. If choice of γ^t always exists, then the total loss satisfies:

$$\sum_t \lambda(\omega^t, \gamma^t) \leq -c \ln \sum_i W_i^{T+1} \leq a L_{\min} + c \ln N$$

- Vovk's result: **yes!** a good choice for γ_t always exists!

Vovk's algorithm is the the highest achiever [Vovk95]

The pair (a, c) is achieved by some algorithm if and only if it is achieved by Vovk's algorithm.

Vovk's algorithm is the the highest achiever [Vovk95]


The pair (a, c) is achieved by some algorithm if and only if it is achieved by Vovk's algorithm.

The separation curve is $\left\{ \left(a(\eta), \frac{a(\eta)}{\eta} \right) \mid \eta \in [0, \infty] \right\}$

Vovk's algorithm is the the highest achiever [Vovk95]

The pair (a, c) is achieved by some algorithm if and only if it is achieved by Vovk's algorithm.

The separation curve is $\left\{ \left(a(\eta), \frac{a(\eta)}{\eta} \right) \mid \eta \in [0, \infty] \right\}$



Mixable Loss Functions

- ▶ A Loss function is **mixable** if a pair of the form $(1, c)$, $c < \infty$ is achievable.

$$L_A \leq L_{\min} + c \ln N$$

Mixable Loss Functions

- ▶ A Loss function is **mixable** if a pair of the form $(1, c)$, $c < \infty$ is achievable.

$$L_A \leq L_{\min} + c \ln N$$

- ▶ Vovk's algorithm with $\eta = 1/c$ achieves this bound.

Mixable Loss Functions

- ▶ A Loss function is **mixable** if a pair of the form $(1, c)$, $c < \infty$ is achievable.

$$L_A \leq L_{\min} + c \ln N$$

- ▶ Vovk's algorithm with $\eta = 1/c$ achieves this bound.
- ▶ $\lambda_{\text{ent}}, \lambda_{\text{sq}}, \lambda_{\text{hel}}$ are **mixable**

Mixable Loss Functions

- ▶ A Loss function is **mixable** if a pair of the form $(1, c)$, $c < \infty$ is achievable.

$$L_A \leq L_{\min} + c \ln N$$

- ▶ Vovk's algorithm with $\eta = 1/c$ achieves this bound.
- ▶ $\lambda_{\text{ent}}, \lambda_{\text{sq}}, \lambda_{\text{hel}}$ are **mixable**
- ▶ $\lambda_{\text{abs}}, \lambda_{\text{dot}}$ are **not mixable**

The convexity condition

- requirement for loss to be $(1, 1/\eta)$ mixable

The convexity condition

- ▶ requirement for loss to be $(1, 1/\eta)$ mixable
- ▶ $\forall \langle (\gamma_1, W_1), \dots, (\gamma_N, W_N) \rangle$
 $\exists \gamma \in \Gamma$
 $\forall \omega \in \Omega:$

$$\lambda(\omega, \gamma) - \frac{1}{\eta} \ln \sum_i W_i \leq -\frac{1}{\eta} \ln \left(\sum_i W_i e^{-\eta \lambda(\omega, \gamma_i)} \right)$$

The convexity condition

- ▶ requirement for loss to be $(1, 1/\eta)$ mixable
- ▶ $\forall \langle (\gamma_1, W_1), \dots, (\gamma_N, W_N) \rangle$
 $\exists \gamma \in \Gamma$
 $\forall \omega \in \Omega:$

$$\lambda(\omega, \gamma) - \frac{1}{\eta} \ln \sum_i W_i \leq -\frac{1}{\eta} \ln \left(\sum_i W_i e^{-\eta \lambda(\omega, \gamma_i)} \right)$$

- ▶ Can be re-written as:

$$e^{-\eta \lambda(\omega, \gamma)} \geq \sum_i \left(\frac{W_i}{\sum_j W_j} \right) e^{-\eta \lambda(\omega, \gamma_i)}$$

The convexity condition

- ▶ requirement for loss to be $(1, 1/\eta)$ mixable
- ▶ $\forall \langle (\gamma_1, W_1), \dots, (\gamma_N, W_N) \rangle$
 $\exists \gamma \in \Gamma$
 $\forall \omega \in \Omega$:

$$\lambda(\omega, \gamma) - \frac{1}{\eta} \ln \sum_i W_i \leq -\frac{1}{\eta} \ln \left(\sum_i W_i e^{-\eta \lambda(\omega, \gamma_i)} \right)$$

- ▶ Can be re-written as:

$$e^{-\eta \lambda(\omega, \gamma)} \geq \sum_i \left(\frac{W_i}{\sum_j W_j} \right) e^{-\eta \lambda(\omega, \gamma_i)}$$

- ▶ Equivalently - the image of the set Γ under the mapping $F(\gamma) = \langle e^{-\eta \lambda(\omega, \gamma)} \rangle_{\omega \in \Omega}$ is concave.

convexity condition: Pictorially


- **Example:** Suppose $\Omega = \{0, 1\}$, $\Gamma = [0, 1]$. then

$$F(\gamma) = \left\langle e^{-\eta\lambda(0,\gamma)}, e^{-\eta\lambda(1,\gamma)} \right\rangle$$

convexity condition: Pictorially

- **Example:** Suppose $\Omega = \{0, 1\}$, $\Gamma = [0, 1]$. then

$$F(\gamma) = \left\langle e^{-\eta\lambda(0,\gamma)}, e^{-\eta\lambda(1,\gamma)} \right\rangle$$



Vovk Algorithm for log loss

- ▶ The log loss is mixable with $\eta = 1$

Vovk Algorithm for log loss

- ▶ The log loss is mixable with $\eta = 1$
- ▶ The image of $[0, 1]$ through $F(\gamma) = \langle e^{-\eta\lambda(0,\gamma)}, e^{-\eta\lambda(1,\gamma)} \rangle$ is a straight line segment.

Vovk Algorithm for log loss

- ▶ The log loss is mixable with $\eta = 1$
- ▶ The image of $[0, 1]$ through $F(\gamma) = \langle e^{-\eta\lambda(0,\gamma)}, e^{-\eta\lambda(1,\gamma)} \rangle$ is a straight line segment.
- ▶ The **only** satisfactory prediction is

$$\gamma = \frac{\sum_i w_i \gamma_i}{\sum_i w_i}$$

Vovk Algorithm for log loss

- ▶ The log loss is mixable with $\eta = 1$
- ▶ The image of $[0, 1]$ through $F(\gamma) = \langle e^{-\eta\lambda(0,\gamma)}, e^{-\eta\lambda(1,\gamma)} \rangle$ is a straight line segment.
- ▶ The **only** satisfactory prediction is

$$\gamma = \frac{\sum_i w_i \gamma_i}{\sum_i w_i}$$

- ▶ We are back to the online Bayes algorithm.

Vovk algorithm for square loss

- ▶ The square loss is mixable with $\eta = 2$.

Vovk algorithm for square loss

- ▶ The square loss is mixable with $\eta = 2$.
- ▶ Prediction must satisfy

$$1 - \sqrt{-\frac{1}{2} \ln \sum_i V_i^t e^{-2(1-p_i^t)^2}} \leq p^t \leq \sqrt{-\frac{1}{2} \ln \sum_i V_i^t e^{-2(p_i^t)^2}}$$

where $V_i^t = \frac{W_i^t}{\sum_s W_i^s}$.

Vovk algorithm for square loss

- ▶ The square loss is mixable with $\eta = 2$.
- ▶ Prediction must satisfy

$$1 - \sqrt{-\frac{1}{2} \ln \sum_i V_i^t e^{-2(1-p_i^t)^2}} \leq p^t \leq \sqrt{-\frac{1}{2} \ln \sum_i V_i^t e^{-2(p_i^t)^2}}$$

where $V_i^t = \frac{W_i^t}{\sum_s W_i^s}$.



$$L_A \leq L_{\min} + \frac{1}{2} \ln N$$

└ Square loss

└ Square loss using simple averaging

Simple prediction for square loss

- We can use the prediction

$$\gamma = \frac{\sum_i w_i \gamma_i}{\sum_i w_i}$$

Simple prediction for square loss

- We can use the prediction

$$\gamma = \frac{\sum_i w_i \gamma_i}{\sum_i w_i}$$

- But in that case we must use $\eta = 1/2$ when updating the weights.

Simple prediction for square loss

- ▶ We can use the prediction

$$\gamma = \frac{\sum_i W_i \gamma_i}{\sum_i W_i}$$

- ▶ But in that case we must use $\eta = 1/2$ when updating the weights.
- ▶ Which yields the bound

$$L_A \leq L_{\min} + 2 \ln N$$

Summary of bounds for mixable losses

TRACKING THE BEST EXPERT

Loss Functions:	c values: ($\eta = 1/c$)	
	$\text{pred}_{\text{wmean}}(v, x)$	$\text{pred}_{\text{Vovk}}(v, x)$
$L_{\text{sq}}(p, q)$	2	$1/2$
$L_{\text{ent}}(p, q)$	1	1
$L_{\text{hel}}(p, q)$	1	$1/\sqrt{2}$

Figure 2. $(c, 1/c)$ -realizability: c values for loss and prediction function pairings

Switching experts setup


- ▶ **Usually:** compare algorithm's total loss to total loss of the best expert.

Switching experts setup

- ▶ **Usually:** compare algorithm's total loss to total loss of the best expert.
- ▶ **Switching experts:** compare algorithm's total loss to total loss of **best expert sequence** with **k switches**.

Switching experts setup

- ▶ **Usually:** compare algorithm's total loss to total loss of the best expert.
- ▶ **Switching experts:** compare algorithm's total loss to total loss of **best expert sequence** with k switches.
- ▶



An inefficient algorithm

► Fix:

An inefficient algorithm

- ▶ Fix:
 - ▶ L - sequence length

An inefficient algorithm

- ▶ Fix:
 - ▶ l - sequence length
 - ▶ k - number of switches

An inefficient algorithm

- ▶ Fix:
 - ▶ l - sequence length
 - ▶ k - number of switches
 - ▶ n - number of experts

An inefficient algorithm

- ▶ Fix:
 - ▶ l - sequence length
 - ▶ k - number of switches
 - ▶ n - number of experts
- ▶ Consider one **partition-expert** per sequence of switching experts.

An inefficient algorithm

- ▶ Fix:
 - ▶ l - sequence length
 - ▶ k - number of switches
 - ▶ n - number of experts
- ▶ Consider one **partition-expert** per sequence of switching experts.
- ▶ No. of **partition-experts** : $\binom{l}{k-1} n(n-1)^k = O\left(n^{k+1} \left(\frac{el}{k}\right)^k\right)$

An inefficient algorithm

- ▶ Fix:
 - ▶ l - sequence length
 - ▶ k - number of switches
 - ▶ n - number of experts
- ▶ Consider one **partition-expert** per sequence of switching experts.
- ▶ No. of **partition-experts** : $\binom{l}{k-1} n(n-1)^k = O\left(n^{k+1} \left(\frac{el}{k}\right)^k\right)$
- ▶ The log-loss regret is at most $(k+1) \log n + k \log \frac{l}{k} + k$

An inefficient algorithm

- ▶ Fix:
 - ▶ l - sequence length
 - ▶ k - number of switches
 - ▶ n - number of experts
- ▶ Consider one **partition-expert** per sequence of switching experts.
- ▶ No. of **partition-experts** : $\binom{l}{k-1} n(n-1)^k = O\left(n^{k+1} \left(\frac{el}{k}\right)^k\right)$
- ▶ The log-loss regret is at most $(k+1) \log n + k \log \frac{l}{k} + k$
- ▶ Requires maintaining $O\left(n^{k+1} \left(\frac{el}{k}\right)^k\right)$ weights.

generalization to mixable losses

- In this lecture we assume loss function is **mixable**.

generalization to mixable losses

- ▶ In this lecture we assume loss function is **mixable**.
- ▶ There is an exponential weights algorithm with learning rate η that achieves (in the non-switching case) a bound

$$L_A \leq \min_i L_i + \frac{1}{\eta} \log n$$

generalization to mixable losses

- ▶ In this lecture we assume loss function is **mixable**.
- ▶ There is an exponential weights algorithm with learning rate η that achieves (in the non-switching case) a bound

$$L_A \leq \min_i L_i + \frac{1}{\eta} \log n$$

- ▶ Then using the **partition-expert** algorithm for the switching-experts case we get a bound on the regret $\frac{1}{\eta} ((k+1) \log n + k \log \frac{l}{k} + k)$

Weight sharing algorithms

- Update weights in two stages: loss update then share update.

Weight sharing algorithms

- ▶ Update weights in two stages: loss update then share update.
- ▶ Prediction uses the normalized s weights $w_{t,i}^s / \sum_j w_{t,j}^s$

Weight sharing algorithms

- ▶ Update weights in two stages: loss update then share update.
- ▶ Prediction uses the normalized **s** weights $w_{t,i}^s / \sum_j w_{t,j}^s$
- ▶ **Loss update** is the same as always, but defines intermediate **m** weights:

$$w_{t,i}^m = w_{t,i}^s e^{-\eta L(y_t, x_{t,i})}$$

Weight sharing algorithms

- ▶ Update weights in two stages: loss update then share update.
- ▶ Prediction uses the normalized s weights $w_{t,i}^s / \sum_j w_{t,j}^s$
- ▶ **Loss update** is the same as always, but defines intermediate m weights:

$$w_{t,i}^m = w_{t,i}^s e^{-\eta L(y_t, x_{t,i})}$$

- ▶ **Share update**: redistribute the weights

Weight sharing algorithms

- Update weights in two stages: loss update then share update.
- Prediction uses the normalized s weights $w_{t,i}^s / \sum_j w_{t,j}^s$
- **Loss update** is the same as always, but defines intermediate m weights:


$$w_{t,i}^m = w_{t,i}^s e^{-\eta L(y_t, x_{t,i})}$$

- **Share update**: redistribute the weights
- **Fixed-share**:


$$pool = \alpha \sum_{i=1}^n w_{t,i}^m$$

$$w_{t+1,i}^s = (1 - \alpha) w_{t,i}^m + \frac{1}{n-1} (pool - \alpha w_{t,i}^m)$$


The fixed-share algorithm



The fixed-share algorithm



The fixed-share algorithm



Proving a bound on the fixed-share

- ▶ The relation between algorithm loss and total weight does not change because share update does not change the total weight.

Proving a bound on the fixed-share

- ▶ The relation between algorithm loss and total weight does not change because share update does not change the total weight.
- ▶ Thus we still have

$$L_A \leq \frac{1}{\eta} \sum_{i=1}^n w_{l+1,i}^s$$

Proving a bound on the fixed-share


- ▶ The relation between algorithm loss and total weight does not change because share update does not change the total weight.
- ▶ Thus we still have

$$L_A \leq \frac{1}{\eta} \sum_{i=1}^n w_{l+1,i}^s$$

- ▶ The harder question is how to lower bound $\sum_{i=1}^n w_{l+1,i}^s$


Lower bounding the final total weight

- Fix some switching experts sequence:



Lower bounding the final total weight


- Fix some switching experts sequence:



- “follow” the weight of the chosen expert i_t .

Lower bounding the final total weight


- Fix some switching experts sequence:



- “follow” the weight of the chosen expert i_t .
- The loss update reduces the weight by a factor of $e^{-\eta \ell_{t,i_t}}$.

Lower bounding the final total weight


- Fix some switching experts sequence:



- “follow” the weight of the chosen expert i_t .
- The loss update reduces the weight by a factor of $e^{-\eta \ell_{t,i_t}}$.
- The share update reduces the weight by a factor larger than:

Lower bounding the final total weight


- Fix some switching experts sequence:



- “follow” the weight of the chosen expert i_t .
- The loss update reduces the weight by a factor of $e^{-\eta \ell_{t,i_t}}$.
- The share update reduces the weight by a factor larger than:
 - $1 - \alpha$ on iterations with no switch.

Lower bounding the final total weight

- Fix some switching experts sequence:



- “follow” the weight of the chosen expert i_t .
- The loss update reduces the weight by a factor of $e^{-\eta \ell_{t,i_t}}$.
- The share update reduces the weight by a factor larger than:
 - $1 - \alpha$ on iterations with no switch.
 - $\frac{\alpha}{n-1}$ on iterations where a switch occurs.

Bound for arbitrary α

- Combining we lower bound the final weight of the last expert in the sequence

$$w_{l+1, e_k}^s \geq \frac{1}{n} e^{-\eta L_*} (1 - \alpha)^{l-k-1} \left(\frac{\alpha}{n-1} \right)^k$$

Where L_* is the cumulative loss of the switching sequence of experts.

Bound for arbitrary α

- Combining we lower bound the final weight of the last expert in the sequence

$$w_{l+1,e_k}^s \geq \frac{1}{n} e^{-\eta L_*} (1 - \alpha)^{l-k-1} \left(\frac{\alpha}{n-1} \right)^k$$

Where L_* is the cumulative loss of the switching sequence of experts.

- Combining the upper and lower bounds we get that for any sequence

$$L_A \leq L_* + \frac{1}{\eta} \left(\ln n + (l - k - 1) \ln \frac{1}{1 - \alpha} + k \left(\ln \frac{1}{\alpha} + \ln(n - 1) \right) \right)$$

Tuning α

- ▶ let k^* be the best number of switches (in hind sight) and
 $\alpha^* = k^*/I$

Tuning α

- ▶ let k^* be the best number of switches (in hind sight) and $\alpha^* = k^*/I$
- ▶ Suppose we use $\alpha \approx \alpha^*$ then the bound that we get is

$$L_A \leq L_* + \frac{1}{\eta} ((k+1) \ln n + (I-1)(H(\alpha^*) + D_{\text{KL}}(\alpha^* || \alpha)))$$

Where

$$H(\alpha^*) = -\alpha^* \ln \alpha^* - (1 - \alpha^*) \ln(1 - \alpha^*)$$

$$D_{\text{KL}}(\alpha^* || \alpha) = \alpha^* \ln \frac{\alpha^*}{\alpha} + (1 - \alpha^*) \ln \frac{1 - \alpha^*}{1 - \alpha}$$

Tuning α

- ▶ let k^* be the best number of switches (in hind sight) and $\alpha^* = k^*/I$
- ▶ Suppose we use $\alpha \approx \alpha^*$ then the bound that we get is

$$L_A \leq L_* + \frac{1}{\eta} ((k+1) \ln n + (I-1)(H(\alpha^*) + D_{\text{KL}}(\alpha^*||\alpha)))$$

Where

$$H(\alpha^*) = -\alpha^* \ln \alpha^* - (1 - \alpha^*) \ln(1 - \alpha^*)$$

$$D_{\text{KL}}(\alpha^*||\alpha) = \alpha^* \ln \frac{\alpha^*}{\alpha} + (1 - \alpha^*) \ln \frac{1 - \alpha^*}{1 - \alpha}$$

- ▶ This is very close to the loss of the computationally inefficient algorithm.

Tuning α

- ▶ let k^* be the best number of switches (in hind sight) and $\alpha^* = k^*/I$
- ▶ Suppose we use $\alpha \approx \alpha^*$ then the bound that we get is

$$L_A \leq L_* + \frac{1}{\eta} ((k+1) \ln n + (I-1)(H(\alpha^*) + D_{\text{KL}}(\alpha^* || \alpha)))$$

Where

$$H(\alpha^*) = -\alpha^* \ln \alpha^* - (1 - \alpha^*) \ln(1 - \alpha^*)$$

$$D_{\text{KL}}(\alpha^* || \alpha) = \alpha^* \ln \frac{\alpha^*}{\alpha} (1 - \alpha^*) \ln \frac{1 - \alpha^*}{1 - \alpha}$$

- ▶ This is very close to the loss of the computationally inefficient algorithm.
- ▶ For the log loss case this is essentially optimal.

Tuning α

- ▶ let k^* be the best number of switches (in hind sight) and $\alpha^* = k^*/I$
- ▶ Suppose we use $\alpha \approx \alpha^*$ then the bound that we get is

$$L_A \leq L_* + \frac{1}{\eta} ((k+1) \ln n + (I-1)(H(\alpha^*) + D_{\text{KL}}(\alpha^* || \alpha)))$$

Where

$$H(\alpha^*) = -\alpha^* \ln \alpha^* - (1 - \alpha^*) \ln(1 - \alpha^*)$$

$$D_{\text{KL}}(\alpha^* || \alpha) = \alpha^* \ln \frac{\alpha^*}{\alpha} + (1 - \alpha^*) \ln \frac{1 - \alpha^*}{1 - \alpha}$$

- ▶ This is very close to the loss of the computationally inefficient algorithm.
- ▶ For the log loss case this is essentially optimal.
- ▶ Not so for square loss!

What can we hope to improve?

- ▶ In the fixed-share algorithm, the weight of a suboptimal expert never decreases below α/n .

What can we hope to improve?

- ▶ In the fixed-share algorithm, the weight of a suboptimal expert never decreases below α/n .
- ▶ The algorithm does not concentrate only on the best expert, even if the last switch is in the distant past.

What can we hope to improve?

- ▶ In the fixed-share algorithm, the weight of a suboptimal expert never decreases below α/n .
- ▶ The algorithm does not concentrate only on the best expert, even if the last switch is in the distant past.
- ▶ The regret depends on the length of the sequence.

The idea of variable-share

- ▶ Let the fraction of the total weight given to the best expert get arbitrarily close to **1**.

The idea of variable-share

- ▶ Let the fraction of the total weight given to the best expert get arbitrarily close to **1**.
- ▶ we can get a regret bound that depends only on the number of switches, not on the length of the sequence.

The idea of variable-share

- ▶ Let the fraction of the total weight given to the best expert get arbitrarily close to **1**.
- ▶ we can get a regret bound that depends only on the number of switches, not on the length of the sequence.
- ▶ Requires that the loss be bounded.

The idea of variable-share

- ▶ Let the fraction of the total weight given to the best expert get arbitrarily close to **1**.
- ▶ we can get a regret bound that depends only on the number of switches, not on the length of the sequence.
- ▶ Requires that the loss be bounded.
- ▶ Works for **square** loss, but not for **log** loss!

Variable-share

$$pool = \sum_{i=1}^n \left(1 - (1 - \alpha)^{\ell_{t,i}}\right) w_{t,i}^m$$

$$w_{t+1,i}^s = (1 - \alpha)^{\ell_{t,i}} w_{t,i}^m + \frac{1}{n-1} \left(pool - \left(1 - (1 - \alpha)^{\ell_{t,i}}\right) w_{t,i}^m \right)$$

Variable-share

$$pool = \sum_{i=1}^n \left(1 - (1 - \alpha)^{\ell_{t,i}}\right) w_{t,i}^m$$

$$w_{t+1,i}^s = (1 - \alpha)^{\ell_{t,i}} w_{t,i}^m + \frac{1}{n-1} \left(pool - (1 - (1 - \alpha)^{\ell_{t,i}}) w_{t,i}^m \right)$$

If $\ell_{t,i} = 0$, then expert i does not contribute to the pool.

Variable-share

$$pool = \sum_{i=1}^n \left(1 - (1 - \alpha)^{\ell_{t,i}}\right) w_{t,i}^m$$

$$w_{t+1,i}^s = (1 - \alpha)^{\ell_{t,i}} w_{t,i}^m + \frac{1}{n-1} \left(pool - (1 - (1 - \alpha)^{\ell_{t,i}}) w_{t,i}^m \right)$$

If $\ell_{t,i} = 0$, then expert i does not contribute to the pool.
Expert can get fraction of the total weight arbitrarily close to 1.

Variable-share

$$pool = \sum_{i=1}^n \left(1 - (1 - \alpha)^{\ell_{t,i}}\right) w_{t,i}^m$$

$$w_{t+1,i}^s = (1 - \alpha)^{\ell_{t,i}} w_{t,i}^m + \frac{1}{n-1} \left(pool - (1 - (1 - \alpha)^{\ell_{t,i}}) w_{t,i}^m \right)$$

If $\ell_{t,i} = 0$, then expert i does not contribute to the pool.
Expert can get fraction of the total weight arbitrarily close to 1.
Shares the weight quickly if $\ell_{t,i} > 0$

Bound for variable share



$$\frac{1}{\eta} \ln n + \left(1 + \frac{1}{(1-\alpha)\eta}\right) L_* + k \left(1 + \frac{1}{\eta} \left(\ln n - 1 + \ln \frac{1}{\alpha} + \ln \frac{1}{1-\alpha}\right)\right)$$

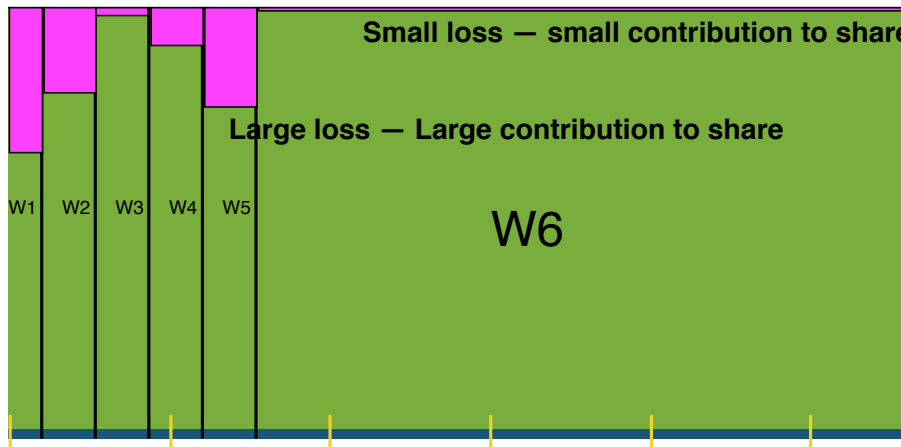
Bound for variable share




$$\frac{1}{\eta} \ln n + \left(1 + \frac{1}{(1-\alpha)\eta}\right) L_* + k \left(1 + \frac{1}{\eta} \left(\ln n - 1 + \ln \frac{1}{\alpha} + \ln \frac{1}{1-\alpha}\right)\right)$$

- α should be tuned so that it is (close to) $\frac{k}{2k+L_*}$

Variable share figure



An experiment using variable share



Next Class

- ▶ Suppose the best switching sequence is repeatedly switching among a small subset of the experts $n' \ll n$

Next Class

- ▶ Suppose the best switching sequence is repeatedly switching among a small subset of the experts $n' \ll n$
- ▶ In the context of speech recognition - the speaker repeatedly uses a small number of phonemes.

Next Class

- ▶ Suppose the best switching sequence is repeatedly switching among a small subset of the experts $n' \ll n$
- ▶ In the context of speech recognition - the speaker repeatedly uses a small number of phonemes.
- ▶ If we know the subset, we can pay $\ln n'$ per switch rather than $\ln n$

Next Class

- ▶ Suppose the best switching sequence is repeatedly switching among a small subset of the experts $n' \ll n$
- ▶ In the context of speech recognition - the speaker repeatedly uses a small number of phonemes.
- ▶ If we know the subset, we can pay $\ln n'$ per switch rather than $\ln n$
- ▶ Can track switches much more closely.

Next Class

- ▶ Suppose the best switching sequence is repeatedly switching among a small subset of the experts $n' \ll n$
- ▶ In the context of speech recognition - the speaker repeatedly uses a small number of phonemes.
- ▶ If we know the subset, we can pay $\ln n'$ per switch rather than $\ln n$
- ▶ Can track switches much more closely.
- ▶ Easy to describe an inefficient algorithm (consider all $\binom{n}{n'}$ subsets.)

Next Class

- ▶ Suppose the best switching sequence is repeatedly switching among a small subset of the experts $n' \ll n$
- ▶ In the context of speech recognition - the speaker repeatedly uses a small number of phonemes.
- ▶ If we know the subset, we can pay $\ln n'$ per switch rather than $\ln n$
- ▶ Can track switches much more closely.
- ▶ Easy to describe an inefficient algorithm (consider all $\binom{n}{n'}$ subsets.)
- ▶ Next class - how to do as well with just one weight per expert.