

Theory and Applications of Boosting

Rob Schapire

Example: “How May I Help You?”

[Gorin et al.]

- goal: automatically categorize type of call requested by phone customer (Collect, CallingCard, PersonToPerson, etc.)
 - yes I'd like to place a collect call long distance please (Collect)
 - operator I need to make a call but I need to bill it to my office (ThirdNumber)
 - yes I'd like to place a call on my master card please (CallingCard)
 - I just called a number in sioux city and I musta rang the wrong number because I got the wrong party and I would like to have that taken off of my bill (BillingCredit)
- observation:
 - easy to find “rules of thumb” that are “often” correct
 - e.g.: “IF ‘card’ occurs in utterance
THEN predict ‘CallingCard’ ”
 - hard to find single highly accurate prediction rule

The Boosting Approach

- devise computer program for deriving rough rules of thumb
- apply procedure to subset of examples
- obtain rule of thumb
- apply to 2nd subset of examples
- obtain 2nd rule of thumb
- repeat T times

Key Details

- how to choose examples on each round?
 - concentrate on “hardest” examples
(those most often misclassified by previous rules of thumb)
- how to combine rules of thumb into single prediction rule?
 - take (weighted) majority vote of rules of thumb

Boosting

- boosting = general method of converting rough rules of thumb into highly accurate prediction rule
- technically:
 - assume given “weak” learning algorithm that can consistently find classifiers (“rules of thumb”) at least slightly better than random, say, accuracy $\geq 55\%$ (in two-class setting) [“weak learning assumption”]
 - given sufficient data, a boosting algorithm can provably construct single classifier with very high accuracy, say, 99%

Outline of Tutorial

- basic algorithm and core theory
- fundamental perspectives
- practical extensions
- advanced topics

Preamble: Early History

Strong and Weak Learnability

- boosting's roots are in "PAC" learning model [Valiant '84]
- get random examples from unknown, arbitrary distribution
- **strong** PAC learning algorithm:
 - for **any** distribution with high probability given polynomially many examples (and polynomial time) can find classifier with **arbitrarily small** generalization error
- **weak** PAC learning algorithm
 - same, but generalization error only needs to be **slightly better than random guessing** ($\frac{1}{2} - \gamma$)
- [Kearns & Valiant '88]:
 - does weak learnability imply strong learnability?

If Boosting Possible, Then...

- can use (fairly) **wild** guesses to produce highly accurate predictions
- if can learn “part way” then can learn “all the way”
- should be able to improve **any** learning algorithm
- for any learning problem:
 - **either** can always learn with nearly **perfect accuracy**
 - **or** there exist cases where **cannot** learn even slightly better than **random guessing**

First Boosting Algorithms

- [Schapire '89]:
 - first provable boosting algorithm
- [Freund '90]:
 - “optimal” algorithm that “boosts by majority”
- [Drucker, Schapire & Simard '92]:
 - first experiments using boosting
 - limited by practical drawbacks
- [Freund & Schapire '95]:
 - introduced “AdaBoost” algorithm
 - strong practical advantages over previous boosting algorithms

Basic Algorithm and Core Theory

- introduction to AdaBoost
- analysis of training error
- analysis of test error
and the margins theory
- experiments and applications

Basic Algorithm and Core Theory

- introduction to AdaBoost
- analysis of training error
- analysis of test error
and the margins theory
- experiments and applications

A Formal Description of Boosting

- given training set $(x_1, y_1), \dots, (x_m, y_m)$
- $y_i \in \{-1, +1\}$ correct label of instance $x_i \in X$
- for $t = 1, \dots, T$:
 - construct distribution D_t on $\{1, \dots, m\}$
 - find weak classifier (“rule of thumb”)

$$h_t : X \rightarrow \{-1, +1\}$$

with small error ϵ_t on D_t :

$$\epsilon_t = \Pr_{i \sim D_t}[h_t(x_i) \neq y_i]$$

- output final classifier H_{final}

AdaBoost

[with Freund]

- constructing D_t :
 - $D_1(i) = 1/m$
 - given D_t and h_t :

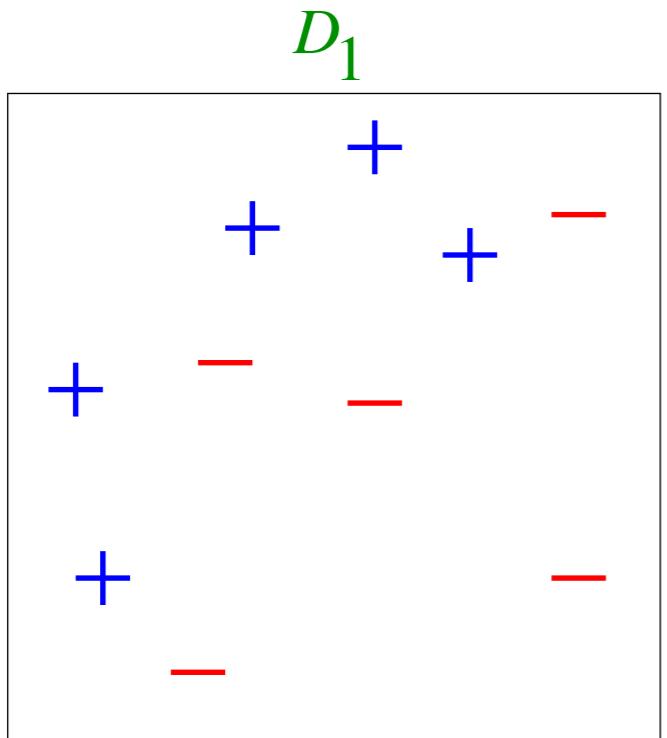
$$\begin{aligned} D_{t+1}(i) &= \frac{D_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } y_i = h_t(x_i) \\ e^{\alpha_t} & \text{if } y_i \neq h_t(x_i) \end{cases} \\ &= \frac{D_t(i)}{Z_t} \exp(-\alpha_t y_i h_t(x_i)) \end{aligned}$$

where Z_t = normalization factor

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right) > 0$$

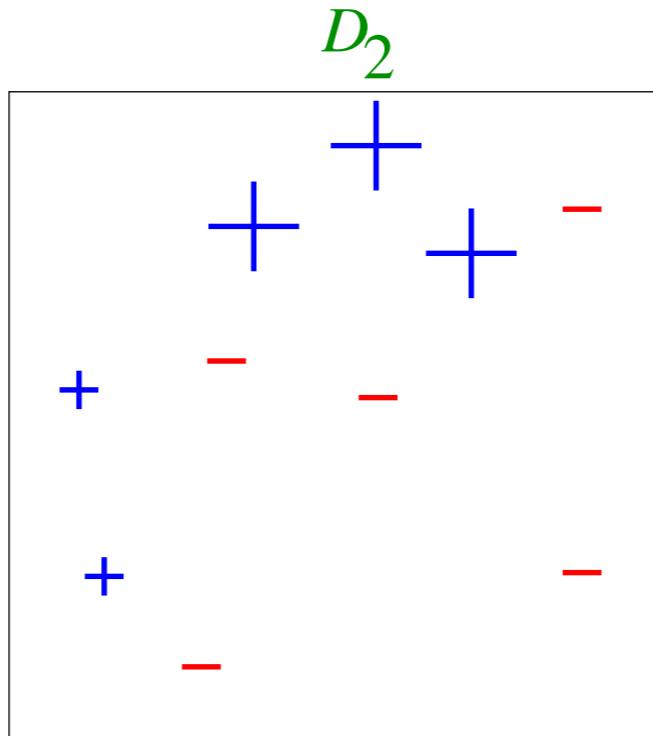
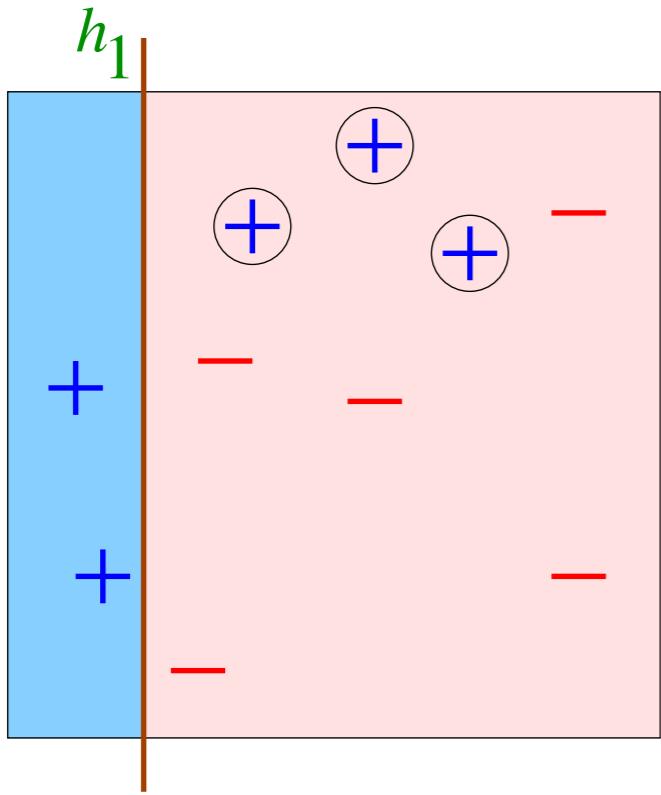
- final classifier:
 - $H_{\text{final}}(x) = \text{sign} \left(\sum_t \alpha_t h_t(x) \right)$

Toy Example



weak classifiers = vertical or horizontal half-planes

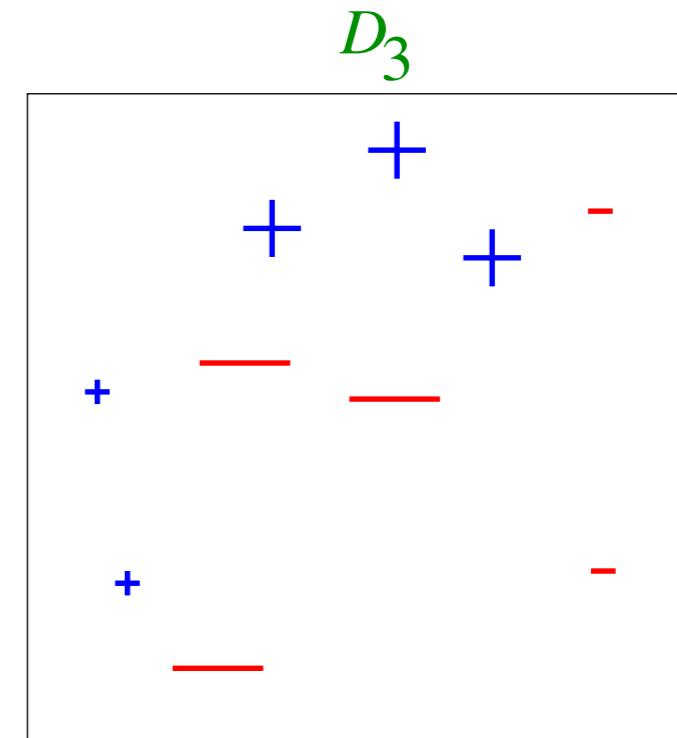
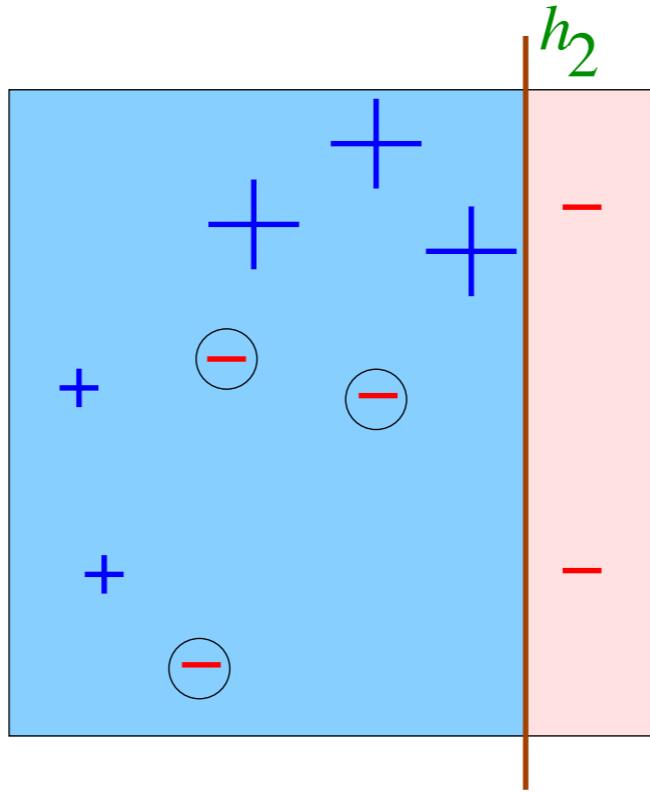
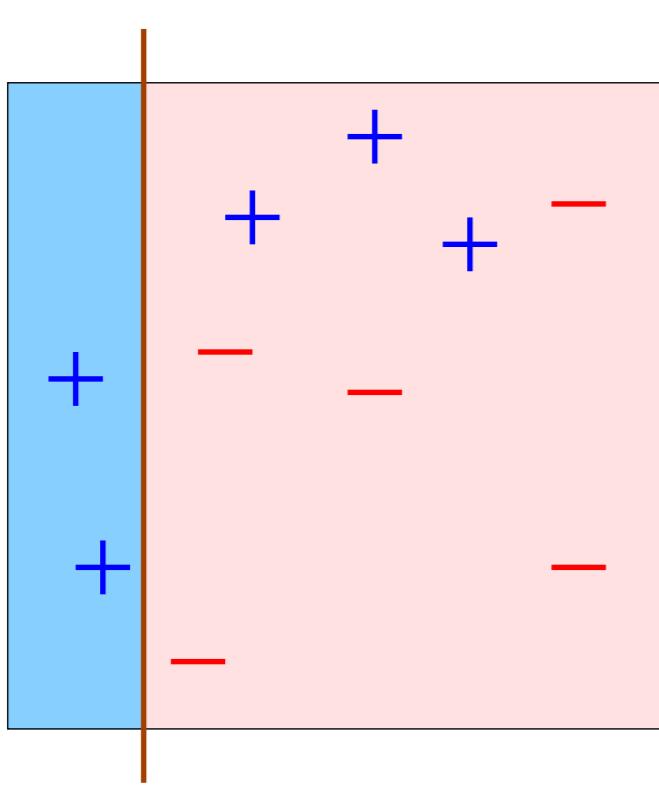
Round 1



$$\varepsilon_1 = 0.30$$

$$\alpha_1 = 0.42$$

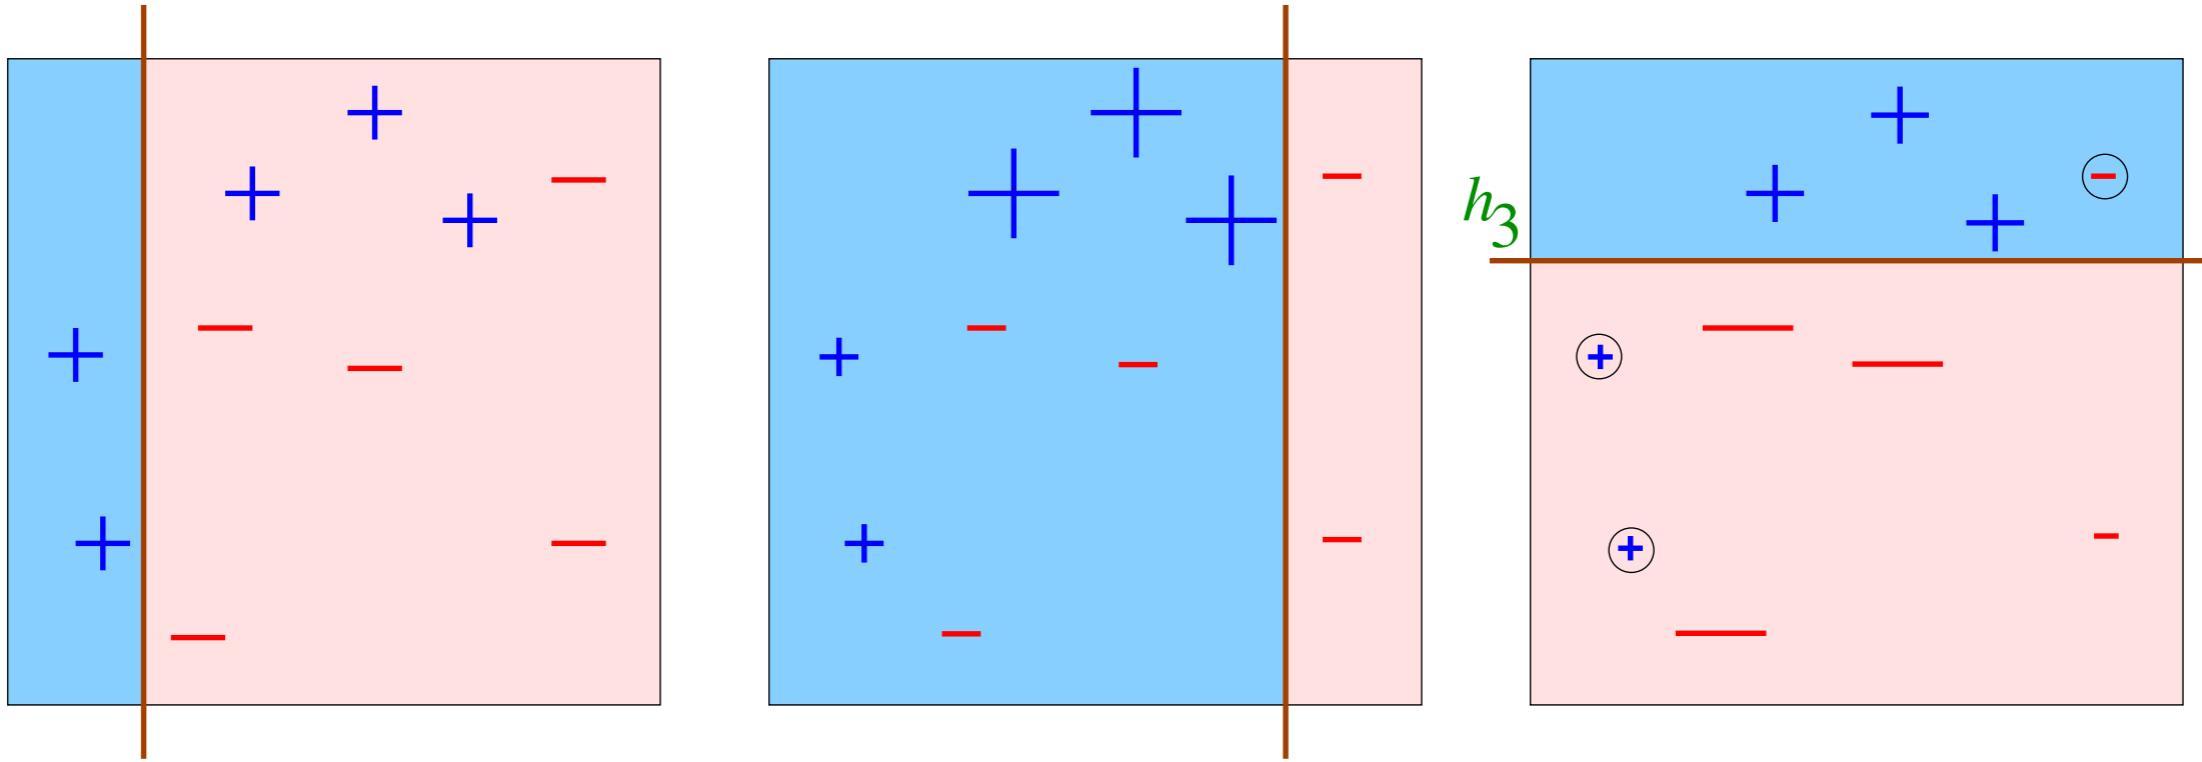
Round 2



$$\varepsilon_2 = 0.21$$

$$\alpha_2 = 0.65$$

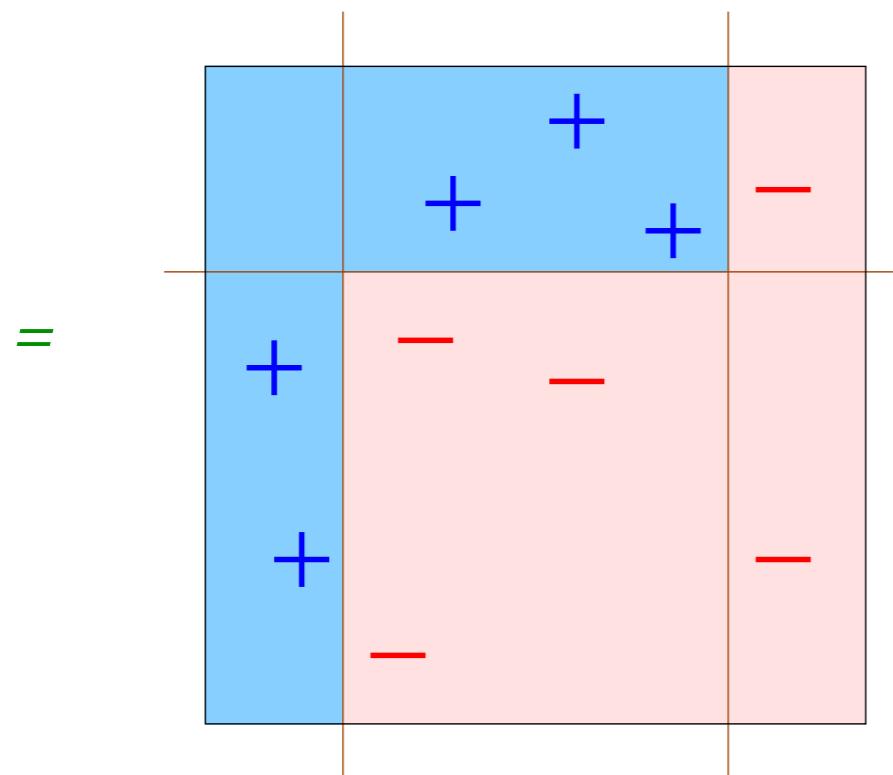
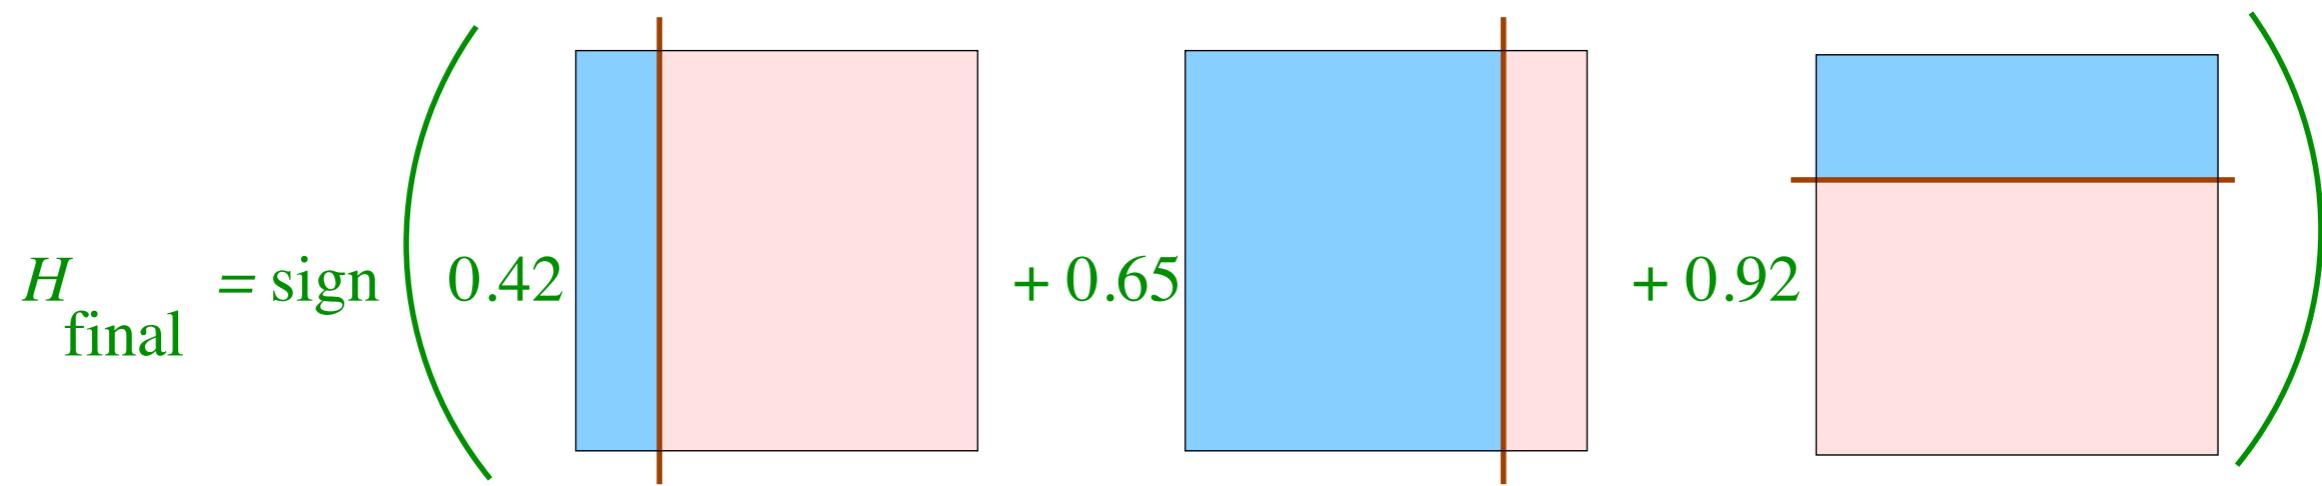
Round 3



$$\varepsilon_3 = 0.14$$

$$\alpha_3 = 0.92$$

Final Classifier



Basic Algorithm and Core Theory

- introduction to AdaBoost
- analysis of training error
- analysis of test error
and the margins theory
- experiments and applications

Analyzing the Training Error

[with Freund]

- **Theorem:**

- write ϵ_t as $\frac{1}{2} - \gamma_t$ [γ_t = “edge”]
- then

$$\begin{aligned}\text{training error}(H_{\text{final}}) &\leq \prod_t \left[2\sqrt{\epsilon_t(1-\epsilon_t)} \right] \\ &= \prod_t \sqrt{1 - 4\gamma_t^2} \\ &\leq \exp \left(-2 \sum_t \gamma_t^2 \right)\end{aligned}$$

- so: if $\forall t : \gamma_t \geq \gamma > 0$
then $\text{training error}(H_{\text{final}}) \leq e^{-2\gamma^2 T}$
- AdaBoost is adaptive:
 - does **not** need to know γ or T a priori
 - can exploit $\gamma_t \gg \gamma$

Proof

- let $F(x) = \sum_t \alpha_t h_t(x) \Rightarrow H_{\text{final}}(x) = \text{sign}(F(x))$
- *Step 1:* unwrapping recurrence:

$$\begin{aligned} D_{\text{final}}(i) &= \frac{1}{m} \frac{\exp \left(-y_i \sum_t \alpha_t h_t(x_i) \right)}{\prod_t Z_t} \\ &= \frac{1}{m} \frac{\exp (-y_i F(x_i))}{\prod_t Z_t} \end{aligned}$$

Proof (cont.)

- *Step 2:* training error(H_{final}) $\leq \prod_t Z_t$

- Proof:

$$\begin{aligned}\text{training error}(H_{\text{final}}) &= \frac{1}{m} \sum_i \begin{cases} 1 & \text{if } y_i \neq H_{\text{final}}(x_i) \\ 0 & \text{else} \end{cases} \\ &= \frac{1}{m} \sum_i \begin{cases} 1 & \text{if } y_i F(x_i) \leq 0 \\ 0 & \text{else} \end{cases} \\ &\leq \frac{1}{m} \sum_i \exp(-y_i F(x_i)) \\ &= \sum_i D_{\text{final}}(i) \prod_t Z_t \\ &= \prod_t Z_t\end{aligned}$$

Proof (cont.)

- *Step 3:* $Z_t = 2\sqrt{\epsilon_t(1 - \epsilon_t)}$

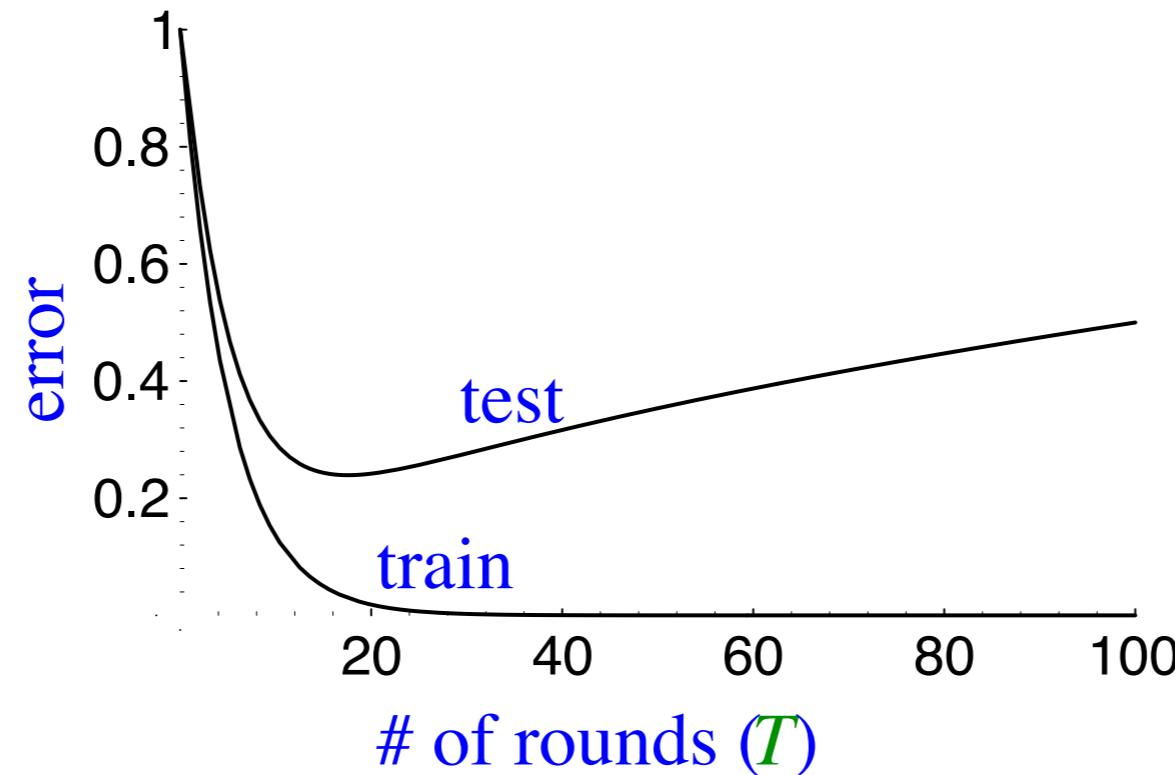
- Proof:

$$\begin{aligned} Z_t &= \sum_i D_t(i) \exp(-\alpha_t y_i h_t(x_i)) \\ &= \sum_{i:y_i \neq h_t(x_i)} D_t(i) e^{\alpha_t} + \sum_{i:y_i = h_t(x_i)} D_t(i) e^{-\alpha_t} \\ &= \epsilon_t e^{\alpha_t} + (1 - \epsilon_t) e^{-\alpha_t} \\ &= 2\sqrt{\epsilon_t(1 - \epsilon_t)} \end{aligned}$$

Basic Algorithm and Core Theory

- introduction to AdaBoost
- analysis of training error
- analysis of test error
and the margins theory
- experiments and applications

How Will Test Error Behave? (A First Guess)



expect:

- training error to continue to drop (or reach zero)
- test error to increase when H_{final} becomes “too complex”
 - “Occam’s razor”
 - overfitting
 - hard to know when to stop training

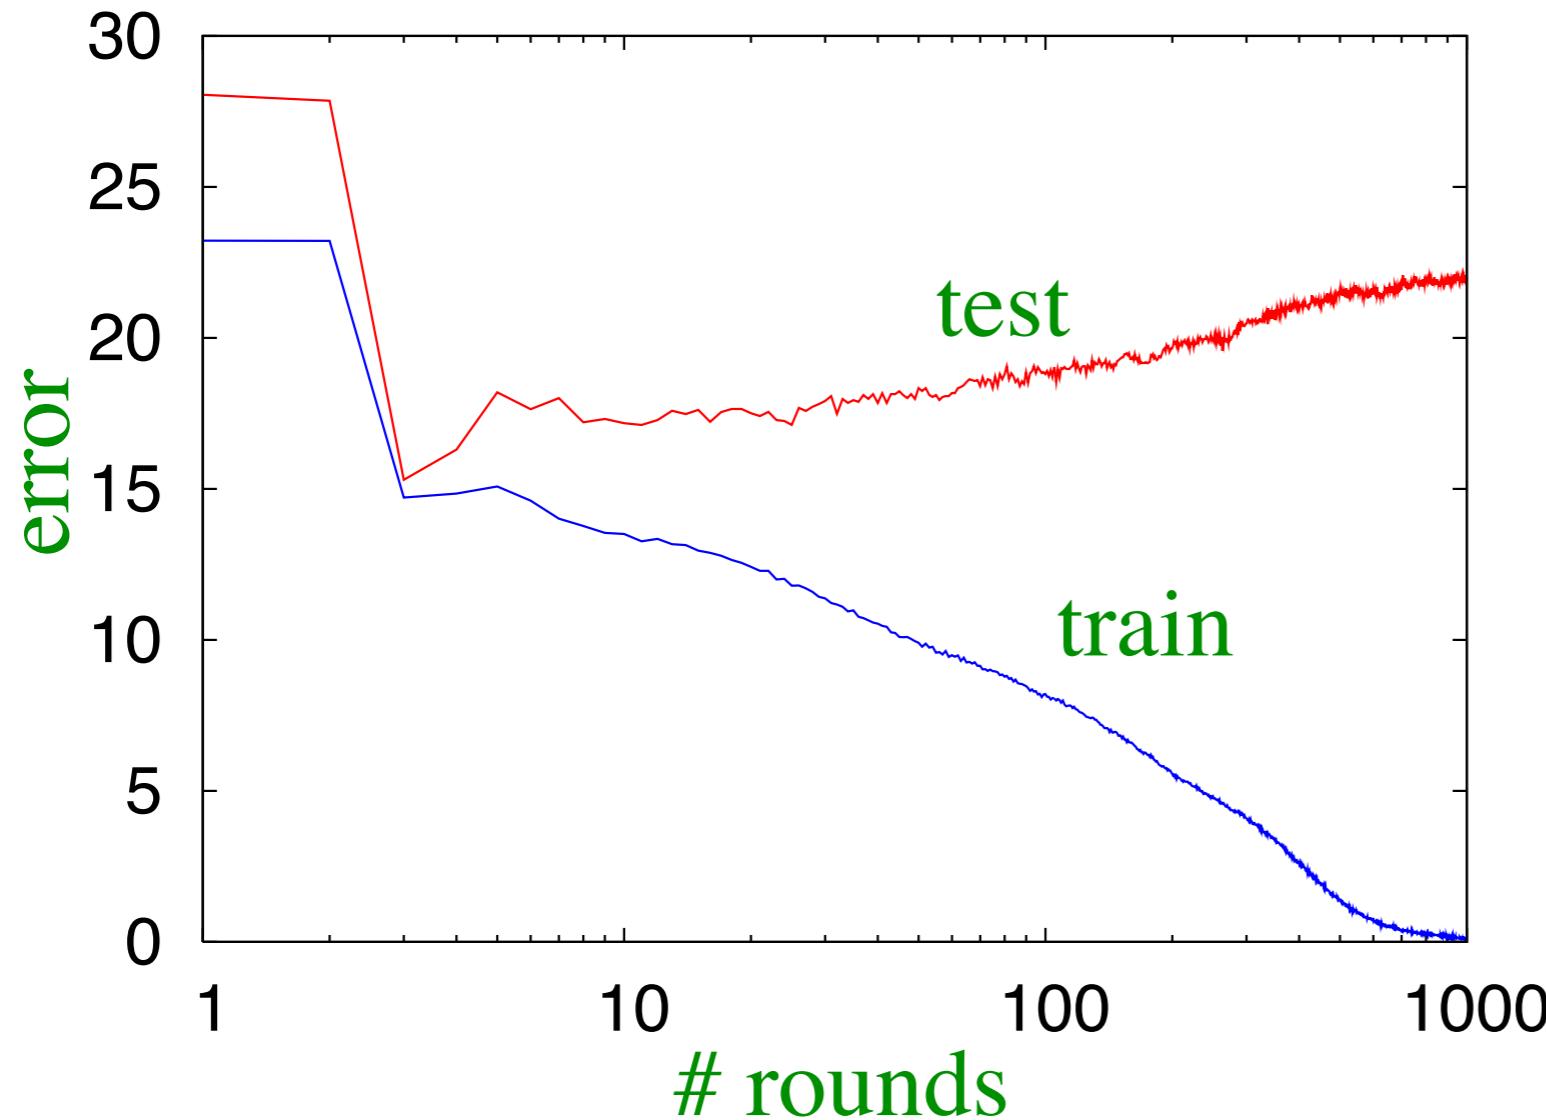
Technically...

- with high probability:

$$\text{generalization error} \leq \text{training error} + \tilde{O}\left(\sqrt{\frac{dT}{m}}\right)$$

- bound depends on
 - m = # training examples
 - d = “complexity” of weak classifiers
 - T = # rounds
- generalization error = $E[\text{test error}]$
- predicts **overfitting**

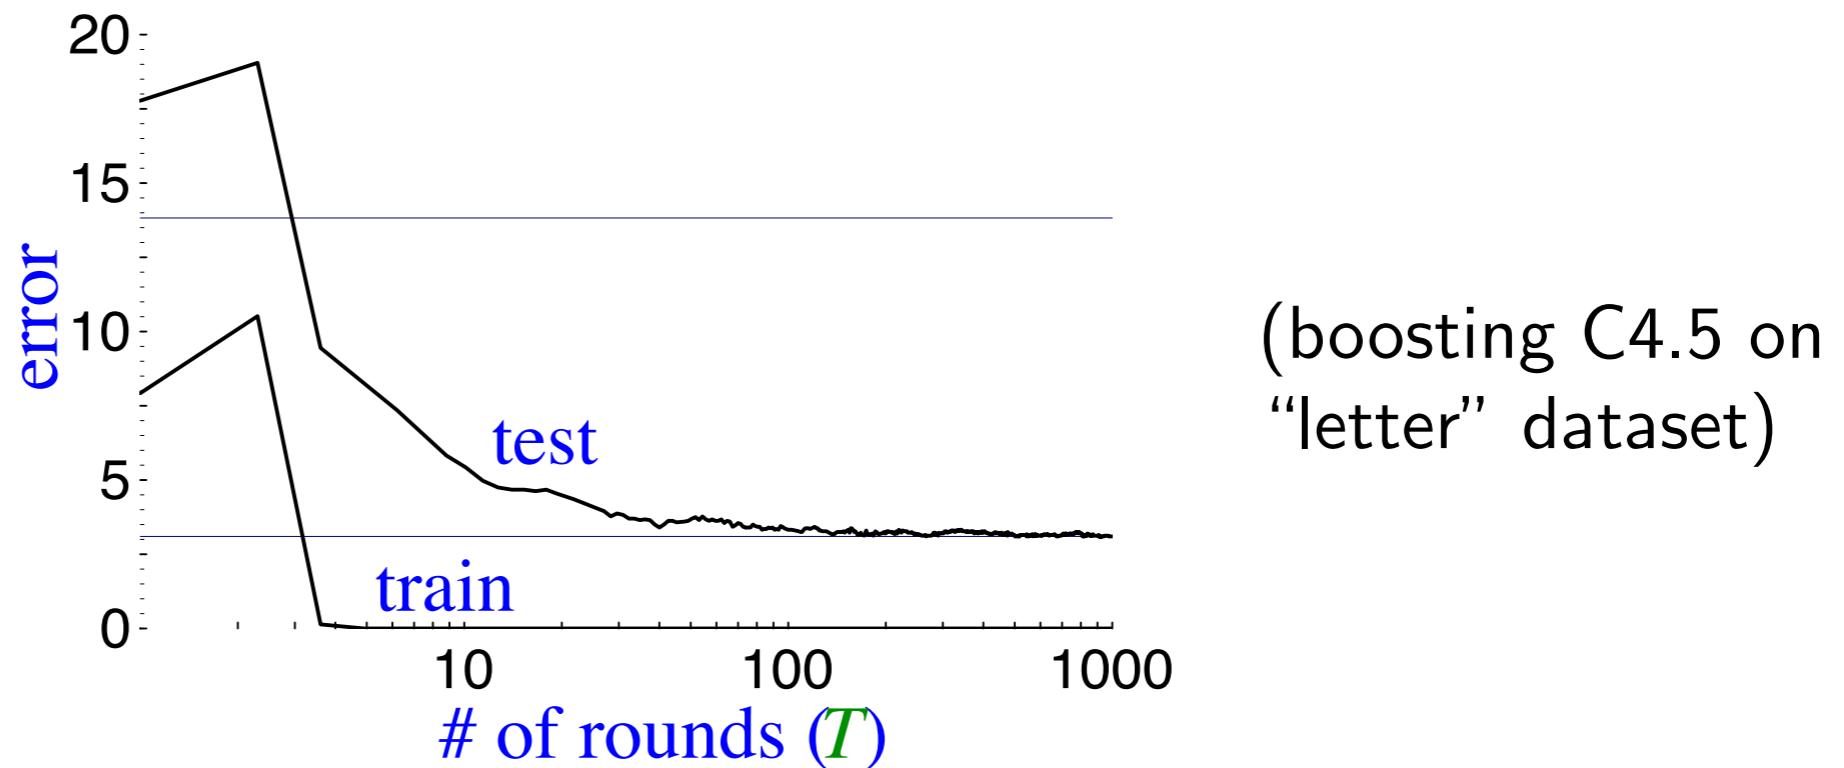
Overfitting Can Happen



(boosting “stumps” on heart-disease dataset)

- but often doesn't...

Actual Typical Run



- test error does **not** increase, even after 1000 rounds
 - (total size > 2,000,000 nodes)
- test error continues to drop even after training error is zero!

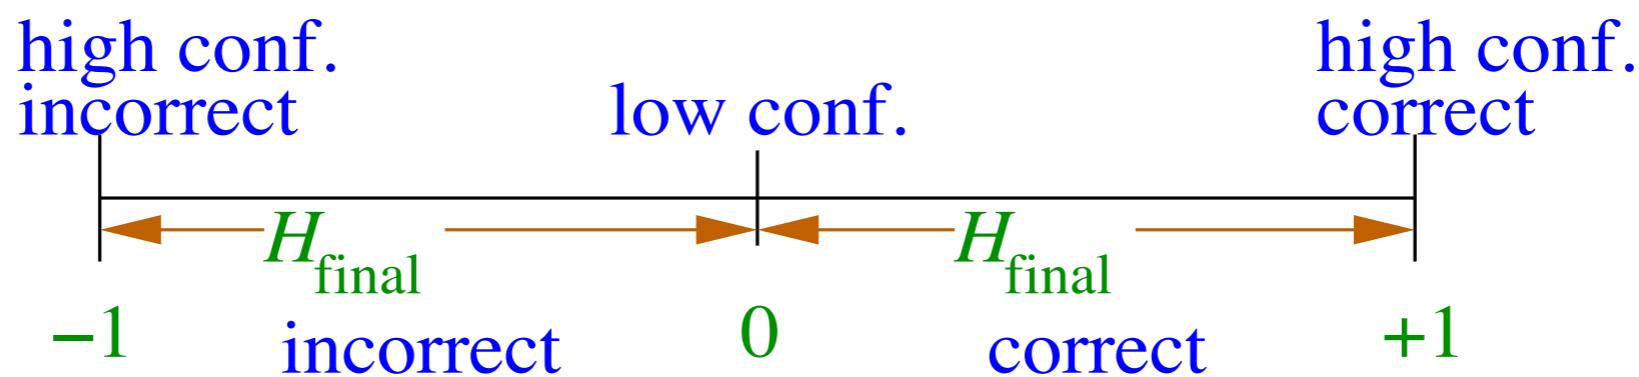
	# rounds	5	100	1000
train error	0.0	0.0	0.0	
test error	8.4	3.3	3.1	

- Occam’s razor **wrongly** predicts “simpler” rule is better

A Better Story: The Margins Explanation

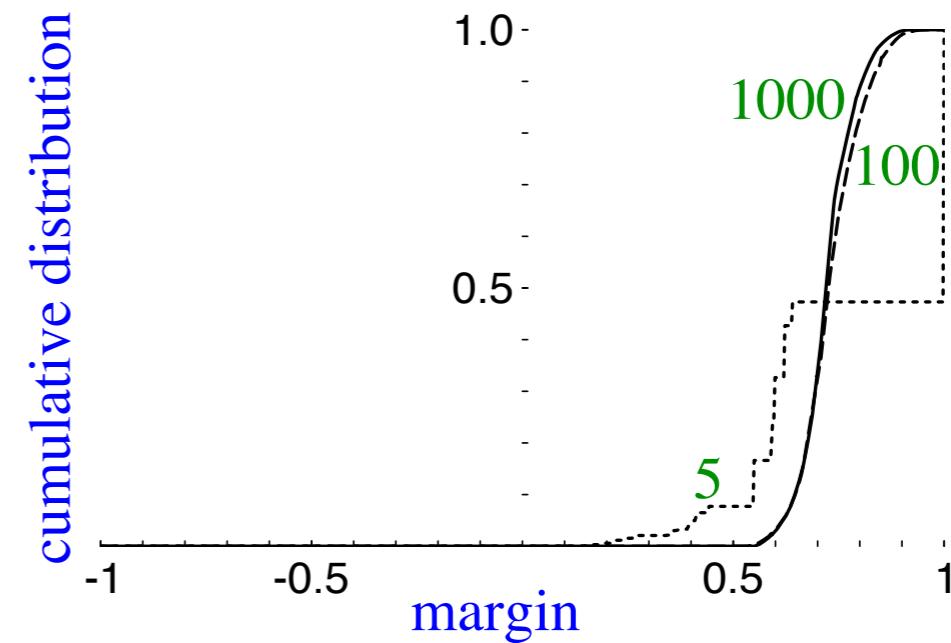
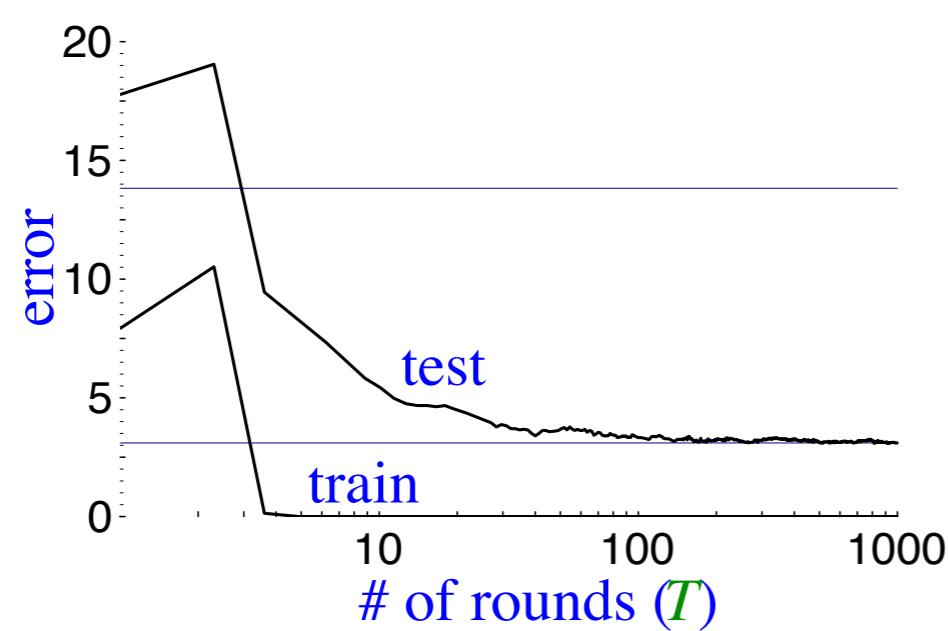
[with Freund, Bartlett & Lee]

- key idea:
 - training error only measures whether classifications are right or wrong
 - should also consider confidence of classifications
- recall: H_{final} is weighted majority vote of weak classifiers
- measure confidence by margin = strength of the vote
= (weighted fraction voting correctly)
–(weighted fraction voting incorrectly)



Empirical Evidence: The Margin Distribution

- margin distribution
= cumulative distribution of margins of training examples



	# rounds	5	100	1000
train error	0.0	0.0	0.0	0.0
test error	8.4	3.3	3.1	3.1
% margins ≤ 0.5	7.7	0.0	0.0	0.0
minimum margin	0.14	0.52	0.55	0.55

Theoretical Evidence: Analyzing Boosting Using Margins

- Theorem: large margins \Rightarrow better bound on generalization error (independent of number of rounds)
 - proof idea: if all margins are large, then can approximate final classifier by a much smaller classifier (just as polls can predict not-too-close election)
- Theorem: boosting tends to increase margins of training examples (given weak learning assumption)
 - moreover, larger edges \Rightarrow larger margins
 - proof idea: similar to training error proof
- so:
 - although final classifier is getting larger, margins are likely to be increasing, so final classifier actually getting close to a simpler classifier, driving down the test error

More Technically...

- with high probability, $\forall \theta > 0$:

$$\text{generalization error} \leq \hat{\Pr}[\text{margin} \leq \theta] + \tilde{O}\left(\frac{\sqrt{d/m}}{\theta}\right)$$

($\hat{\Pr}[\cdot]$ = empirical probability)

- bound depends on
 - m = # training examples
 - d = “complexity” of weak classifiers
 - entire distribution of margins of training examples
- $\hat{\Pr}[\text{margin} \leq \theta] \rightarrow 0$ exponentially fast (in T)
if $\epsilon_t < \frac{1}{2} - \theta$ ($\forall t$)
 - so: if weak learning assumption holds, then all examples will quickly have “large” margins

Consequences of Margins Theory

- predicts good generalization with no overfitting if:
 - weak classifiers have **large edges** (implying **large margins**)
 - weak classifiers not too complex relative to size of training set
- e.g., boosting decision trees resistant to overfitting since trees often have large edges and limited complexity
- overfitting **may** occur if:
 - small edges (underfitting), or
 - overly complex weak classifiers
- e.g., heart-disease dataset:
 - stumps yield small edges
 - also, small dataset

Improved Boosting with Better Margin-Maximization?

- can design algorithms **more effective** than AdaBoost at maximizing the minimum margin
- in practice, often perform **worse** [Breiman]
- **why??**
- more aggressive margin maximization seems to lead to:
 - more **complex** weak classifiers
(even using same weak learner); or
 - **higher minimum margins**,
but margin **distributions** that are **lower overall** [with Reyzin]

Comparison to SVM's

- both AdaBoost and SVM's:
 - work by maximizing “margins”
 - find linear threshold function in high-dimensional space
- differences:
 - margin measured slightly differently
(using different norms)
 - SVM's handle high-dimensional space using kernel trick;
AdaBoost uses weak learner to search over space

Basic Algorithm and Core Theory

- introduction to AdaBoost
- analysis of training error
- analysis of test error
and the margins theory
- experiments and applications

Practical Advantages of AdaBoost

- fast
- simple and easy to program
- no parameters to tune (except T)
- flexible — can combine with any learning algorithm
- no prior knowledge needed about weak learner
- provably effective, provided can consistently find rough rules of thumb
 - shift in mind set — goal now is merely to find classifiers barely better than random guessing
- versatile
 - can use with data that is textual, numeric, discrete, etc.
 - has been extended to learning problems well beyond binary classification

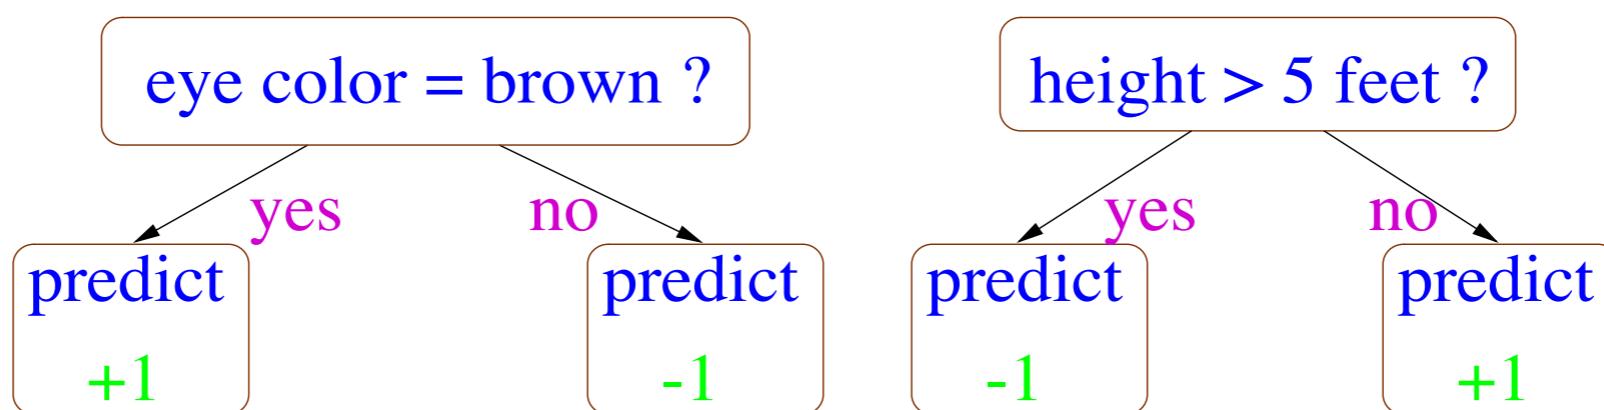
Caveats

- performance of AdaBoost depends on **data** and **weak learner**
- consistent with theory, AdaBoost can **fail** if
 - weak classifiers too **complex**
 - overfitting
 - weak classifiers too **weak** ($\gamma_t \rightarrow 0$ too quickly)
 - underfitting
 - low margins → overfitting
- empirically, AdaBoost seems especially susceptible to uniform noise

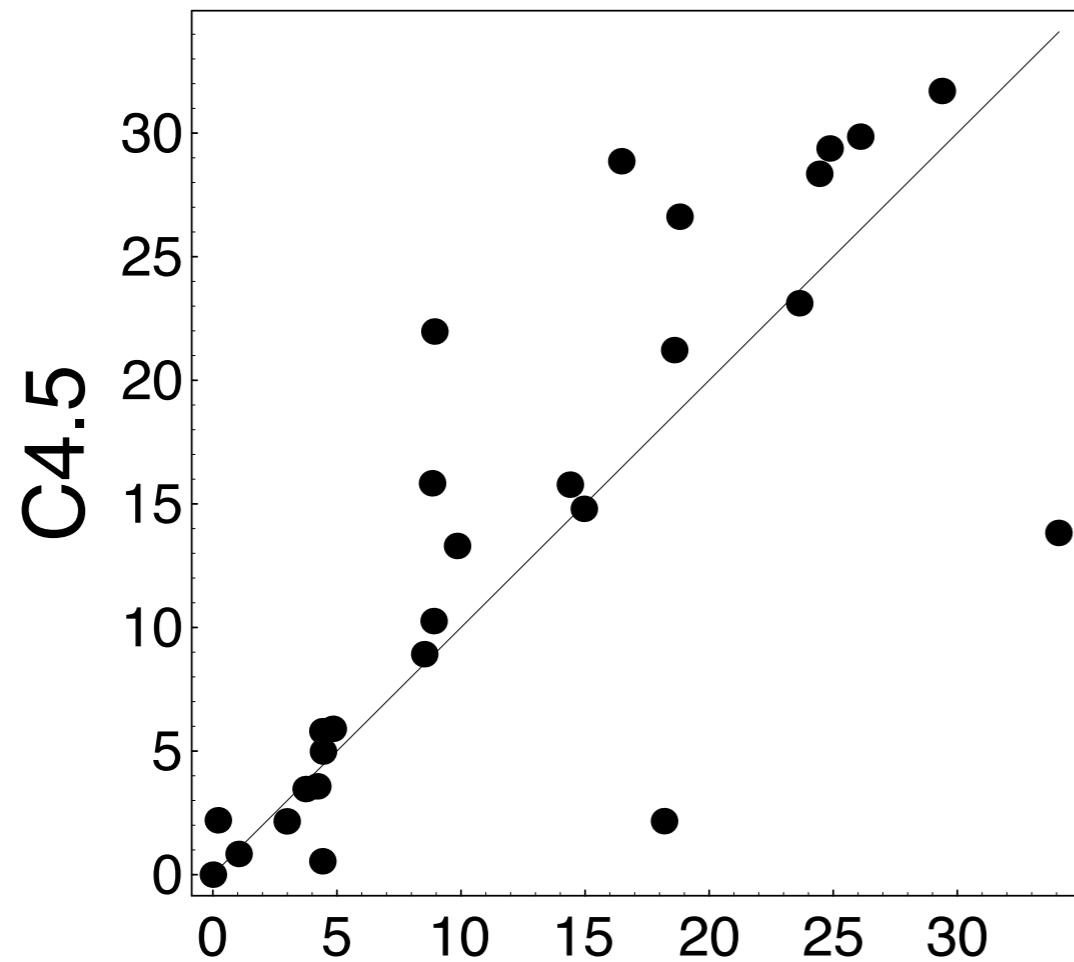
UCI Experiments

[with Freund]

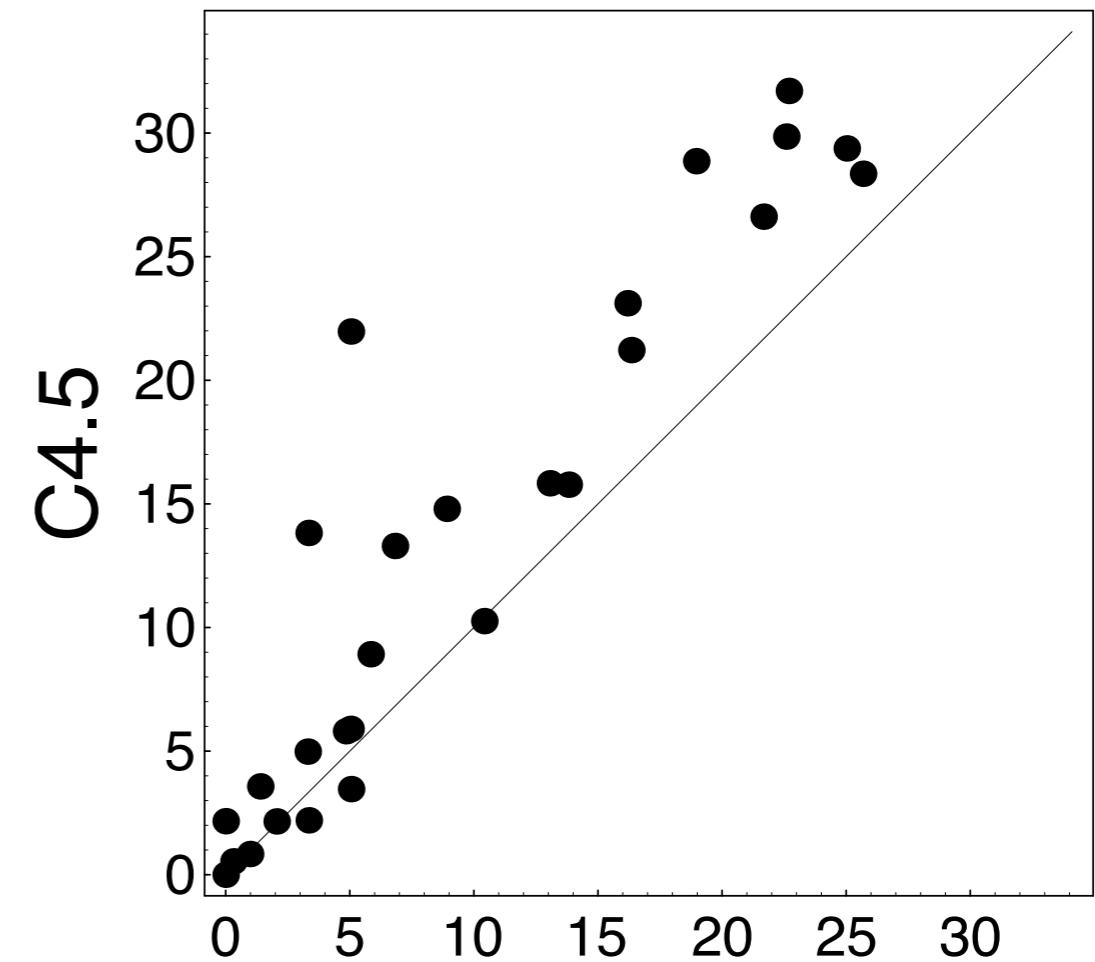
- tested AdaBoost on UCI benchmarks
- used:
 - C4.5 (Quinlan's decision tree algorithm)
 - “decision stumps”: very simple rules of thumb that test on single attributes



UCI Results



boosting Stumps

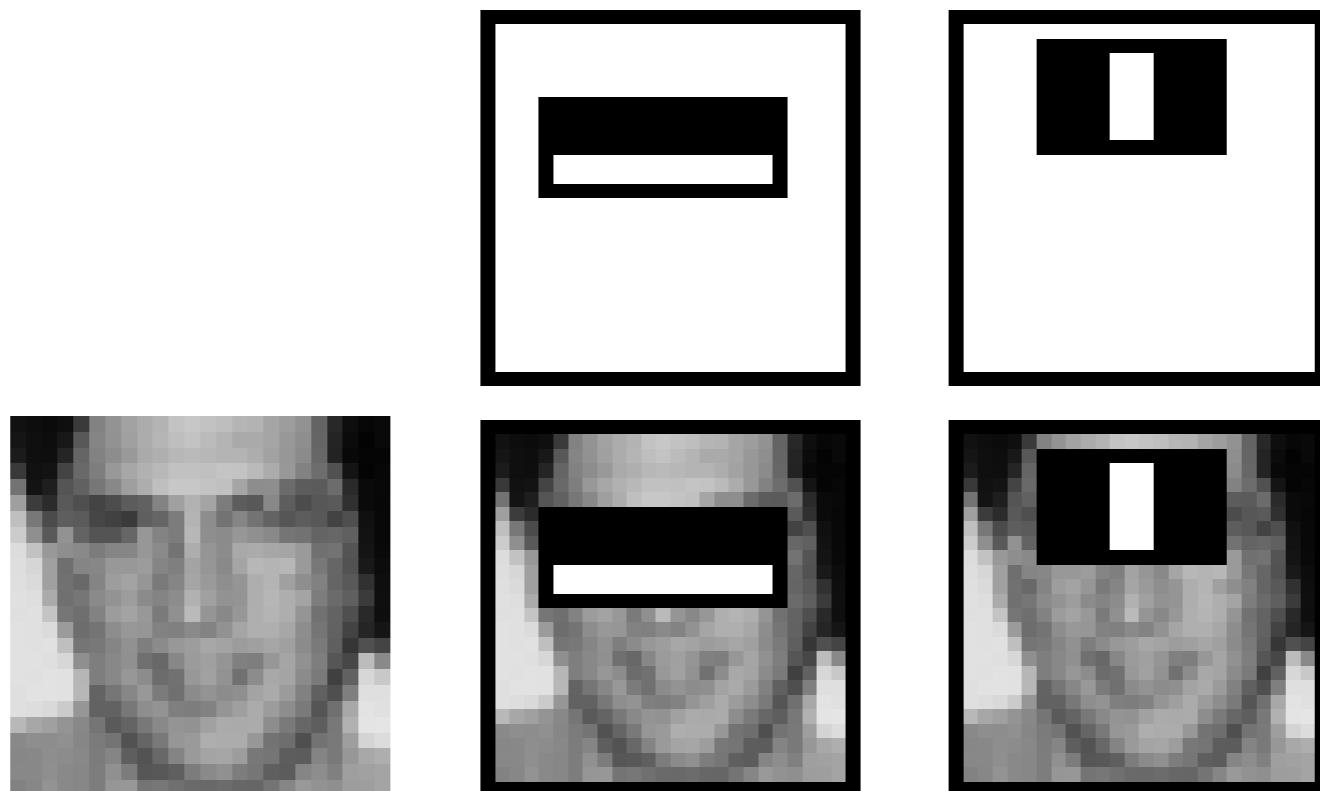


boosting C4.5

Application: Detecting Faces

[Viola & Jones]

- problem: find faces in photograph or movie
- weak classifiers: detect light/dark rectangles in image



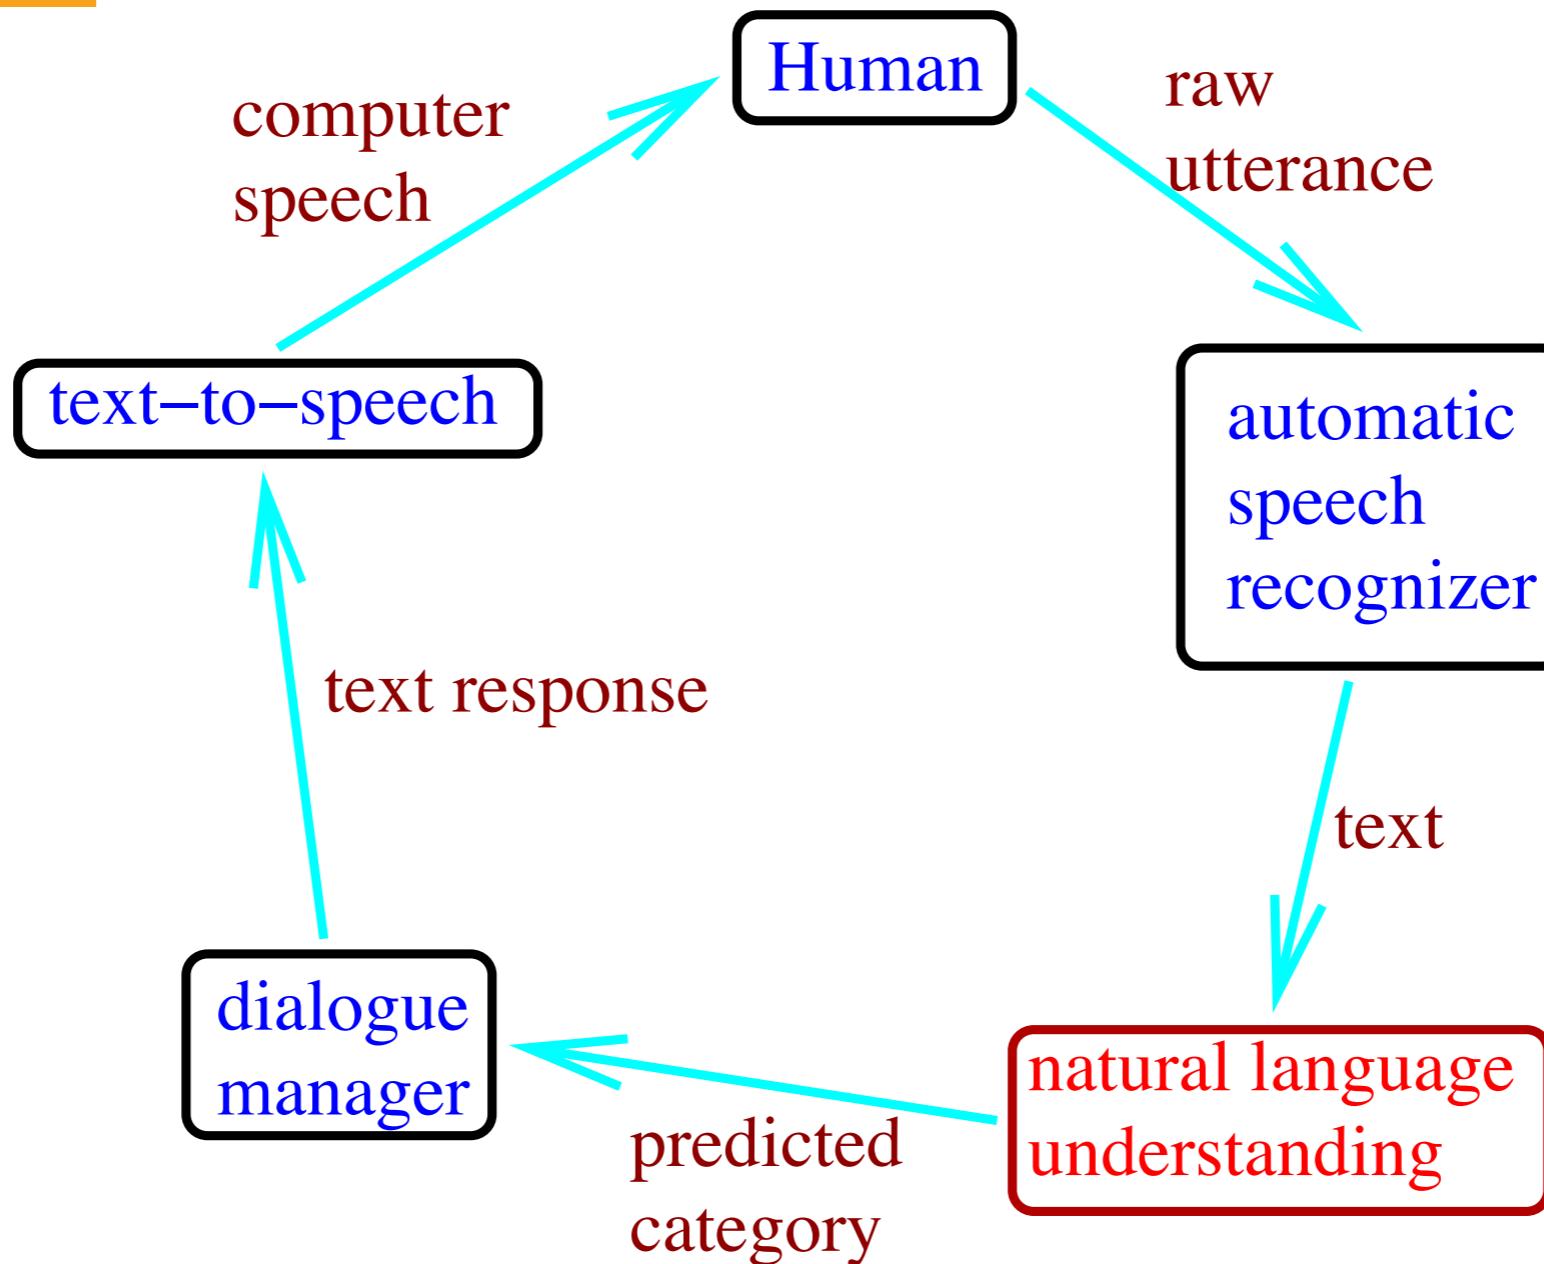
- many clever tricks to make extremely fast and accurate

Application: Human-computer Spoken Dialogue

[with Rahim, Di Fabrizio, Dutton, Gupta, Hollister & Riccardi]

- application: automatic “store front” or “help desk” for AT&T Labs’ Natural Voices business
- caller can request demo, pricing information, technical support, sales agent, etc.
- interactive dialogue

How It Works



- NLU's job: classify caller utterances into 24 categories (demo, sales rep, pricing info, yes, no, etc.)
- weak classifiers: test for presence of word or phrase

Problem: Labels are Expensive

- for spoken-dialogue task
 - getting examples is cheap
 - getting **labels** is expensive
 - must be annotated by humans
- how to reduce number of **labels** needed?

Active Learning

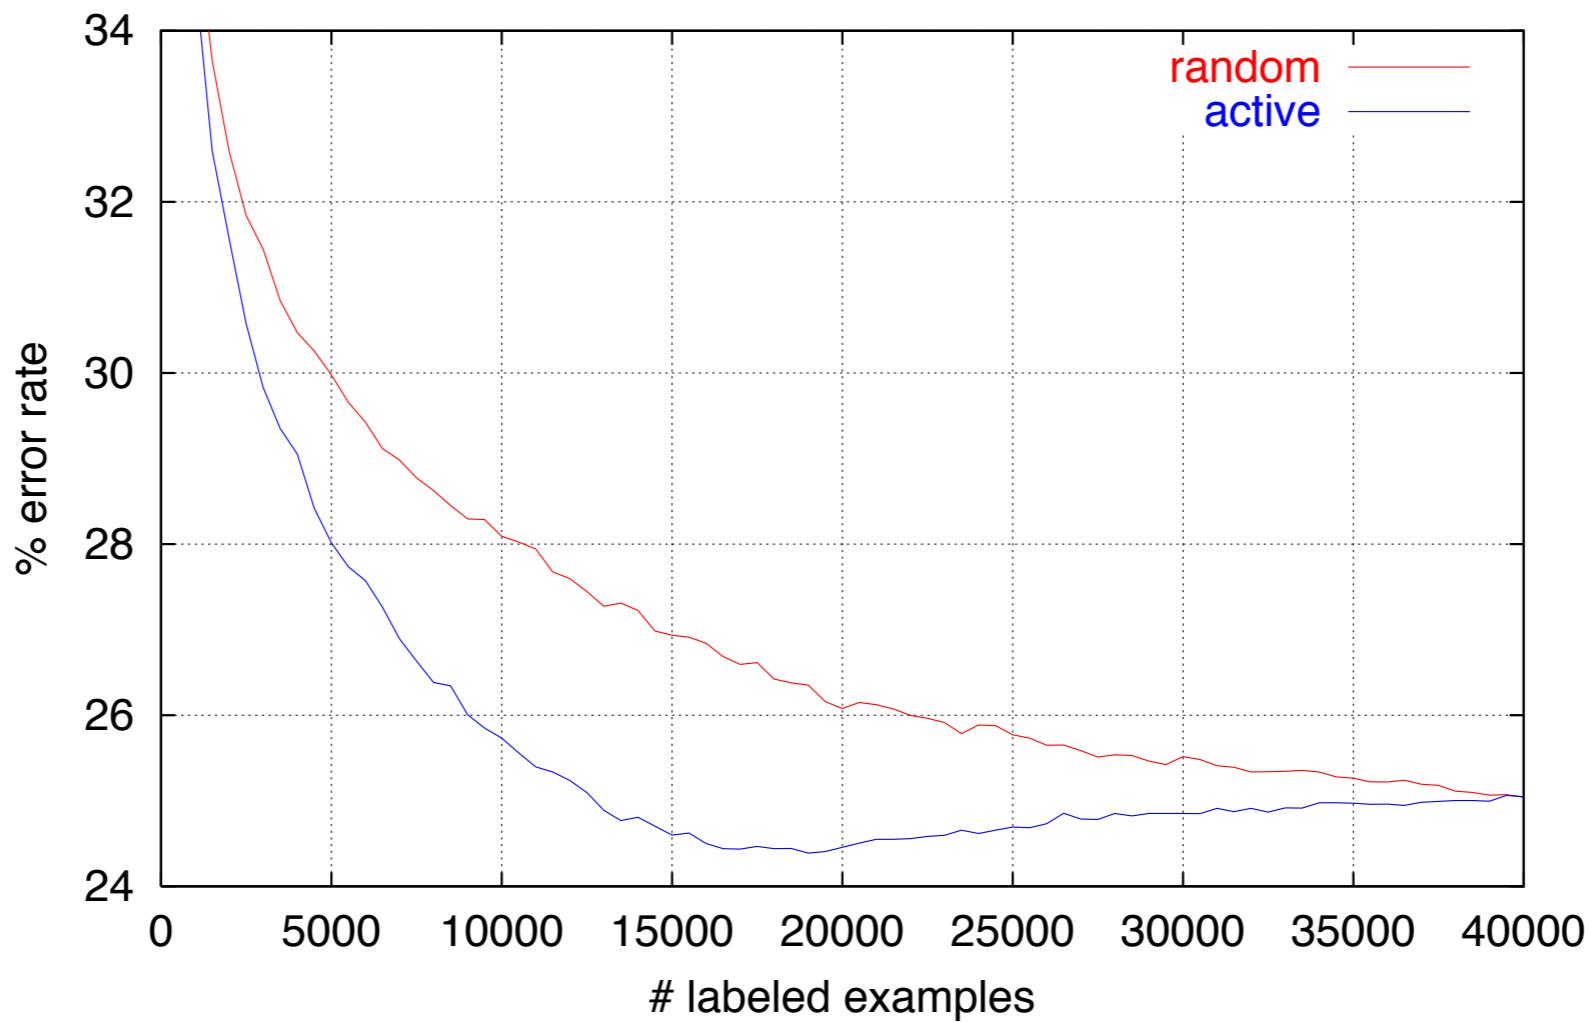
[with Tur & Hakkani-Tür]

- idea:
 - use **selective sampling** to choose which examples to label
 - focus on **least confident** examples [Lewis & Gale]
- for boosting, use (absolute) margin as natural confidence measure [Abe & Mamitsuka]

Labeling Scheme

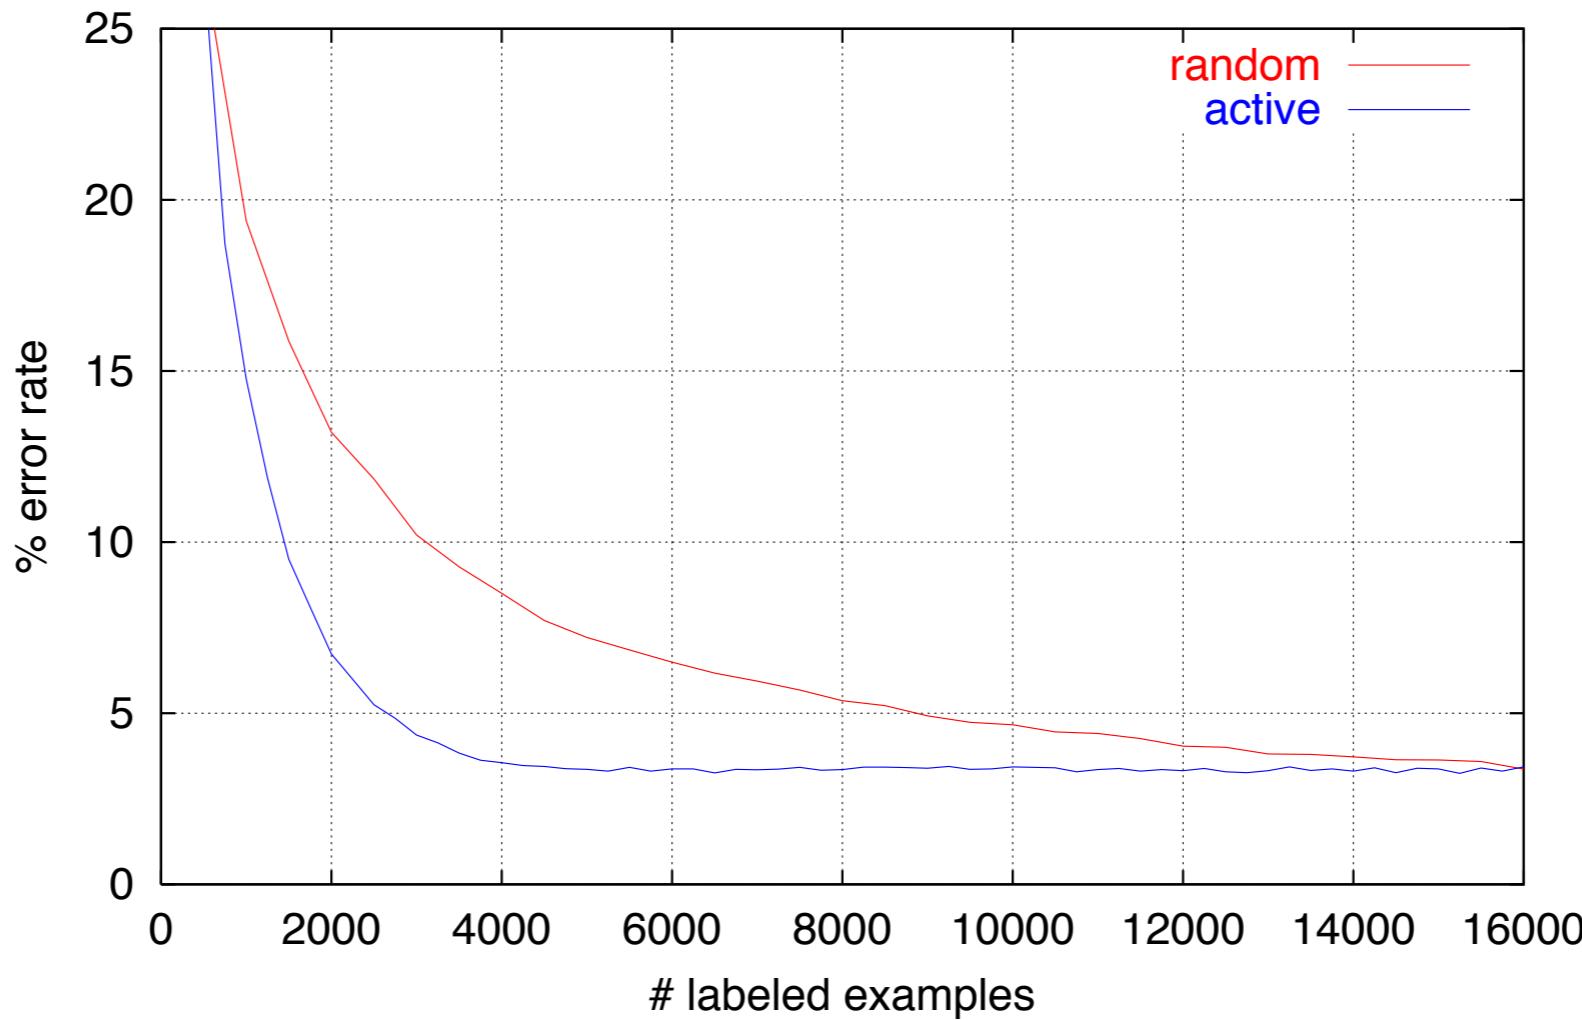
- start with pool of **unlabeled** examples
- choose (say) 500 examples at random for labeling
- run boosting on **all labeled** examples
 - get combined classifier F
- pick (say) 250 additional examples from pool for labeling
 - choose examples with minimum $|F(x)|$
(proportional to absolute margin)
- repeat

Results: How-May-I-Help-You?



% error	first reached		% label savings
	random	active	
28	11,000	5,500	50
26	22,000	9,500	57
25	40,000	13,000	68

Results: Letter



% error	first reached		% label savings
	random	active	
10	3,500	1,500	57
5	9,000	2,750	69
4	13,000	3,500	73

Fundamental Perspectives

- game theory
- loss minimization
- an information-geometric view

Fundamental Perspectives

- game theory
- loss minimization
- an information-geometric view

Just a Game

[with Freund]

- can view boosting as a **game**, a formal interaction between **booster** and **weak learner**
- on each round t :
 - booster chooses distribution D_t
 - weak learner responds with weak classifier h_t
- **game theory**: studies interactions between all sorts of “players”

Games

- game defined by matrix \mathbf{M} :

	Rock	Paper	Scissors
Rock	1/2	1	0
Paper	0	1/2	1
Scissors	1	0	1/2

- row player (“Mindy”) chooses row i
- column player (“Max”) chooses column j (simultaneously)
- Mindy’s goal: minimize her loss $\mathbf{M}(i, j)$
- assume (wlog) all entries in $[0, 1]$

Randomized Play

- usually allow randomized play:
 - Mindy chooses distribution \mathbf{P} over rows
 - Max chooses distribution \mathbf{Q} over columns (simultaneously)
- Mindy's (expected) loss

$$\begin{aligned} &= \sum_{i,j} \mathbf{P}(i) \mathbf{M}(i,j) \mathbf{Q}(j) \\ &= \mathbf{P}^\top \mathbf{M} \mathbf{Q} \equiv \mathbf{M}(\mathbf{P}, \mathbf{Q}) \end{aligned}$$

- i, j = “pure” strategies
- \mathbf{P}, \mathbf{Q} = “mixed” strategies
- m = # rows of \mathbf{M}
- also write $\mathbf{M}(i, \mathbf{Q})$ and $\mathbf{M}(\mathbf{P}, j)$ when one side plays pure and other plays mixed

Sequential Play

- say Mindy plays **before** Max
- if Mindy chooses **P** then Max will pick **Q** to maximize $M(P, Q)$ \Rightarrow loss will be

$$L(P) \equiv \max_Q M(P, Q)$$

- so Mindy should pick **P** to minimize $L(P)$
 \Rightarrow loss will be

$$\min_P L(P) = \min_P \max_Q M(P, Q)$$

- similarly, if **Max** plays first, loss will be

$$\max_Q \min_P M(P, Q)$$

Minmax Theorem

- playing **second** (with knowledge of other player's move) cannot be **worse** than playing **first**, so:

$$\underbrace{\min_P \max_Q M(P, Q)}_{\text{Mindy plays first}} \geq \underbrace{\max_Q \min_P M(P, Q)}_{\text{Mindy plays second}}$$

- von Neumann's minmax theorem:

$$\min_P \max_Q M(P, Q) = \max_Q \min_P M(P, Q)$$

- in words: **no** advantage to playing second

Optimal Play

- minmax theorem:

$$\min_P \max_Q M(P, Q) = \max_Q \min_P M(P, Q) = \text{value } v \text{ of game}$$

- optimal strategies:

- $P^* = \arg \min_P \max_Q M(P, Q)$ = minmax strategy
- $Q^* = \arg \max_Q \min_P M(P, Q)$ = maxmin strategy

- in words:

- Mindy's minmax strategy P^* guarantees loss $\leq v$ (regardless of Max's play)
- optimal because Max has maxmin strategy Q^* that can force loss $\geq v$ (regardless of Mindy's play)
- e.g.: in RPS, $P^* = Q^* = \text{uniform}$
- solving game = finding minmax/maxmin strategies

Weaknesses of Classical Theory

- seems to fully answer how to play games — just compute minmax strategy (e.g., using linear programming)
 - weaknesses:
 - game **M** may be **unknown**
 - game **M** may be **extremely large**
 - opponent may **not** be fully adversarial
 - may be possible to do **better** than value **v**
 - e.g.:
Lisa (thinks):
Poor predictable Bart, always takes Rock.
- Bart (thinks):**
- Good old Rock, nothing beats that.*

Repeated Play

- if only playing **once**, hopeless to overcome ignorance of game **M** or opponent
- but if game played **repeatedly**, may be possible to **learn** to play well
- **goal**: play (almost) as well as if **knew** game and how opponent would play ahead of time

Repeated Play (cont.)

- \mathbf{M} unknown
- for $t = 1, \dots, T$:
 - Mindy chooses \mathbf{P}_t
 - Max chooses \mathbf{Q}_t (possibly depending on \mathbf{P}_t)
 - Mindy's loss = $\mathbf{M}(\mathbf{P}_t, \mathbf{Q}_t)$
 - Mindy observes loss $\mathbf{M}(i, \mathbf{Q}_t)$ of each pure strategy i
- want:

$$\underbrace{\frac{1}{T} \sum_{t=1}^T \mathbf{M}(\mathbf{P}_t, \mathbf{Q}_t)}_{\text{actual average loss}} \leq \underbrace{\min_{\mathbf{P}} \frac{1}{T} \sum_{t=1}^T \mathbf{M}(\mathbf{P}, \mathbf{Q}_t)}_{\text{best loss (in hindsight)}} + [\text{"small amount"}]$$

Multiplicative-weights Algorithm (MW)

[with Freund]

- choose $\eta > 0$
- initialize: $\mathbf{P}_1 = \text{uniform}$
- on round t :

$$\mathbf{P}_{t+1}(i) = \frac{\mathbf{P}_t(i) \exp(-\eta \mathbf{M}(i, \mathbf{Q}_t))}{\text{normalization}}$$

- idea: decrease weight of strategies suffering the most loss
- directly generalizes [Littlestone & Warmuth]
- other algorithms:
 - [Hannan'57]
 - [Blackwell'56]
 - [Foster & Vohra]
 - [Fudenberg & Levine]
 - ⋮

Analysis

- **Theorem:** can choose η so that, for any game \mathbf{M} with m rows, and any opponent,

$$\underbrace{\frac{1}{T} \sum_{t=1}^T \mathbf{M}(\mathbf{P}_t, \mathbf{Q}_t)}_{\text{actual average loss}} \leq \underbrace{\min_{\mathbf{P}} \frac{1}{T} \sum_{t=1}^T \mathbf{M}(\mathbf{P}, \mathbf{Q}_t)}_{\text{best average loss } (\leq v)} + \Delta_T$$

where $\Delta_T = O\left(\sqrt{\frac{\ln m}{T}}\right) \rightarrow 0$

- regret Δ_T is:
 - logarithmic in # rows m
 - independent of # columns
- therefore, can use when working with very large games

Solving a Game

[with Freund]

- suppose game \mathbf{M} played repeatedly
 - Mindy plays using MW
 - on round t , Max chooses best response:

$$\mathbf{Q}_t = \arg \max_{\mathbf{Q}} \mathbf{M}(\mathbf{P}_t, \mathbf{Q})$$

- let

$$\bar{\mathbf{P}} = \frac{1}{T} \sum_{t=1}^T \mathbf{P}_t, \quad \bar{\mathbf{Q}} = \frac{1}{T} \sum_{t=1}^T \mathbf{Q}_t$$

- can prove that $\bar{\mathbf{P}}$ and $\bar{\mathbf{Q}}$ are Δ_T -approximate minmax and maxmin strategies:

$$\max_{\mathbf{Q}} \mathbf{M}(\bar{\mathbf{P}}, \mathbf{Q}) \leq v + \Delta_T$$

and

$$\min_{\mathbf{P}} \mathbf{M}(\mathbf{P}, \bar{\mathbf{Q}}) \geq v - \Delta_T$$

Boosting as a Game

- Mindy (row player) \leftrightarrow booster
- Max (column player) \leftrightarrow weak learner
- matrix \mathbf{M} :
 - row \leftrightarrow training example
 - column \leftrightarrow weak classifier
 - $\mathbf{M}(i, j) = \begin{cases} 1 & \text{if } j\text{-th weak classifier correct on } i\text{-th training example} \\ 0 & \text{else} \end{cases}$
 - encodes which weak classifiers correct on which examples
 - huge # of columns — one for every possible weak classifier

Boosting and the Minmax Theorem

- γ -weak learning assumption:
 - for every distribution on examples
 - can find weak classifier with weighted error $\leq \frac{1}{2} - \gamma$

- equivalent to:

$$(\text{value of game } \mathbf{M}) \geq \frac{1}{2} + \gamma$$

- by minmax theorem, implies that:
 - \exists some weighted majority classifier that correctly classifies all training examples with margin $\geq 2\gamma$
 - further, weights are given by maxmin strategy of game \mathbf{M}

Idea for Boosting

- maxmin strategy of \mathbf{M} has perfect (training) accuracy and large margins
- find approximately using earlier algorithm for solving a game
 - i.e., apply MW to \mathbf{M}
- yields (variant of) AdaBoost

AdaBoost and Game Theory

- summarizing:
 - weak learning assumption implies maxmin strategy for \mathbf{M} defines large-margin classifier
 - AdaBoost finds maxmin strategy by applying general algorithm for solving games through repeated play
- consequences:
 - weights on weak classifiers converge to (approximately) maxmin strategy for game \mathbf{M}
 - (average) of distributions D_t converges to (approximately) minmax strategy
 - margins and edges connected via minmax theorem
 - explains why AdaBoost maximizes margins
- different instantiation of game-playing algorithm gives online learning algorithms (such as weighted majority algorithm)

Fundamental Perspectives

- game theory
- loss minimization
- an information-geometric view

AdaBoost and Loss Minimization

- many (most?) learning and statistical methods can be viewed as minimizing **loss** (a.k.a. cost or objective) function measuring **fit to data**:
 - e.g. least squares regression $\sum_i(F(x_i) - y_i)^2$
- AdaBoost also minimizes a loss function
- **helpful** to understand because:
 - clarifies goal of algorithm and useful in proving convergence properties
 - decoupling of algorithm from its objective means:
 - faster algorithms possible for same objective
 - same algorithm may generalize for new learning challenges

What AdaBoost Minimizes

- recall proof of training error bound:
 - training error(H_{final}) $\leq \prod_t Z_t$
 - $Z_t = \epsilon_t e^{\alpha_t} + (1 - \epsilon_t) e^{-\alpha_t} = 2\sqrt{\epsilon_t(1 - \epsilon_t)}$
- closer look:
 - α_t chosen to minimize Z_t
 - h_t chosen to minimize ϵ_t
 - same as minimizing Z_t
(since increasing in ϵ_t on $[0, 1/2]$)
- so: both AdaBoost and weak learner minimize Z_t on round t
 - equivalent to **greedily** minimizing $\prod_t Z_t$

AdaBoost and Exponential Loss

- so AdaBoost is greedy procedure for minimizing exponential loss

$$\prod_t Z_t = \frac{1}{m} \sum_i \exp(-y_i F(x_i))$$

where

$$F(x) = \sum_t \alpha_t h_t(x)$$

- why exponential loss?
 - intuitively, strongly favors $F(x_i)$ to have same sign as y_i
 - upper bound on training error
 - smooth and convex (but very loose)
- how does AdaBoost minimize it?

Coordinate Descent

[Breiman]

- $\{g_1, \dots, g_N\}$ = space of all weak classifiers

- then can write $F(x) = \sum_t \alpha_t h_t(x) = \sum_{j=1}^N \lambda_j g_j(x)$

- want to find $\lambda_1, \dots, \lambda_N$ to minimize

$$L(\lambda_1, \dots, \lambda_N) = \sum_i \exp \left(-y_i \sum_j \lambda_j g_j(x_i) \right)$$

- AdaBoost is actually doing coordinate descent on this optimization problem:
 - initially, all $\lambda_j = 0$
 - each round: choose one coordinate λ_j (corresponding to h_t) and update (increment by α_t)
 - choose update causing biggest decrease in loss
- powerful technique for minimizing over huge space of functions

Functional Gradient Descent

[Mason et al.][Friedman]

- want to minimize

$$\mathcal{L}(F) = \mathcal{L}(F(x_1), \dots, F(x_m)) = \sum_i \exp(-y_i F(x_i))$$

- say have current estimate F and want to improve
- to do gradient descent, would like update

$$F \leftarrow F - \alpha \nabla_F \mathcal{L}(F)$$

- but update restricted in class of weak classifiers

$$F \leftarrow F + \alpha h_t$$

- so choose h_t “closest” to $-\nabla_F \mathcal{L}(F)$
- equivalent to AdaBoost

Estimating Conditional Probabilities

[Friedman, Hastie & Tibshirani]

- often want to estimate probability that $y = +1$ given x
- AdaBoost minimizes (empirical version of):

$$E_{x,y} \left[e^{-yF(x)} \right] = E_x \left[\Pr [y = +1|x] e^{-F(x)} + \Pr [y = -1|x] e^{F(x)} \right]$$

where x, y random from true distribution

- over all F , minimized when

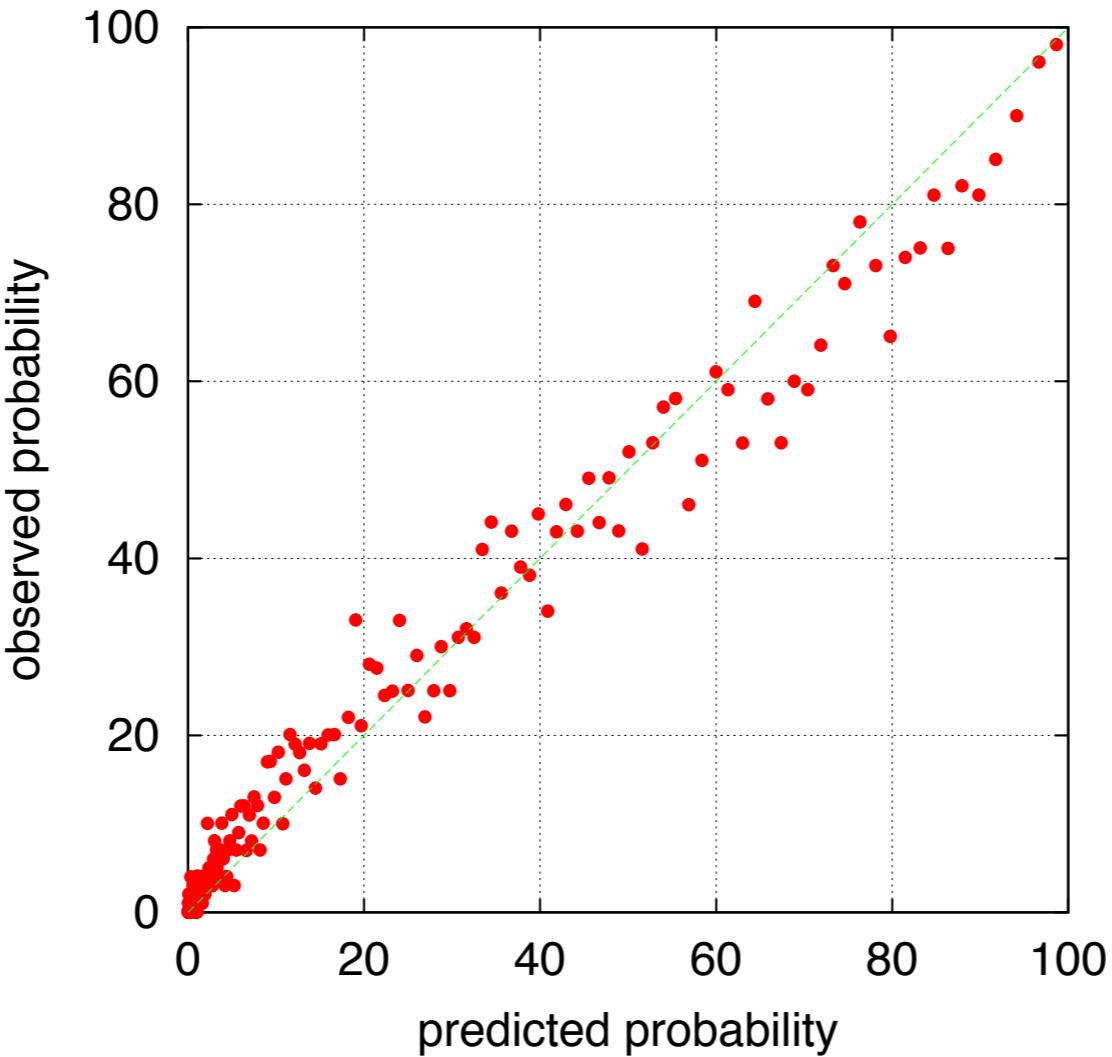
$$F(x) = \frac{1}{2} \cdot \ln \left(\frac{\Pr [y = +1|x]}{\Pr [y = -1|x]} \right)$$

or

$$\Pr [y = +1|x] = \frac{1}{1 + e^{-2F(x)}}$$

- so, to convert F output by AdaBoost to probability estimate, use same formula

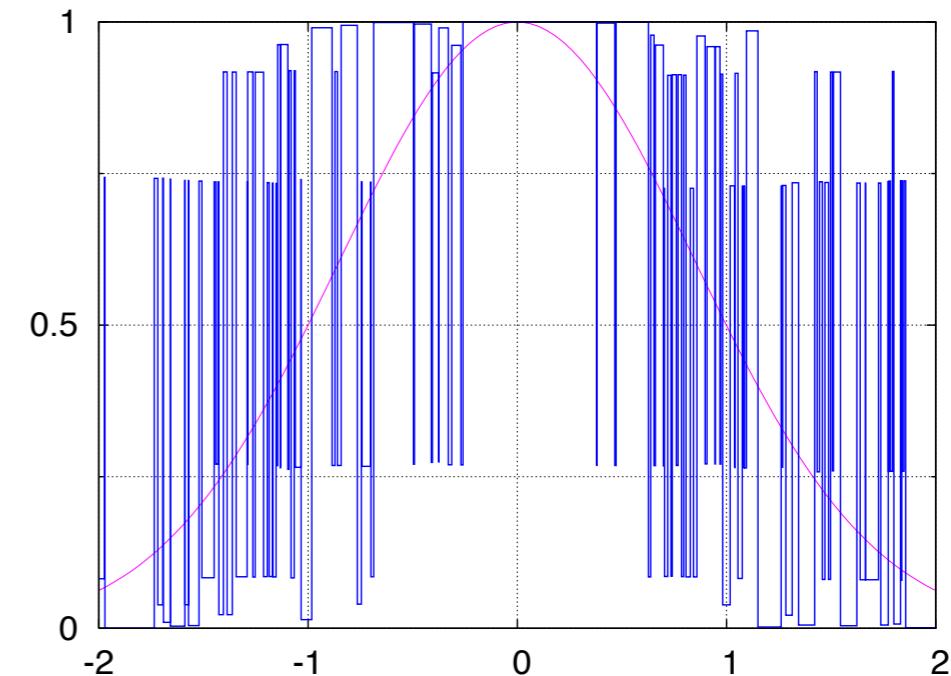
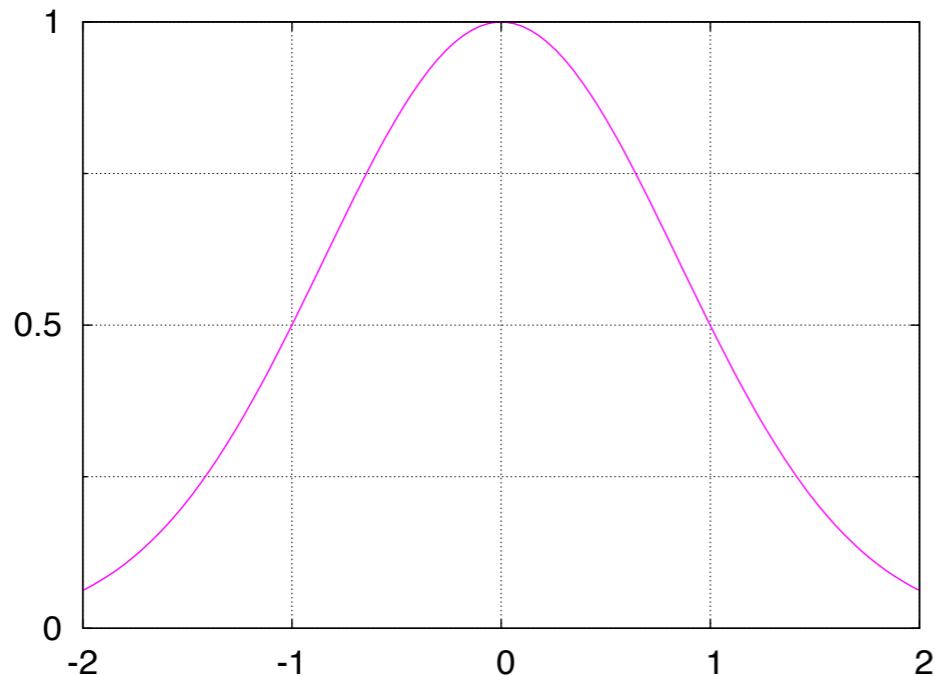
Calibration Curve



- order examples by F value output by AdaBoost
- break into bins of fixed size
- for each bin, plot a point:
 - x -value: average estimated probability of examples in bin
 - y -value: actual fraction of positive examples in bin

A Synthetic Example

- $x \in [-2, +2]$ uniform
- $\Pr [y = +1|x] = 2^{-x^2}$
- $m = 500$ training examples



- if run AdaBoost with stumps and convert to probabilities, result is poor
 - extreme **overfitting**

Regularization

- AdaBoost minimizes

$$L(\boldsymbol{\lambda}) = \sum_i \exp \left(-y_i \sum_j \lambda_j g_j(x_i) \right)$$

- to avoid overfitting, want to constrain $\boldsymbol{\lambda}$ to make solution “smoother”
- (ℓ_1) regularization:

minimize: $L(\boldsymbol{\lambda})$

subject to: $\|\boldsymbol{\lambda}\|_1 \leq B$

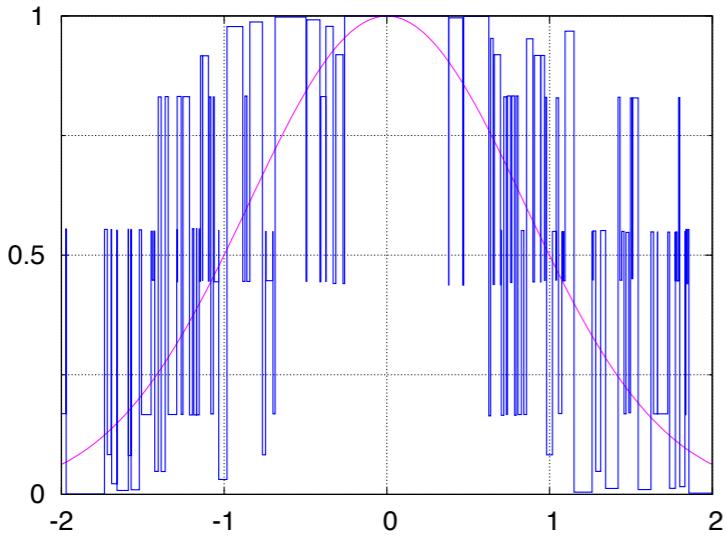
- or:

minimize: $L(\boldsymbol{\lambda}) + \beta \|\boldsymbol{\lambda}\|_1$

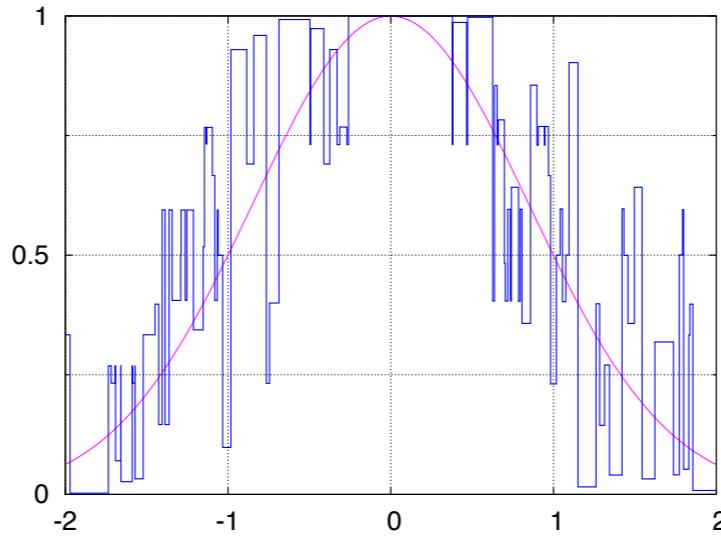
- other norms possible
 - ℓ_1 (“lasso”) currently popular since encourages sparsity

[Tibshirani]

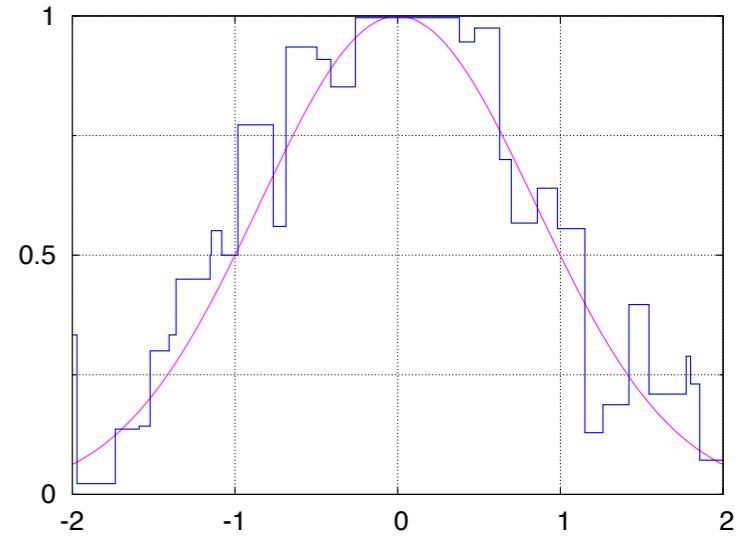
Regularization Example



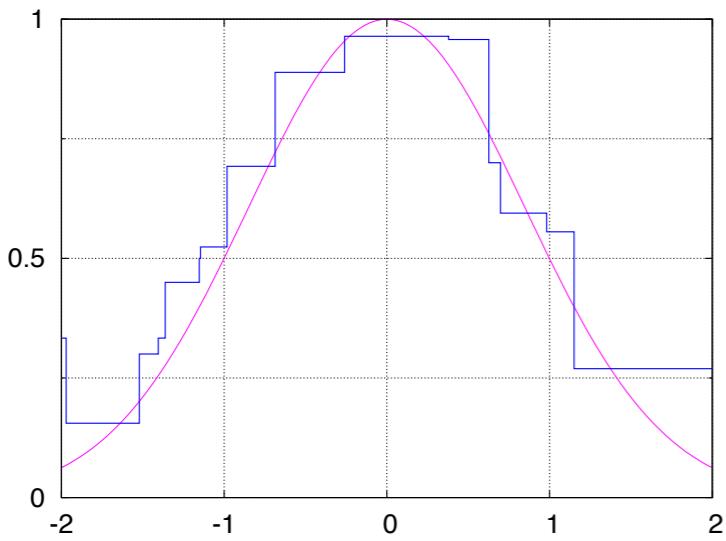
$$\beta = 10^{-3}$$



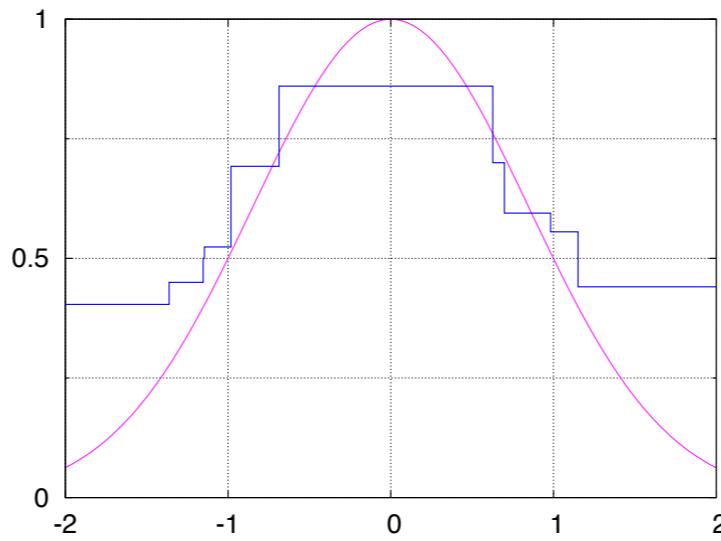
$$\beta = 10^{-2.5}$$



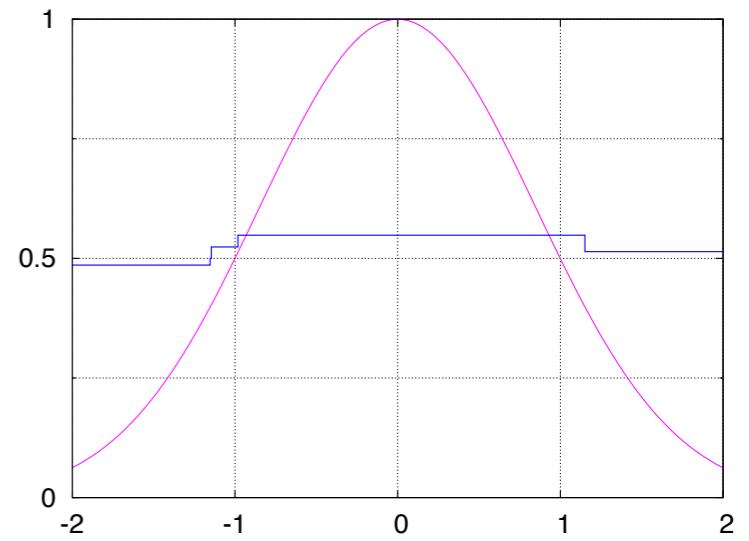
$$\beta = 10^{-2}$$



$$\beta = 10^{-1.5}$$



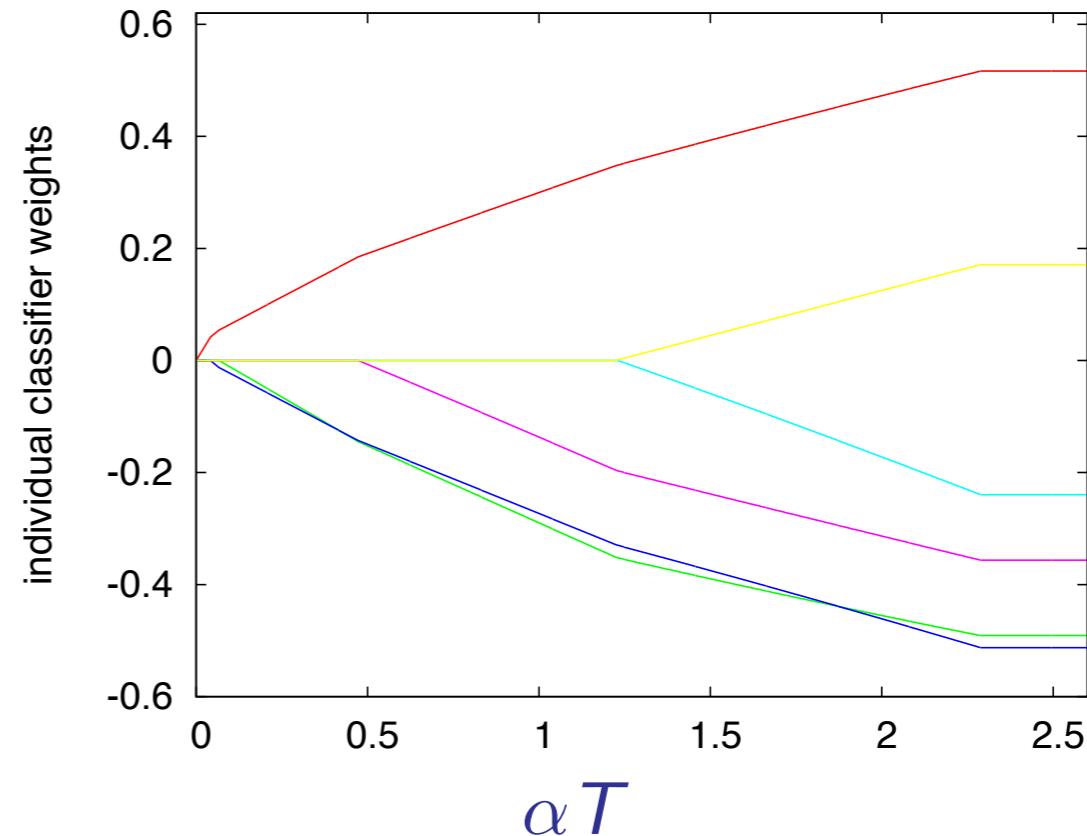
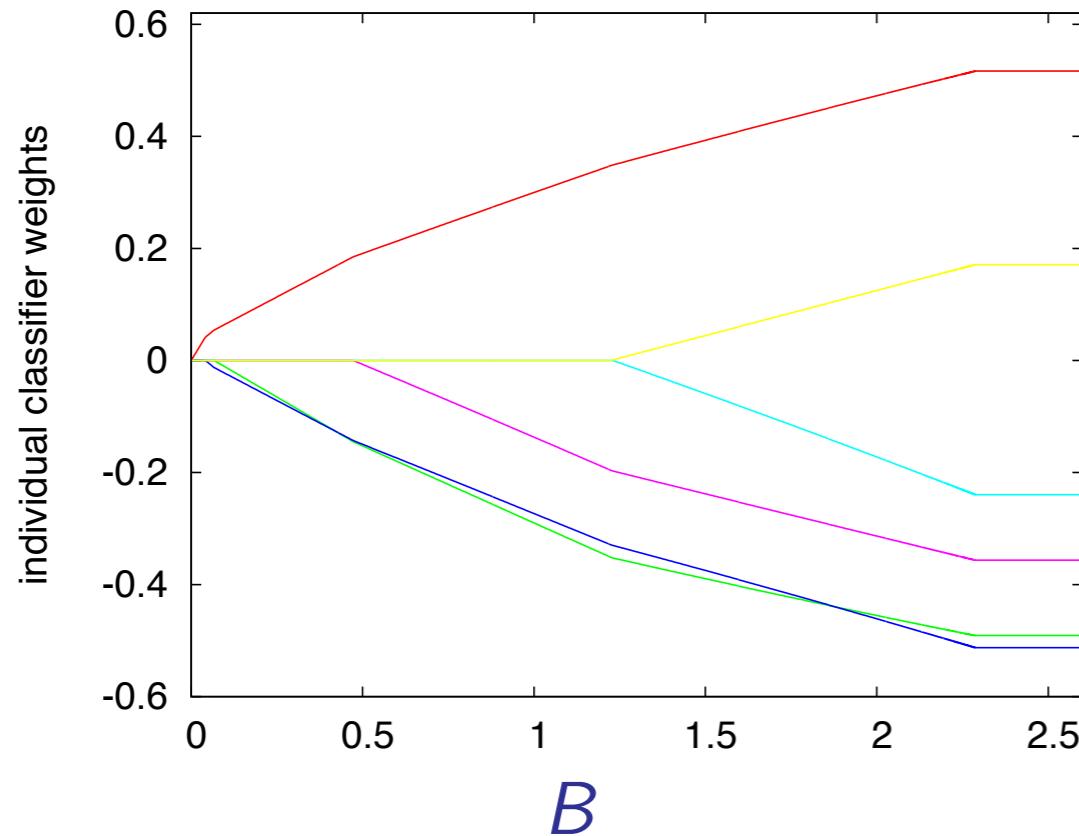
$$\beta = 10^{-1}$$



$$\beta = 10^{-0.5}$$

Regularization and AdaBoost

[Hastie, Tibshirani & Friedman; Rosset, Zhu & Hastie]



- **Experiment 1:** regularized solution vectors λ plotted as function of B
- plots are identical!
- can prove under certain (but not all) conditions that results will be the same (as $\alpha \rightarrow 0$)
- **Experiment 2:** AdaBoost run with α_t fixed to (small) α
 - solution vectors λ plotted as function of αT

[Zhao & Yu]

Regularization and AdaBoost

- suggests stopping AdaBoost early is akin to applying ℓ_1 -regularization
- caveats:
 - does **not** strictly apply to AdaBoost (only variant)
 - not helpful when boosting run “**to convergence**” (would correspond to very weak regularization)
- in fact, in limit of vanishingly **weak** regularization ($B \rightarrow \infty$), solution converges to maximum margin solution

[Rosset, Zhu & Hastie]

Benefits of Loss-Minimization View

- immediate generalization to other loss functions and learning problems
 - e.g. squared error for regression
 - e.g. logistic regression
(by only changing one line of AdaBoost)
- sensible approach for converting output of boosting into conditional probability estimates
- helpful connection to regularization
- basis for proving AdaBoost is statistically “consistent”
 - i.e., under right assumptions, converges to best possible classifier

[Bartlett & Traskin]

A Note of Caution

- tempting (but **incorrect!**) to conclude:
 - AdaBoost is **just** an algorithm for minimizing exponential loss
 - AdaBoost works **only because** of its loss function
∴ more powerful optimization techniques for same loss should work even better
- **incorrect** because:
 - **other** algorithms that minimize exponential loss can give very **poor** generalization performance compared to AdaBoost
- for example...

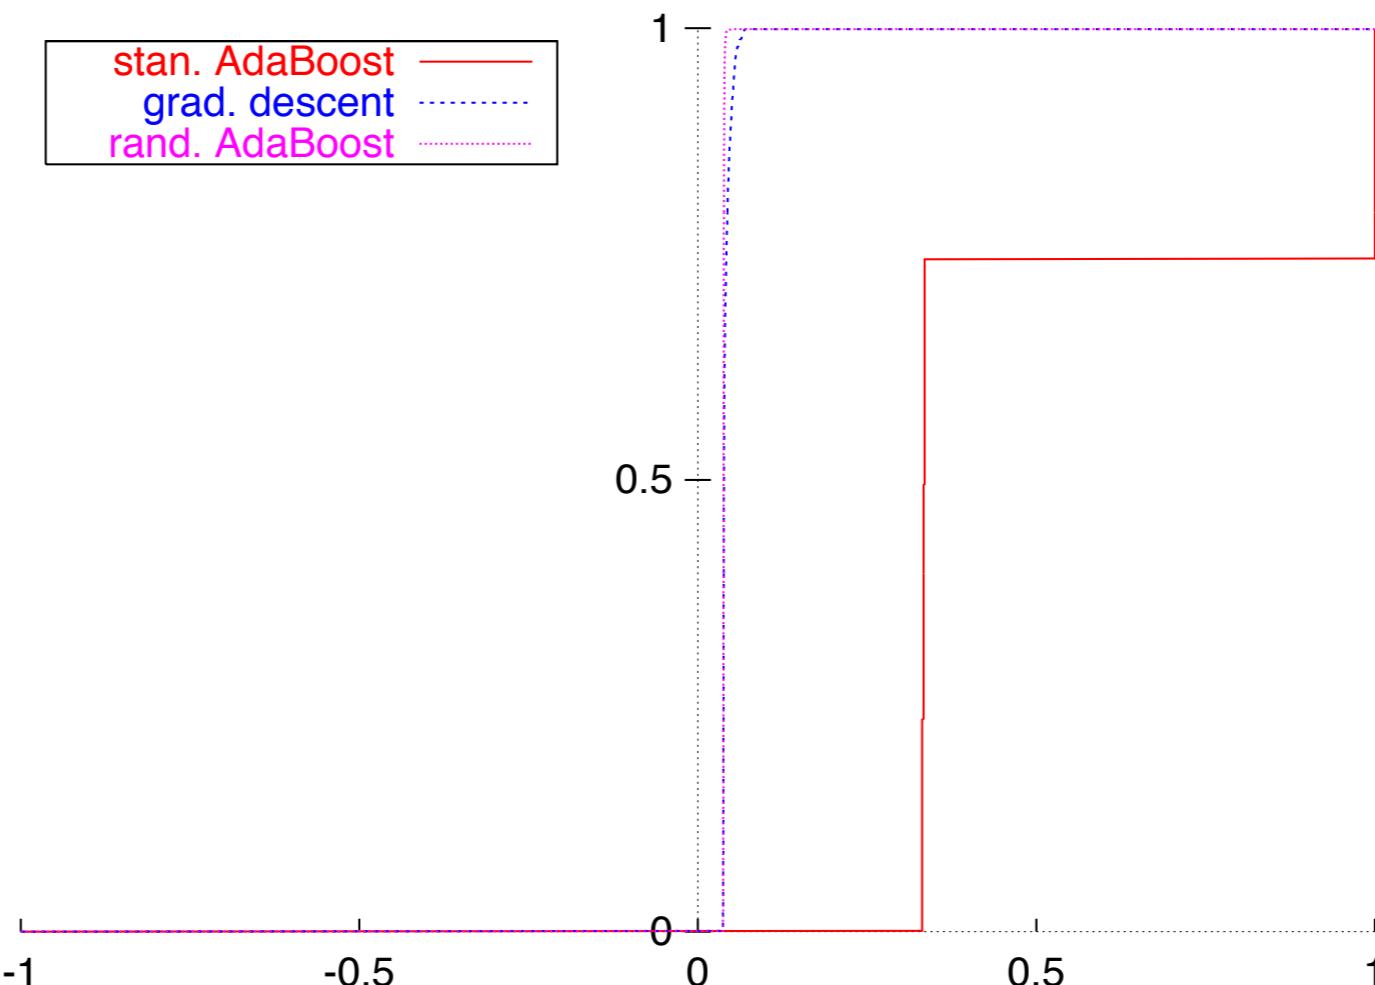
An Experiment

- data:
 - instances x uniform from $\{-1, +1\}^{10,000}$
 - label y = majority vote of three coordinates
 - weak classifier = single coordinate (or its negation)
 - training set size $m = 1000$
- algorithms (all provably minimize exponential loss):
 - standard AdaBoost
 - gradient descent on exponential loss
 - AdaBoost, but in which weak classifiers chosen at random
- results:

exp. loss	% test error [# rounds]				
	stand. AdaB.	grad. desc.	random AdaB.		
10^{-10}	0.0 [94]	40.7 [5]	44.0	[24,464]	
10^{-20}	0.0 [190]	40.8 [9]	41.6	[47,534]	
10^{-40}	0.0 [382]	40.8 [21]	40.9	[94,479]	
10^{-100}	0.0 [956]	40.8 [70]	40.3	[234,654]	

An Experiment (cont.)

- conclusions:
 - not just **what** is being minimized that matters, but **how** it is being minimized
 - loss-minimization view has benefits and is fundamental to understanding AdaBoost
 - but is **limited** in what it says about generalization
- results **are** consistent with **margins theory**



Fundamental Perspectives

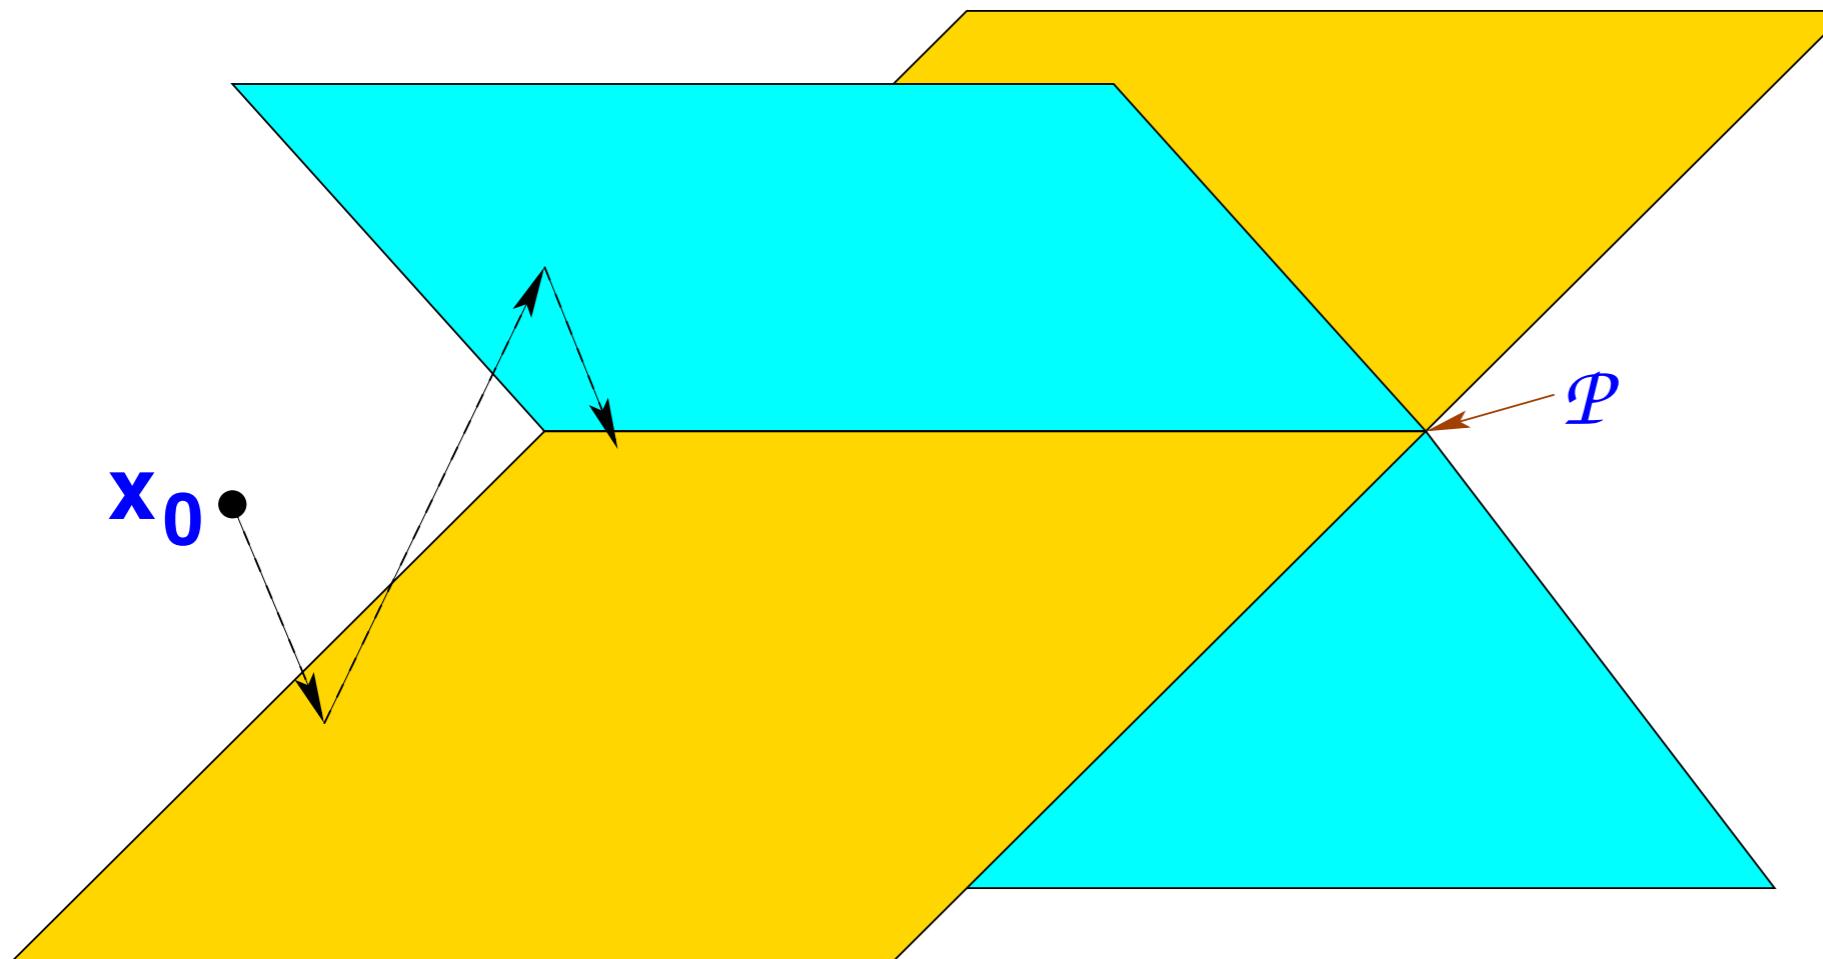
- game theory
- loss minimization
- an information-geometric view

A Dual Information-Geometric Perspective

- loss minimization focuses on **function** computed by AdaBoost
(i.e., weights on **weak classifiers**)
- **dual view**: instead focus on **distributions** D_t
(i.e., weights on **examples**)
- dual perspective combines **geometry** and information theory
- exposes underlying mathematical **structure**
- basis for proving **convergence**

An Iterative-Projection Algorithm

- say want to find point closest to \mathbf{x}_0 in set
 $\mathcal{P} = \{ \text{intersection of } N \text{ hyperplanes} \}$
- algorithm: [Bregman; Censor & Zenios]
 - start at \mathbf{x}_0
 - repeat: pick a hyperplane and project onto it



- if $\mathcal{P} \neq \emptyset$, under general conditions, will converge correctly

AdaBoost is an Iterative-Projection Algorithm

[Kivinen & Warmuth]

- points = distributions D_t over training examples
- distance = relative entropy:

$$\text{RE}(P \parallel Q) = \sum_i P(i) \ln \left(\frac{P(i)}{Q(i)} \right)$$

- reference point \mathbf{x}_0 = uniform distribution
- hyperplanes defined by all possible weak classifiers g_j :

$$\sum_i D(i) y_i g_j(x_i) = 0 \Leftrightarrow \Pr_{i \sim D} [g_j(x_i) \neq y_i] = \frac{1}{2}$$

- intuition: looking for “hardest” distribution

AdaBoost as Iterative Projection (cont.)

- algorithm:
 - start at $D_1 = \text{uniform}$
 - for $t = 1, 2, \dots$:
 - pick hyperplane/weak classifier $h_t \leftrightarrow g_j$
 - $D_{t+1} = (\text{entropy}) \text{ projection of } D_t \text{ onto hyperplane}$
 $= \arg \min_{D: \sum_i D(i) y_i g_j(x_i) = 0} \text{RE}(D \parallel D_t)$
- claim: equivalent to AdaBoost
- further: choosing h_t with minimum error \equiv choosing farthest hyperplane

Boosting as Maximum Entropy

- corresponding optimization problem:

$$\min_{D \in \mathcal{P}} \text{RE}(D \parallel \text{uniform}) \leftrightarrow \max_{D \in \mathcal{P}} \text{entropy}(D)$$

- where

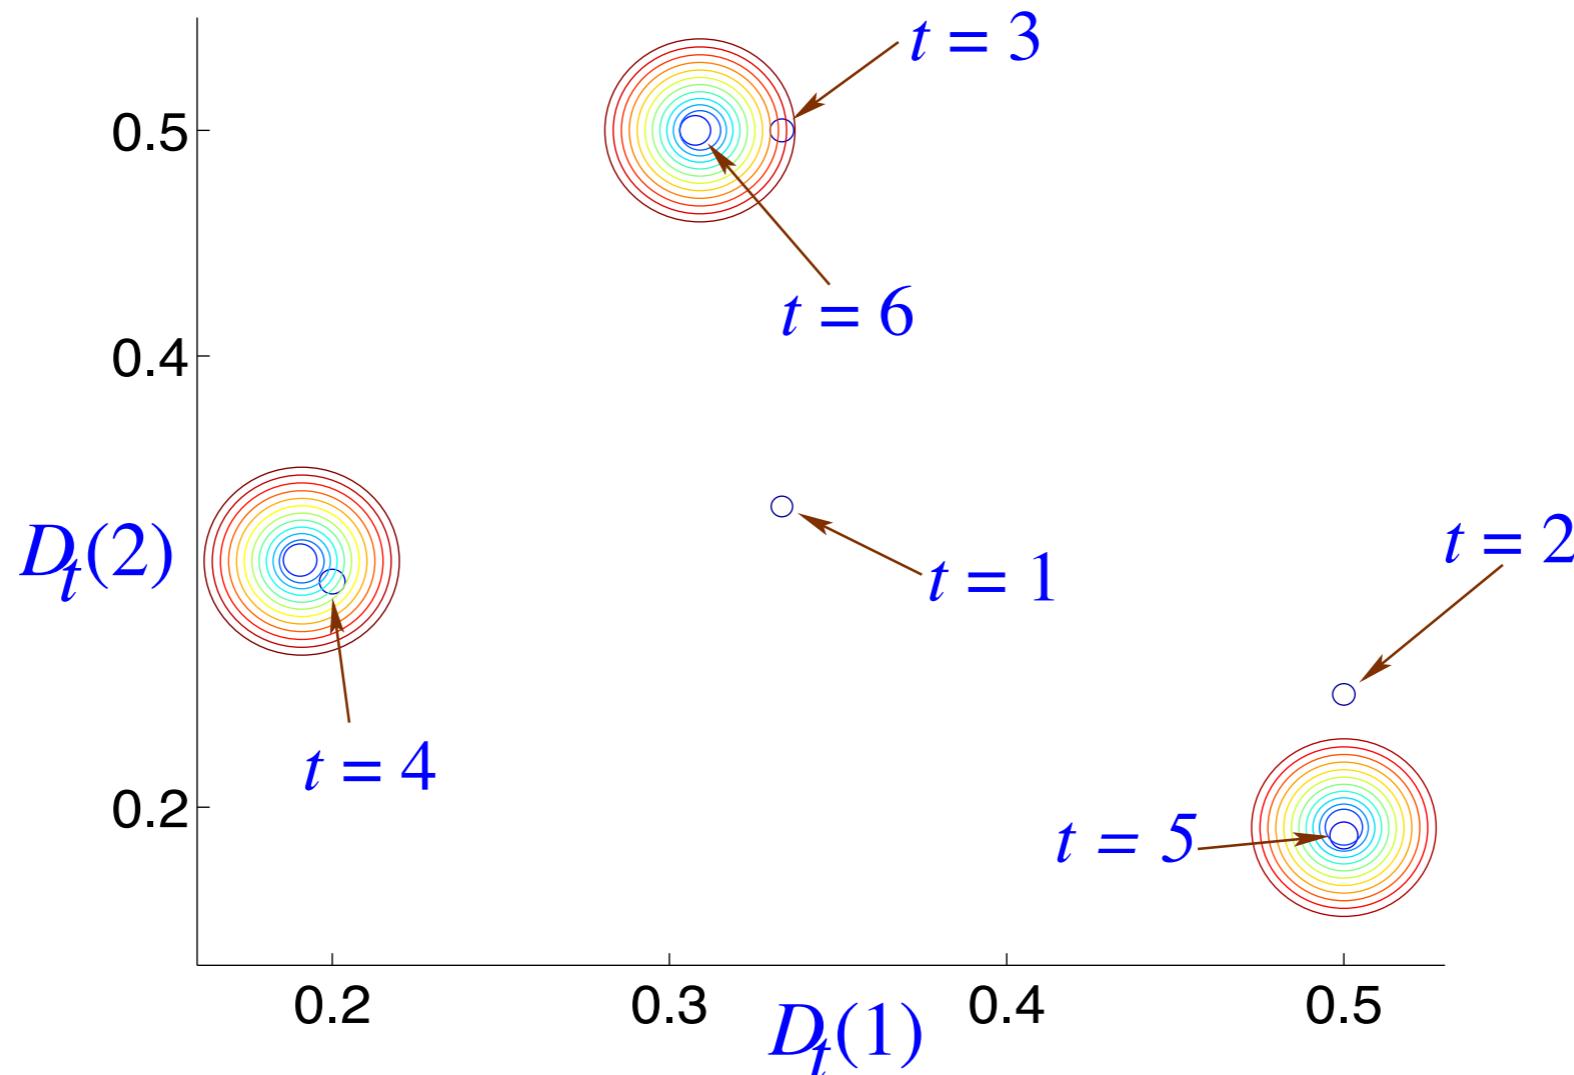
$$\begin{aligned}\mathcal{P} &= \text{feasible set} \\ &= \left\{ D : \sum_i D(i) y_i g_j(x_i) = 0 \quad \forall j \right\}\end{aligned}$$

- $\mathcal{P} \neq \emptyset \Leftrightarrow$ weak learning assumption does **not** hold
 - in this case, $D_t \rightarrow$ (unique) solution
- if weak learning assumption **does** hold then
 - $\mathcal{P} = \emptyset$
 - D_t can **never** converge
 - dynamics are fascinating but unclear in this case

Visualizing Dynamics

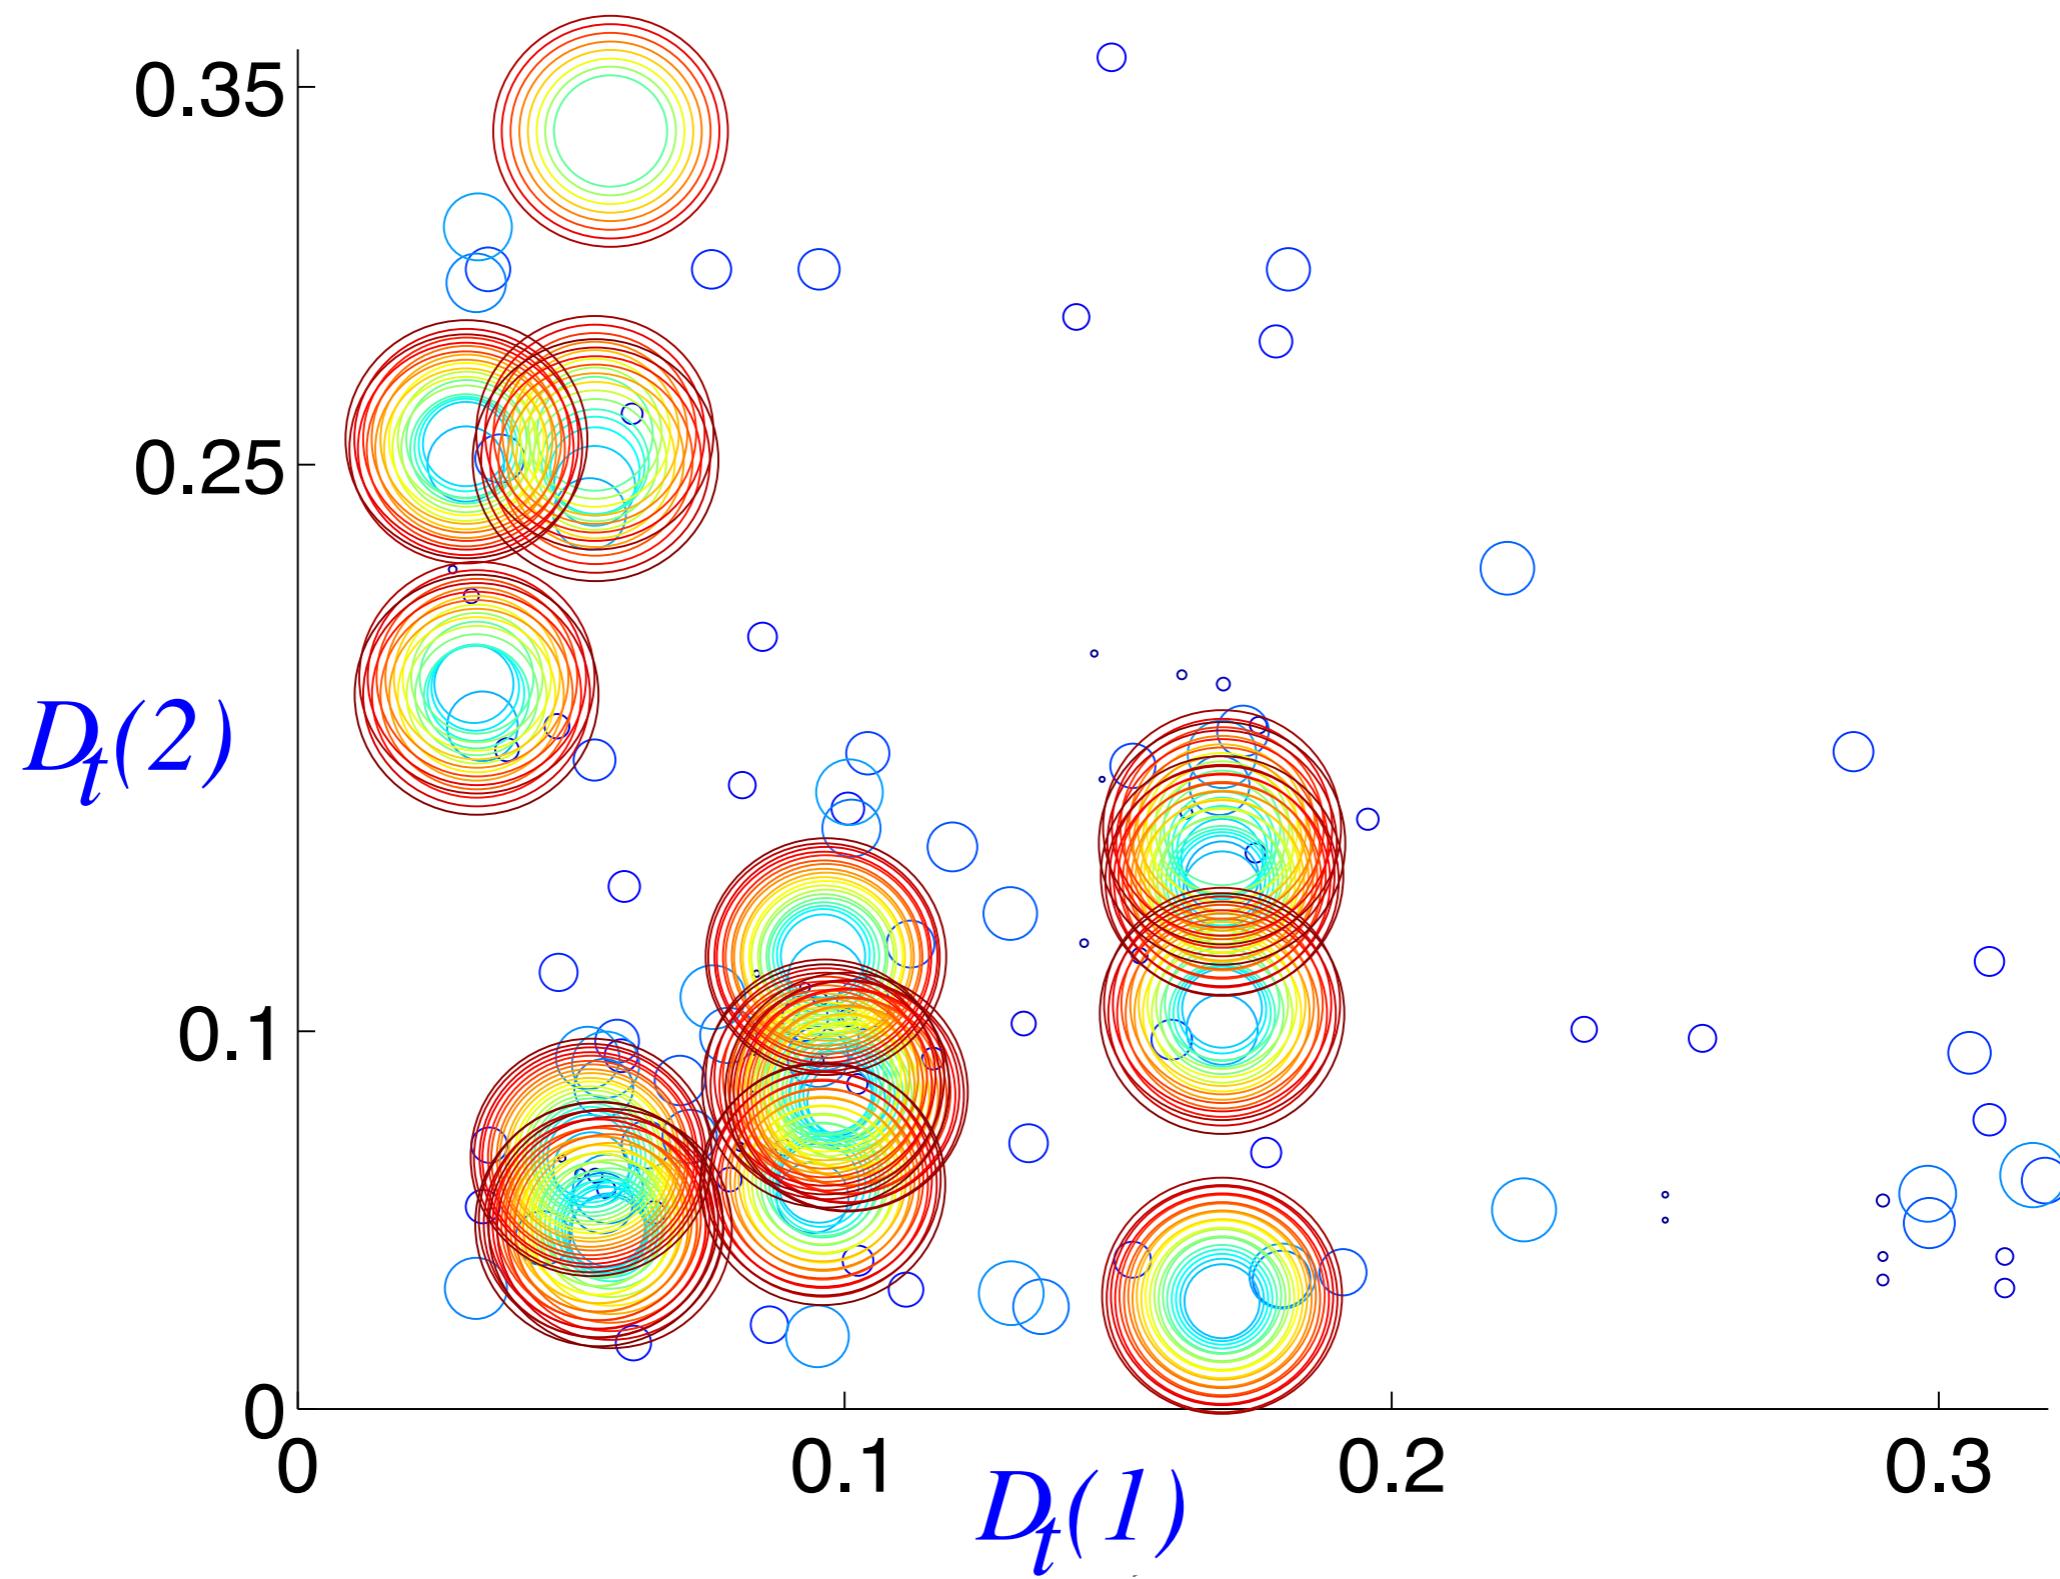
[with Rudin & Daubechies]

- plot one circle for each round t :
 - center at $(D_t(1), D_t(2))$
 - radius $\propto t$ (color also varies with t)

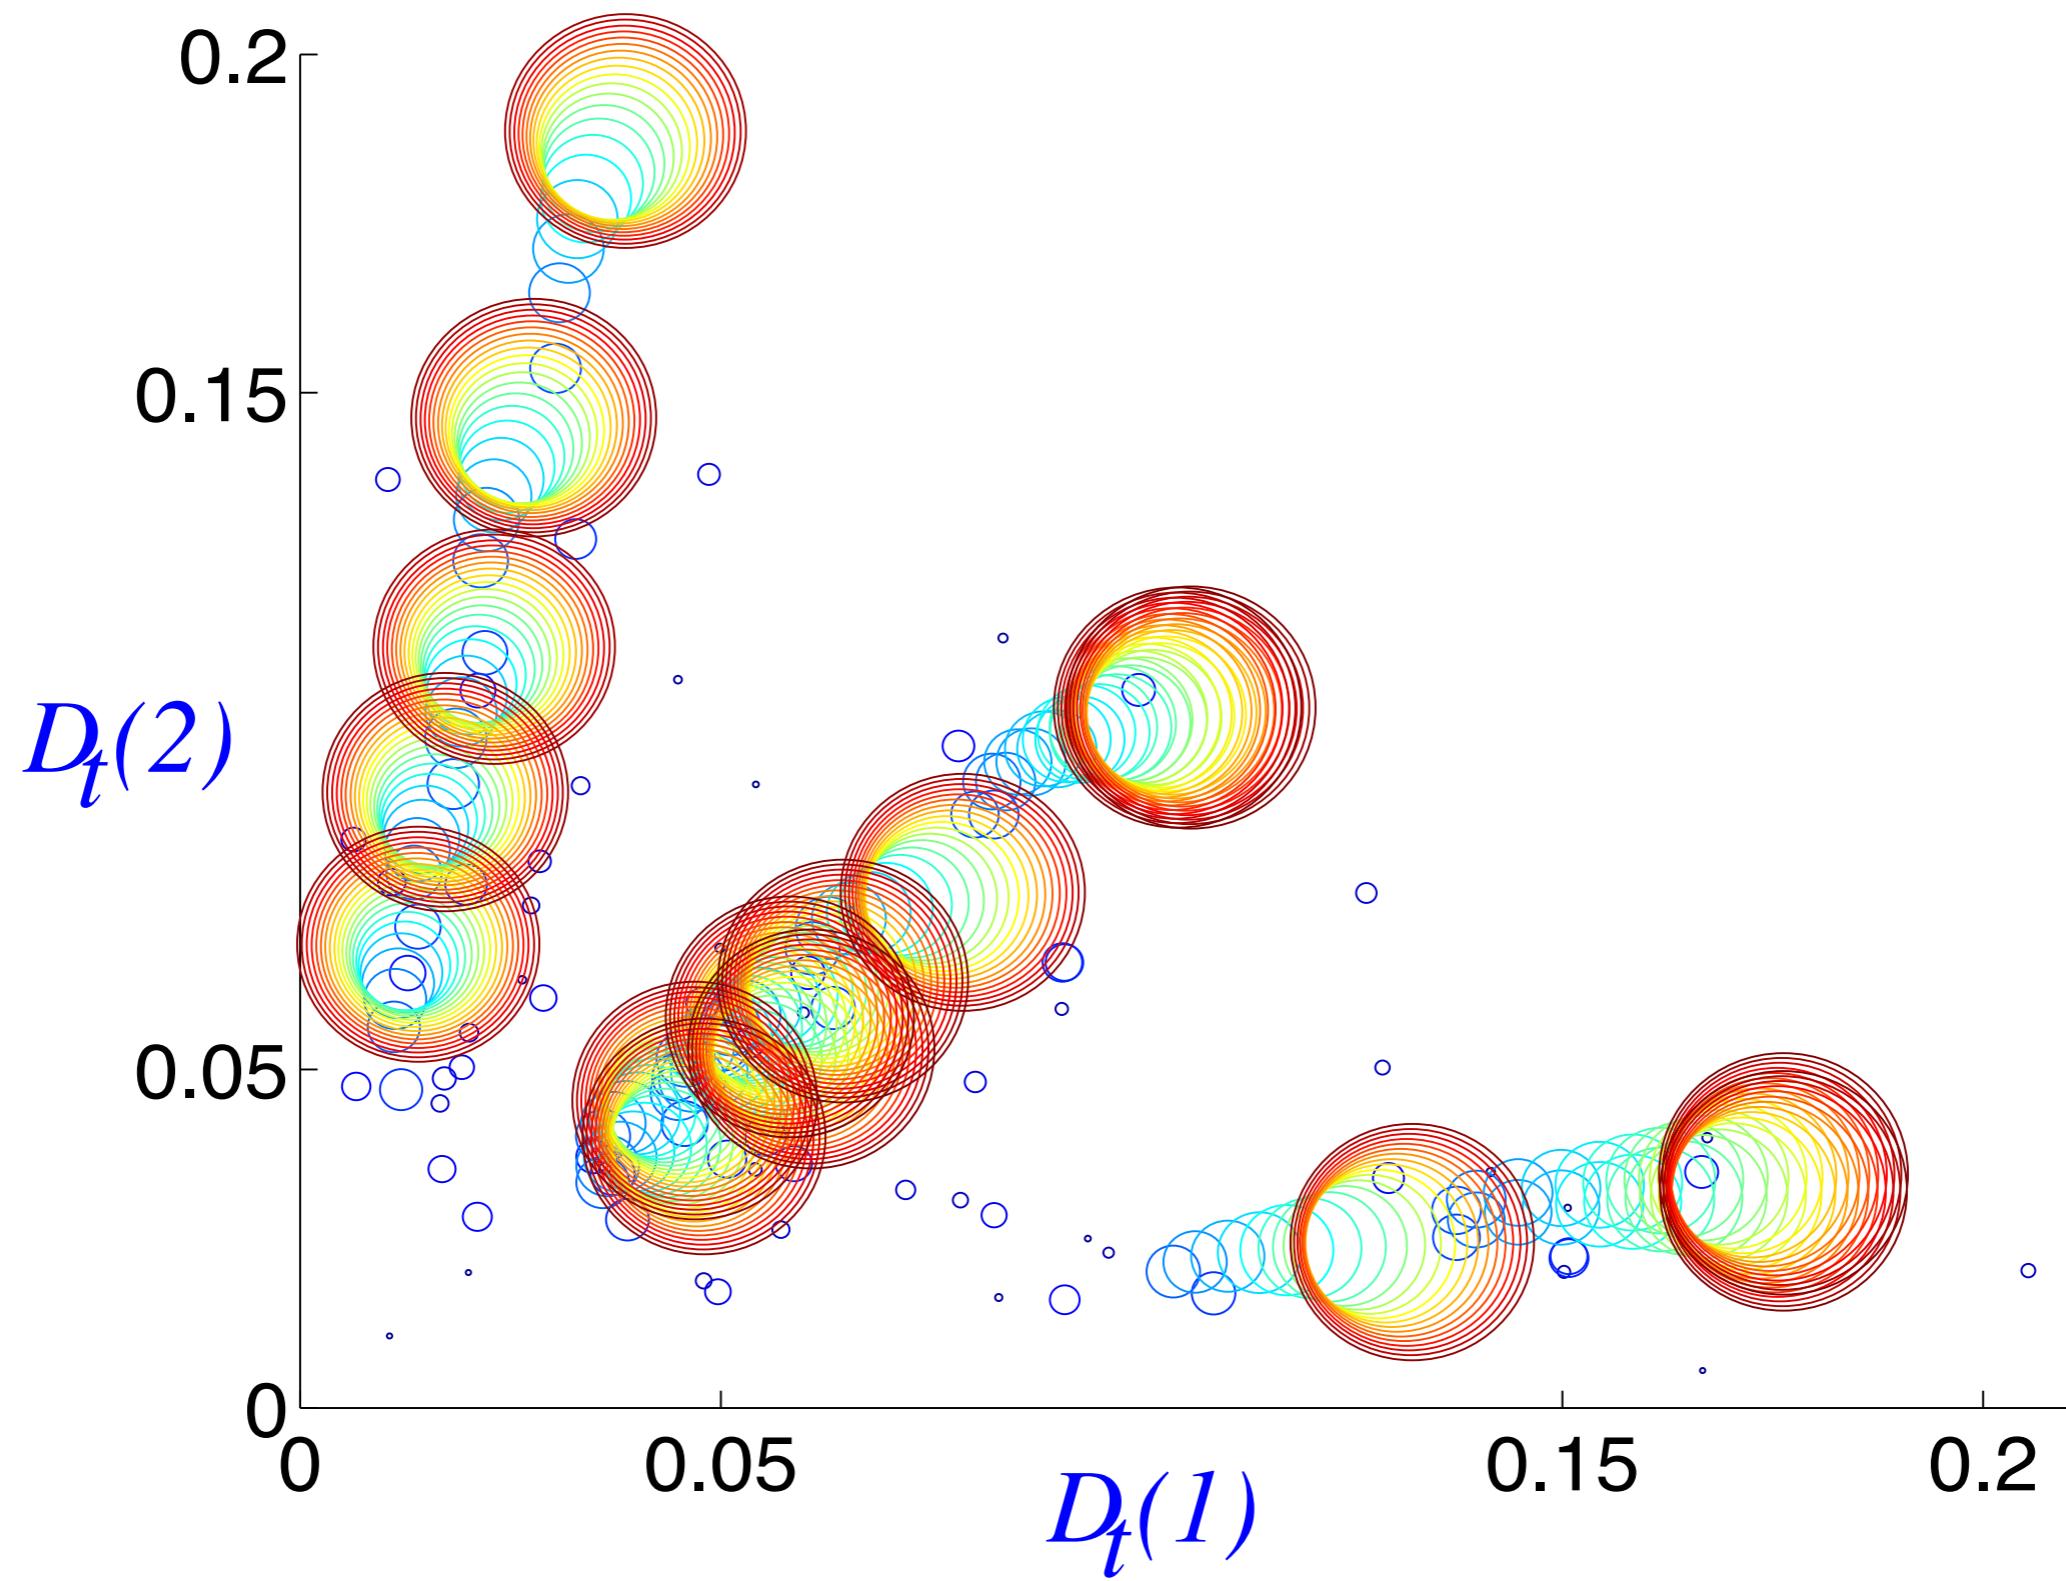


- in all cases examined, appears to converge eventually to cycle
 - open if always true

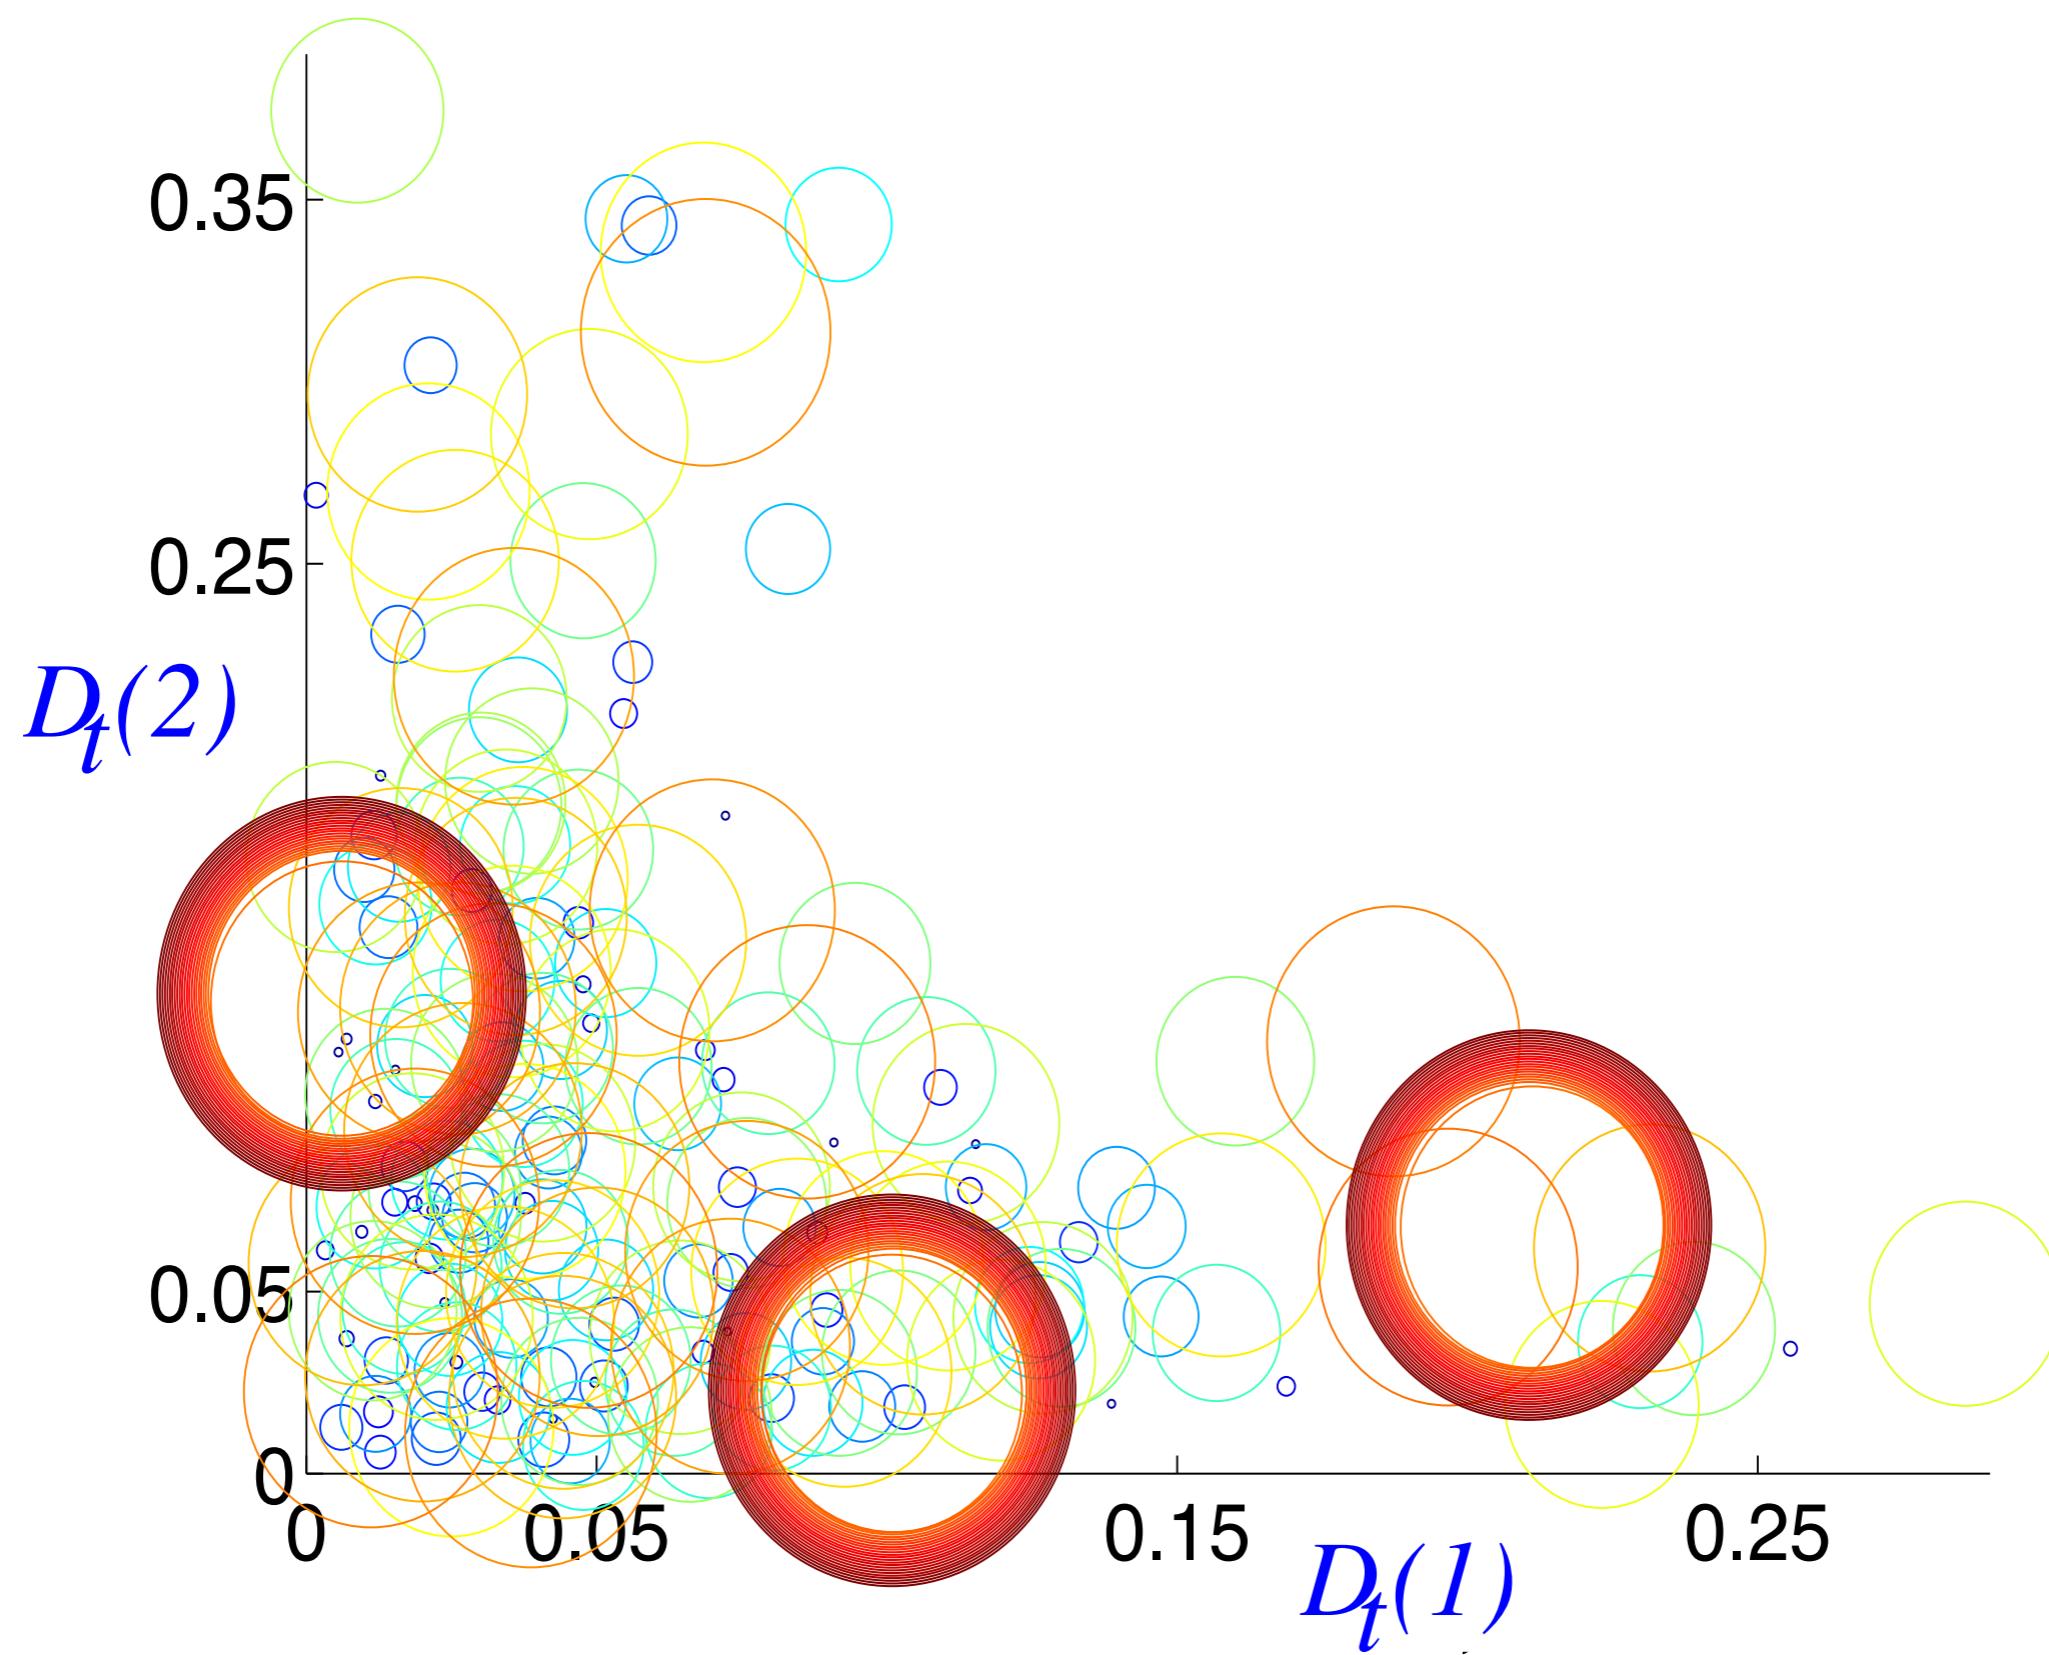
More Examples



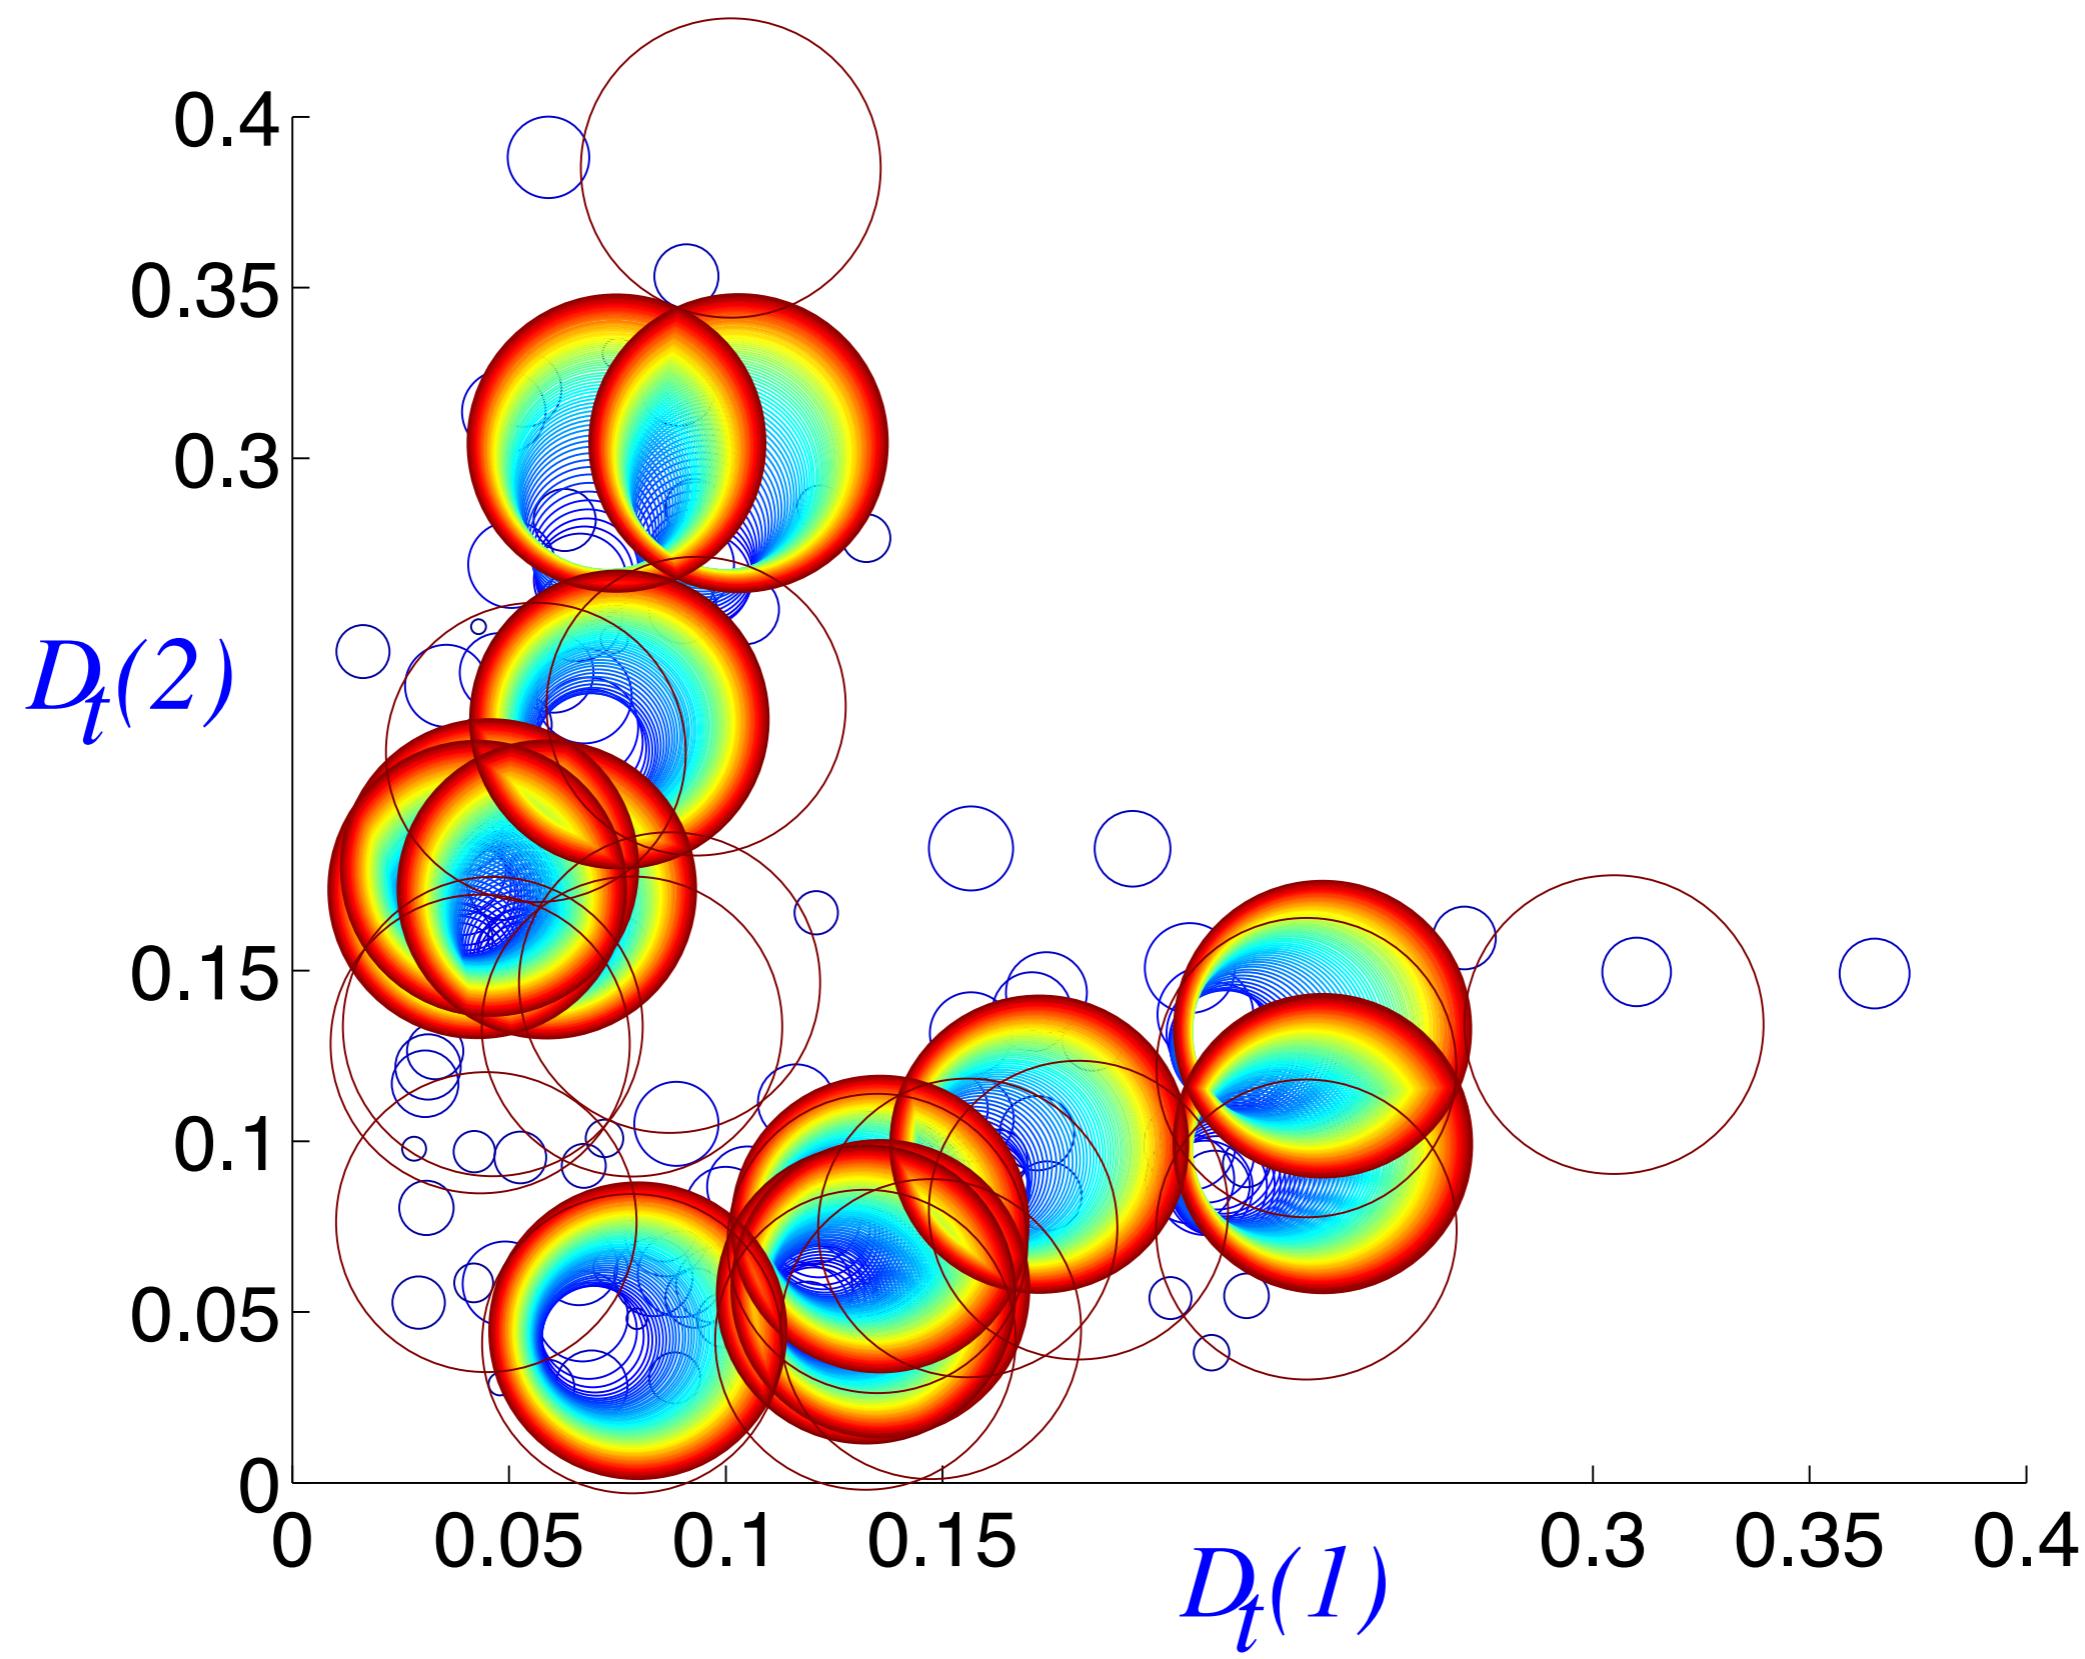
More Examples



More Examples



More Examples



Unifying the Two Cases

[with Collins & Singer]

- two distinct cases:
 - weak learning assumption holds
 - $\mathcal{P} = \emptyset$
 - dynamics unclear
 - weak learning assumption does **not** hold
 - $\mathcal{P} \neq \emptyset$
 - can prove convergence of D_t 's
- to **unify**: work instead with **unnormalized** versions of D_t 's
 - standard AdaBoost: $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{\text{normalization}}$
 - instead:

$$d_{t+1}(i) = d_t(i) \exp(-\alpha_t y_i h_t(x_i))$$

$$D_{t+1}(i) = \frac{d_{t+1}(i)}{\text{normalization}}$$

- algorithm is **unchanged**

Reformulating AdaBoost as Iterative Projection

- points = nonnegative vectors \mathbf{d}_t
- distance = unnormalized relative entropy:

$$\text{RE}(\mathbf{p} \parallel \mathbf{q}) = \sum_i \left[p(i) \ln \left(\frac{p(i)}{q(i)} \right) + q(i) - p(i) \right]$$

- reference point $\mathbf{x}_0 = \mathbf{1}$ (all 1's vector)
- hyperplanes defined by weak classifiers g_j :

$$\sum_i d(i) y_i g_j(x_i) = 0$$

- resulting iterative-projection algorithm is again equivalent to AdaBoost

Reformulated Optimization Problem

- optimization problem:

$$\min_{\mathbf{d} \in \mathcal{P}} \text{RE}(\mathbf{d} \parallel \mathbf{1})$$

- where

$$\mathcal{P} = \left\{ \mathbf{d} : \sum_i d(i) y_i g_j(x_i) = 0 \quad \forall j \right\}$$

- note: feasible set \mathcal{P} **never** empty (since $\mathbf{0} \in \mathcal{P}$)

Exponential Loss as Entropy Optimization

- all vectors \mathbf{d}_t created by AdaBoost have form:

$$d(i) = \exp \left(-y_i \sum_j \lambda_j g_j(x_i) \right)$$

- let $\mathcal{Q} = \{ \text{all vectors } \mathbf{d} \text{ of this form} \}$
- can rewrite exponential loss:

$$\begin{aligned} \inf_{\boldsymbol{\lambda}} \sum_i \exp \left(-y_i \sum_j \lambda_j g_j(x_i) \right) &= \inf_{\mathbf{d} \in \mathcal{Q}} \sum_i d(i) \\ &= \min_{\mathbf{d} \in \overline{\mathcal{Q}}} \sum_i d(i) \\ &= \min_{\mathbf{d} \in \overline{\mathcal{Q}}} \text{RE}(\mathbf{0} \parallel \mathbf{d}) \end{aligned}$$

- $\overline{\mathcal{Q}} = \text{closure of } \mathcal{Q}$

Duality

[Della Pietra, Della Pietra & Lafferty]

- presented two optimization problems:
 - $\min_{\mathbf{d} \in \mathcal{P}} \text{RE}(\mathbf{d} \parallel \mathbf{1})$
 - $\min_{\mathbf{d} \in \overline{\mathcal{Q}}} \text{RE}(\mathbf{0} \parallel \mathbf{d})$
- which is AdaBoost solving? Both!
- problems have **same** solution
- moreover: solution given by **unique** point in $\mathcal{P} \cap \overline{\mathcal{Q}}$
- problems are **convex duals** of each other

Convergence of AdaBoost

- can use to prove AdaBoost **converges** to common solution of both problems:
 - can argue that $\mathbf{d}^* = \lim \mathbf{d}_t$ is in \mathcal{P}
 - vectors \mathbf{d}_t are in \mathcal{Q} always $\Rightarrow \mathbf{d}^* \in \overline{\mathcal{Q}}$
 - ∴ $\mathbf{d}^* \in \mathcal{P} \cap \overline{\mathcal{Q}}$
 - ∴ \mathbf{d}^* solves both optimization problems
- so:
 - AdaBoost **minimizes** exponential loss
 - exactly **characterizes** limit of unnormalized “distributions”
 - likewise for normalized distributions when weak learning assumption does not hold
- also, provides additional link to **logistic regression**
 - only need slight change in optimization problem
[with Collins & Singer; Lebannon & Lafferty]

Practical Extensions

- multiclass classification
- ranking problems
- confidence-rated predictions

Practical Extensions

- multiclass classification
- ranking problems
- confidence-rated predictions

Multiclass Problems

[with Freund]

- say $y \in Y = \{1, \dots, k\}$
- direct approach (AdaBoost.M1):

$$h_t : X \rightarrow Y$$

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \cdot \begin{cases} e^{-\alpha_t} & \text{if } y_i = h_t(x_i) \\ e^{\alpha_t} & \text{if } y_i \neq h_t(x_i) \end{cases}$$

$$H_{\text{final}}(x) = \arg \max_{y \in Y} \sum_{t: h_t(x) = y} \alpha_t$$

- can prove same bound on error if $\forall t : \epsilon_t \leq 1/2$
 - in practice, not usually a problem for “strong” weak learners (e.g., C4.5)
 - significant problem for “weak” weak learners (e.g., decision stumps)
- instead, reduce to binary

Reducing Multiclass to Binary

[with Singer]

- say possible labels are $\{a, b, c, d, e\}$
- each training example replaced by five $\{-1, +1\}$ -labeled examples:

$$x, c \rightarrow \left\{ \begin{array}{ll} (x, a), & -1 \\ (x, b), & -1 \\ (x, c), & +1 \\ (x, d), & -1 \\ (x, e), & -1 \end{array} \right.$$

- predict with label receiving most (weighted) votes

AdaBoost.MH

- can prove:

$$\text{training error}(H_{\text{final}}) \leq \frac{k}{2} \cdot \prod Z_t$$

- reflects fact that small number of errors in binary predictors can cause overall prediction to be incorrect
- extends immediately to **multi-label** case
(more than one correct label per example)

Using Output Codes

[with Allwein & Singer][Dietterich & Bakiri]

- alternative: choose “codeword” for each label

	π_1	π_2	π_3	π_4
a	-	+	-	+
b	-	+	+	-
c	+	-	-	+
d	+	-	+	+
e	-	-	-	-

- each training example mapped to one example per column

$$x, c \rightarrow \begin{cases} (x, \pi_1), +1 \\ (x, \pi_2), -1 \\ (x, \pi_3), -1 \\ (x, \pi_4), +1 \end{cases}$$

- to classify new example x :
 - evaluate classifier on $(x, \pi_1), \dots, (x, \pi_4)$
 - choose label “most consistent” with results

Output Codes (cont.)

- training error bounds **independent** of # of classes
- overall prediction robust to large number of errors in binary predictors
- **but:** binary problems may be harder

Practical Extensions

- multiclass classification
- ranking problems
- confidence-rated predictions

Ranking Problems

[with Freund, Iyer & Singer]

- other problems can also be handled by reducing to binary
- e.g.: want to learn to **rank** objects (say, movies) from examples
- can **reduce** to multiple **binary** questions of form:
“is or is not object A preferred to object B?”
- now apply (binary) AdaBoost \Rightarrow “**RankBoost**”

Application: Finding Cancer Genes

[Agarwal & Sengupta]

- examples are genes (described by microarray vectors)
- want to rank genes from most to least relevant to leukemia
- data sizes:
 - 7129 genes total
 - 10 known relevant
 - 157 known irrelevant

Top-Ranked Cancer Genes

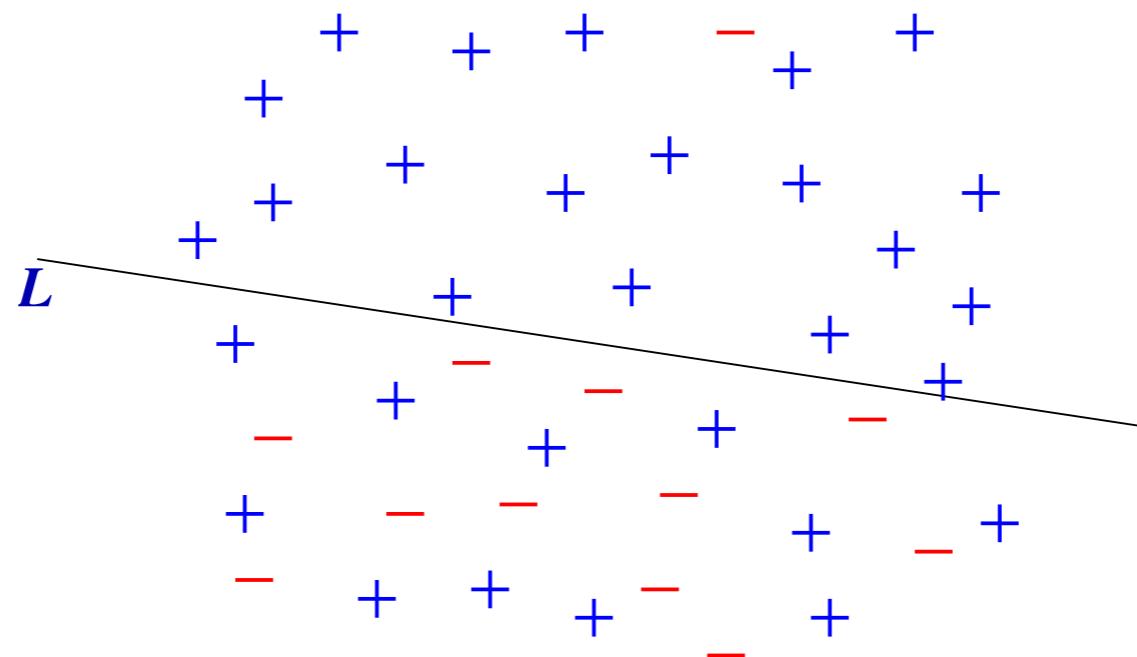
Gene	Relevance Summary
1. KIAA0220	□
2. G-gamma globin	◆
3. Delta-globin	◆
4. Brain-expressed HCPA78 homolog	□
5. Myeloperoxidase	◆
6. Probable protein disulfide isomerase ER-60 precursor	□
7. NPM1 Nucleophosmin	◆
8. CD34	◆
9. Elongation factor-1-beta	×
10. CD24	◆

- = known therapeutic target
- = potential therapeutic target
- ◆ = known marker
- ◇ = potential marker
- × = no link found

Practical Extensions

- multiclass classification
- ranking problems
- confidence-rated predictions

“Hard” Predictions Can Slow Learning



- ideally, want weak classifier that says:

$$h(x) = \begin{cases} +1 & \text{if } x \text{ above } L \\ \text{“don't know”} & \text{else} \end{cases}$$

- problem: cannot express using “hard” predictions
- if must predict ± 1 below L , will introduce many “bad” predictions
 - need to “clean up” on later rounds
- dramatically increases time to convergence

Confidence-Rated Predictions

[with Singer]

- useful to allow weak classifiers to assign **confidences** to predictions
- formally, allow $h_t : X \rightarrow \mathbb{R}$

$$\begin{aligned}\text{sign}(h_t(x)) &= \text{prediction} \\ |h_t(x)| &= \text{"confidence"}\end{aligned}$$

- use identical update:

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \cdot \exp(-\alpha_t y_i h_t(x_i))$$

and identical rule for combining weak classifiers

- question: how to choose α_t and h_t on each round

Confidence-Rated Predictions (cont.)

- saw earlier:

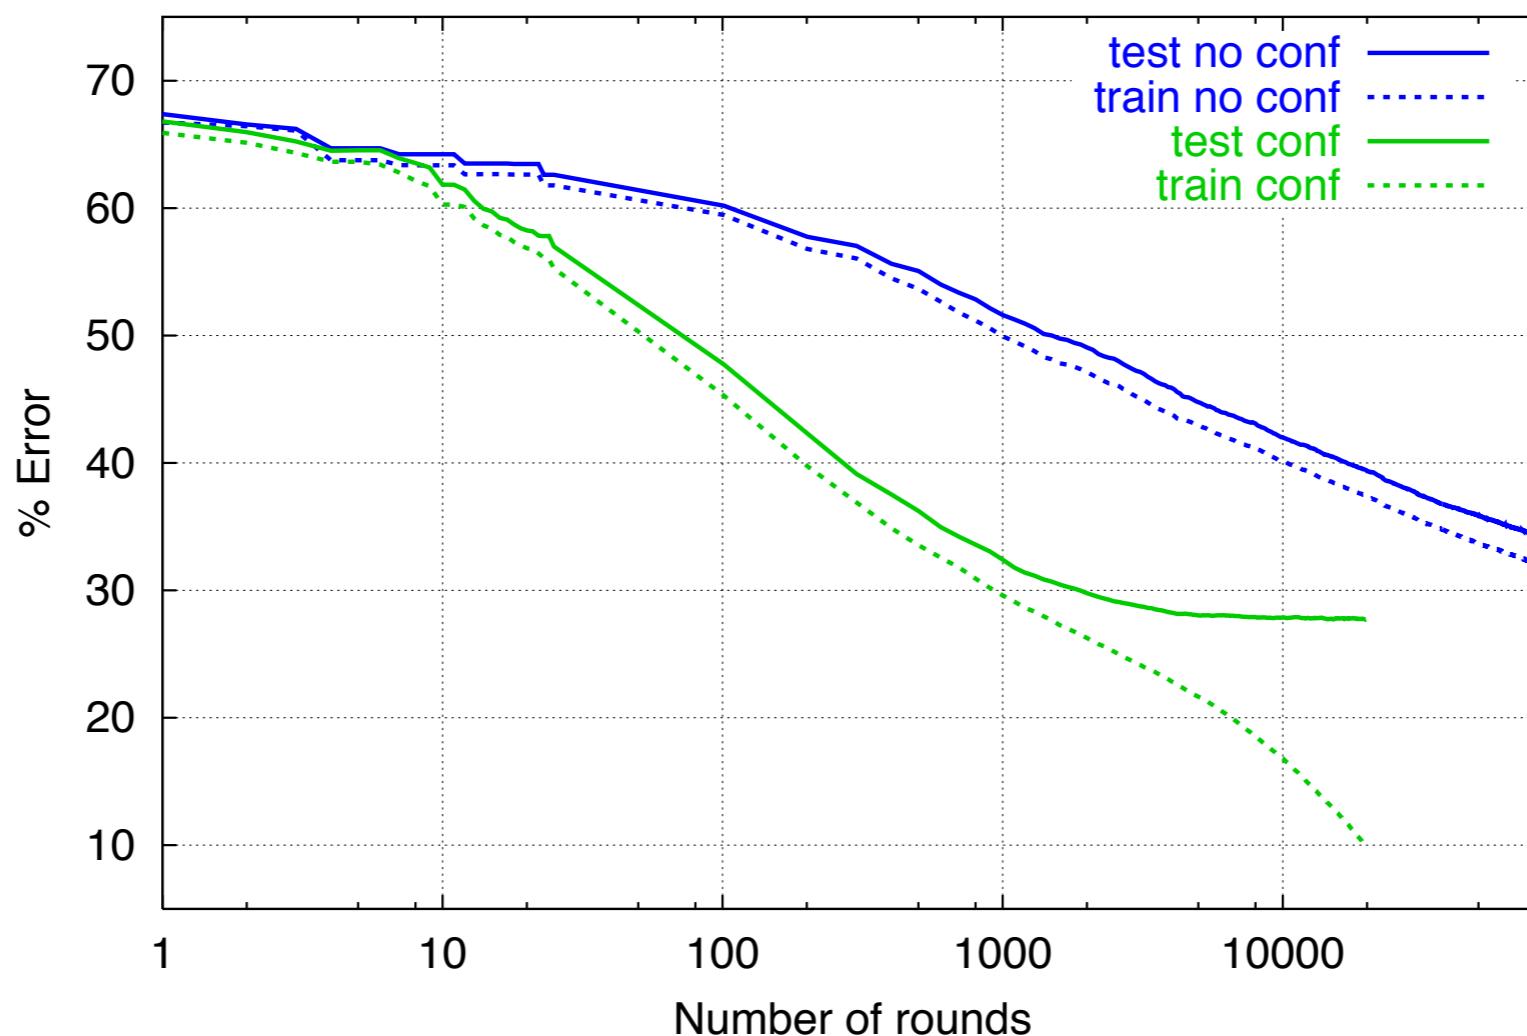
$$\text{training error}(H_{\text{final}}) \leq \prod_t Z_t = \frac{1}{m} \sum_i \exp \left(-y_i \sum_t \alpha_t h_t(x_i) \right)$$

- therefore, on each round t , should choose $\alpha_t h_t$ to minimize:

$$Z_t = \sum_i D_t(i) \exp(-\alpha_t y_i h_t(x_i))$$

- in many cases (e.g., decision stumps), best confidence-rated weak classifier has simple form that can be found efficiently

Confidence-Rated Predictions Help a Lot



% error	round first reached			speedup
	conf.	no conf.		
40	268	16,938		63.2
35	598	65,292		109.2
30	1,888	>80,000		-

Application: Boosting for Text Categorization

[with Singer]

- **weak classifiers:** very simple weak classifiers that test on simple patterns, namely, (sparse) n -grams
 - find parameter α_t and rule h_t of given form which minimize Z_t
 - use efficiently implemented exhaustive search
- “How may I help you” data:
 - 7844 training examples
 - 1000 test examples
 - categories: AreaCode, AttService, BillingCredit, CallingCard, Collect, Competitor, DialForMe, Directory, HowToDial, PersonToPerson, Rate, ThirdNumber, Time, TimeCharge, Other.

Weak Classifiers

More Weak Classifiers

More Weak Classifiers

Finding Outliers

examples with most weight are often **outliers** (mislabeled and/or ambiguous)

- I'm trying to make a credit card call (**Collect**)
- hello (**Rate**)
- yes I'd like to make a long distance collect call please (**CallingCard**)
- calling card please (**Collect**)
- yeah I'd like to use my calling card number (**Collect**)
- can I get a collect call (**CallingCard**)
- yes I would like to make a long distant telephone call and have the charges billed to another number
(**CallingCard DialForMe**)
- yeah I can not stand it this morning I did oversea call is so bad (**BillingCredit**)
- yeah special offers going on for long distance
(**AttService Rate**)
- mister allen please william allen (**PersonToPerson**)
- yes ma'am I I'm trying to make a long distance call to a non dialable point in san miguel philippines
(**AttService Other**)

Advanced Topics

- optimal accuracy
- optimal efficiency
- boosting in continuous time

Advanced Topics

- optimal accuracy
- optimal efficiency
- boosting in continuous time

Optimal Accuracy

[Bartlett & Traskin]

- usually, impossible to get perfect accuracy due to intrinsic noise or uncertainty
- Bayes optimal error = best possible error of any classifier
 - usually > 0
- can prove AdaBoost's classifier converges to Bayes optimal if:
 - enough data
 - run for many (but not too many) rounds
 - weak classifiers "sufficiently rich"
- "universally consistent"
- related results: [Jiang], [Lugosi & Vayatis], [Zhang & Yu], ...
- means:
 - AdaBoost can (theoretically) learn "optimally" even in noisy settings
 - but: does not explain why works when run for very many rounds

Boosting and Noise

[Long & Servedio]

- can construct data source that “breaks” AdaBoost with even **tiny** amount of noise (say, 1%)
 - Bayes optimal error $= 1\%$
(obtainable by classifier of same form as AdaBoost)
 - AdaBoost provably has error $\geq 50\%$
- holds even if:
 - given **unlimited** training data
 - use any method for minimizing exponential loss
(holds for most other convex losses as well)
- shows:
 - consistency result can fail badly if weak classifiers “not rich enough”
 - boosting susceptible to noise
- on “real-world” datasets, AdaBoost often works anyway
- various theoretical algorithms proposed for handling noise
(e.g., [Kalai & Servedio], [Long & Servedio])

Advanced Topics

- optimal accuracy
- **optimal efficiency**
- boosting in continuous time

Optimal Efficiency

[Freund]

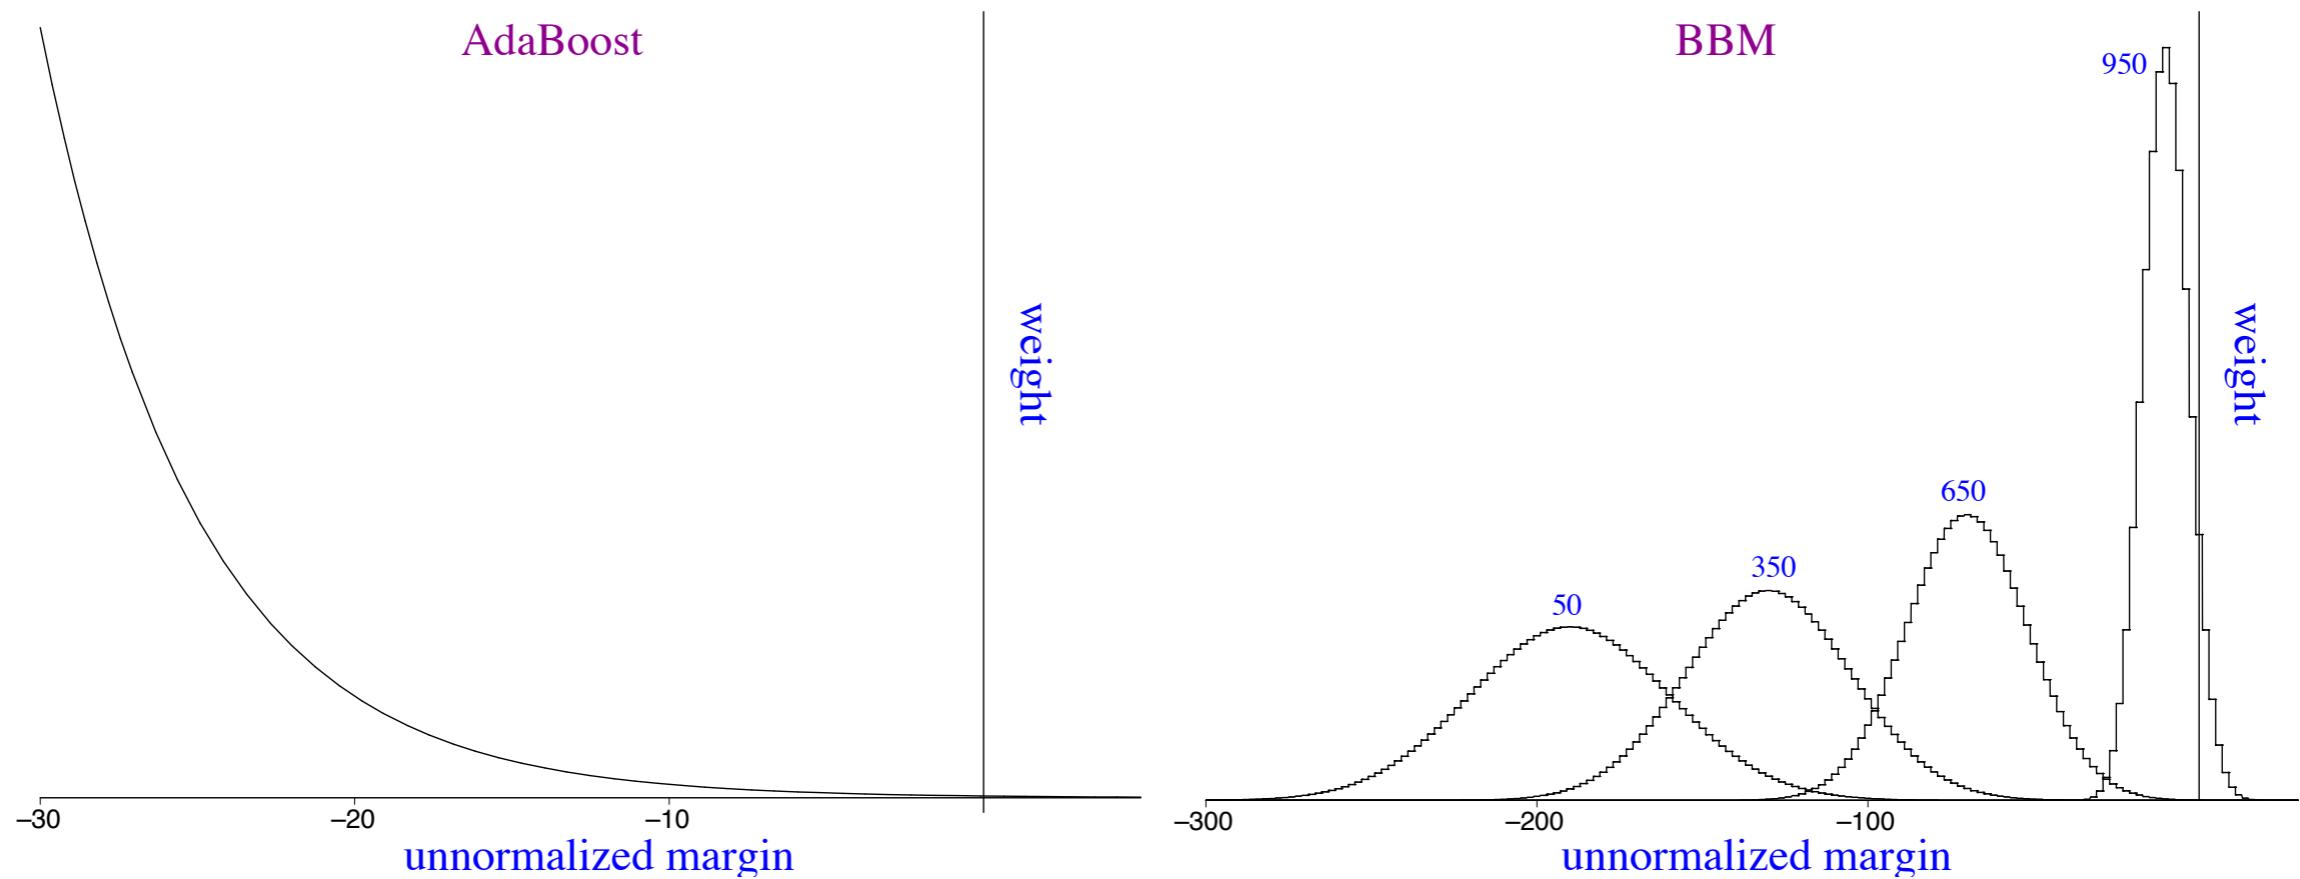
- for AdaBoost, saw: training error $\leq e^{-2\gamma^2 T}$
- is AdaBoost most efficient boosting algorithm?
no!
- given T rounds and γ -weak learning assumption,
boost-by-majority (BBM) algorithm is provably **exactly** best
possible:

$$\text{training error} \leq \sum_{j=0}^{\lfloor T/2 \rfloor} \binom{T}{j} \left(\frac{1}{2} + \gamma\right)^j \left(\frac{1}{2} - \gamma\right)^{T-j}$$

(probability of $\leq T/2$ heads in T coin flips if probability of
heads $= \frac{1}{2} + \gamma$)

- AdaBoost's training error is like Chernoff approximation of
BBM's

Weighting Functions: AdaBoost versus BBM



- both put more weight on harder examples, but BBM “gives up” on **very hardest** examples
 - may make more robust to noise
- **problem:** BBM not adaptive
 - need to know γ and T a priori

Advanced Topics

- optimal accuracy
- optimal efficiency
- boosting in continuous time

Boosting in Continuous Time

[Freund]

- idea: let γ get very small so that γ -weak learning assumption eventually satisfied
- need to make T correspondingly large
- if scale “time” to begin at $\tau = 0$ and end at $\tau = 1$, then each boosting round takes time $1/T$
- in limit $T \rightarrow \infty$, boosting is happening in **continuous time**

BrownBoost

- algorithm has sensible limit called “**BrownBoost**” (due to connection to Brownian motion)
- harder to implement, but potentially more resistant to noise and outliers, e.g.:

dataset	noise	AdaBoost	BrownBoost
letter	0%	3.7	4.2
	10%	10.8	7.0
	20%	15.7	10.5
satimage	0%	4.9	5.2
	10%	12.1	6.2
	20%	21.3	7.4

[Cheamanunkul, Ettinger & Freund]

Conclusions

- from different perspectives, AdaBoost can be interpreted as:
 - a method for **boosting** the accuracy of a weak learner
 - a procedure for **maximizing margins**
 - an algorithm for playing **repeated games**
 - a numerical method for **minimizing exponential loss**
 - an **iterative-projection** algorithm based on an information-theoretic geometry
- none is entirely satisfactory by itself, but each useful in its own way
- taken together, create rich theoretical understanding
 - connect boosting to other learning problems and techniques
 - provide foundation for versatile set of methods with many extensions, variations and applications

References

Coming soon:

- Robert E. Schapire and Yoav Freund.
Boosting: Foundations and Algorithms.
MIT Press, 2012.

Survey articles:

- Ron Meir and Gunnar Rätsch.
An Introduction to Boosting and Leveraging.
In *Advanced Lectures on Machine Learning (LNAI2600)*, 2003.
<http://www.boosting.org/papers/MeiRae03.pdf>
- Robert E. Schapire.
The boosting approach to machine learning: An overview.
In *MSRI Workshop on Nonlinear Estimation and Classification*, 2002.
<http://www.cs.princeton.edu/~schapire/boost.html>