







Learning Games and Brownian Motion

Yoav Freund

Center for Computational Learning
Systems (CLASS)

Columbia University


Drifting games (in 2d)



$$\sum_i \|\vec{w}_i\| = 1$$

$$\sum_i \vec{w}_i \vec{x}_i \geq \delta > 0$$

Drifting games (in 2d)



Target
region




Plan of talk

- The one dimensional case
- Applications (in learning and elsewhere)
- Solution of the one-dimensional case
- Generalization to higher dimensions
- The continuous time limit (Brownian motion)
- Other applications and open problems

Simple 1d drifting games

- Chips (formerly sheep) start from origin.
- Goal of shepherd: move all chips above origin.
- Sheep strategist divided into Splitter and Chooser.
- **GAME:** repeat
 - 1) **Shepherd** set a weight for each chip, weights sum to 1.
 - 2) **Splitter** splits the chips into two sets.
 - 3) **Chooser** selects one of the two sets which is then moved one step up.

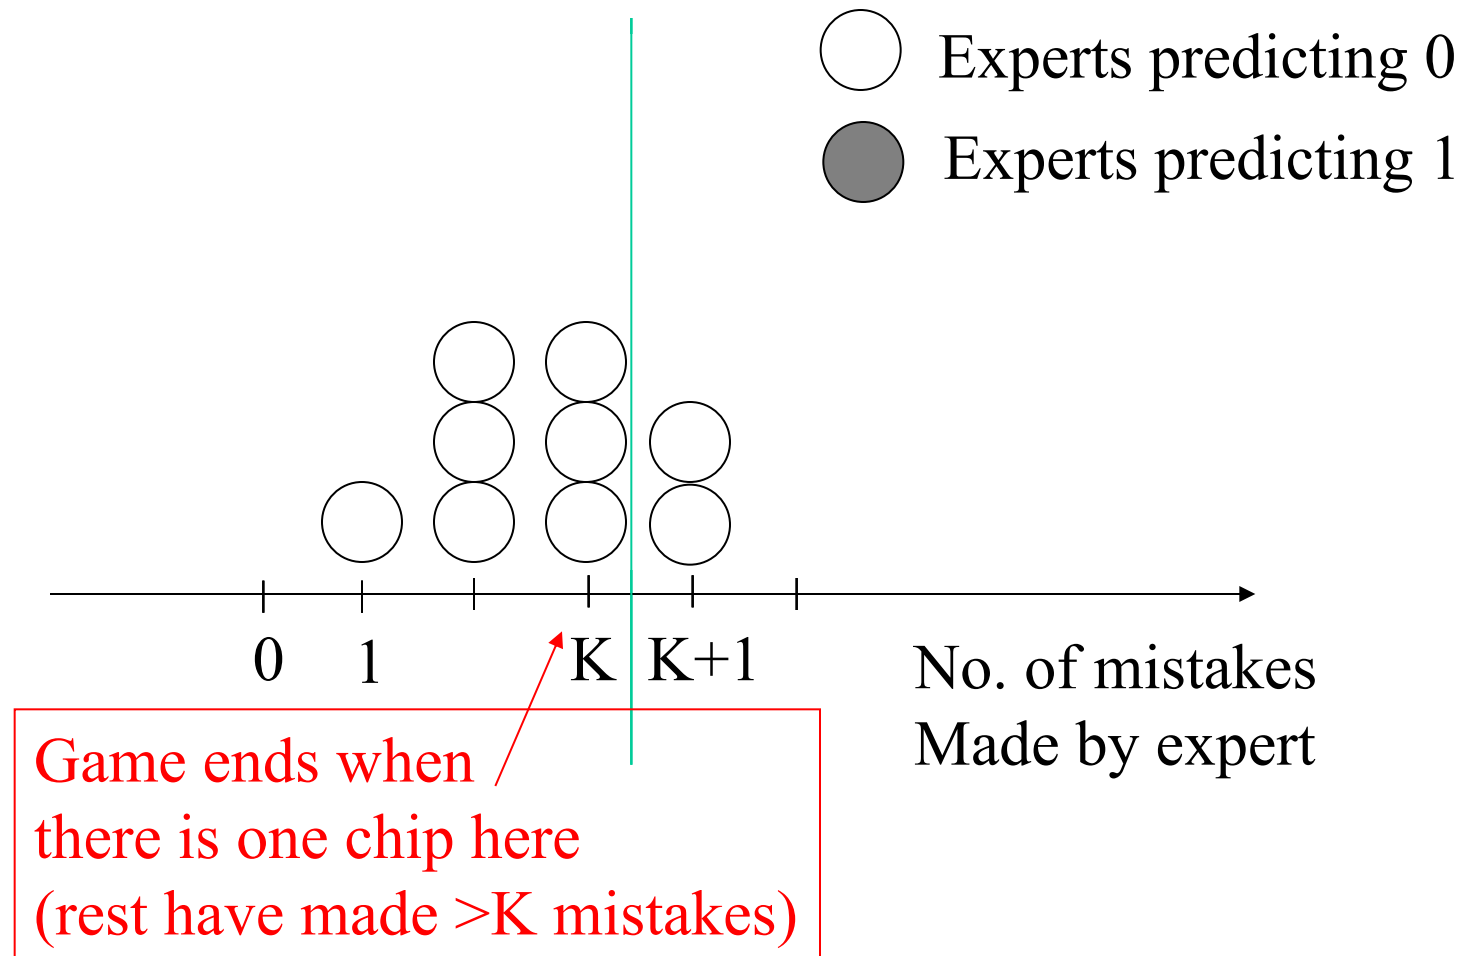
The 1d Chip game



Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:
 - expert 1: 1,0,1,0,1,1,1
 - expert 2: 0,0,0,1,1,0,1
 - ...
 - expert N: 1,0,0,1,1,1,0
- Algorithm makes prediction: 1,0,1,1,0,1,1
- Assumption: there exists an expert which makes at most k mistakes
- Goal: make least mistakes under the assumption (no statistics!)
- Chip = expert, bin = number of mistakes made
- Game step = algorithm's prediction is incorrect.

Chip game for expert advice



Boosting

- A method for improving accuracy of learning algorithms for classification problems.
- **Weak Learner**: a learning algorithm generating rules that are only slightly better than random guessing.
- Basic idea: **re-weight training examples** to force weak learner into different parts of the space.
- Combine weak rules by a majority vote.

batch learning for binary classification

Data distribution:


Generalization error:

Training set:

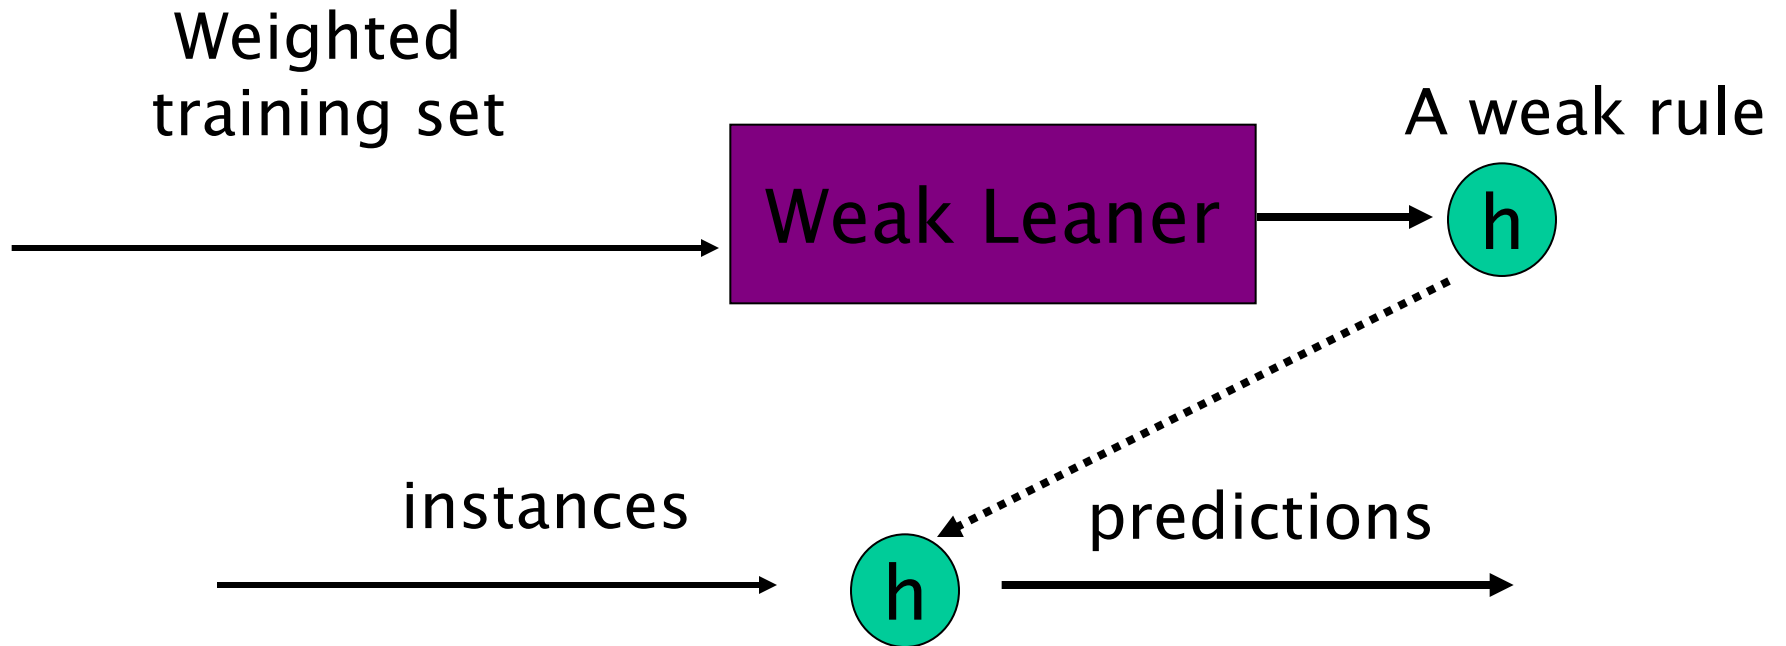
Training error:

A weighted training set

Feature vectors

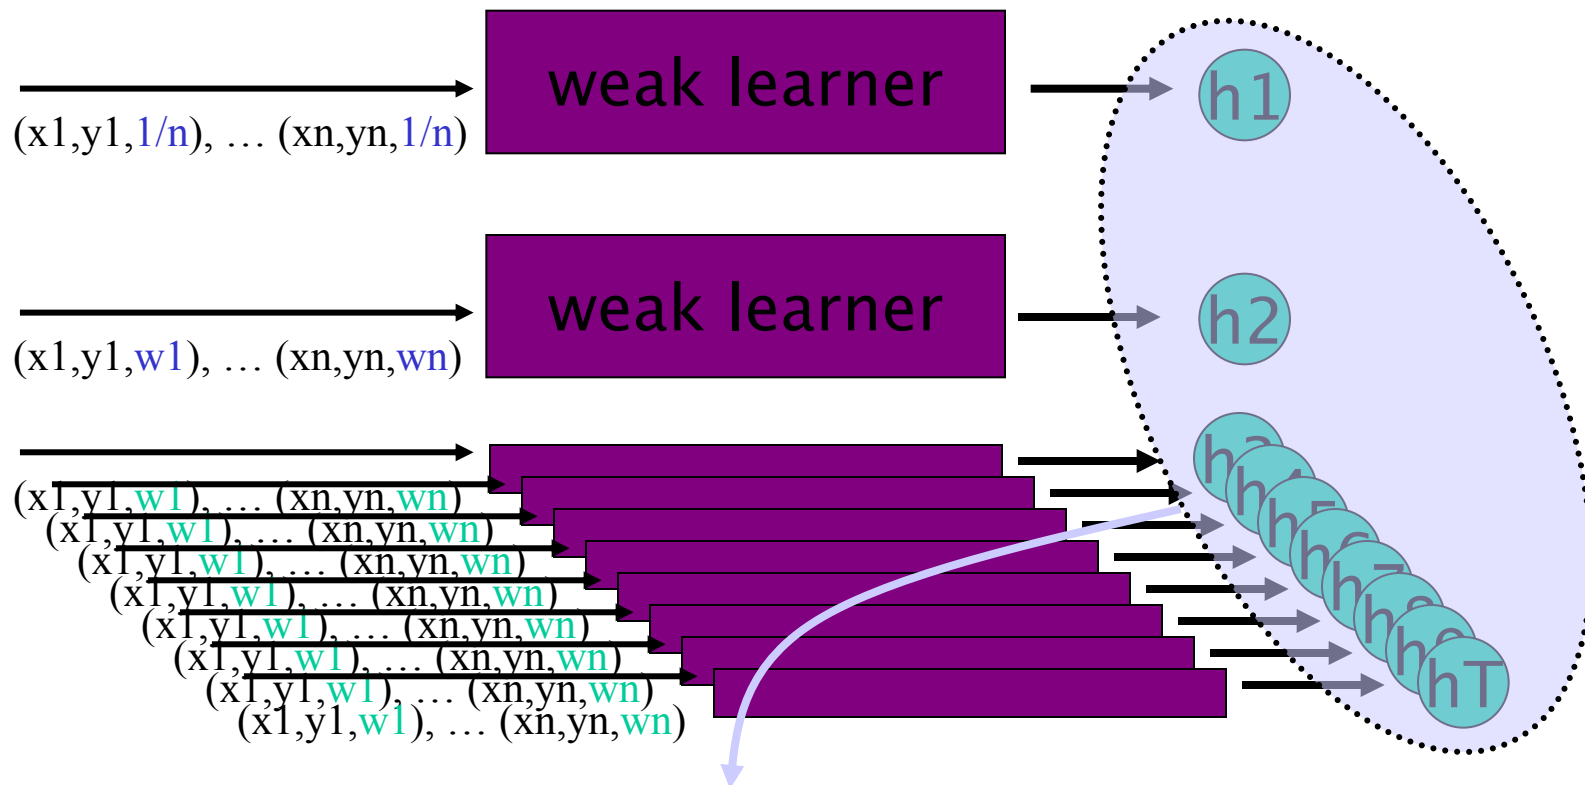


A weak learner



The weak requirement:

The boosting process




Final rule:

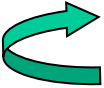
Boosting as a chip game

- Chip = training example
- Booster assigns a weight to each chip, weights sum to 1.
- Learner selects a subset with weight $\geq \frac{1}{2} + \gamma$
 - Selected set moves a step right (correct)
 - Unselected set moves a step left (incorrect)
- Booster wins examples on right of origin.
 - Implies that majority vote is correct.

The boosting chip game



Ulam's game with k lies


- Player 1 chooses secret x from S_1, S_2, \dots, S_N
- • Player 2 asks “is x in the set A ?”
- Player 1 can lie k times.
- Game ends when player 2 can determine x .
- A smart player 1 aims to have more than one consistent x for as long as possible.
- Chip = element (S_i)
- Bin = number of lies made regarding element

Ron Rivest's game

- You are given a bag with 1000 fair coins.
- Your goal is to have at least 700 “heads”
- You repeatedly choose one of two actions:
 - 1) Take a single coin out of the bag and flip it.
 - 2) Put all the coins back into the bag.
- Find strategy that minimizes the expected time to success?

Taking number of chips to infinity

Replace individual chips by chip **mass**



Each bin can be divided arbitrarily

Notation

loss of unit mass in bin i at game's end

Simple case:

Fraction of mass that is in bin i at time t

Goal of shepherd is to minimize

More Notation

Fraction of mass in bin i at time t that moves up

Weights assigned by Shepherd:

allowed movement given weighting

Backwards recursion

The Potential of a unit mass in bin i at time t

Given we express the min/max value for $t-1$

Matching min/max strategies

Shepherd's strategy

implies

Sheep's strategy

Recursive definition of potential

Definition:

Yields:

min-max value of the game

if at $t=0$ mass is concentrated in origin

Probabilistic interpretation


Simplifying assumption (not critical)

Each sheep moves up independently with probability

The expected final loss of a chip performing
a biased random walk starting from bin i at time t

Solution for boost-by-majority

Example potential and weight



Generalization to general vector spaces

General drifting game notation

Sheep mass
in $1d$

Discrete Sheep
in $d>1$

location	Bin no. i	Sheep j location
Sheep steps		bounded set
Shepherd's choice		

The min/max solution

[Schapire99]

- A potential defined by a min/max recursion

$$\phi_T(\mathbf{s}) = L(\mathbf{s})$$

$$\phi_{t-1}(\mathbf{s}) = \min_{\mathbf{w} \in \mathbb{R}^d} \sup_{\mathbf{z} \in B} (\phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w} \cdot \mathbf{z} - \delta \|\mathbf{w}\|)$$

$$\phi_0(0) = \text{the value of the game}$$

- Shepherd's strategy

$$\mathbf{w}_i^t = \arg \min_{\mathbf{w}} \sup_{\mathbf{z} \in B} (\phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w} \cdot \mathbf{z} - \delta \|\mathbf{w}\|)$$


Restricting allowed steps

B = the set of all allowable step

Normal B = minimal set that spans the space. (~basis)

Regular B = a symmetric regular set. (~orthonormal basis)

Regular step set in \mathbb{R}^2



The solution simplifies when $\delta \rightarrow 0$


If B is normal, and δ is sufficiently small then $\exists \mathbf{w}^*$ such that

$$\phi_{t-1}(\mathbf{s}) = \phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w}^* \cdot \mathbf{z} - \delta \|\mathbf{w}^*\|$$

for all $\mathbf{z} \in B$ (and all $t = 1, 2, \dots, \mathbf{s} \in \mathbb{R}^d$)

Implies that: \mathbf{w}^* is the “local slope” at $\phi_t(\mathbf{s})$, i.e.

$$\phi_t(\mathbf{s} + \mathbf{z}_i) = C + \mathbf{w}^* \mathbf{z}_i; \quad C \doteq \frac{\sum_{j=0}^d \phi_t(\mathbf{s} + \mathbf{z}_j)}{d+1}$$



and that

$$\phi_{t-1}(\mathbf{s}) = C - \delta \|\mathbf{w}^*\|$$

Increasing the number of steps

- Consider T steps in a unit time
- Drift d should scale like $1/T$
- Step size $O(1/T)$ gives game to shepherd
- Step size $O(1/\sqrt{T})$ keeps game balanced

The solution for

The local slope becomes the **gradient**

The potential is a solution of a **PDE**

$$\frac{\partial \phi(\mathbf{s}, \tau)}{\partial \tau} = -\frac{1}{2} \sum_{k,l} D_{kl} \frac{\partial^2 \phi(\mathbf{s}, t)}{\partial s^k \partial s^l} + \delta \|\nabla \phi(\mathbf{s}, \tau)\|_p$$

Same PDE describes time evolution of
Brownian motion with drift proportional to gradient

Applications of continuous drifting games


Some known solutions for 1d

Adaboost [\[Freund, Schapire 97\]](#)

Brownboost [\[Freund 01\]](#)


Boosting for more than 2 labels


Predict according to the most popular prediction




$\phi(x,y,t=1)$


$\phi(x,y,t=0)$







Boosting for variational optimization







Figure 2: The potential $\phi(\mathbf{s}, t)$ for the square loss $L(y) = y^2$.

solution for $L(s) = I_{|s|>1}$






Figure 3: The potential $\phi(\mathbf{s}, t)$ for the loss $L(y) = \min(y^2, 1)$.

Normalized drifting

- Boosting tends not to over-fit
- Explanation: “boosting the margins”
[schapire et al. 98]
- Goal is to minimize

Normalized drifting game in 1d

Regular steps

Leads to Brownian motion

Normalized steps

What is the appropriate limit for

?

1d problem with 3 allowed steps.

- Allowed steps: $\{-1, 0, +1\}$
- Importance: allows weak rules that **abstain**.
- Worst case choice of sheep depends on whether potential is convex or concave.

Open question

- In the continuous time limit, assuming that there are two regimes what characterizes the boundary condition between them?
- “The one-phase Stefan problem with temperature-boundary specification”?
- “The least sub-parabolic majorant of a function u ”?
- The boundary between solid and fluid?