


Chip games
Online Learning, Boosting
and
why continuous is easier than discrete

Yoav Freund
UCSD

The Chip game




Plan of the talk

- Why is the chip game is interesting?
- Letting the number of **chips** go to infinity.
- The boosting game.
- Drifting games (let chip=sheep).
- Letting the number of **rounds** go to infinity
- Brownian motion and generalized boosting (nice pictures!)


Combining expert advice

- Binary sequence: 1,0,0,1,1,0,?
- N experts make predictions:
 - expert 1: 1,0,1,0,1,1,1
 - expert 2: 0,0,0,1,1,0,1
 - ...
 - expert N : 1,0,0,1,1,1,0
- Algorithm makes prediction: 1,0,1,1,0,1,1
- Assumption: there exists an expert which makes at most k mistakes
- Goal: make least mistakes under the assumption (no statistics!)
- Chip = expert, bin = number of mistakes made
- Game step = algorithm's prediction is incorrect.

Chip game for expert advice




20 questions with k lies



- Player 1 chooses secret x from S_1, S_2, \dots, S_N
- Player 2 asks “is x in the set A ?”
- Player 1 can lie k times.
- Game ends when player 2 can determine x .
- A smart player 1 aims to have more than one consistent x for as long as possible.
- **Chip** = element (S_i)
- **Bin** = number of lies made regarding element

Chip game for 20 questions




Simple case – all chips in bin 1

- 21 questions without lies
- Combining experts where one is perfect
- Splitting a cookie
- Optimal strategies:
 - Player 1: split chips into two equal parts
 - Player 2: choose larger part
- Note problem when number of chips is odd

Number of chips to infinity

Replace individual chips by chip **mass**



Optimal splitter strategy:

Split **each** bin into **two equal** parts

Binomial weights strategies

- What should **chooser** do if parts are not equal?
- Assume that on **following** iterations splitter will play optimally = split each bin into 2 equal parts.
- Future configuration independent of chooser's choice
- **Potential**: Fraction of bin **i** that will remain in bins 1..k when **m** iterations remain

$$\binom{m}{\leq k-i} \doteq \frac{1}{2^m} \sum_{j=0}^{k-i} \binom{m}{j}$$

- Choose part so that next configuration will have maximal (or minimal) **potential**.

- Solve for **m**:
$$m := \max \left\{ q \in \mathbb{N} : \sum_{E \in \mathcal{E}} \binom{q}{\leq k - i_E} > 1 \right\}$$

Optimality of strategy

- If the chips are infinitely divisible then the solution is min/max optimal
- [Spencer95] Enough for optimality if number of chips is at least


$$\Omega\left(2^{2^k}\right)$$

Equivalence to a random walk

- Both sides playing optimally is equivalent to each chip performing an independent random walk.
- **Potential** = **probability** of a chip in bin **i** ending in bins **1..K** after **m** iteration
- **Weight** = difference between the potentials of a chip in its two possible locations on the following iteration.
- Chooser's optimal strategy: choose set with smaller (larger) **weight**

Example potential and weight

$m=5; k=10$




Boosting

- A method for improving classifier accuracy
- **Weak Learner**: a learning algorithm generating rules **slightly** better than random guessing.
- **Basic idea**: **re-weight** training examples to force weak learner into different parts of the space.
- **Combine** weak rules by a **majority vote**.

Boosting as a chip game

- Chip = training example
- **Booster** assigns a weight to each chip, weights sum to 1.
- **Learner** selects a subset with weight $\geq \frac{1}{2} + \gamma$
 - Selected set moves a step right (correct)
 - Unselected set moves a step left (incorrect)
- Booster wins examples on right of origin.
 - Implies that majority vote is correct.

The boosting chip game









Binomial strategies for boosting

- **Learner** moves chip right independently with probability $\frac{1}{2} + \gamma$
- The **potential** of example that is correctly classified **r** times after **i** out of **k** iterations

$$\sum_{j=0}^{\lfloor \frac{k}{2} \rfloor - r} \binom{k-i}{j} \left(\frac{1}{2} + \gamma\right)^j \left(\frac{1}{2} - \gamma\right)^{k-i-j}$$

- **Booster** assigns to each example a **weight** proportional to the **difference** between its possible next-step potentials.


Drifting games (in 2d)



$$\sum_i \|\vec{w}_i\| = 1$$

$$\sum_i \vec{w}_i \vec{x}_i \geq \delta > 0$$

Drifting games (in 2d)



Target
region




The allowable steps

B = the set of all allowable step

Normal B = minimal set that spans the space. (~basis)

Regular B = a symmetric regular set. (~orthonormal basis)

Regular step set in \mathbb{R}^2



The min/max solution

[Schapire99]

- A potential defined by a min/max recursion

$$\phi_T(\mathbf{s}) = L(\mathbf{s})$$

$$\phi_{t-1}(\mathbf{s}) = \min_{\mathbf{w} \in \mathbb{R}^d} \sup_{\mathbf{z} \in B} (\phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w} \cdot \mathbf{z} - \delta \|\mathbf{w}\|)$$

$$\phi_0(0) = \text{the value of the game}$$

- Shepherd's strategy

$$\mathbf{w}_i^t = \arg \min_{\mathbf{w}} \sup_{\mathbf{z} \in B} (\phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w} \cdot \mathbf{z} - \delta \|\mathbf{w}\|)$$

The solution simplifies when $\delta \rightarrow 0$


If B is normal, and δ is sufficiently small then $\exists \mathbf{w}^*$ such that

$$\phi_{t-1}(\mathbf{s}) = \phi_t(\mathbf{s} + \mathbf{z}) + \mathbf{w}^* \cdot \mathbf{z} - \delta \|\mathbf{w}^*\|$$

for all $\mathbf{z} \in B$ (and all $t = 1, 2, \dots, \mathbf{s} \in \mathbb{R}^d$)

Implies that: \mathbf{w}^* is the “local slope” at $\phi_t(\mathbf{s})$, i.e.

$$\phi_t(\mathbf{s} + \mathbf{z}_i) = C + \mathbf{w}^* \mathbf{z}_i ; \quad C \doteq \frac{\sum_{j=0}^d \phi_t(\mathbf{s} + \mathbf{z}_j)}{d+1}$$



and that

$$\phi_{t-1}(\mathbf{s}) = C - \delta \|\mathbf{w}^*\|$$

Increasing the number of steps

- Consider T steps in a unit time
- Drift δ should scale like $1/T$
- Step size $O(1/T)$ gives game to shepherd
- Step size $O(1/\sqrt{T})$ keeps game balanced

The continuous time limit

For $t = 0, 1, 2, \dots, T$ instantaneous loss is $l_i^t \in \{-1, +1\}$

Instead of $t = 0, 1, 2, \dots, T$ let $t = 0, \frac{1}{T}, \frac{2}{T}, \dots, 1$


Let $T \rightarrow \infty$

How should we set the instantaneous loss?


Natural choice: $l_i^t = \pm \frac{1}{T}$?

For optimal adversary, cumulative loss of each expert defines a random walk. What is random walk in continuous time?


The game lattice




Using step $l_i^t = \pm \frac{1}{T}$



Using step $l_i^t = \pm \frac{1}{T}$




Using step $l_i^t = \pm \frac{1}{T}$




Looks fine but $\text{var}(L) = T \frac{1}{T^2} = \frac{1}{T} \xrightarrow{T \rightarrow \infty} 0$

Using step $l_i^t = \pm \frac{1}{\sqrt{T}}$




Using step $l_i^t = \pm \frac{1}{\sqrt{T}}$

L $T = 3, \Delta t = \frac{1}{3}, l_i^t = \pm \frac{1}{\sqrt{3}}$



Using step $l_i^t = \pm \frac{1}{\sqrt{T}}$



The solution when $T \rightarrow \infty$

The local slope becomes the gradient

$$\mathbf{w}^* = \nabla \phi_\tau(\mathbf{s})$$

The recursion becomes a PDE

$$\begin{aligned} \frac{\partial \phi_\tau(\mathbf{s})}{\partial \tau} &= -\frac{1}{2} \sum_{k=1}^d \frac{\partial^2 \phi_\tau(\mathbf{s})}{\partial^2 s_k} + \delta \|\mathbf{w}^*\| \\ &= -\frac{1}{2} \Delta \phi_\tau(\mathbf{s}) + \delta \|\nabla \phi_\tau(\mathbf{s})\| \end{aligned}$$

Same PDE describes time development of
Brownian motion with drift


Special Cases

- Target func. = PDE boundary condition at time=1.
- Target func. = **exponential**
 - potential and weight are exponential at all times
 - Adaboost.
 - Exponential weights algs for online learning.
- Target func. = **step function**
 - Potential is the error function, weight is the normal distribution.
 - Potential and weight change with time.
 - Brownboost.


Potential applications

- Generalized boosting
 - Classification for >2 labels
 - Regression
 - Density estimation
 - Variational function fitting
- Generalized online learning
 - Continuous predictions (instead of binary)

solution for $L(s) = I_{|s|>1}$



Boosting for variational optimization






Figure 2: The potential $\phi(\mathbf{s}, t)$ for the square loss $L(y) = y^2$.




Figure 3: The potential $\phi(\mathbf{s}, t)$ for the loss $L(y) = \min(y^2, 1)$.

Plan of talk

- How to use expert advice,
A quick tour of some old results.
- Decision Theoretic Online Learning
 - Hedge
 - Binomial Weights
 - Normal Hedge
 - Open problem

The Hedge Algorithm

[Freund & Schapire 1997]

based on [Littlestone and Warmuth 1989],[Cesa-Bianchi et al 1997]

Initial weights: $w^1 = \left\langle \frac{1}{N}, \dots, \frac{1}{N} \right\rangle$

Weights update rule: $w_i^{t+1} = w_i^t e^{-\eta l_i^t} = e^{-\eta L_i^{t+1}}$ **Learning rate**

Alternatively: $w_i^{t+1} = \frac{1}{N} e^{-\eta L_i^t}$

Posterior
probability
(un-normalized)

Prior
probability

Potential-based bound

Potential:
$$W^{t+1} = \sum_{i=1}^N w_i^{t+1} = \sum_{i=1}^N e^{-\eta L_i^t}$$

Large loss of algorithm \Rightarrow small potential
$$L_A^t \leq \frac{-\log W^{t+1}}{1 - e^{-\eta}}$$

Good expert \Rightarrow large potential
$$W^{t+1} \geq w_i^{t+1} = e^{-\eta L_i^t}$$

Combining, we get:
$$\forall i, L_A^T \leq \frac{\eta L_i^T + \ln N}{1 - e^{-\eta}}$$

Tuning the learning rate

$$\forall i, L_A^T \leq \frac{\eta L_i^T + \ln N}{1 - e^{-\eta}}$$

If we set $\eta = \sqrt{\frac{2 \ln N}{T}}$

Then we guarantee $L_A^T \leq \min_i L_i^T + \sqrt{2T \ln N} + \ln N$

Equivalently $\forall i, R_i^T \leq \sqrt{2T \ln N} + \ln N; \quad \lim_{T \rightarrow \infty} \frac{\sqrt{2T \ln N} + \ln N}{T} = 0$


Tuning the learning rate

$$\forall i, L_A^T \leq \frac{\eta L_i^T + \ln N}{1 - e^{-\eta}}$$

If we set $\eta = \sqrt{\frac{2 \ln N}{T}}$

Then we guarantee $L_A^T \leq \min_i L_i^T + \sqrt{2T \ln N} + \ln N$

Equivalently $\forall i, R_i^T \leq \sqrt{2T \ln N} + \ln N; \quad \lim_{T \rightarrow \infty} \frac{\sqrt{2T \ln N} + \ln N}{T} = 0$



Is it possible to do better?

Plan of talk

- How to use expert advice,
A quick tour of some old results.
- Decision Theoretic Online Learning
 - Hedge
 - Binomial Weights
 - **Normal Hedge**
 - Open problem

Design of NormalHedge

- BW: potential function depends on **loss** and **number of remaining mistakes**
- Normal-Hedge: Potential function based on **regret** and **variance of the positive cumulative regrets**

The NormalHedge potential

$$\text{Potential: } \psi(r, c) = \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$

$$\text{Weight: } w(R, c) = \frac{\partial}{\partial R} \psi(R, c) = \begin{cases} \frac{R}{c} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 0 & \text{if } R \leq 0 \end{cases}$$

Intuition: If we play against the random walk player, then

the probability that the cumulative regret is R is approximately $e^{-\frac{R^2}{2t}}$,
to have any hope of keeping the potential constant the potential
function cannot increase faster than $1/\text{probability}$.

NormalHedge algorithm

for $t=0,1,2,\dots$

if $\forall i, R_i^t \leq 0$ then $w_i^t = 1 / N$

else

set $c(t)$ so that $\frac{1}{N} \sum_{i=1}^N \psi(R_i^t, c(t)) = e$

$$w_i^t = w(R_i^t, c(t))$$

Incur instantaneous losses: $\langle l_1^t, l_2^t, \dots, l_N^t \rangle$

$$\text{Algorithm loss: } l_A^t = \frac{\sum_{i=1}^N w_i^t l_i^t}{\sum_{i=1}^N w_i^t}$$

Update regrets: $R_i^{t+1} = R_i^t + l_A^t - l_i^t$


Illustrative Example

● Expert

● Algorithm

— $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




Illustrative Example

● Expert

● Algorithm

— $\exp(\eta G)$

$$\text{— } \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




Illustrative Example

● Expert

● Algorithm

— $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$




Illustrative Example

● Expert

● Algorithm

— $\exp(\eta G)$

$$\text{—} \begin{cases} \exp\left(\frac{R^2}{2c}\right) & \text{if } R \geq 0 \\ 1 & \text{if } R \leq 0 \end{cases}$$



Normal-Hedge Performance bound

[Chaudhuri, Freund & Hsu 2009]

Main Lemma: $c(t) \leq t + C_N$

The regret of NormalHedge is upper bounded by

$$\sqrt{3t(1 + \ln N) + o(t)}$$

The regret to the top ϵ -percentile is upper bounded by


$$O\left(\sqrt{t \ln(1/\epsilon)} + \ln^2 N\right)$$

We failed to get rid of the dependence on N

BW depends only on $1/\epsilon$ not on N

Performance on flip-flop

- Worst case for follow the leader



NormalHedge.DT / Luo and Schapire, 2014

$$\text{Potential: } \psi(r, t) = \begin{cases} \exp\left(\frac{r^2}{3t}\right) & \text{if } r \geq 0 \\ 1 & \text{if } r \leq 0 \end{cases}$$

$$\text{Weight: } w(r, t) = \psi(r + 1, t + 1) - \psi(r + 1, t - 1)$$

$$\text{Recall NormalHedge weight: } w(r, c) = \frac{\partial}{\partial r} \psi(r, c)$$

A very clean application of drifting games analysis

NormalHedge.DT / analysis

Instead of keeping potential constant,
allow potential to increase like $c \log(t)$

The regret of AdaNormalHedge is bounded by

$$\sqrt{3t \ln \left(\frac{1}{2\epsilon} (e^{4/3} - 1) (\ln t + 1) + 1 \right)} = O \left(\sqrt{t \ln(1/\epsilon)} + t \ln \ln t \right)$$

for all t and for all $\epsilon > 0$

Algorithm does not depend on ϵ !!!

Plan of talk

- How to use expert advice,
A quick tour of some old results.
- Decision Theoretic Online Learning
 - Hedge
 - Binomial Weights
 - Normal Hedge
 - Open problem

Open Problem

- Luo and Schapire got rid of the dependence on N
- What about the dependence on T ?
- Suppose that, although loss is bounded in $[0, 1]$, it turns out that the loss is always in $[0, 1/2]$.
- If alg. has information in advance, bound is halved
- Equivalently: $T \rightarrow \frac{T}{4}$
- Can we achieve this performance without a -priori knowledge?
- Hypothesis: The increase $c(t+1) - c(t)$ for normal hedge can upper bounded by a function of the form

$$\int_{\theta} \alpha_{\theta} (r_{\theta}^t)^2 d\mu(\theta) \quad (\text{current bound is } 1)$$

- Without degrading current performance.

Recent Progress

Schapire and Liu 2015:

[AdaNormalHedge](#), an adaptive version of NormalHedge.DT

The regret of AdaNormalHedge relative to the ϵ best experts at time T

$$\forall T, \epsilon \quad R(\epsilon, T) \leq O\left(\sqrt{T \log\left(\frac{\log T}{\epsilon}\right)}\right)$$

Open Problem

- Luo and Schapire got rid of the dependence on N
- What about the dependence on T ?
- Suppose that, although loss is bounded in $[0, 1]$, it turns out that the loss is always in $[0, 1/2]$.
- If alg. has information in advance, bound is halved
- Equivalently: $T \rightarrow \frac{T}{4}$
- Can we achieve this performance without a -priori knowledge?
- Hypothesis: The increase $c(t+1) - c(t)$ for normal hedge can upper bounded by a function of the form


$$\int_{\theta} \alpha_{\theta} (r_{\theta}^t)^2 d\mu(\theta) \quad (\text{current bound is } 1)$$

- Without degrading current performance.


Hypothesis


- In the bound on NormalHedge, T can be replaced by cumulative variance.
- Similar to measuring return vs. volatility in the stock market.

Hedging differentiable processes is trivial



Intersections of Stochastic processes are complex





Thank You!

<http://www.mindreaderpro.appspot.com/>

$\phi(x,y,t=1)$

$\phi(x,y,t=0)$

