

Exponential Weights Algorithms for Online Learning

Yoav Freund

January 11, 2018

Outline

The hedging problem

- ▶ N possible actions
- ▶ At each time step $t = 1, 2, \dots, T$:
 - ▶ Algorithm chooses a distribution \mathbf{p}^t over actions.
 - ▶ Losses $0 \leq \ell_i^t \leq 1$ of all actions $i = 1, \dots, N$ are revealed.
 - ▶ Algorithm suffers **expected** loss $\mathbf{p}^t \cdot \boldsymbol{\ell}^t$
- ▶ **Goal:** minimize total expected loss
- ▶ Here we have stochasticity - but only in **algorithm**, not in **outcome**
- ▶ Fits nicely in game theory

Hedging vs. Halving

- ▶ Like halving - we want to zoom into best action (expert).
- ▶ Unlike halving - no action is perfect.
- ▶ Basic idea - reduce probability of lossy actions, but **not all the way to zero**.
- ▶ **Modified Goal:** minimize **difference between** expected total loss and minimal total loss of repeating one action.

$$\sum_{t=1}^T \mathbf{p}^t \cdot \ell^t - \min_i \left(\sum_{t=1}^T \ell_i^t \right)$$

Using hedge to generalize halving alg.

- ▶ Suppose that there is no perfect expert.
- ▶ action i = use prediction of expert i
- ▶ Now each iteration of game consistst of **three** steps:
 - ▶ Experts make predictions $e_i^t \in \{0, 1\}$
 - ▶ Algorithm predicts **1** with probability $\sum_{i: e_i^t = 1} p_i^t$.
 - ▶ outcome o_i^t is revealed. $\ell_i^t = 0$ if $e_i^t = o_i^t$, $\ell_i^t = 1$ otherwise.

The Hedge(η) Algorithm

Consider action i at time t

- ▶ Total loss:

$$L_i^t = \sum_{s=1}^{t-1} \ell_i^s$$

- ▶ Weight:

$$w_i^t = w_i^1 e^{-\eta L_i^t}$$

Note freedom to choose initial weight (w_i^1) $\sum_{i=1}^n w_i^1 = 1$.

- ▶ $\eta > 0$ is the learning rate parameter. Halving: $\eta \rightarrow \infty$
- ▶ Probability:

$$p_i^t = \frac{w_i^t}{\sum_{j=1}^N w_j^t}, \quad \mathbf{p}^t = \frac{\mathbf{w}^t}{\sum_{j=1}^N w_j^t}$$

Choosing the initial weights

- ▶ Giving an action high initial weight makes alg perform well **if** that action performs well.
- ▶ If good action has low initial weight, our total loss will be larger.
- ▶ As $\sum_{i=1}^n w_i^1 = 1$ increasing one weight implies decreasing some others.
- ▶ Plays a similar role to prior distribution in Bayesian algorithms.

Bound on the loss of **Hedge**(η) Algorithm

► Theorem (main theorem)

For any sequence of loss vectors ℓ^1, \dots, ℓ^T , and for any $i \in \{1, \dots, N\}$, we have

$$L_{\text{Hedge}(\eta)} \leq \frac{-\ln(w_i^1) + \eta L_i}{1 - e^{-\eta}}.$$

- **Proof:** by combining upper and lower bounds on $\sum_{i=1}^N w_i^{T+1}$

Hedge(η)

└ Bound on total loss

└ Upper bound on $\sum_{i=1}^N w_i^{T+1}$

Upper bound on $\sum_{i=1}^N w_i^{T+1}$


Lemma (upper bound)

For any sequence of loss vectors ℓ^1, \dots, ℓ^T we have

$$\ln \left(\sum_{i=1}^N w_i^{T+1} \right) \leq -(1 - e^{-\eta}) L_{\text{Hedge}(\eta)}.$$

Proof of upper bound (slide 1)

- ▶ If $a \geq 0$ then a^r is convex.
- ▶ For $r \in [0, 1]$, $a^r \leq 1 - (1 - a)r$



Hedge(η)

└ Bound on total loss

└ Upper bound on $\sum_{i=1}^N w_i^{T+1}$

Proof of upper bound (slide 2)

Applying $a^r \leq 1 - (1 - a)^r$ where $a = e^{-\eta}, r = \ell_i^t$

$$\begin{aligned}\sum_{i=1}^N w_i^{t+1} &= \sum_{i=1}^N w_i^t e^{-\eta \ell_i^t} \\ &\leq \sum_{i=1}^N w_i^t (1 - (1 - e^{-\eta}) \ell_i^t) \\ &= \left(\sum_{i=1}^N w_i^t \right) \left(1 - (1 - e^{-\eta}) \frac{\mathbf{w}^t}{\sum_{i=1}^N w_i^t} \cdot \boldsymbol{\ell}^t \right) \\ &= \left(\sum_{i=1}^N w_i^t \right) (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \boldsymbol{\ell}^t)\end{aligned}$$

Proof of upper bound (slide 3)

► Combining

$$\sum_{i=1}^N w_i^{t+1} \leq \left(\sum_{i=1}^N w_i^t \right) (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t)$$

► for $t = 1, \dots, T$

► yields

$$\begin{aligned} \sum_{i=1}^N w_i^{T+1} &\leq \prod_{t=1}^T (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t) \\ &\leq \exp \left(-(1 - e^{-\eta}) \sum_{t=1}^T \mathbf{p}^t \cdot \ell^t \right) \end{aligned}$$

since $1 + x \leq e^x$ for $x = -(1 - e^{-\eta})$.

Hedge(η)

└ Bound on total loss

└ Lower bound on $\sum_{i=1}^N w_i^{T+1}$

Lower bound on $\sum_{i=1}^N w_i^{T+1}$

For any $j = 1, \dots, N$:

$$\sum_{i=1}^N w_i^{T+1} \geq w_j^{T+1} = w_j^1 e^{-\eta L_j}$$

Combining Upper and Lower bounds

- ▶ Combining bounds on $\ln \left(\sum_{i=1}^N w_i^{T+1} \right)$

$$\ln w_j^1 - \eta L_j \leq \ln \sum_{i=1}^N w_i^{T+1} \leq -(1 - e^{-\eta}) \sum_{t=1}^T \mathbf{p}^t \cdot \ell^t$$


- ▶ Reversing signs, using $L_{\text{Hedge}(\eta)} = \sum_{t=1}^T \mathbf{p}^t \cdot \ell^t$ and reorganizing we get

$$L_{\text{Hedge}(\eta)} \leq \frac{-\ln(w_i^1) + \eta L_i}{1 - e^{-\eta}}$$

Tuning η

How to Use Expert Advice

451



Tuning η

- ▶ Suppose $\min_i L_i \leq \tilde{L}$
- ▶ set

$$\eta = \ln \left(1 + \sqrt{\frac{2 \ln N}{\tilde{L}}} \right) \approx \sqrt{\frac{2 \ln N}{\tilde{L}}}$$

- ▶ use uniform initial weights $\mathbf{w}^1 = \langle 1/N, \dots, 1/N \rangle$
- ▶ Then

$$L_{\text{Hedge}(\eta)} \leq \frac{-\ln(w_i^1) + \eta L_i}{1 - e^{-\eta}} \leq \min_i L_i + \sqrt{2\tilde{L} \ln N} + \ln N$$

Tuning η as a function of T

- ▶ trivially $\min_i L_i \leq T$, yielding

$$L_{\text{Hedge}(\eta)} \leq \min_i L_i + \sqrt{2T \ln N} + \ln N$$

- ▶ per iteration we get:

$$\frac{L_{\text{Hedge}(\eta)}}{T} \leq \min_i \frac{L_i}{T} + \sqrt{\frac{2 \ln N}{T}} + \frac{\ln N}{T}$$

How good is this bound?

- ▶ **Very good!** There is a closely matching lower bound!
- ▶ There exists a stochastic adversarial strategy such that with high probability for **any** hedging strategy **S** after **T** trials

$$L_S - \min_i L_i \geq (1 - o(1))\sqrt{2T \ln N}$$

- ▶ The adversarial strategy is random, extremely simple, and does not depend on the hedging strategy!

The adversarial strategy

- ▶ Adversary sets each loss ℓ_i^t independently at random to 0 or 1 with equal probabilities (1/2, 1/2).
- ▶ Obviously, nothing to learn !
 $L_S \approx T/2$.
- ▶ On the other hand $\min_i L_i \approx T/2 - \sqrt{2T \ln N}$
- ▶ The difference $L_S - \min_i L_i$ is due to unlearnable random fluctuations!
- ▶ Detailed proof quite involved. See games paper.

Summary

- ▶ Given learning rate η the **Hedge**(η) algorithm satisfies

$$L_{\text{Hedge}(\eta)} \leq \frac{\ln N + \eta L_i}{1 - e^{-\eta}}$$

- ▶ Setting $\eta \approx \sqrt{\frac{2 \ln N}{T}}$ guarantees

$$L_{\text{Hedge}(\eta)} \leq \min_i L_i + \sqrt{2T \ln N} + \ln N$$

- ▶ A trivial random data, in which there is nothing to be learned forces **any** algorithm to suffer this total loss

Some loose threads

- ▶ Total Loss of best action usually scales linearly with time, but we can't change η on the fly. (I think El Yaniv proposed a reasonable solution).
- ▶ Observing only the loss of chosen action - the multi-armed bandit problem. Will get to that later in the course.
- ▶ Register to the class on the google drive.