

# Exponential Weights Algorithms for Online Learning

Yoav Freund

January 11, 2018

# Outline

## The Halving Algorithm

# Outline

The Halving Algorithm

**Hedge( $\eta$ )**Algorithm

Hedging vs. Halving

# Outline

## The Halving Algorithm

## Hedge( $\eta$ ) Algorithm

Hedging vs. Halving

## Bound on total loss

Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

Lower bound on  $\sum_{i=1}^N w_i^{T+1}$

Combining Upper and Lower bounds

# Outline

## The Halving Algorithm

## Hedge( $\eta$ ) Algorithm

Hedging vs. Halving

## Bound on total loss

Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

Lower bound on  $\sum_{i=1}^N w_i^{T+1}$

Combining Upper and Lower bounds

tuning  $\eta$

# Outline

The Halving Algorithm

**Hedge( $\eta$ )**Algorithm

Hedging vs. Halving

Bound on total loss

Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

Lower bound on  $\sum_{i=1}^N w_i^{T+1}$

Combining Upper and Lower bounds

tuning  $\eta$

Lower Bounds

# The hedging problem

- ▶  $N$  possible actions

## The hedging problem

- ▶  $N$  possible actions
- ▶ At each time step  $t = 1, 2, \dots, T$ :



## The hedging problem

- ▶  $N$  possible actions
- ▶ At each time step  $t = 1, 2, \dots, T$ :
  - ▶ Algorithm chooses a distribution  $\mathbf{p}^t$  over actions.

## The hedging problem

- ▶  $N$  possible actions
- ▶ At each time step  $t = 1, 2, \dots, T$ :
  - ▶ Algorithm chooses a distribution  $\mathbf{p}^t$  over actions.
  - ▶ Losses  $0 \leq \ell_i^t \leq 1$  of all actions  $i = 1, \dots, N$  are revealed.

## The hedging problem

- ▶  $N$  possible actions
- ▶ At each time step  $t = 1, 2, \dots, T$ :
  - ▶ Algorithm chooses a distribution  $\mathbf{p}^t$  over actions.
  - ▶ Losses  $0 \leq \ell_i^t \leq 1$  of all actions  $i = 1, \dots, N$  are revealed.
  - ▶ Algorithm suffers **expected** loss  $\mathbf{p}^t \cdot \boldsymbol{\ell}^t$

## The hedging problem

- ▶  $N$  possible actions
- ▶ At each time step  $t = 1, 2, \dots, T$ :
  - ▶ Algorithm chooses a distribution  $\mathbf{p}^t$  over actions.
  - ▶ Losses  $0 \leq \ell_i^t \leq 1$  of all actions  $i = 1, \dots, N$  are revealed.
  - ▶ Algorithm suffers **expected** loss  $\mathbf{p}^t \cdot \boldsymbol{\ell}^t$
- ▶ **Goal:** minimize total expected loss

## The hedging problem

- ▶  $N$  possible actions
- ▶ At each time step  $t = 1, 2, \dots, T$ :
  - ▶ Algorithm chooses a distribution  $\mathbf{p}^t$  over actions.
  - ▶ Losses  $0 \leq \ell_i^t \leq 1$  of all actions  $i = 1, \dots, N$  are revealed.
  - ▶ Algorithm suffers **expected** loss  $\mathbf{p}^t \cdot \boldsymbol{\ell}^t$
- ▶ **Goal:** minimize total expected loss
- ▶ Here we have stochasticity - but only in **algorithm**, not in **outcome**

## The hedging problem

- ▶  $N$  possible actions
- ▶ At each time step  $t = 1, 2, \dots, T$ :
  - ▶ Algorithm chooses a distribution  $\mathbf{p}^t$  over actions.
  - ▶ Losses  $0 \leq \ell_i^t \leq 1$  of all actions  $i = 1, \dots, N$  are revealed.
  - ▶ Algorithm suffers **expected** loss  $\mathbf{p}^t \cdot \boldsymbol{\ell}^t$
- ▶ **Goal:** minimize total expected loss
- ▶ Here we have stochasticity - but only in **algorithm**, not in **outcome**
- ▶ Fits nicely in game theory

## Hedging vs. Halving

- ▶ Like halving - we want to zoom into best action (expert).

## Hedging vs. Halving

- ▶ Like halving - we want to zoom into best action (expert).
- ▶ Unlike halving - no action is perfect.



## Hedging vs. Halving

- ▶ Like halving - we want to zoom into best action (expert).
- ▶ Unlike halving - no action is perfect.
- ▶ Basic idea - reduce probability of lossy actions, but **not all the way to zero**.

## Hedging vs. Halving

- ▶ Like halving - we want to zoom into best action (expert).
- ▶ Unlike halving - no action is perfect.
- ▶ Basic idea - reduce probability of lossy actions, but **not all the way to zero**.
- ▶ **Modified Goal:** minimize **difference between** expected total loss and minimal total loss of repeating one action.

$$\sum_{t=1}^T \mathbf{p}^t \cdot \ell^t - \min_i \left( \sum_{t=1}^T \ell_i^t \right)$$

## Using hedge to generalize halving alg.

- Suppose that there is no perfect expert.

## Using hedge to generalize halving alg.

- ▶ Suppose that there is no perfect expert.
- ▶ action  $i$  = use prediction of expert  $i$

## Using hedge to generalize halving alg.

- ▶ Suppose that there is no perfect expert.
- ▶ action  $i$  = use prediction of expert  $i$
- ▶ Now each iteration of game consistst of **three** steps:

## Using hedge to generalize halving alg.

- ▶ Suppose that there is no perfect expert.
- ▶ action  $i$  = use prediction of expert  $i$
- ▶ Now each iteration of game consistst of **three** steps:
  - ▶ Experts make predictions  $e_i^t \in \{0, 1\}$

## Using hedge to generalize halving alg.

- ▶ Suppose that there is no perfect expert.
- ▶ action  $i$  = use prediction of expert  $i$
- ▶ Now each iteration of game consistst of **three** steps:
  - ▶ Experts make predictions  $e_i^t \in \{0, 1\}$
  - ▶ Algorithm predicts **1** with probability  $\sum_{i: e_i^t=1} p_i^t$ .

## Using hedge to generalize halving alg.

- ▶ Suppose that there is no perfect expert.
- ▶ action  $i$  = use prediction of expert  $i$
- ▶ Now each iteration of game consistst of **three** steps:
  - ▶ Experts make predictions  $e_i^t \in \{0, 1\}$
  - ▶ Algorithm predicts **1** with probability  $\sum_{i: e_i^t = 1} p_i^t$ .
  - ▶ outcome  $o_i^t$  is revealed.  $\ell_i^t = 0$  if  $e_i^t = o_i^t$ ,  $\ell_i^t = 1$  otherwise.



## The Hedge( $\eta$ ) Algorithm

Consider action  $i$  at time  $t$

- ▶ Total loss:

$$L_i^t = \sum_{s=1}^{t-1} \ell_i^s$$

# The Hedge( $\eta$ ) Algorithm

Consider action  $i$  at time  $t$

- ▶ Total loss:

$$L_i^t = \sum_{s=1}^{t-1} \ell_i^s$$

- ▶ Weight:

$$w_i^t = w_i^1 e^{-\eta L_i^t}$$

Note freedom to choose initial weight ( $w_i^1$ )  $\sum_{i=1}^n w_i^1 = 1$ .

# The Hedge( $\eta$ ) Algorithm

Consider action  $i$  at time  $t$

- ▶ Total loss:

$$L_i^t = \sum_{s=1}^{t-1} \ell_i^s$$

- ▶ Weight:

$$w_i^t = w_i^1 e^{-\eta L_i^t}$$

Note freedom to choose initial weight ( $w_i^1$ )  $\sum_{i=1}^n w_i^1 = 1$ .

- ▶  $\eta > 0$  is the learning rate parameter. Halving:  $\eta \rightarrow \infty$

# The Hedge( $\eta$ ) Algorithm

Consider action  $i$  at time  $t$

- ▶ Total loss:

$$L_i^t = \sum_{s=1}^{t-1} \ell_i^s$$

- ▶ Weight:

$$w_i^t = w_i^1 e^{-\eta L_i^t}$$

Note freedom to choose initial weight ( $w_i^1$ )  $\sum_{i=1}^n w_i^1 = 1$ .

- ▶  $\eta > 0$  is the learning rate parameter. Halving:  $\eta \rightarrow \infty$
- ▶ Probability:

$$p_i^t = \frac{w_i^t}{\sum_{j=1}^N w_j^t},$$

# The Hedge( $\eta$ ) Algorithm

Consider action  $i$  at time  $t$

- ▶ Total loss:

$$L_i^t = \sum_{s=1}^{t-1} \ell_i^s$$

- ▶ Weight:

$$w_i^t = w_i^1 e^{-\eta L_i^t}$$

Note freedom to choose initial weight ( $w_i^1$ )  $\sum_{i=1}^n w_i^1 = 1$ .

- ▶  $\eta > 0$  is the learning rate parameter. Halving:  $\eta \rightarrow \infty$
- ▶ Probability:

$$p_i^t = \frac{w_i^t}{\sum_{j=1}^N w_j^t},$$

# The Hedge( $\eta$ ) Algorithm

Consider action  $i$  at time  $t$

- ▶ Total loss:

$$L_i^t = \sum_{s=1}^{t-1} \ell_i^s$$

- ▶ Weight:

$$w_i^t = w_i^1 e^{-\eta L_i^t}$$

Note freedom to choose initial weight ( $w_i^1$ )  $\sum_{i=1}^n w_i^1 = 1$ .

- ▶  $\eta > 0$  is the learning rate parameter. Halving:  $\eta \rightarrow \infty$
- ▶ Probability:

$$p_i^t = \frac{w_i^t}{\sum_{j=1}^N w_j^t}, \quad \mathbf{p}^t = \frac{\mathbf{w}^t}{\sum_{j=1}^N w_j^t}$$

## Choosing the initial weights

- ▶ Giving an action high initial weight makes alg perform well  
if that action performs well.

## Choosing the initial weights

- ▶ Giving an action high initial weight makes alg perform well **if** that action performs well.
- ▶ If good action has low initial weight, our total loss will be larger.



## Choosing the initial weights

- ▶ Giving an action high initial weight makes alg perform well **if** that action performs well.
- ▶ If good action has low initial weight, our total loss will be larger.
- ▶ As  $\sum_{i=1}^n w_i^1 = 1$  increasing one weight implies decreasing some others.

## Choosing the initial weights

- ▶ Giving an action high initial weight makes alg perform well **if** that action performs well.
- ▶ If good action has low initial weight, our total loss will be larger.
- ▶ As  $\sum_{i=1}^n w_i^1 = 1$  increasing one weight implies decreasing some others.
- ▶ Plays a similar role to prior distribution in Bayesian algorithms.

## Bound on the loss of **Hedge**( $\eta$ ) Algorithm

## Bound on the loss of **Hedge**( $\eta$ ) Algorithm

► Theorem (main theorem)

For any sequence of loss vectors  $\ell^1, \dots, \ell^T$ , and for any  $i \in \{1, \dots, N\}$ , we have

$$L_{\text{Hedge}(\eta)} \leq \frac{-\ln(w_i^1) + \eta L_i}{1 - e^{-\eta}}.$$

## Bound on the loss of **Hedge**( $\eta$ ) Algorithm

► Theorem (main theorem)

For any sequence of loss vectors  $\ell^1, \dots, \ell^T$ , and for any  $i \in \{1, \dots, N\}$ , we have

$$L_{\text{Hedge}(\eta)} \leq \frac{-\ln(w_i^1) + \eta L_i}{1 - e^{-\eta}}.$$

- **Proof:** by combining upper and lower bounds on  $\sum_{i=1}^N w_i^{T+1}$

## Hedge( $\eta$ )

└ Bound on total loss

└ Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

# Upper bound on $\sum_{i=1}^N w_i^{T+1}$

## Lemma (upper bound)

For any sequence of loss vectors  $\ell^1, \dots, \ell^T$  we have

$$\ln \left( \sum_{i=1}^N w_i^{T+1} \right) \leq -(1 - e^{-\eta}) L_{\text{Hedge}(\eta)}.$$

## Hedge( $\eta$ )

└ Bound on total loss

└ Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

## Proof of upper bound (slide 1)

- ▶ If  $a \geq 0$  then  $a^r$  is convex.

## Hedge( $\eta$ )

└ Bound on total loss

└ Upper bound on  $\sum_{i=1}^N w_i^{T+1}$


## Proof of upper bound (slide 1)

- ▶ If  $a \geq 0$  then  $a^r$  is convex.
- ▶ For  $r \in [0, 1]$ ,  $a^r \leq 1 - (1 - a)r$



## Proof of upper bound (slide 1)

- ▶ If  $a \geq 0$  then  $a^r$  is convex.
- ▶ For  $r \in [0, 1]$ ,  $a^r \leq 1 - (1 - a)r$



## Hedge( $\eta$ )

└ Bound on total loss

└ Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

## Proof of upper bound (slide 2)

Applying  $a^r \leq 1 - (1 - a)^r$  where  $a = e^{-\eta}$ ,  $r = \ell_i^t$

$$\sum_{i=1}^N w_i^{t+1} = \sum_{i=1}^N w_i^t e^{-\eta \ell_i^t}$$

## Hedge( $\eta$ )

└ Bound on total loss

└ Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

## Proof of upper bound (slide 2)

Applying  $a^r \leq 1 - (1 - a)^r$  where  $a = e^{-\eta}, r = \ell_i^t$

$$\begin{aligned}\sum_{i=1}^N w_i^{t+1} &= \sum_{i=1}^N w_i^t e^{-\eta \ell_i^t} \\ &\leq \sum_{i=1}^N w_i^t (1 - (1 - e^{-\eta}) \ell_i^t)\end{aligned}$$

## Hedge( $\eta$ )

└ Bound on total loss

└ Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

## Proof of upper bound (slide 2)

Applying  $a^r \leq 1 - (1 - a)^r$  where  $a = e^{-\eta}, r = \ell_i^t$

$$\begin{aligned}\sum_{i=1}^N w_i^{t+1} &= \sum_{i=1}^N w_i^t e^{-\eta \ell_i^t} \\ &\leq \sum_{i=1}^N w_i^t (1 - (1 - e^{-\eta}) \ell_i^t) \\ &= \left( \sum_{i=1}^N w_i^t \right) \left( 1 - (1 - e^{-\eta}) \frac{\sum_{i=1}^N w_i^t \ell_i^t}{\sum_{i=1}^N w_i^t} \right)\end{aligned}$$

## Proof of upper bound (slide 2)

Applying  $a^r \leq 1 - (1 - a)^r$  where  $a = e^{-\eta}, r = \ell_i^t$

$$\begin{aligned}\sum_{i=1}^N w_i^{t+1} &= \sum_{i=1}^N w_i^t e^{-\eta \ell_i^t} \\ &\leq \sum_{i=1}^N w_i^t (1 - (1 - e^{-\eta}) \ell_i^t) \\ &= \left( \sum_{i=1}^N w_i^t \right) \left( 1 - (1 - e^{-\eta}) \frac{\mathbf{w}^t}{\sum_{i=1}^N w_i^t} \cdot \boldsymbol{\ell}^t \right) \\ &= \left( \sum_{i=1}^N w_i^t \right) (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \boldsymbol{\ell}^t)\end{aligned}$$

## Hedge( $\eta$ )

└ Bound on total loss

└ Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

## Proof of upper bound (slide 3)

### ► Combining

$$\sum_{i=1}^N w_i^{t+1} \leq \left( \sum_{i=1}^N w_i^t \right) (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t)$$

## Hedge( $\eta$ )

└ Bound on total loss

└ Upper bound on  $\sum_{i=1}^N w_i^{T+1}$

## Proof of upper bound (slide 3)

### ► Combining

$$\sum_{i=1}^N w_i^{t+1} \leq \left( \sum_{i=1}^N w_i^t \right) (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t)$$

### ► for $t = 1, \dots, T$

## Proof of upper bound (slide 3)

► Combining

$$\sum_{i=1}^N w_i^{t+1} \leq \left( \sum_{i=1}^N w_i^t \right) (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t)$$

► for  $t = 1, \dots, T$

► yields

$$\sum_{i=1}^N w_i^{T+1} \leq \prod_{t=1}^T (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t)$$



## Proof of upper bound (slide 3)

► Combining

$$\sum_{i=1}^N w_i^{t+1} \leq \left( \sum_{i=1}^N w_i^t \right) (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t)$$

► for  $t = 1, \dots, T$

► yields

$$\sum_{i=1}^N w_i^{T+1} \leq \prod_{t=1}^T (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t)$$

## Proof of upper bound (slide 3)

► Combining

$$\sum_{i=1}^N w_i^{t+1} \leq \left( \sum_{i=1}^N w_i^t \right) (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t)$$

► for  $t = 1, \dots, T$

► yields

$$\begin{aligned} \sum_{i=1}^N w_i^{T+1} &\leq \prod_{t=1}^T (1 - (1 - e^{-\eta}) \mathbf{p}^t \cdot \ell^t) \\ &\leq \exp \left( -(1 - e^{-\eta}) \sum_{t=1}^T \mathbf{p}^t \cdot \ell^t \right) \end{aligned}$$

since  $1 + x \leq e^x$  for  $x = -(1 - e^{-\eta})$ .

## Hedge( $\eta$ )

└ Bound on total loss

└ Lower bound on  $\sum_{i=1}^N w_i^{T+1}$

## Lower bound on $\sum_{i=1}^N w_i^{T+1}$

For any  $j = 1, \dots, N$ :

$$\sum_{i=1}^N w_i^{T+1} \geq w_j^{T+1} = w_j^1 e^{-\eta L_j}$$

## Combining Upper and Lower bounds

- Combining bounds on  $\ln \left( \sum_{i=1}^N w_i^{T+1} \right)$

$$\ln w_j^1 - \eta L_j \leq \ln \sum_{i=1}^N w_i^{T+1} \leq -(1 - e^{-\eta}) \sum_{t=1}^T \mathbf{p}^t \cdot \ell^t$$

## Combining Upper and Lower bounds

- Combining bounds on  $\ln \left( \sum_{i=1}^N w_i^{T+1} \right)$

$$\ln w_j^1 - \eta L_j \leq \ln \sum_{i=1}^N w_i^{T+1} \leq -(1 - e^{-\eta}) \sum_{t=1}^T \mathbf{p}^t \cdot \ell^t$$

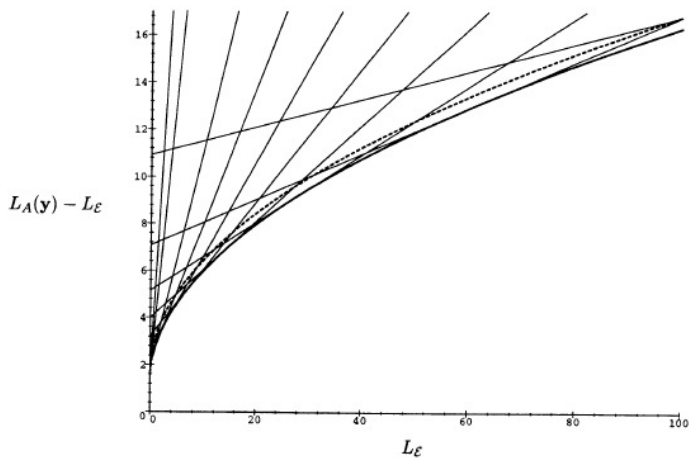
- Reversing signs, using  $L_{\text{Hedge}(\eta)} = \sum_{t=1}^T \mathbf{p}^t \cdot \ell^t$  and reorganizing we get

$$L_{\text{Hedge}(\eta)} \leq \frac{-\ln(w_i^1) + \eta L_i}{1 - e^{-\eta}}$$

# Tuning $\eta$

*How to Use Expert Advice*

451



## Tuning $\eta$

- ▶ Suppose  $\min_i L_i \leq \tilde{L}$

## Tuning $\eta$

- ▶ Suppose  $\min_i L_i \leq \tilde{L}$
- ▶ set

$$\eta = \ln \left( 1 + \sqrt{\frac{2 \ln N}{\tilde{L}}} \right) \approx \sqrt{\frac{2 \ln N}{\tilde{L}}}$$



## Tuning $\eta$

- ▶ Suppose  $\min_i L_i \leq \tilde{L}$
- ▶ set

$$\eta = \ln \left( 1 + \sqrt{\frac{2 \ln N}{\tilde{L}}} \right) \approx \sqrt{\frac{2 \ln N}{\tilde{L}}}$$

- ▶ use uniform initial weights  $\mathbf{w}^1 = \langle 1/N, \dots, 1/N \rangle$

## Tuning $\eta$

- ▶ Suppose  $\min_i L_i \leq \tilde{L}$
- ▶ set

$$\eta = \ln \left( 1 + \sqrt{\frac{2 \ln N}{\tilde{L}}} \right) \approx \sqrt{\frac{2 \ln N}{\tilde{L}}}$$

- ▶ use uniform initial weights  $\mathbf{w}^1 = \langle 1/N, \dots, 1/N \rangle$
- ▶ Then

$$L_{\text{Hedge}(\eta)} \leq \frac{-\ln(w_i^1) + \eta L_i}{1 - e^{-\eta}} \leq \min_i L_i + \sqrt{2\tilde{L} \ln N} + \ln N$$

## Tuning $\eta$ as a function of $T$

- ▶ trivially  $\min_i L_i \leq T$ , yielding

$$L_{\text{Hedge}(\eta)} \leq \min_i L_i + \sqrt{2T \ln N} + \ln N$$

## Tuning $\eta$ as a function of $T$

- ▶ trivially  $\min_i L_i \leq T$ , yielding

$$L_{\text{Hedge}(\eta)} \leq \min_i L_i + \sqrt{2T \ln N} + \ln N$$

- ▶ per iteration we get:

$$\frac{L_{\text{Hedge}(\eta)}}{T} \leq \min_i \frac{L_i}{T} + \sqrt{\frac{2 \ln N}{T}} + \frac{\ln N}{T}$$

## How good is this bound?

- ▶ **Very good!** There is a closely matching lower bound!

## How good is this bound?

- ▶ **Very good!** There is a closely matching lower bound!
- ▶ There exists a stochastic adversarial strategy such that with high probability for **any** hedging strategy **S** after **T** trials

$$L_S - \min_i L_i \geq (1 - o(1))\sqrt{2T \ln N}$$

## How good is this bound?

- ▶ **Very good!** There is a closely matching lower bound!
- ▶ There exists a stochastic adversarial strategy such that with high probability for **any** hedging strategy **S** after **T** trials

$$L_S - \min_i L_i \geq (1 - o(1))\sqrt{2T \ln N}$$

- ▶ The adversarial strategy is random, extremely simple, and does not depend on the hedging strategy!

## The adversarial strategy

- ▶ Adversary sets each loss  $\ell_i^t$  independently at random to 0 or 1 with equal probabilities (1/2, 1/2).



## The adversarial strategy

- ▶ Adversary sets each loss  $\ell_i^t$  independently at random to 0 or 1 with equal probabilities (1/2, 1/2).
- ▶ Obviously, nothing to learn !  
 $L_S \approx T/2$ .

## The adversarial strategy

- ▶ Adversary sets each loss  $\ell_i^t$  independently at random to 0 or 1 with equal probabilities (1/2, 1/2).
- ▶ Obviously, nothing to learn !  
 $L_S \approx T/2.$
- ▶ On the other hand  $\min_i L_i \approx T/2 - \sqrt{2T \ln N}$

## The adversarial strategy

- ▶ Adversary sets each loss  $\ell_i^t$  independently at random to 0 or 1 with equal probabilities (1/2, 1/2).
- ▶ Obviously, nothing to learn !  
 $L_S \approx T/2$ .
- ▶ On the other hand  $\min_i L_i \approx T/2 - \sqrt{2T \ln N}$
- ▶ The difference  $L_S - \min_i L_i$  is due to unlearnable random fluctuations!

## The adversarial strategy

- ▶ Adversary sets each loss  $\ell_i^t$  independently at random to 0 or 1 with equal probabilities (1/2, 1/2).
- ▶ Obviously, nothing to learn !  
 $L_S \approx T/2$ .
- ▶ On the other hand  $\min_i L_i \approx T/2 - \sqrt{2T \ln N}$
- ▶ The difference  $L_S - \min_i L_i$  is due to unlearnable random fluctuations!
- ▶ Detailed proof quite involved. See games paper.

## Summary

- ▶ Given learning rate  $\eta$  the **Hedge( $\eta$ )** algorithm satisfies

$$L_{\text{Hedge}(\eta)} \leq \frac{\ln N + \eta L_i}{1 - e^{-\eta}}$$

## Summary

- ▶ Given learning rate  $\eta$  the **Hedge**( $\eta$ ) algorithm satisfies

$$L_{\text{Hedge}(\eta)} \leq \frac{\ln N + \eta L_i}{1 - e^{-\eta}}$$

- ▶ Setting  $\eta \approx \sqrt{\frac{2 \ln N}{T}}$  guarantees

$$L_{\text{Hedge}(\eta)} \leq \min_i L_i + \sqrt{2T \ln N} + \ln N$$

## Summary

- ▶ Given learning rate  $\eta$  the **Hedge**( $\eta$ ) algorithm satisfies

$$L_{\text{Hedge}(\eta)} \leq \frac{\ln N + \eta L_i}{1 - e^{-\eta}}$$

- ▶ Setting  $\eta \approx \sqrt{\frac{2 \ln N}{T}}$  guarantees

$$L_{\text{Hedge}(\eta)} \leq \min_i L_i + \sqrt{2T \ln N} + \ln N$$

- ▶ A trivial random data, in which there is nothing to be learned forces **any** algorithm to suffer this total loss

## Some loose threads

- ▶ Total Loss of best action usually scales linearly with time  $T$ , but we need to know the **horizon**  $T$  ahead of time to choose  $\eta$  correctly.



## Some loose threads

- ▶ Total Loss of best action usually scales linearly with time  $T$ , but we need to know the **horizon**  $T$  ahead of time to choose  $\eta$  correctly.
- ▶  $T$  is tight only when the loss of experts at each iteration is either 0 or 1. If the loss of the best expert is  $o(T)$  then we would like to have a tighter bound.

## Some loose threads

- ▶ Total Loss of best action usually scales linearly with time  $T$ , but we need to know the **horizon**  $T$  ahead of time to choose  $\eta$  correctly.
- ▶  $T$  is tight only when the loss of experts at each iteration is either 0 or 1. If the loss of the best expert is  $o(T)$  then we would like to have a tighter bound.
- ▶ Observing only the loss of chosen action - the multi-armed bandit problem. Will get to that later in the course.

## Some loose threads

- ▶ Total Loss of best action usually scales linearly with time  $T$ , but we need to know the **horizon**  $T$  ahead of time to choose  $\eta$  correctly.
- ▶  $T$  is tight only when the loss of experts at each iteration is either 0 or 1. If the loss of the best expert is  $o(T)$  then we would like to have a tighter bound.
- ▶ Observing only the loss of chosen action - the multi-armed bandit problem. Will get to that later in the course.
- ▶ Register to the class on the google drive.