

# On-Line Learning of Non-Stationary Data

---

Manfred K. Warmuth  
UC, Santa Cruz

Joint work with:

Olivier Bousquet

Mark Herbster

---

- Motivate on-line learning
- Motivate relative loss bounds
- Halving Algorithm as example
- Loss Update
- Flavor of proof techniques
- Comparator on-line as well
- How to adapt the algs
- Future work

Loop

Get next instance

Predict

Get label

Incur loss

- Choose comparison class of predictors  
e.g. linear
- Goal:  
Do well compared to best off-line comparator
- No statistical assumptions on the data

	experts				predic tion	<i>true</i> <i>label</i>	loss
	$E_1$	$E_2$	$E_3$	$E_n$			
day 1	1	1	0	0	0	1	1
day 2	1	0	1	0	1	0	1
day 3	0	1	1	1	1	1	0
day $t$	$x_{t,1}$	$x_{t,2}$	$x_{t,3}$	$x_{t,n}$	$\hat{y}_t$	$y_t$	$ y_t - \hat{y}_t $

## Protocol of the Master Algorithm

For  $t = 1$  To  $T$  Do

Get instance  $\mathbf{x}_t \in \{0, 1\}^n$

Predict  $\hat{y}_t \in \{0, 1\}$

Get label  $y_t \in \{0, 1\}$

Incur loss  $|y_t - \hat{y}_t|$

- Learner against adversary

$$\sup_{\mathbf{x}_1} \inf_{\hat{y}_1} \sup_{y_1} \quad \sup_{\mathbf{x}_2} \inf_{\hat{y}_2} \sup_{y_2} \quad \dots \quad \sup_{\mathbf{x}_T} \inf_{\hat{y}_T} \sup_{y_T}$$

$$\sum_{t=1}^T L(y_t, \hat{y}_t) \quad - \quad \inf_{\mathbf{f} \in \mathbf{C}} \left( \sum_{t=1}^T L(f(\mathbf{x}_t), y_t) \right)$$

total loss of  
on-line  
algorithm

total loss of  
off-line  
algorithm

- $\mathbf{C}$  is comparison class
- Minimax algorithm usually intractable

# What kind of performance can we expect ?

---

5

- $L_{1..T,A}$  be the total loss of algorithm  $A$
- $L_{1..T,i}$  be the total loss of  $i$ -th expert  $E_i$

- Form of bounds

$$\forall S : \quad L_{1..T,A} \leq \min_i (L_{1..T,i} + c \log n)$$

where  $c$  is constant

- Bounds the loss of the algorithm **relative to** the loss of best expert

- Master algorithm predicts with weighted average

$$\hat{y}_t = \mathbf{v}_t \cdot \mathbf{x}_t$$

- The weights are updated according to the Loss Update

$$v_{t+1,i} := \frac{v_{t,i} e^{-\eta L_{t,i}}}{\text{normaliz.}}$$


where  $L_{t,i}$  is loss of expert  $i$  in trial  $t$

→ Weighted Majority Algorithm

[LW89]

→ Generalized by Vovk

[Vovk90]



- Off-line alg. **partitions** sequence into sections and chooses best expert in each section
- Goal:  
Do well compared to best off-line partition
- Problem:  
Loss Update **learns too well**  
and does **not recover fast enough**



- Predict  $\hat{y}_t = \mathbf{v}_t \cdot \mathbf{x}_t$
- Loss Update

$$v_{t,i}^m := \frac{v_{t,i} e^{-\eta L_{t,i}}}{\text{normaliz.}}$$

- Share Update


- Static Expert

$$\mathbf{v}_{t+1} = \mathbf{v}_t^m$$

- Fixed Share to Start Vector ( $\alpha \in [0, 1)$ )


$$\mathbf{v}_{t+1} = (1 - \alpha) \mathbf{v}_t^m + \alpha \mathbf{v}_0$$

where  $\mathbf{v}_0 = (\frac{1}{n}, \frac{1}{n}, \dots, \frac{1}{n})$




- Square loss, target outcome always 0, experts have predictions between 0 and 1/2 uniform for typical experts and restricted to  $[0, 0.12]$  for current best expert
- $T = 1400$  trials,  $n = 20000$  experts,  $k = 6$  shifts


- Tracks the best expert



# Weights of Fixed Share Alg.

17






- Variable Share to Start Vector

- Replace

$$\mathbf{v}_{t+1,i} = (1 - \alpha) \mathbf{v}_{t,i}^m + \alpha \frac{1}{n}$$

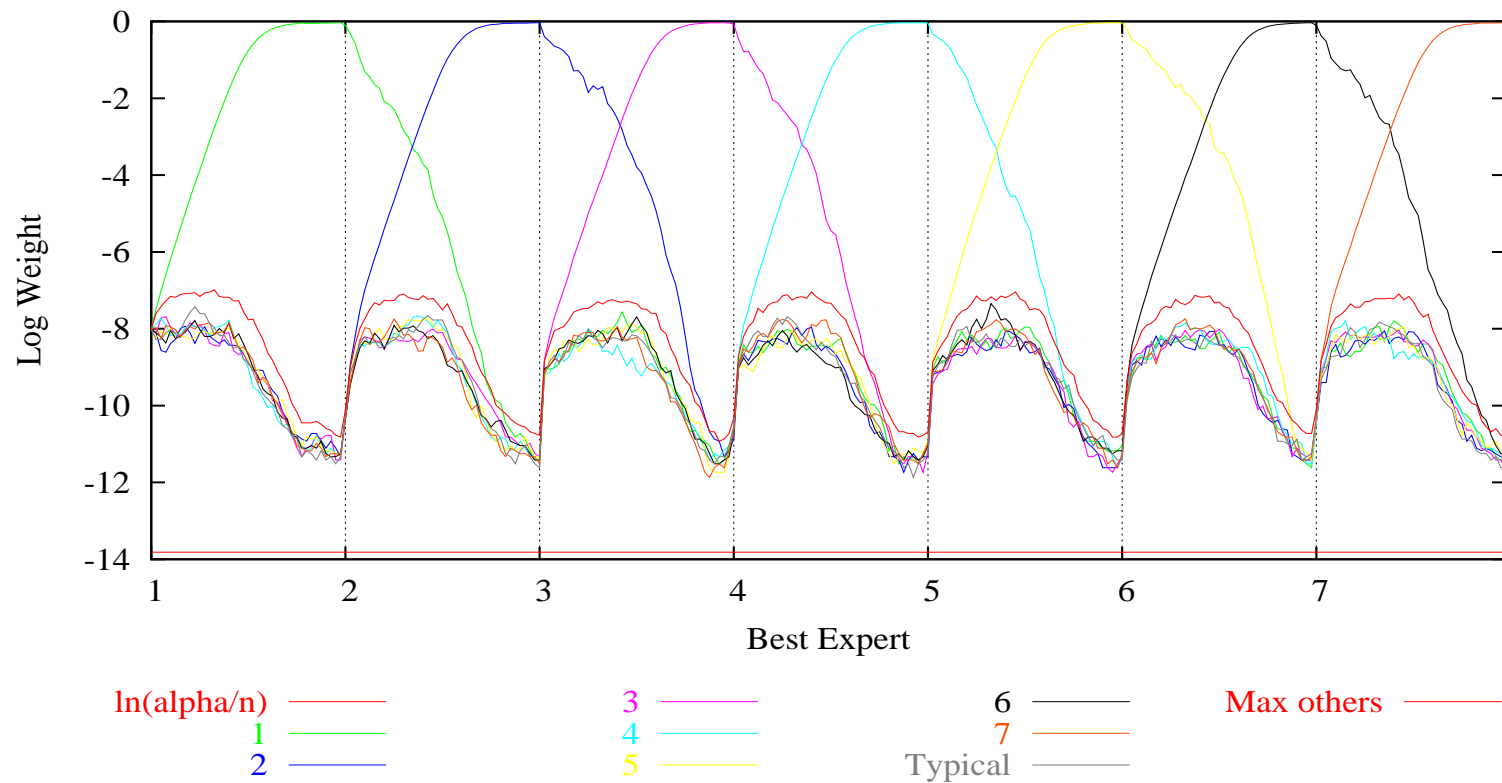
- by

$$\mathbf{v}_{t+1,i} = (1 - \alpha)^{L_{t,i}} \mathbf{v}_{t,i}^m + \left(1 - \sum_{i=1}^n (1 - \alpha)^{L_{t,i}}\right) \frac{1}{n} \text{ where } L_{t,i} \in [0, 1]$$



# Weights of Variable Share Alg.

21





- Recall Static Expert bound

$$L_{\text{Alg}}(S) \leq \min_i (L_i(S) + O(\log n))$$

- Comparison class: set of experts

- Bounds for Share Algs.

$$L_{\text{Alg}}(S) \leq \min_P (L_P(S) + O(\# \text{ of bits for } P))$$


- Comparison class: set of partitions
- # of bits for partitions with  $k$  shifts:

$$k \log n + \log \binom{T}{k}$$


- Number of possible experts  $n$  is large  $n \approx 10^6$
- Experts in partition from small subset of size  $m$   $m \approx 10$
- # of bits for partitions with  $k$  shifts:

$$\log \binom{n}{m} + k \log m + \log \binom{T}{k}$$


- Predict  $\hat{y}_t = \mathbf{v}_t \cdot \mathbf{x}_t$
- Loss Update  $v_{t,i}^m := \frac{v_{t,i} e^{-\eta L_{t,i}}}{\text{normaliz.}}$
- Mixing Update:  $\mathbf{v}_{t+1} = \sum_{q=0}^t \beta_{t+1,q} \mathbf{v}_q^m$
- Mixing scheme




FS to Start Vector



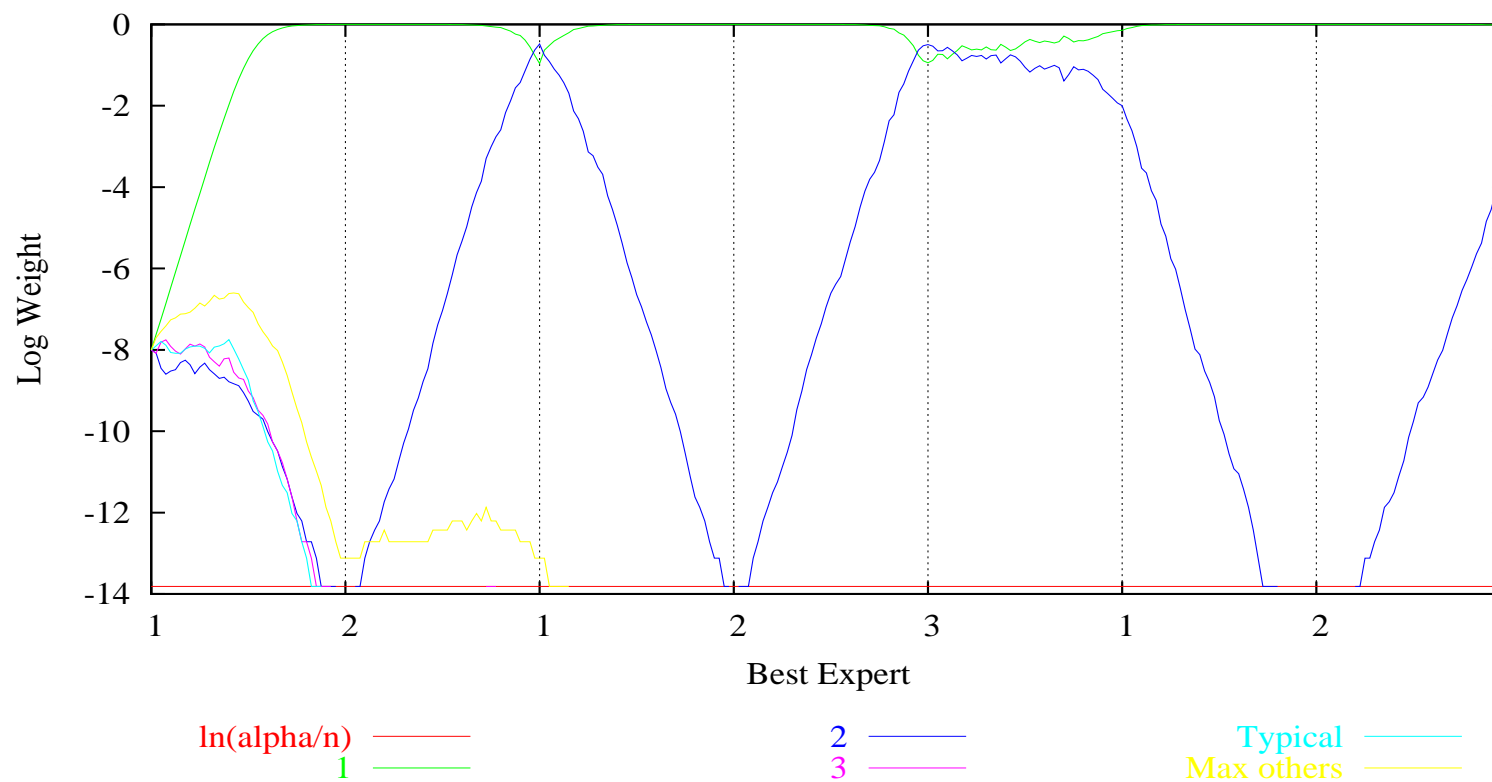
FS to Uniform Past




FS to Decaying Past




- Square loss, target outcome always 0, experts have predictions between 0 and 1/2 uniform for typical experts and restricted to  $[0, 0.12]$  for current best expert
- $T = 1400$  trials,  $n = 20000$  experts,  $k = 6$  shifts,  $m = 3$  experts in the small subset






# Fixed Share to Decaying Past - Log Weights

29




# More Experts Remembered

30







- Weights past good expert remain at higher level



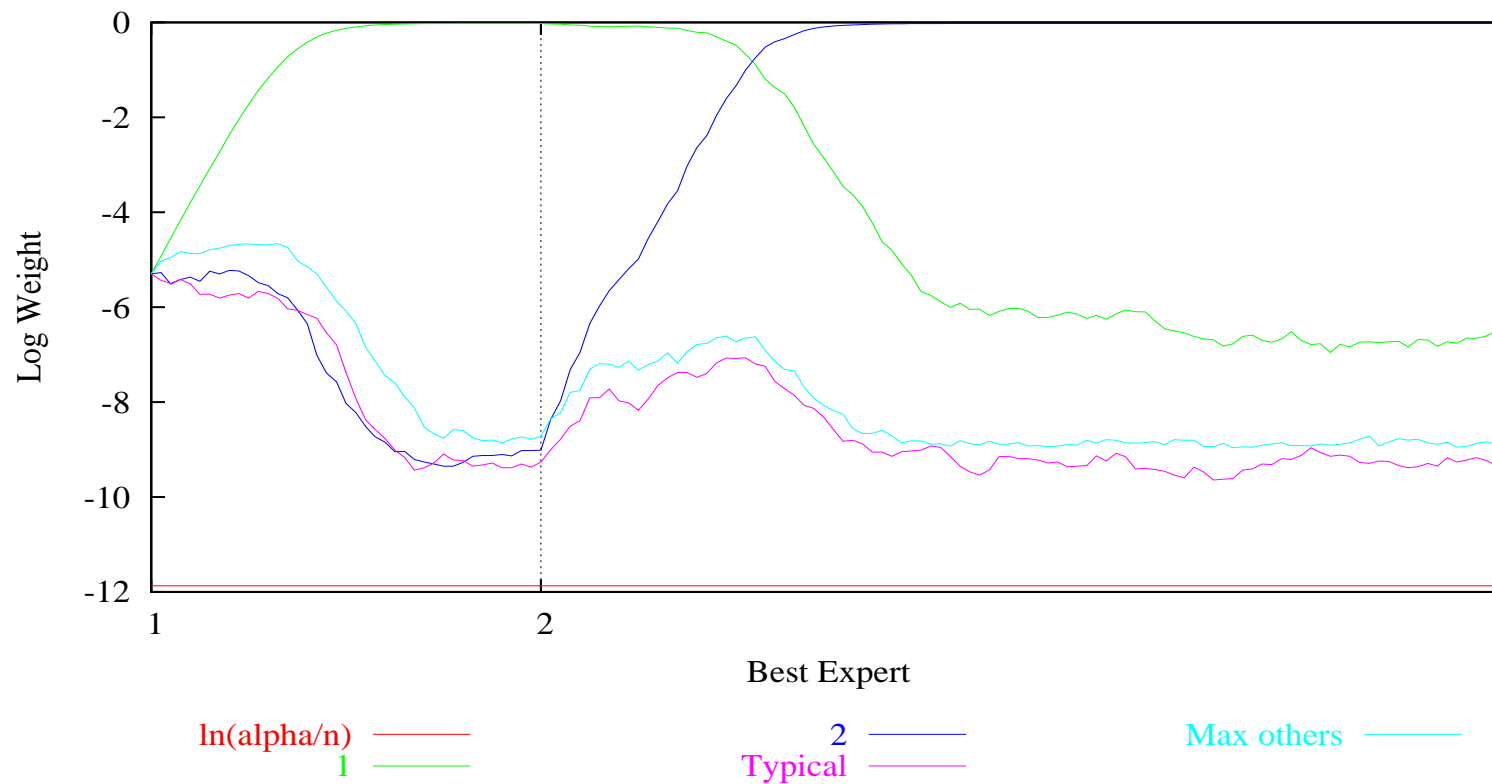
- No memory

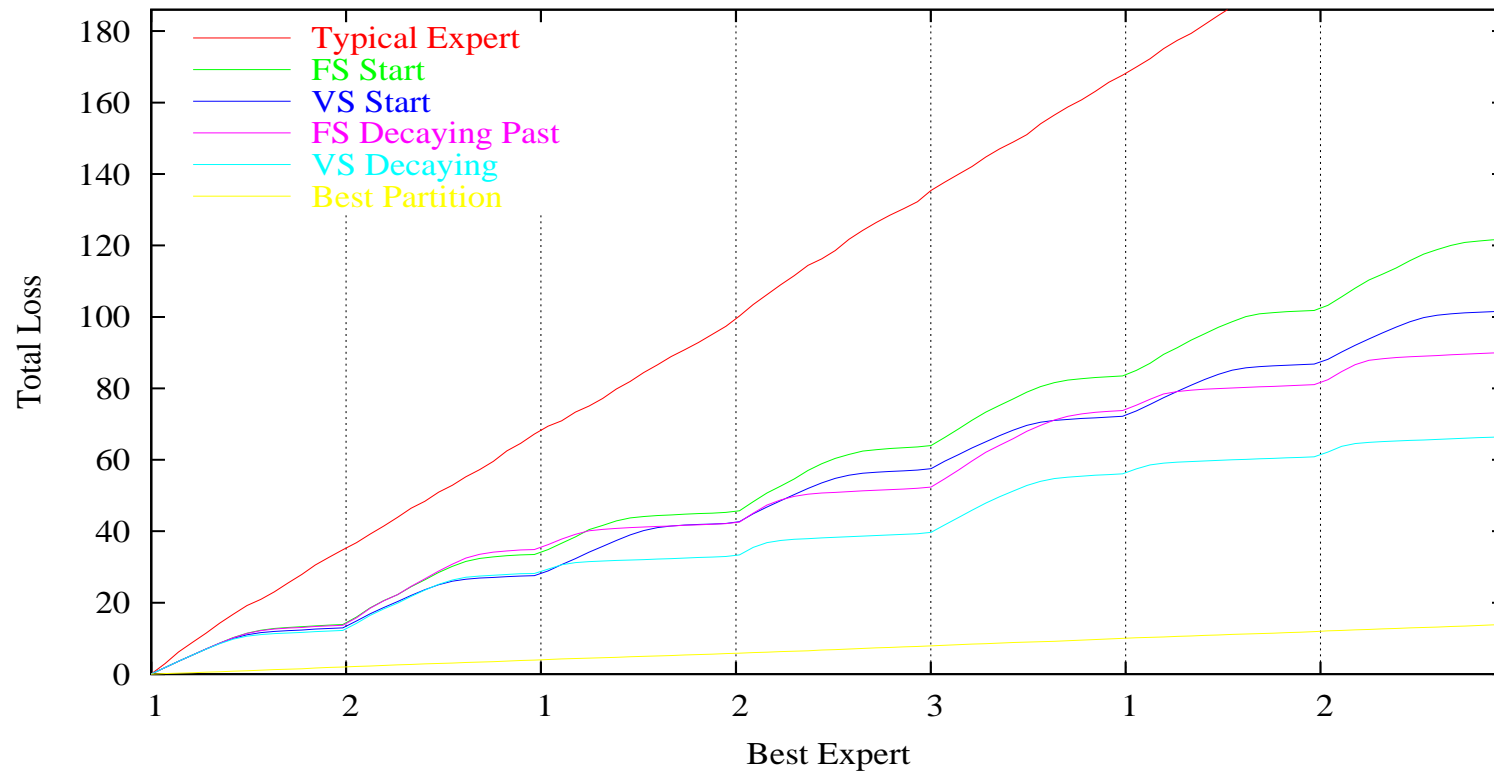




# Fixed Share to Decaying Past

34





- Bounds still have the form

$$L_{1..T,A} \leq \min_P (L_{1..T,P} + O(\# \text{ of bits for } P))$$

- Excess loss for naive alg.

$$O(\log \binom{n}{m} + k \log m + \log \binom{T}{k})$$

- Excess loss for Fixed Share to Decaying Past

$$O\left(m \log n + k \log m + 2 \log \binom{T}{k}\right)$$

→ Boundaries are encoded twice

→ Off-line problem NP-complete

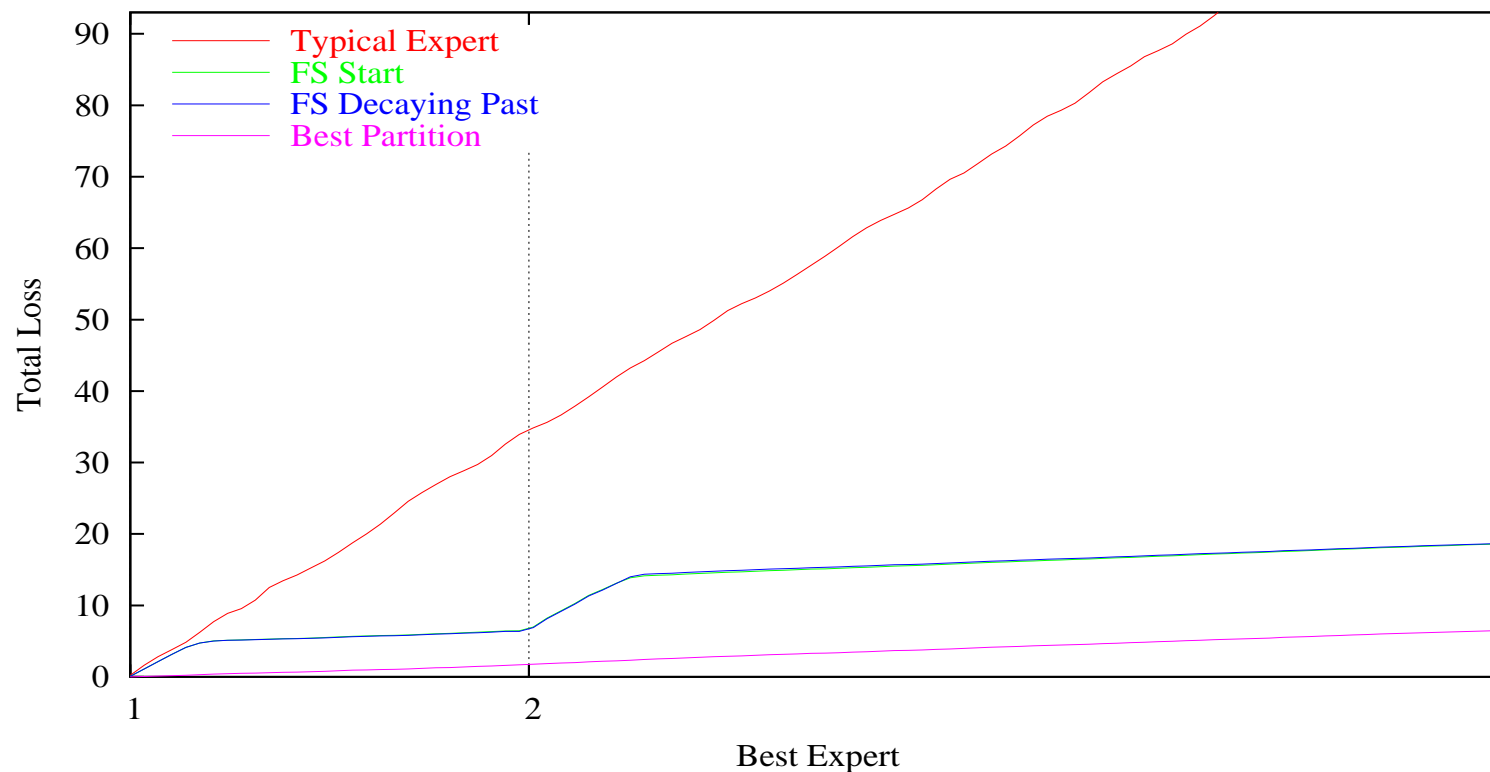
- Naive alg. has optimal bound - exponential storage
  - Fixed Share to Uniform Past -  $O(n)$  weights
  - Fixed Share to Decaying Past -  $O(nT)$  weights and better bound
- With tricks  $O(n \ln T)$  weights and essentially same bound

- What we need for bounds


$$\mathbf{v}_{t+1} = \beta_{t+1,q} \mathbf{v}_q^m, \text{ for } 0 \leq q \leq t \quad (*)$$

Mixing Update	$\mathbf{v}_{t+1} = \sum_{q=0}^t \beta_{t+1,q} \mathbf{v}_q^m$
Max Update	$\mathbf{v}_{t+1} = \frac{1}{\text{normaliz.}} \max_{q=0,\dots,t} \beta_{t+1,q} \mathbf{v}_q^m$
Projection Update	$\mathbf{v}_{t+1} = \arg \min_{\mathbf{v} \in (*)} \Delta(\mathbf{v}, \mathbf{v}_t^m)$






- $T = 1400$  trials,  $n = 20000$  experts
- $k = 1$  shift (at trial 400),  $m = 2$  experts in the small subset




- Larger alpha gives better long-term memory

# Fixed Share to Start Vector - Log Weights


28





- $T = 3200$  trials,  $n = 20000$  experts
- $k = 30$  shifts (every 200 and 50 trials),  $m = 2$  experts in the small subset

- The memory from many short sections accumulates



- Bayesian interpretation
- Variable share
- Lower bounds
- Automatic tuning
- Mixing Update works for EG family
- Connections to Universal Coding
- Applications
  - Load balancing
  - Switching between a few users
  - Segmentation